



计算机视觉表征与识别

Chapter 2: Images and Filtering

王利民

媒体计算课题组

<http://mcg.nju.edu.cn/>



Overview



- What is an image?
- Image formation: light and color
- Image transformation and filtering
- Image noise and image smoothing
- Convolution operation
- Non-linear filter: Media filter, Bilateral Filter



Images as matrices



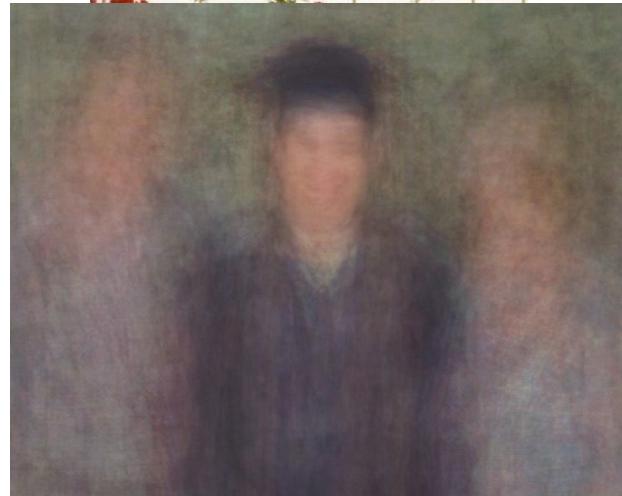
Result of averaging 100 similar snapshots



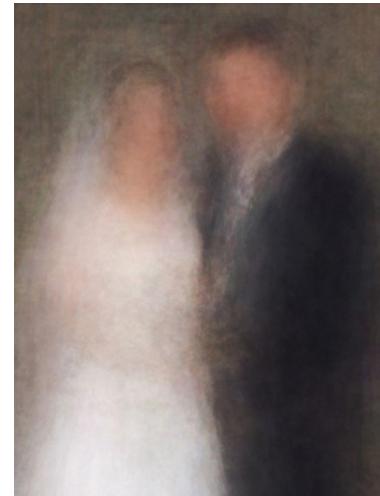
Little Leaguer



Kids with Santa



The Graduate

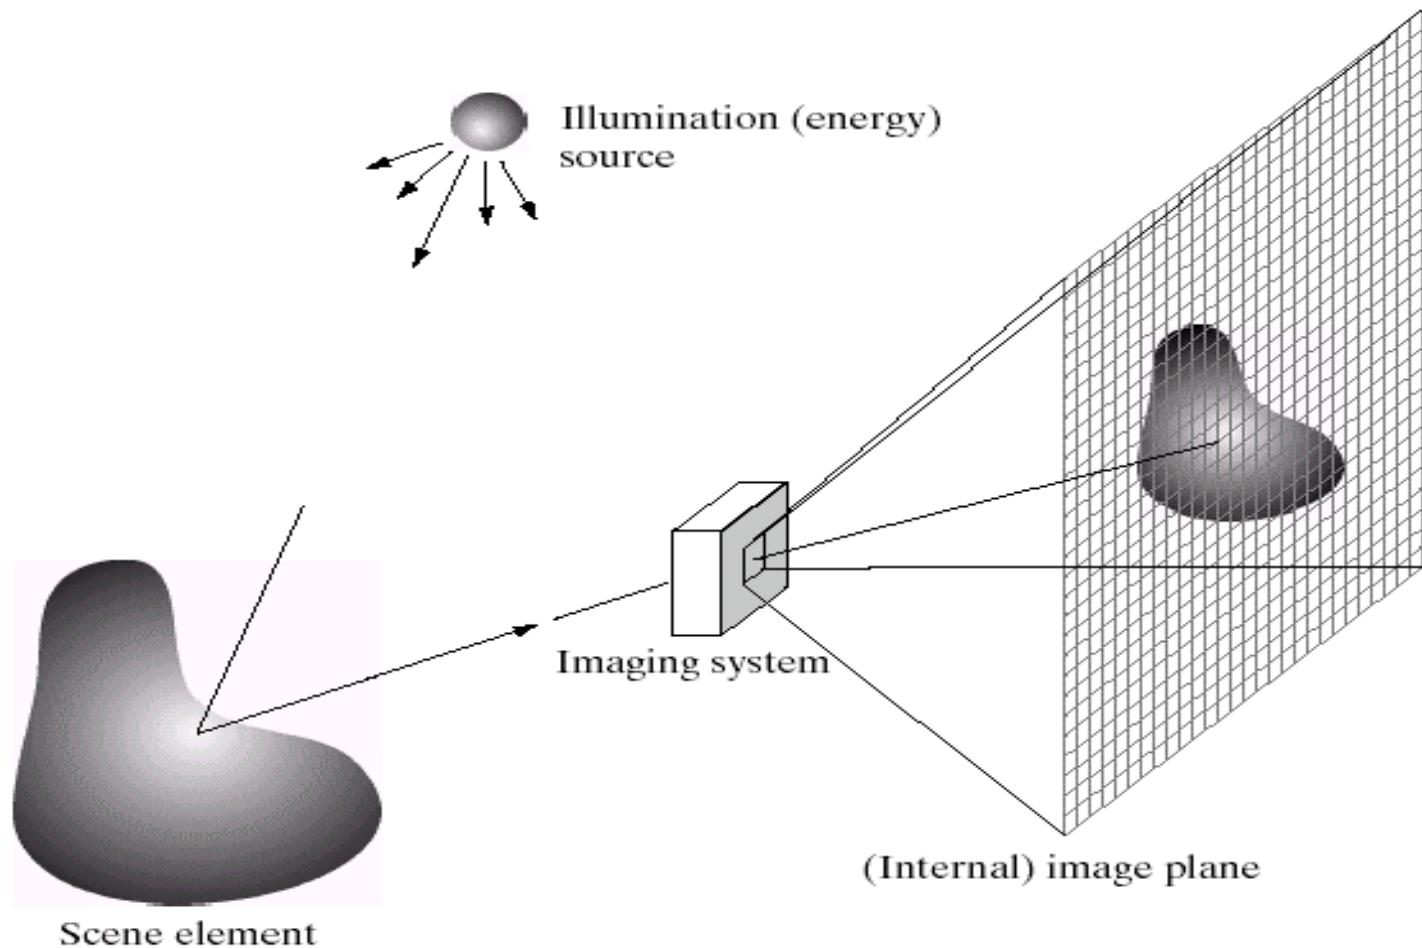


Newlyweds

From: *100 Special Moments*, by Jason Salavon (2004)
<http://salavon.com/SpecialMoments/SpecialMoments.shtml>



How light is recorded





Digital camera

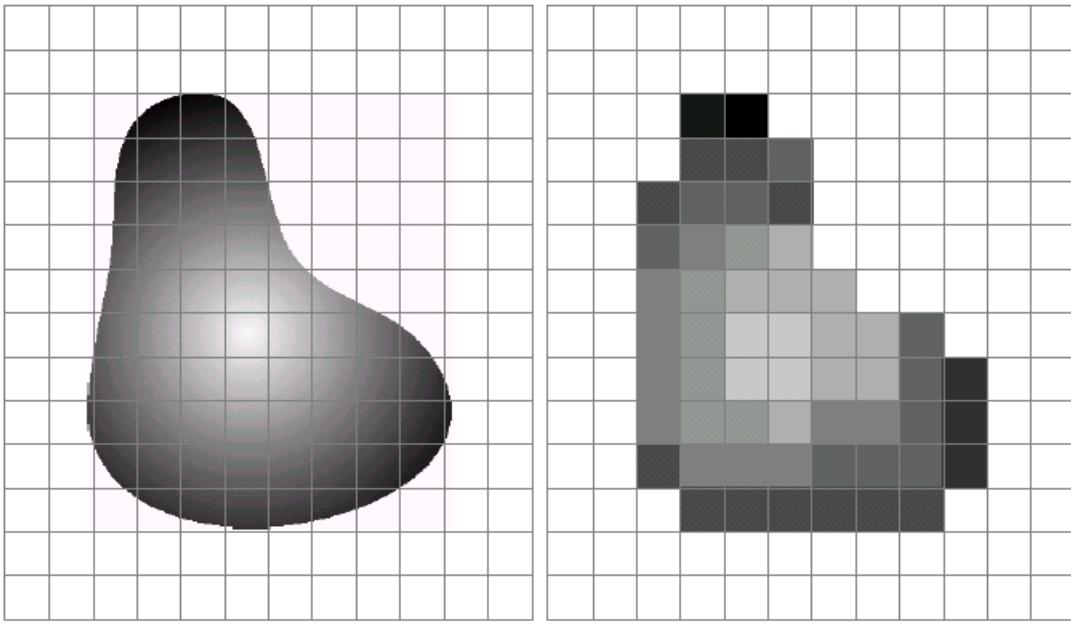


A digital camera replaces film with a sensor array

- Each cell in the array is light-sensitive diode that converts photons to electrons
- <http://electronics.howstuffworks.com/digital-camera.htm>



Sensor array



a b

FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

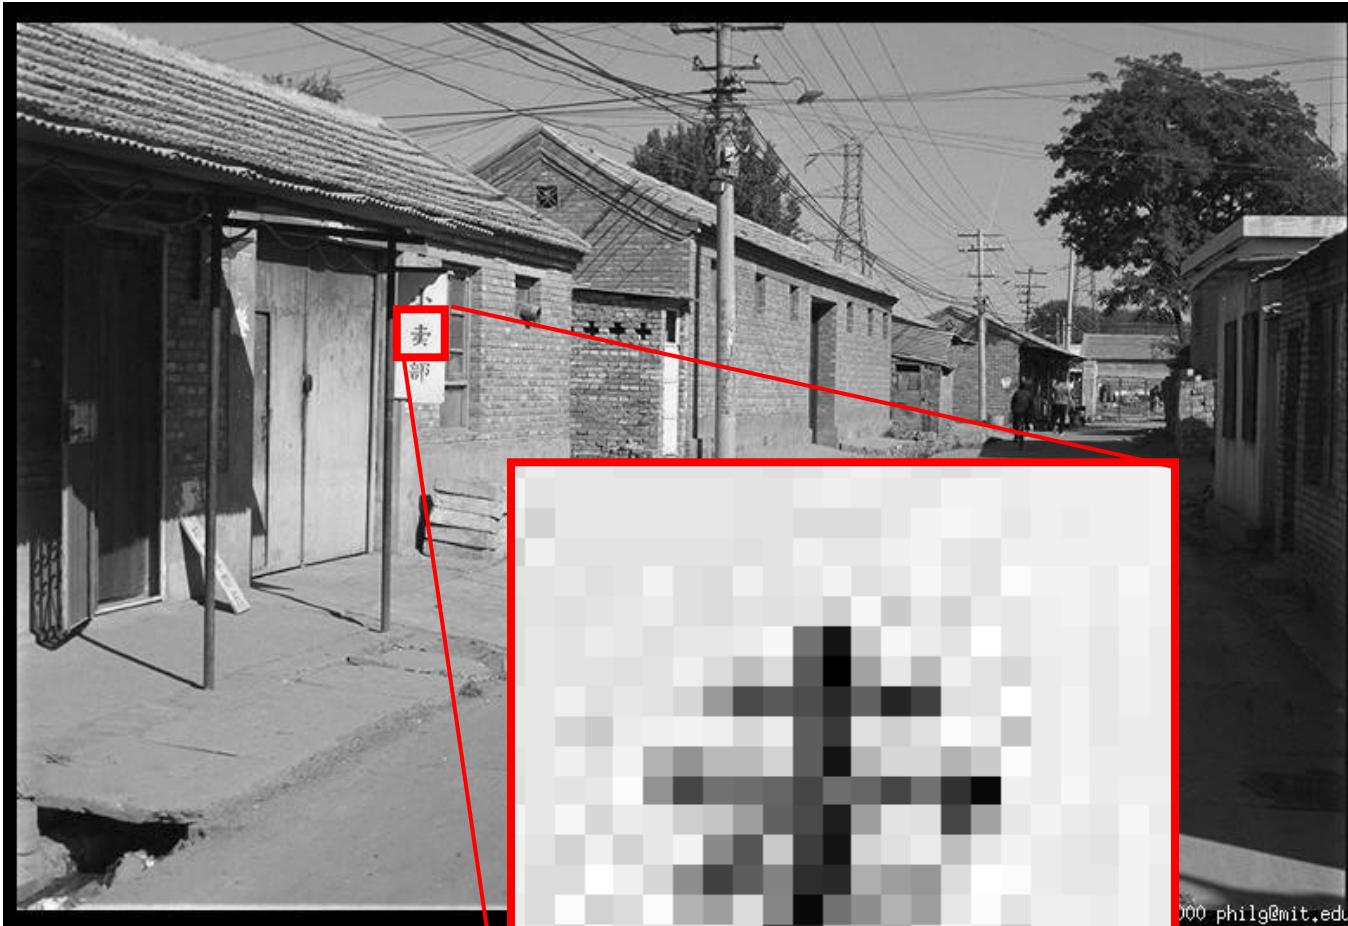


CMOS sensor

Each sensor cell records amount of light coming in at a small range of orientations



The raster image (pixel matrix)

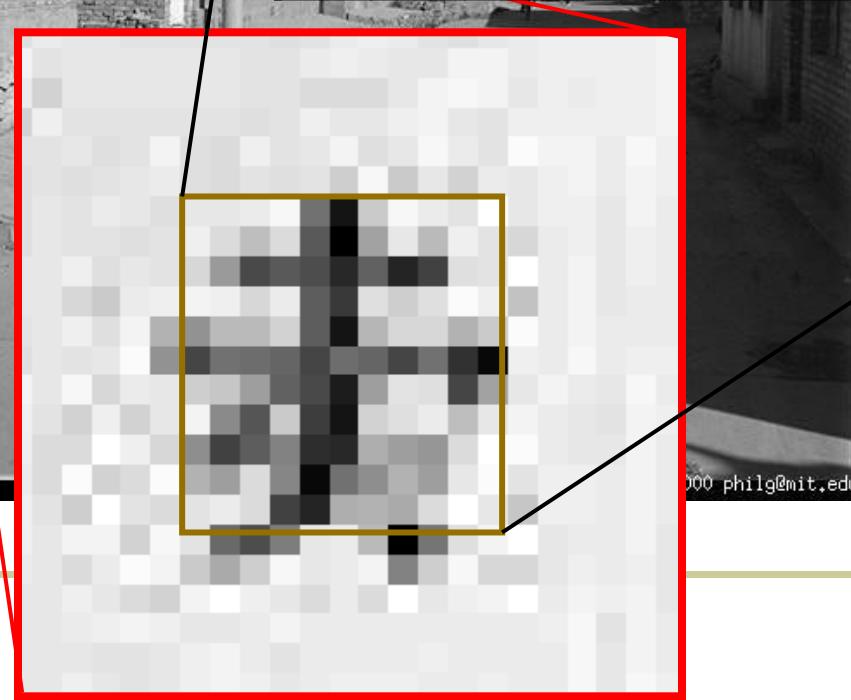




The raster image (pixel matrix)



0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93





What determines a pixel's intensity





Overview



- What is an image?
- **Image formation: light and color**
- Image transformation
- Image noise and image smoothing
- Convolution operation
- Media filter

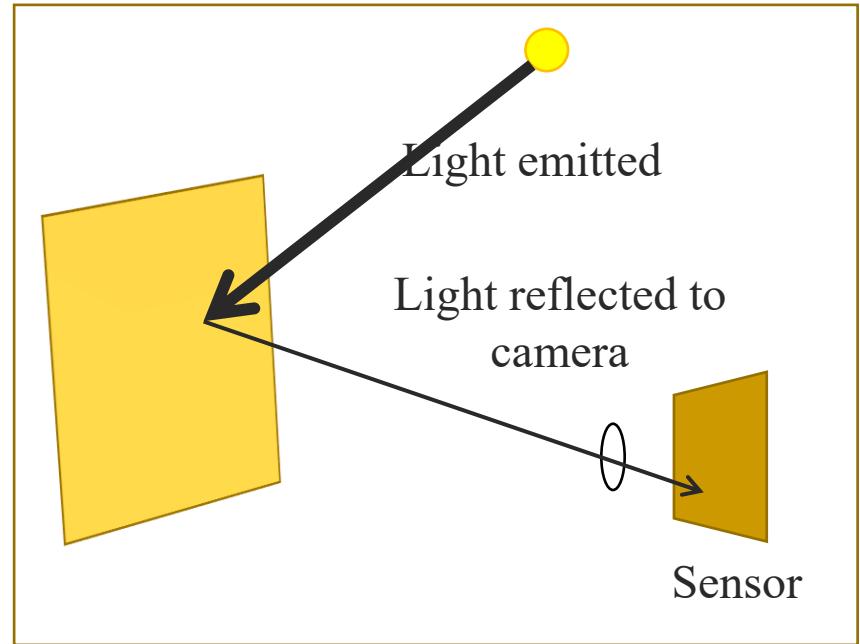


How much light is recorded



■ Major factors

- Illumination strength and direction
- Surface geometry
- Surface material
- Nearby surfaces
- Camera gain/exposure

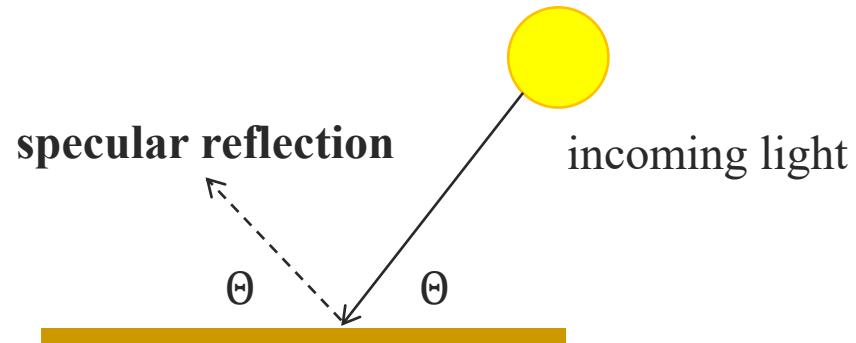




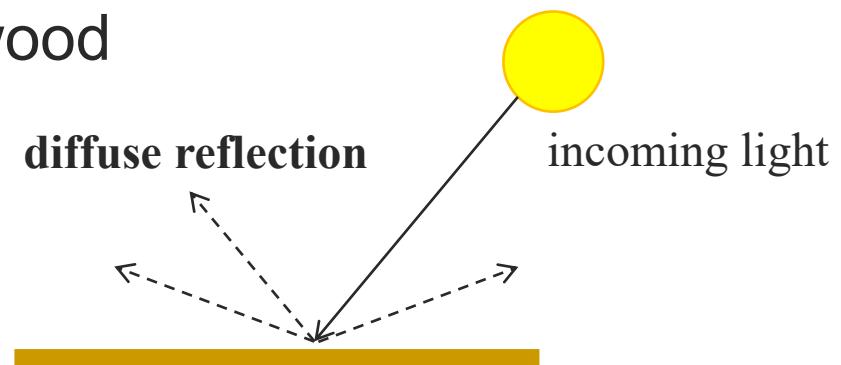
Basic models of reflection



- Specular: light bounces off at the incident angle
 - E.g., mirror



- Diffuse: light scatters in all directions
 - E.g., brick, cloth, rough wood

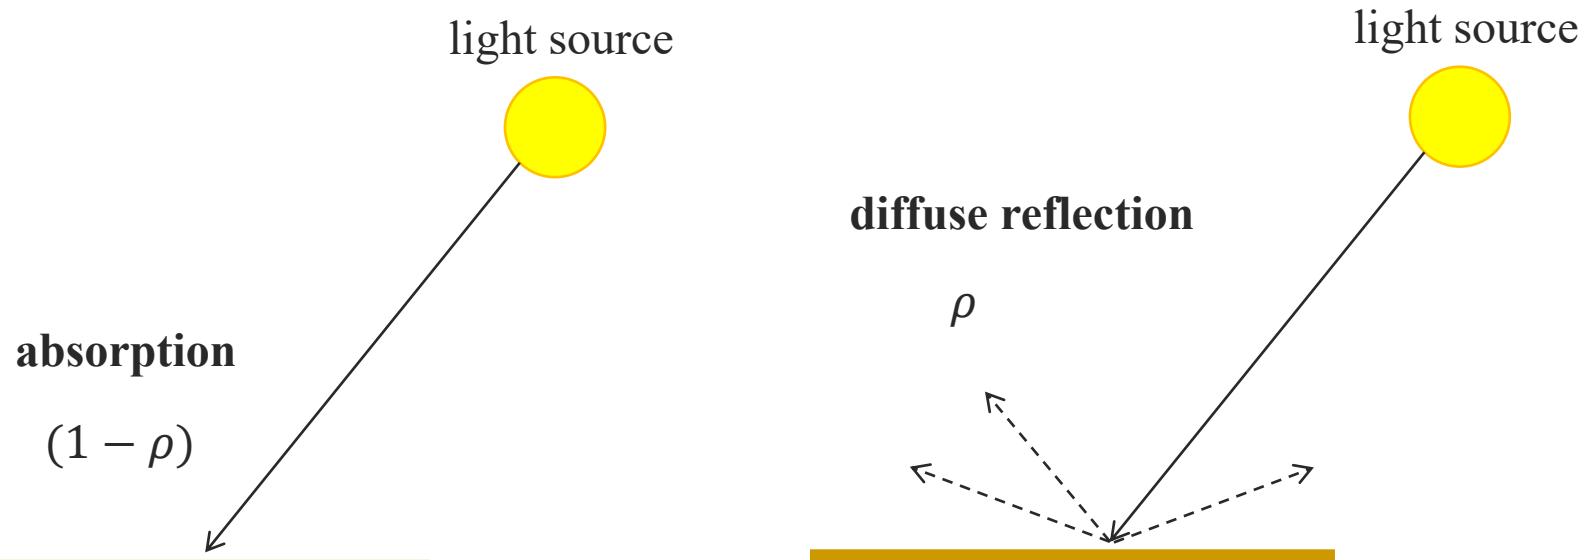




Lambertian reflectance model (理想散射)



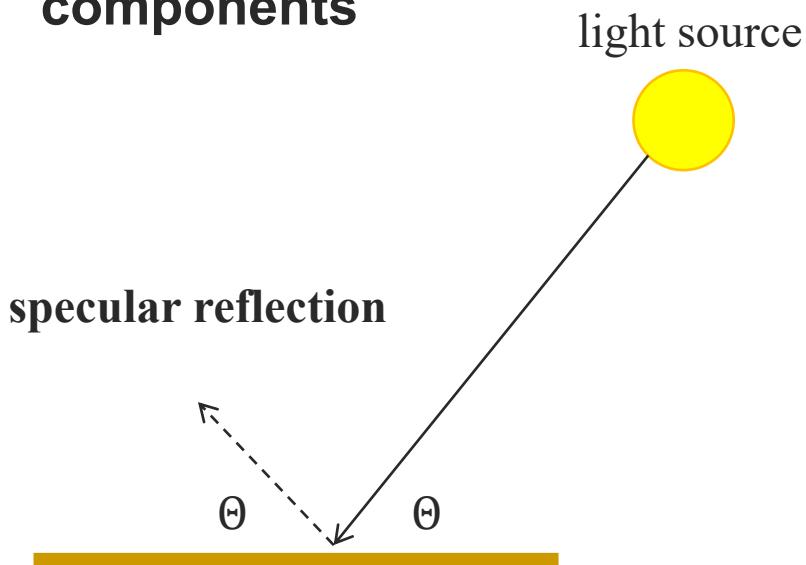
- Some light is absorbed (function of albedo ρ)
- Remaining light is scattered (diffuse reflection)
- Examples: soft cloth, concrete, matte paints





Specular Reflection

- Reflected direction depends on light orientation and surface normal
 - E.g., mirrors are fully specular
 - **Most surfaces can be modeled with a mixture of diffuse and specular components**



Flickr, by suzysputnik



Flickr, by piratejohnny



Most surfaces have both specular and diffuse components



- Specularity (高光) = spot where specular reflection dominates (typically reflects light source)

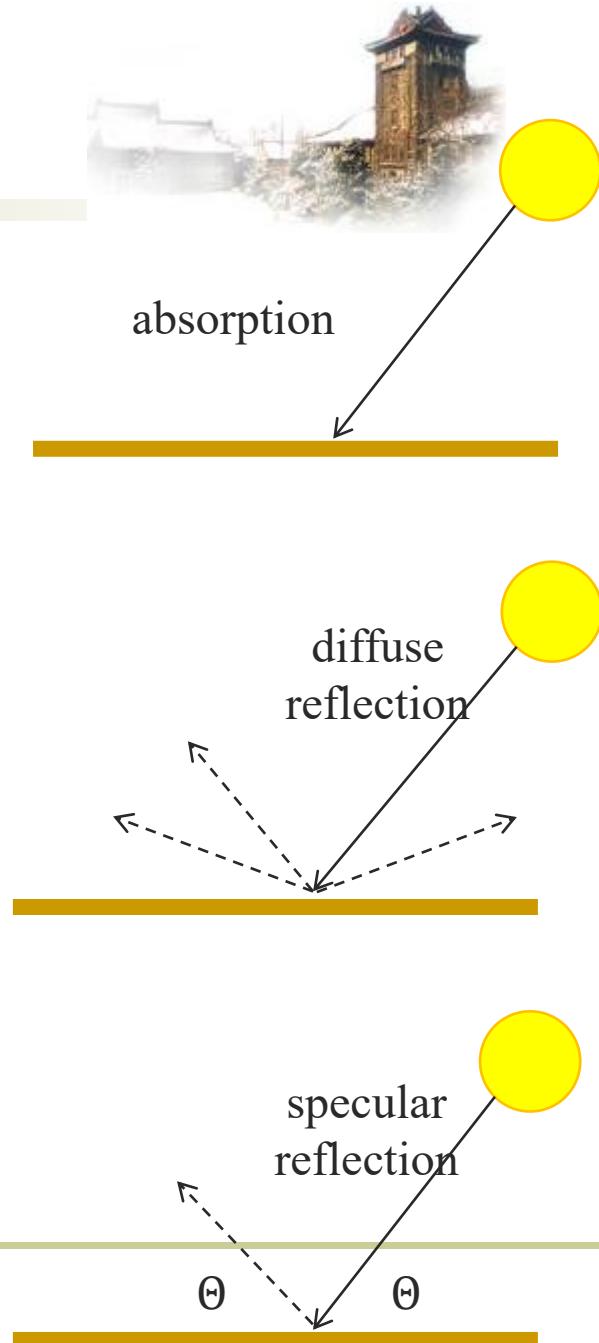


Typically, specular component is small



Summary

- When light hits a typical surface
 - Some light is absorbed ($1-\rho$)
 - More absorbed for low albedos
 - Some light is reflected diffusely
 - Independent of viewing direction
 - Some light is reflected specularly
 - Light bounces off (like a mirror), depends on viewing direction





Intensity and Surface Orientation



Intensity depends on illumination angle because less light comes in at oblique angles.

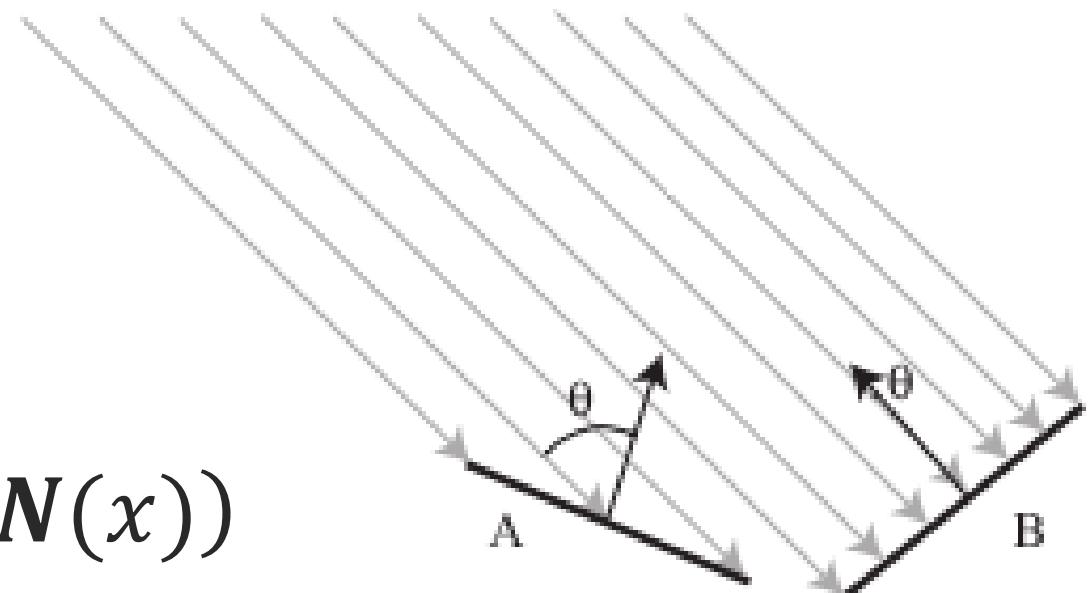
ρ = albedo

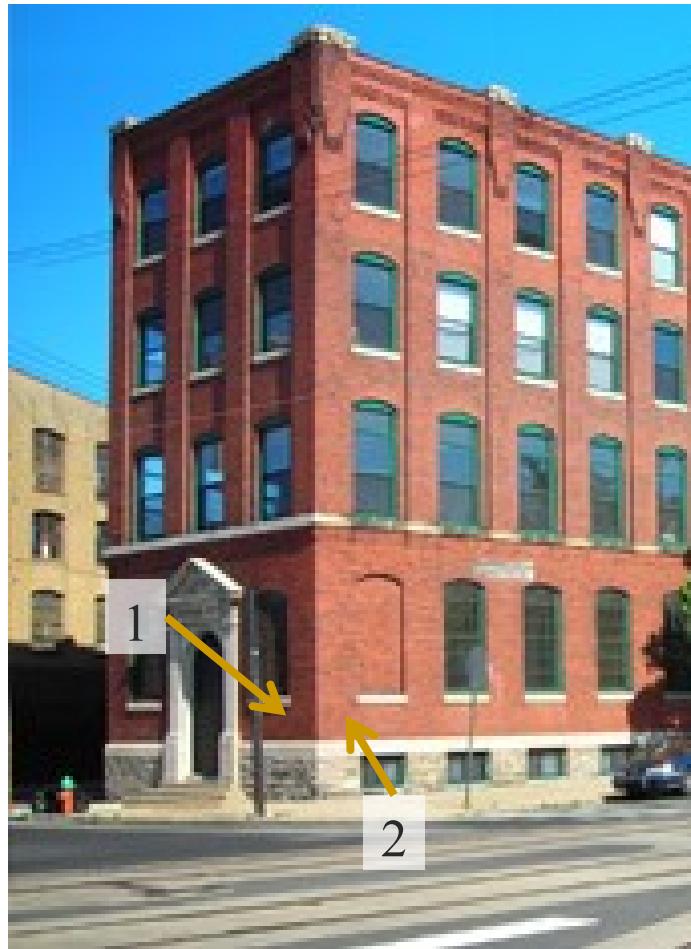
S = directional source

N = surface normal

I = reflected intensity

$$I(x) = \rho(x)(S \cdot N(x))$$





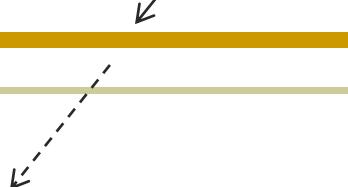
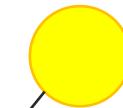


Other possible effects



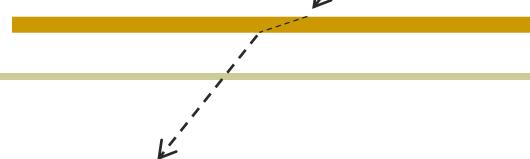
Transparency (透视)

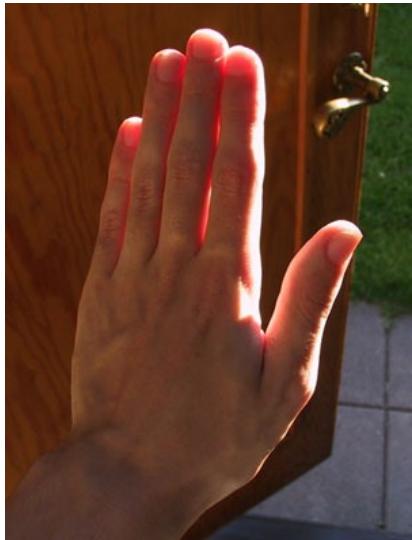
light source



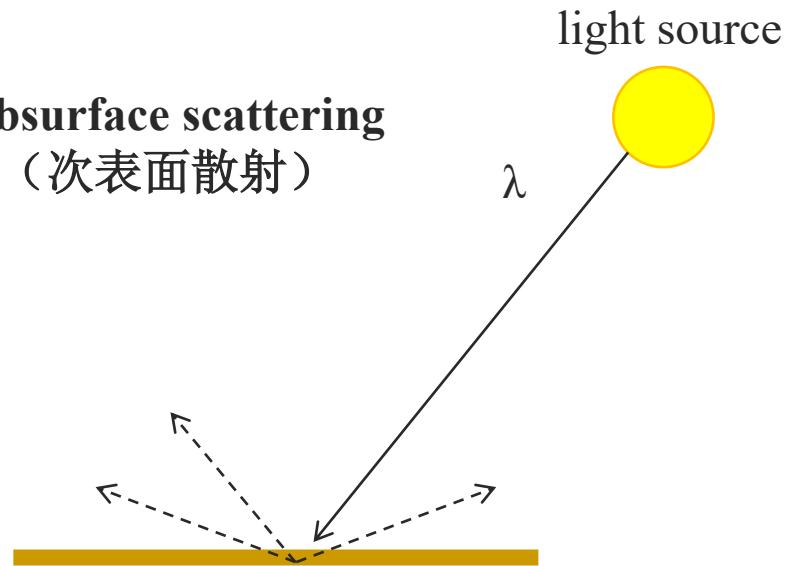
Refraction (折射)

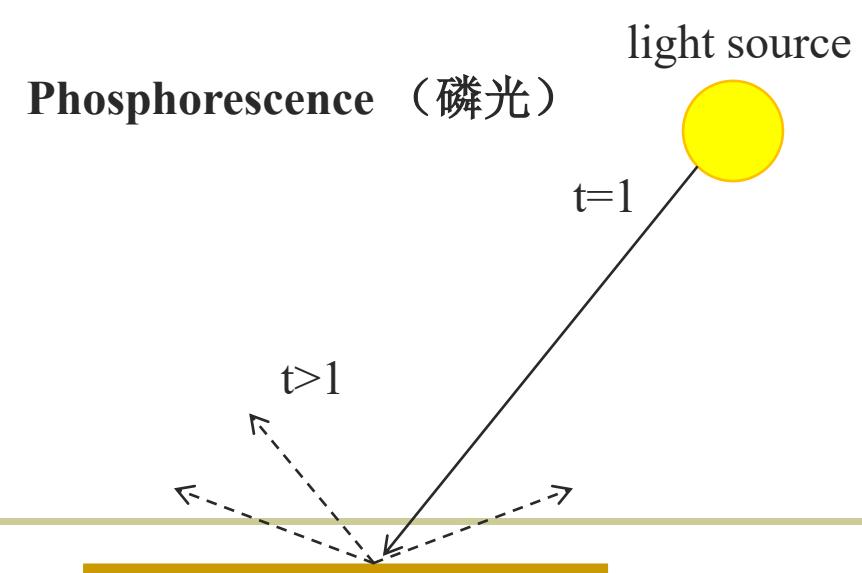
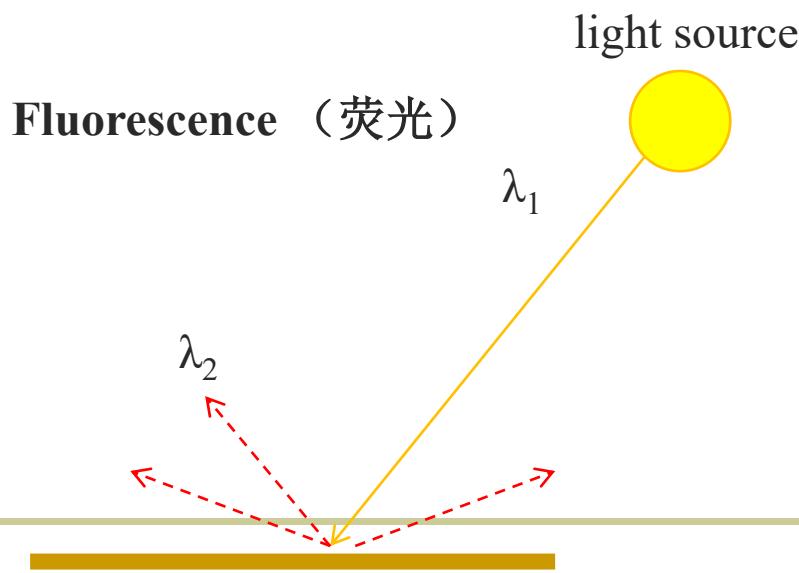
light source





subsurface scattering (次表面散射)



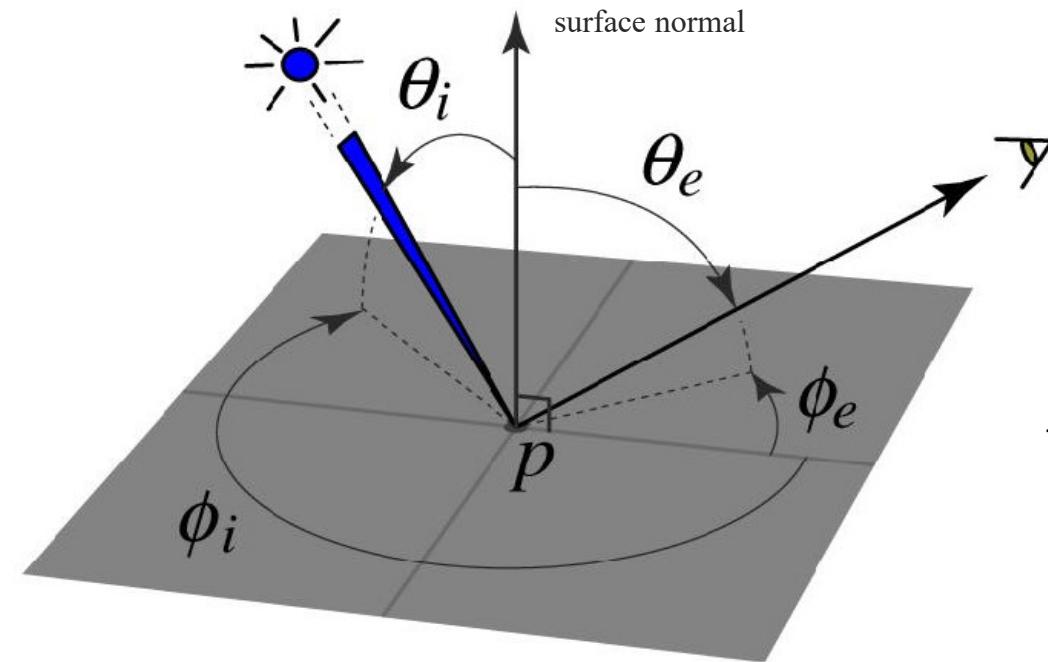




BRDF: Bidirectional Reflectance Distribution Function (双向反射分布函数)



- Model of local reflection
 - How bright a surface appears when viewed from one direction when light falls on it from another

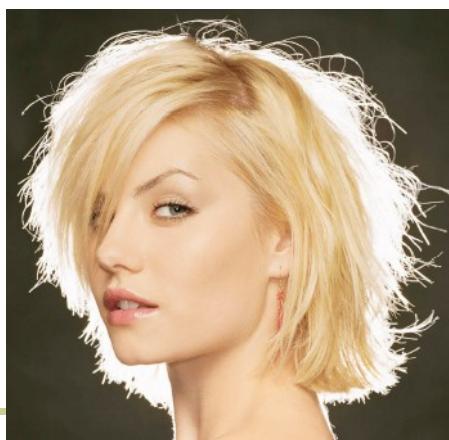


$$\rho(\theta_i, \phi_i, \theta_e, \phi_e; \lambda) =$$

$$\frac{L_e(\theta_e, \phi_e)}{E_i(\theta_i, \phi_i)} = \frac{L_e(\theta_e, \phi_e)}{L_i(\theta_i, \phi_i) \cos \theta_i d\omega}$$



BRDFs can be incredibly complicated

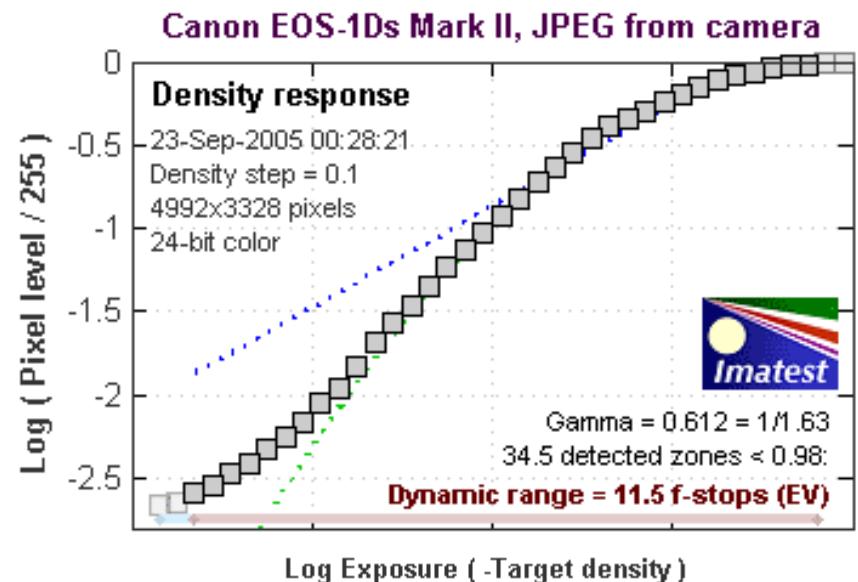
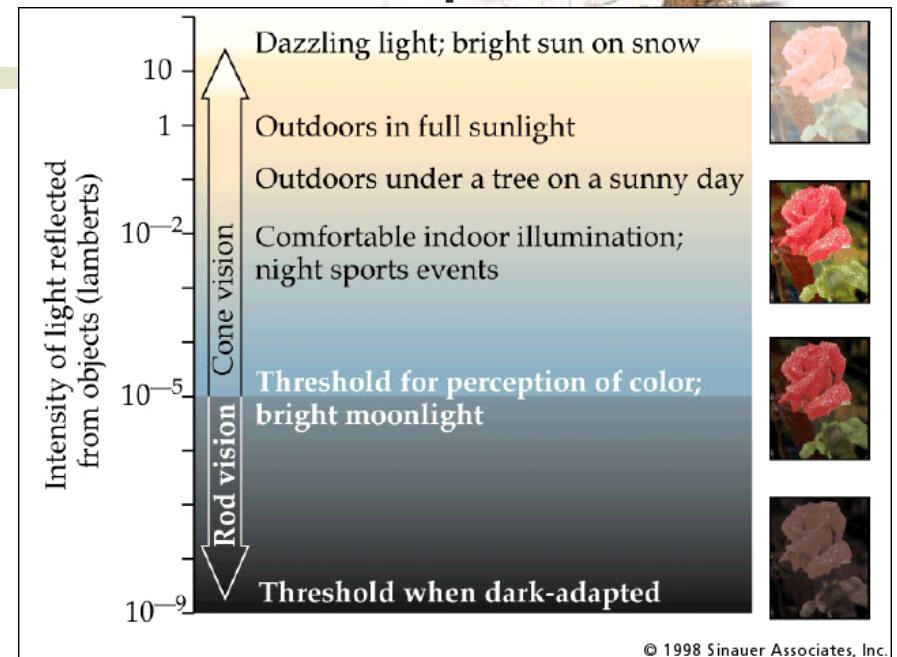




Dynamic range and camera response



- Typical scenes have a huge dynamic range
- Camera response is roughly linear in the mid range (15 to 240) but non-linear at the extremes
 - called saturation or undersaturation





Theorem of Perception



Perceived brightness $\sim =$ # photons



Theorem of Perception



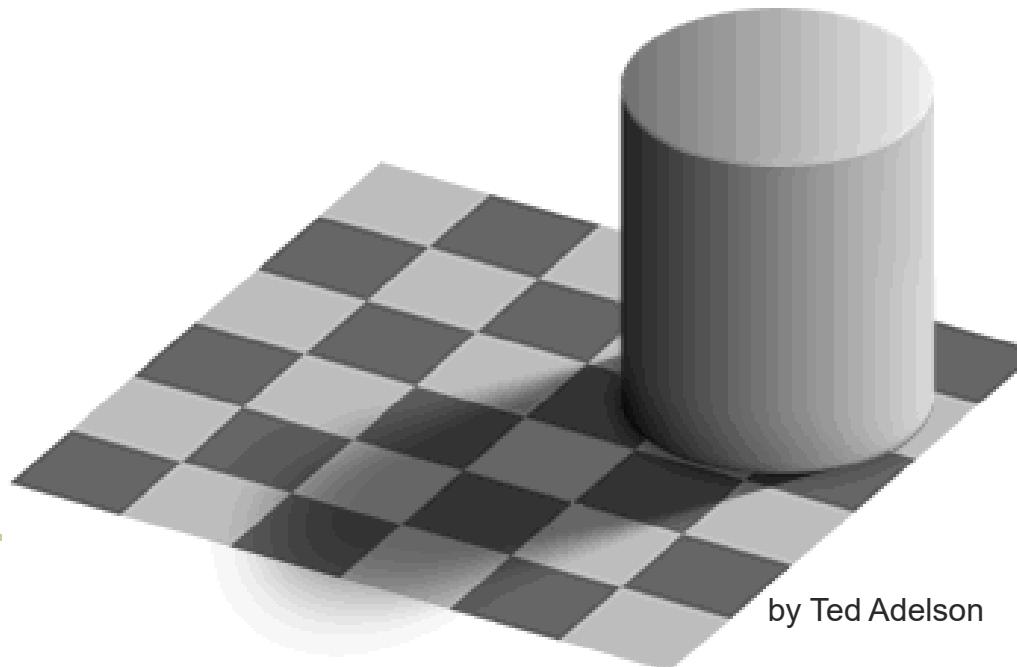
Perceived brightness $\sim =$ # photons



Brightness contrast



The apparent brightness depends on the surrounding region



by Ted Adelson



Brightness constancy

A surface looks the same under widely varying lighting conditions

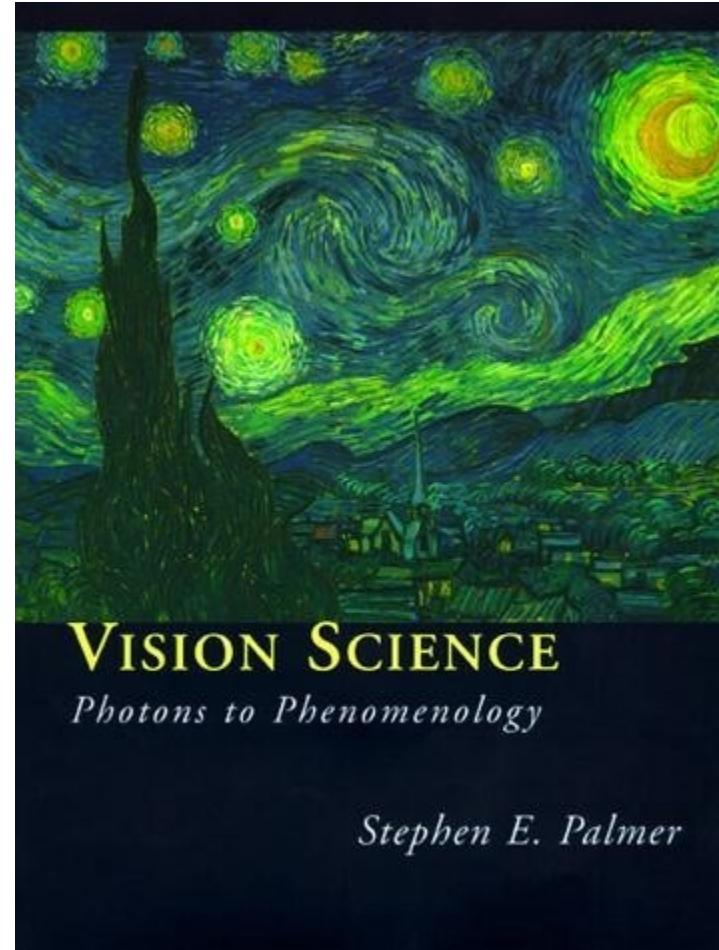




What is color?



- Color is the result of interaction between physical light in the environment and our visual system
- Color is a psychological property of our visual experiences when we look at objects and lights, *not* a physical property of those objects or lights
(S. Palmer, *Vision Science: Photons to Phenomenology*)

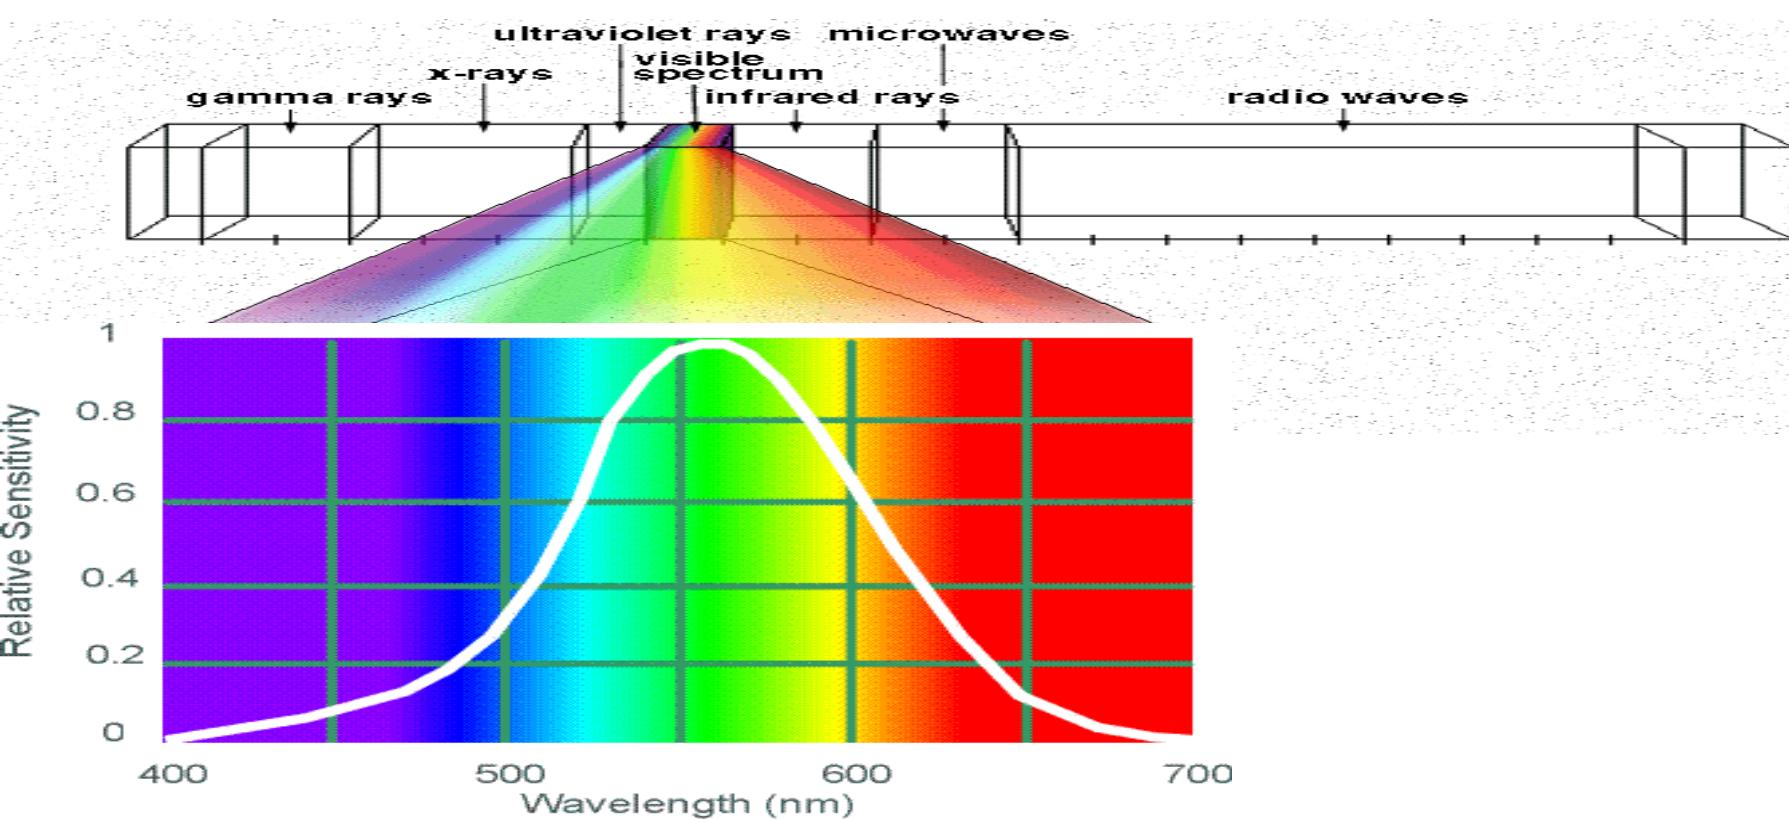




Color



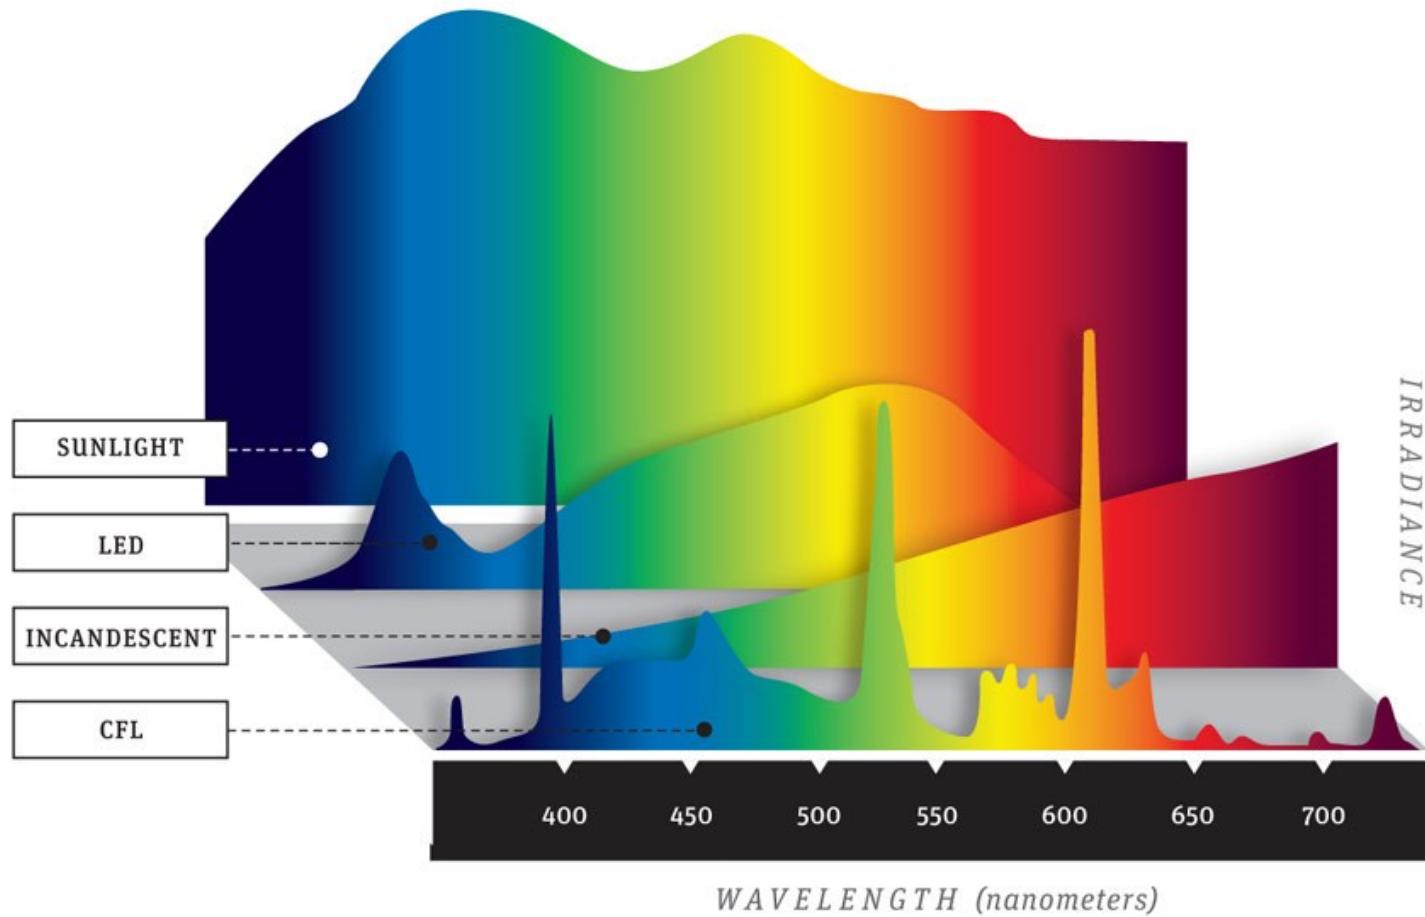
Light is composed of a spectrum of wavelengths



Human Luminance Sensitivity Function



Spectra of light sources

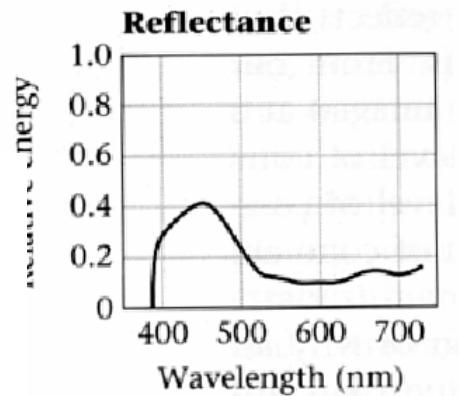
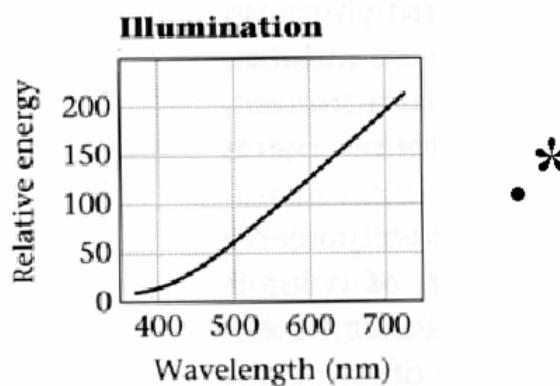




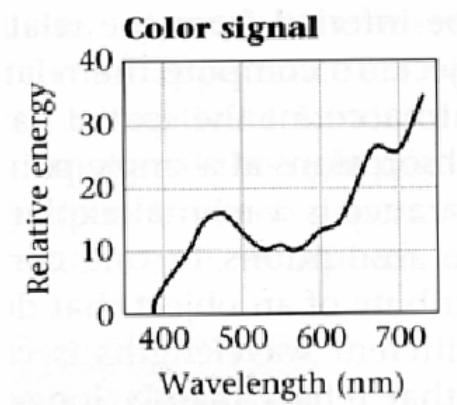
Interaction of light and surfaces



- Reflected color is the result of interaction of light source spectrum with surface reflectance

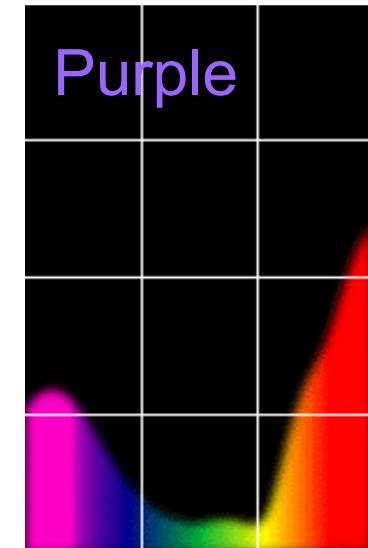
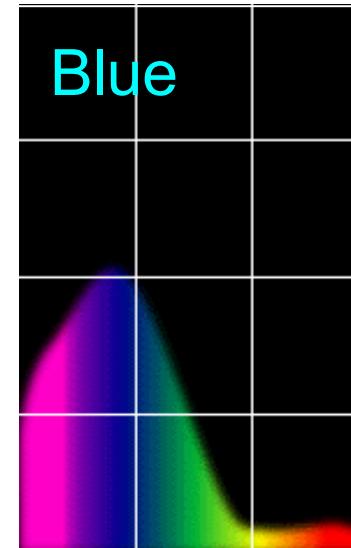
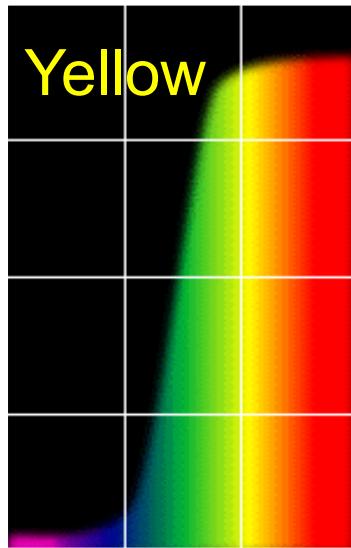
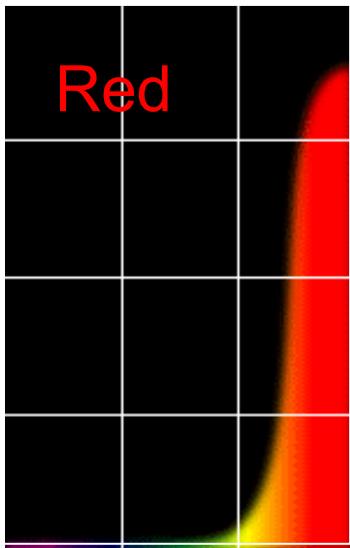


=





Some examples of the reflectance spectra of surfaces

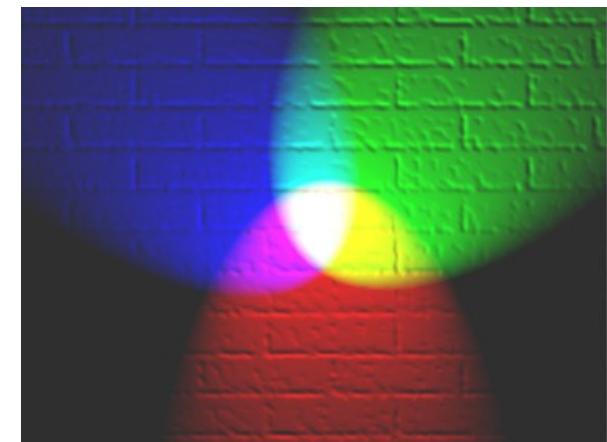
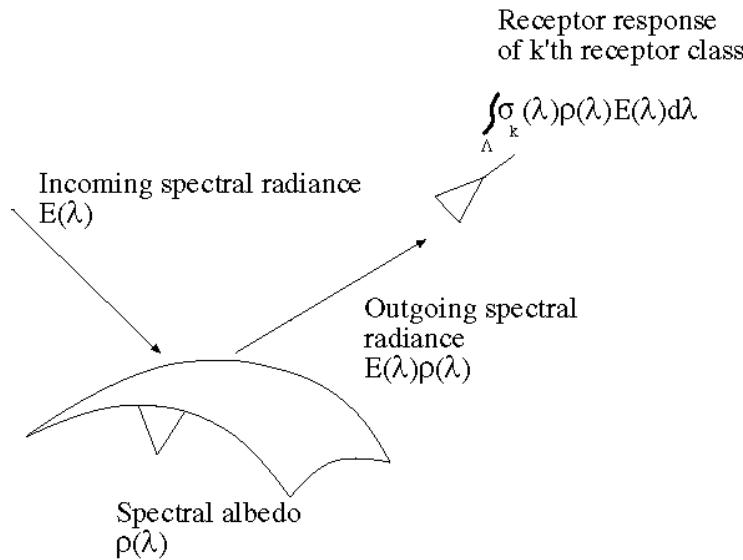




The color of objects



- Colored light arriving at the camera involves two effects
 - The color of the light source (illumination + inter-reflections)
 - The color of the surface





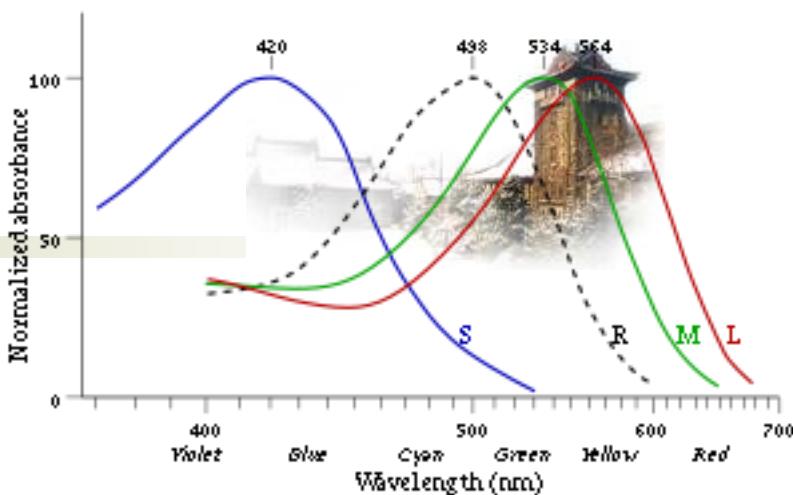
Why RGB?



If light is a spectrum, why are images RGB?



Human color receptors

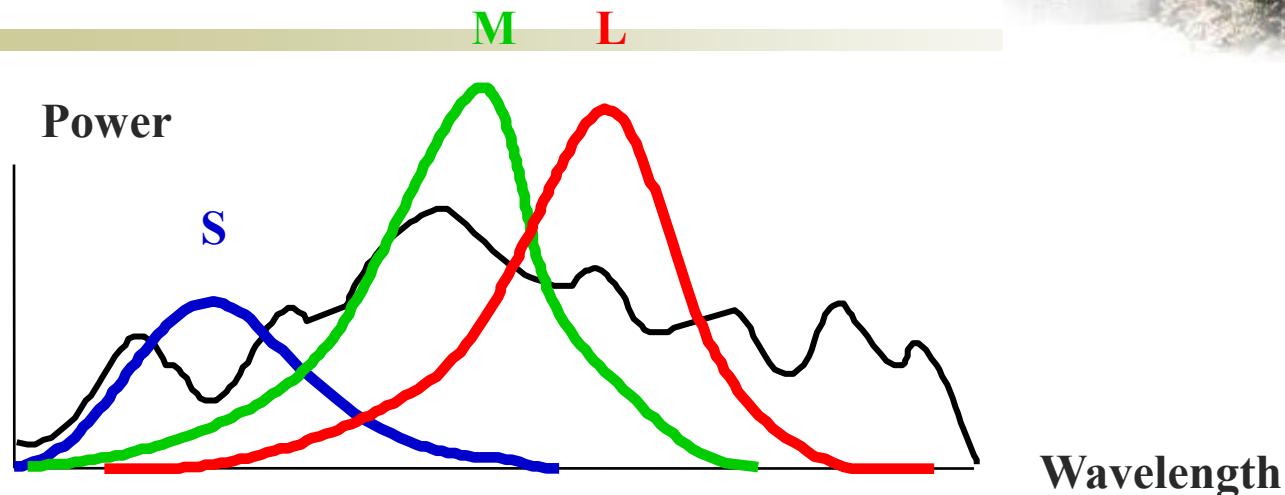


Long (red), Medium (green), and Short (blue) cones, plus intensity rods

- Fun facts
 - “M” and “L” on the X-chromosome
 - That’s why men are more likely to be color blind (see what it’s like:
<http://www.vischeck.com/vischeck/vischeckImage.php>)
 - “L” has high variation, so some women are tetrachromatric
 - Ratio of L to M to S cones: approx. 10:5:1
 - Some animals have 1 (night animals), 2 (e.g., dogs), 4 (fish, birds), 5 (pigeons, some reptiles/amphibians), or even 12 (mantis shrimp) types of cones



Color perception



Rods and cones act as *filters* on the spectrum

- To get the output of a filter, multiply its response curve by the spectrum, integrate over all wavelengths
 - Each cone yields one number

How can we represent an entire spectrum with three numbers?

We can't! Most of the information is lost

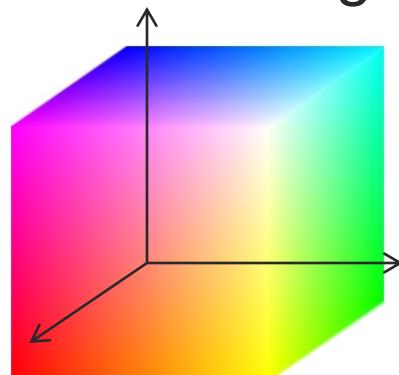
As a result, two different spectra may appear indistinguishable
such spectra are known as **metamers** (同色光)



Linear color spaces

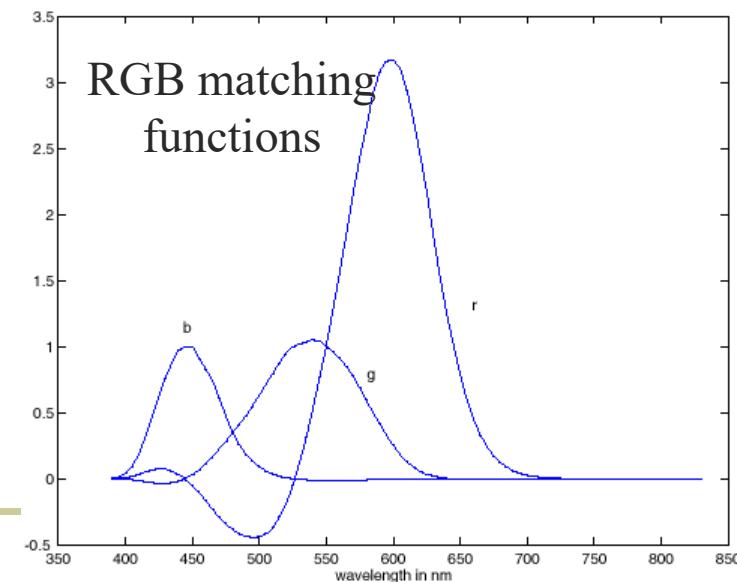


- Defined by a choice of three **primaries**
- The **coordinates** of a color are given by the weights of the primaries used to match it
- In addition to primaries, need to specify **matching functions**: the amount of each primary needed to match a monochromatic light source at each wavelength



RGB primaries

- $p_1 = 645.2 \text{ nm}$
- $p_2 = 525.3 \text{ nm}$
- $p_3 = 444.4 \text{ nm}$



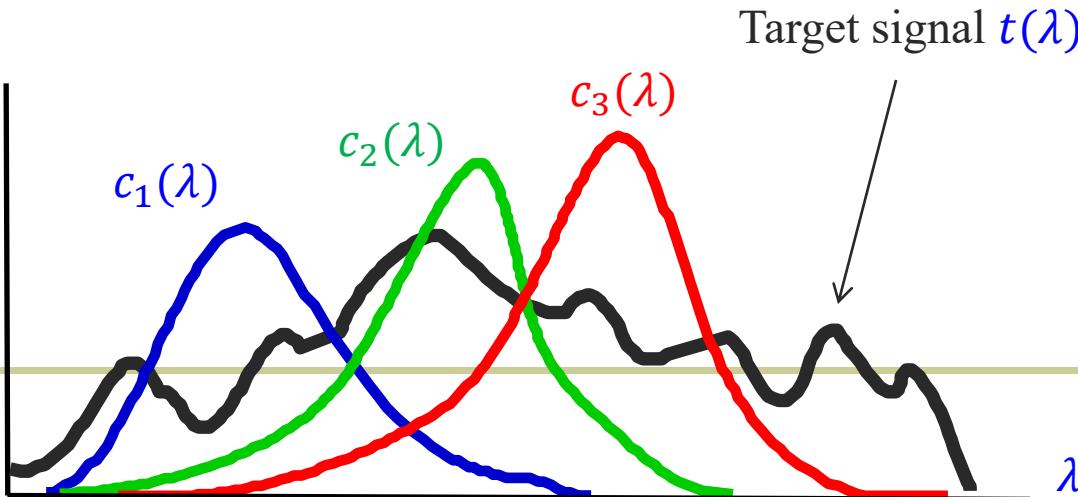


Matching functions



- Let $t(\lambda)$ be the spectrum of the target signal
- Let $c_1(\lambda)$, $c_2(\lambda)$, and $c_3(\lambda)$ be the *matching functions*, or the amounts of each primary needed to match monochromatic sources with wavelengths λ
- Then the coordinates of t in the corresponding linear space are given by

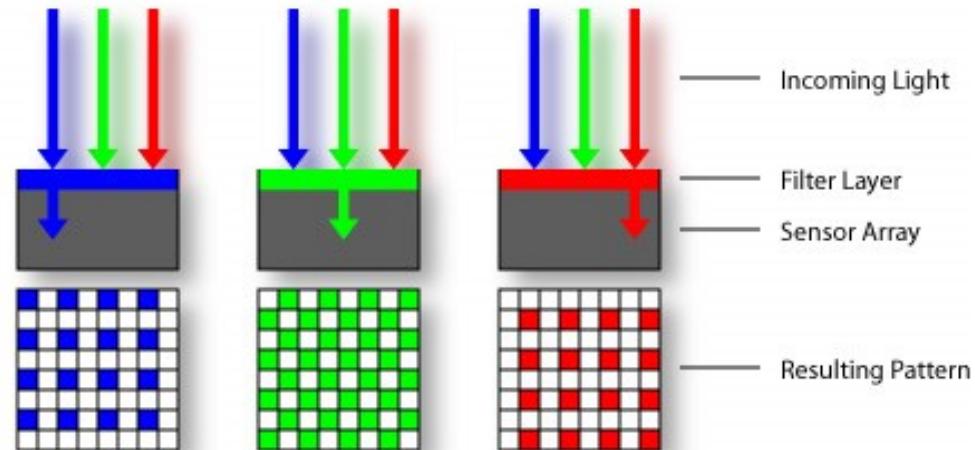
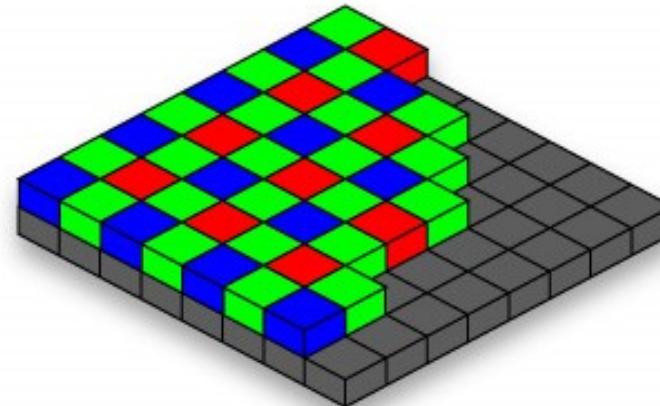
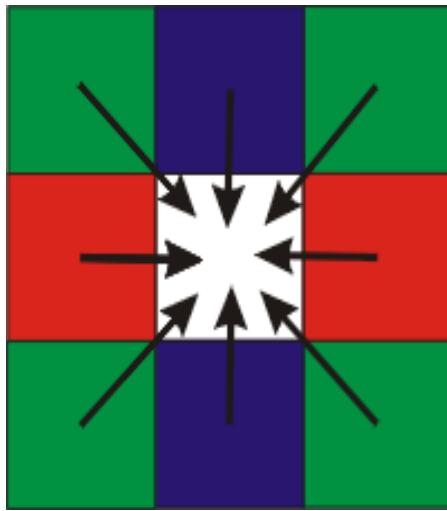
$$w_p = \int t(\lambda) c_p(\lambda) d\lambda$$



Matching functions act as filters on the target spectrum, like response curves of color receptors!



Color Sensing: Bayer Grid



- Estimate RGB at each cell from neighboring values

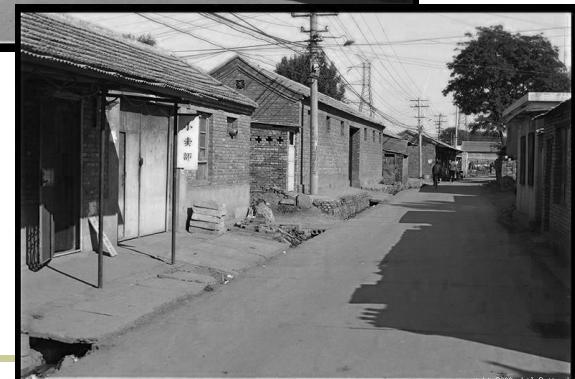
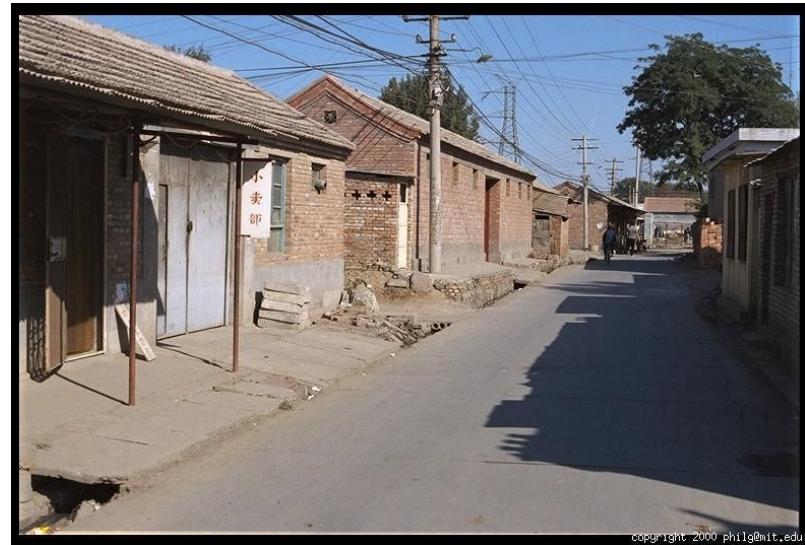


Color Image

R

G

B



copyright 2000 phil@mit.edu



Images in Matlab



- Images represented as a matrix
- Suppose we have a NxM RGB image called “im”
 - $im(1,1,1)$ = top-left pixel value in R-channel
 - $im(y, x, b)$ = y pixels down, x pixels to right in the bth channel
 - $im(N, M, 3)$ = bottom-right pixel in B-channel
- `imread(filename)` returns a uint8 image (values 0 to 255)
 - Convert to double format (values 0 to 1) with `im2double`

row ↓ column →

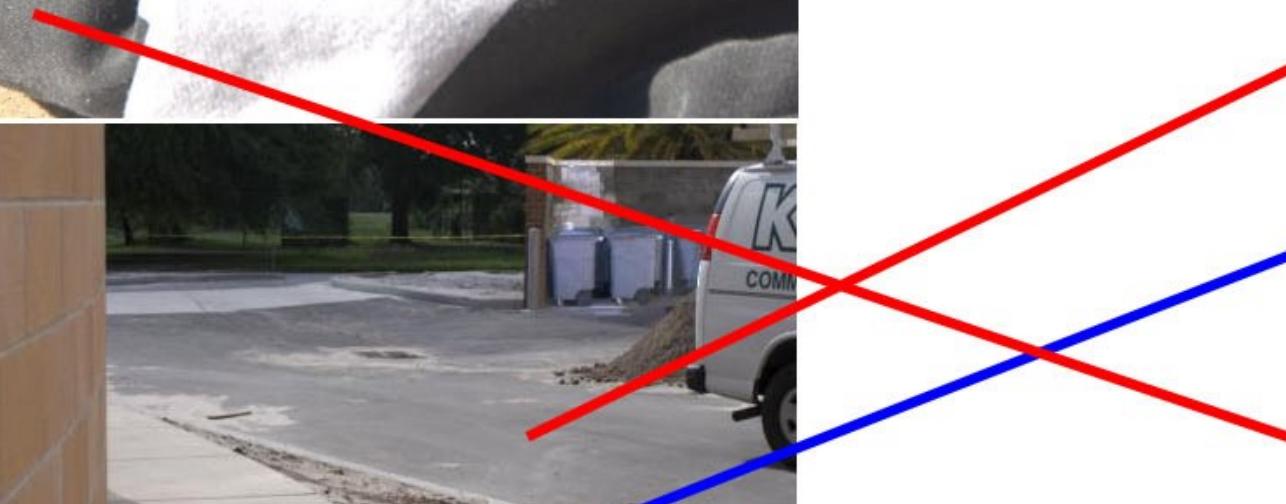
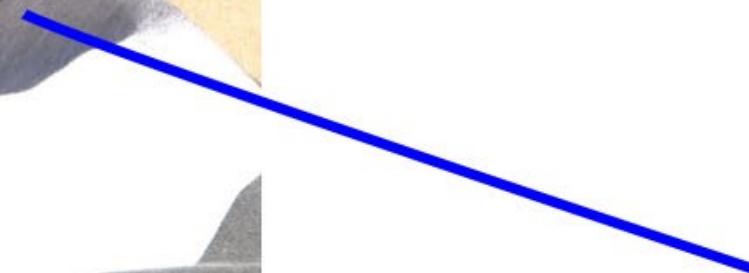
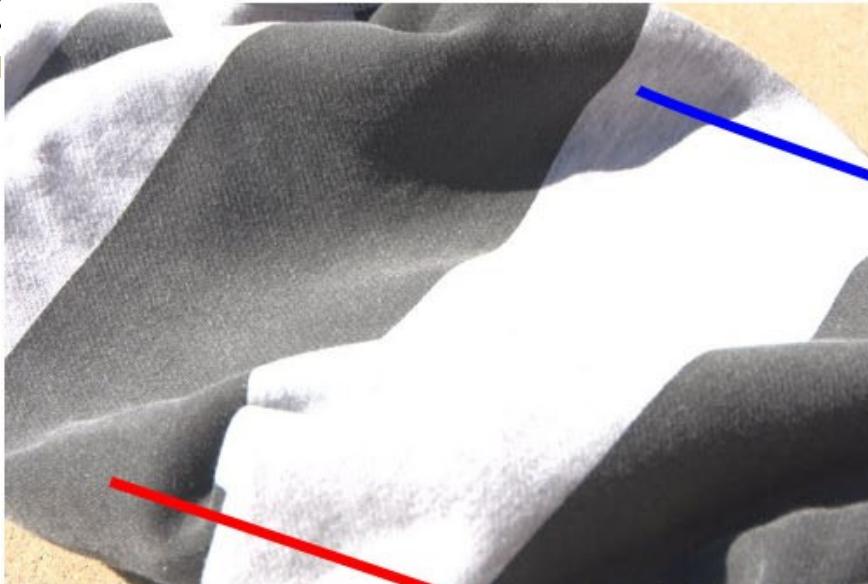
0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99	R
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91	G
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92	B
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95	
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85	
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33	
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74	
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93	
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99	
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	
0.69	0.45	0.56	0.58	0.58	0.45	0.42	0.77	0.75	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	0.49
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	0.82
0.69	0.45	0.56	0.58	0.58	0.45	0.42	0.77	0.75	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	0.90
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	0.93
0.69	0.45	0.56	0.58	0.58	0.45	0.42	0.77	0.75	0.71	0.90	0.99
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	0.93
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	0.97



The plight of the poor pixel



- A pixel's brightness is determined by
 - Light source (strength, direction, color)
 - Surface orientation
 - Surface material and albedo
 - Reflected light and shadows from surrounding surfaces
 - Gain on the sensor
- A pixel's brightness tells us nothing by itself





And yet we can interpret images...



copyright 2000 philg@mit.edu

- Key idea: for nearby scene points, most factors do not change much
- The information is mainly contained in *local differences* of brightness



Darkness = Large Difference in Neighboring Pixels





What is this?







Overview



- What is an image?
- Image formation: light and color
- **Image transformation: filtering**
- Image noise and image smoothing
- Convolution operation
- Media filter



Image transformations

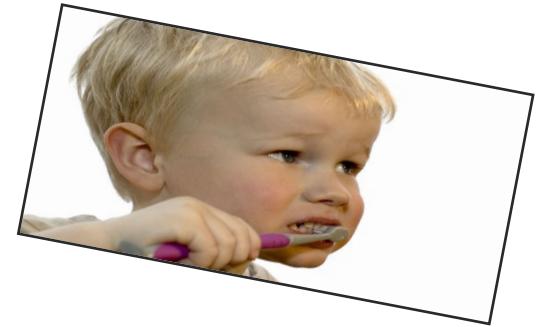


Filtering



changes pixel *values*

Warping



changes pixel *locations*



Image transformations



F



Filtering

$$G(\mathbf{x}) = h\{F(\mathbf{x})\}$$

G



changes *range* of image function

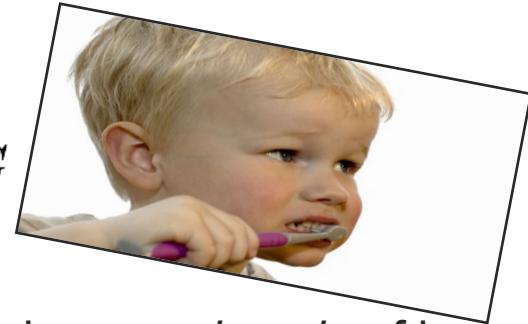
F



Warping

$$G(\mathbf{x}) = F(h\{\mathbf{x}\})$$

G



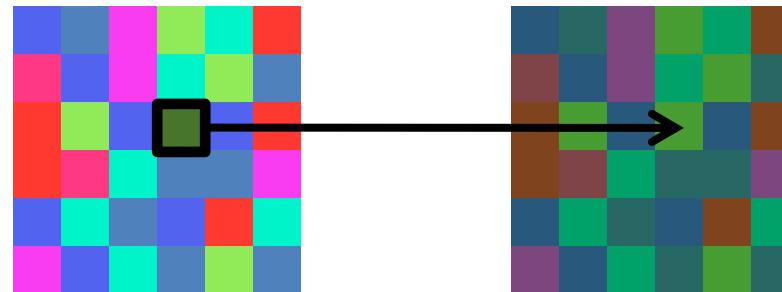
changes *domain* of image function



Image filtering

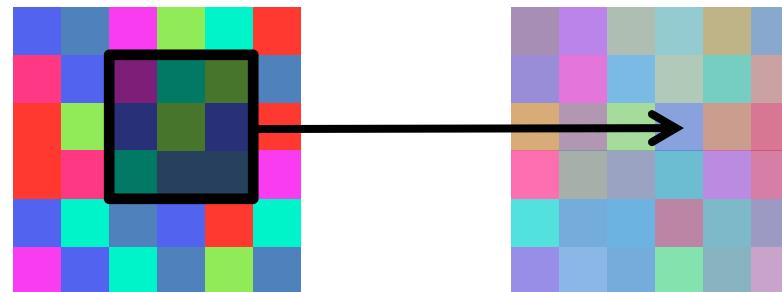


Point Operation



point processing

Neighborhood Operation



“filtering”



Examples of point processing



original



darken



lower contrast



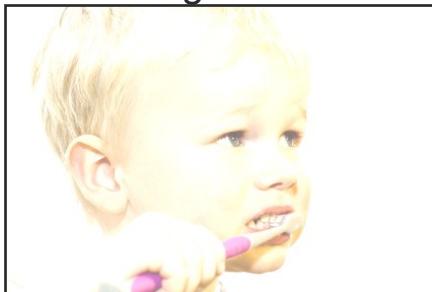
non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



darker



lower contrast



non-linear lower contrast

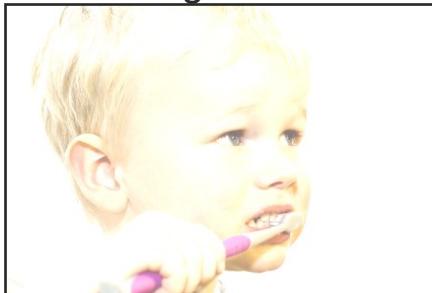


x

invert



lighten



raise contrast



non-linear raise contrast





Examples of point processing

How would you implement these?

original



x

darker



$x - 128$

lower contrast



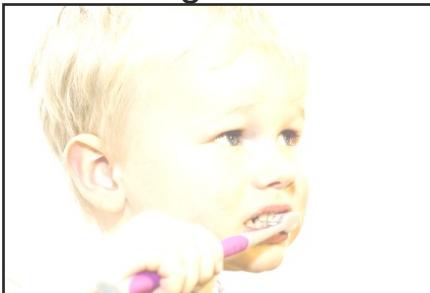
non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



invert



lighten



raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast

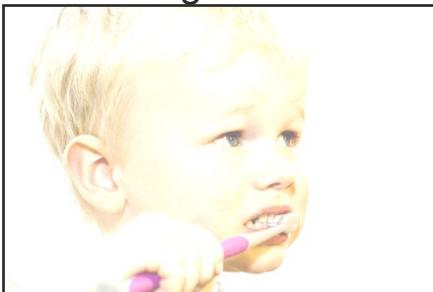


$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

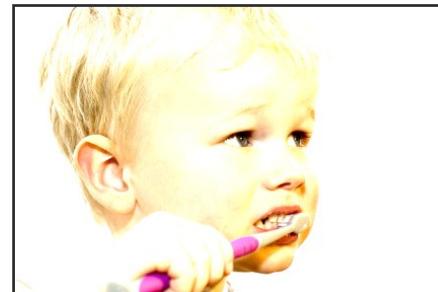
invert



lighten



raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



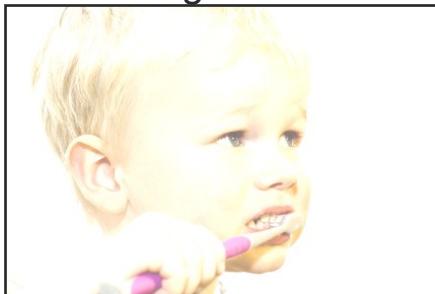
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



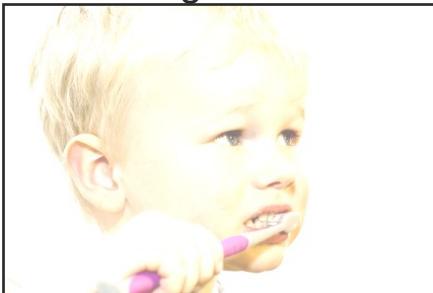
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



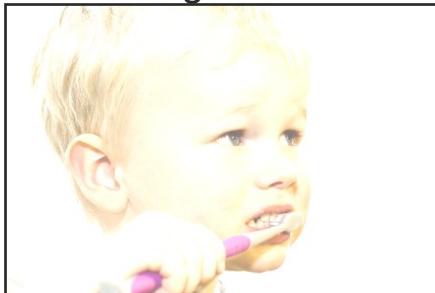
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear raise contrast





Examples of point processing



How would you implement these?

original



$$x$$

darker



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



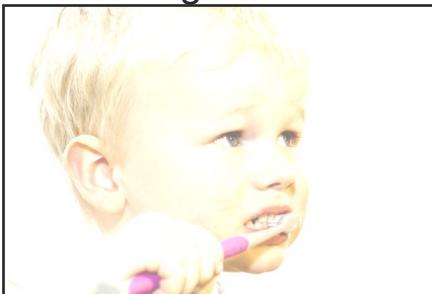
$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$



$$\left(\frac{x}{255}\right)^2 \times 255$$



Image filtering



- Compute a function of the local neighborhood at each pixel in the image
 - Function specified by a “filter” or mask saying how to combine values from neighbors.
- Uses of filtering:
 - Enhance an image (denoise, resize, etc)
 - Extract information (texture, edges, etc)
 - Detect patterns (template matching)



Three views of filtering



- Image filters in spatial domain
 - Filter is a mathematical operation on values of each patch
 - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
 - Filtering is a way to modify the frequencies of images
 - Denoising, sampling, image compression
- Templates and Image Pyramids
 - Filtering is a way to match a template to the image
 - Detection, coarse-to-fine registration



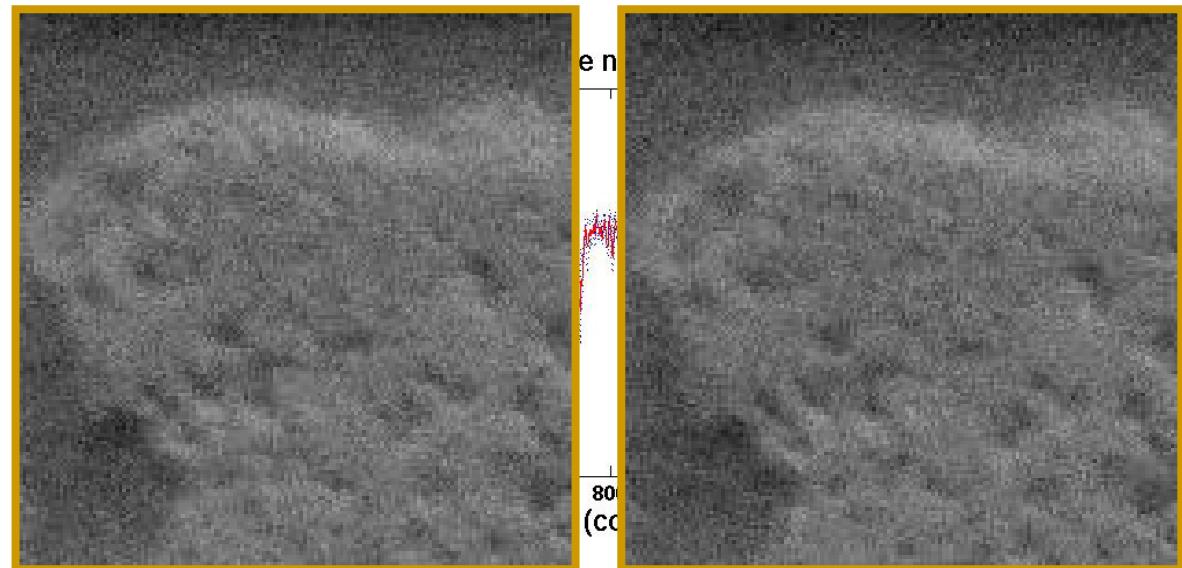
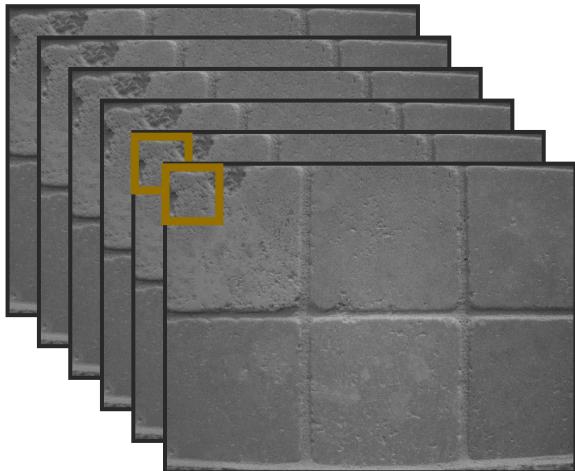
Overview



- What is an image?
- Image formation: light and color
- Image transform
- **Image noise and image smoothing**
- Convolution operation
- Non-linear filter: Media filter, Bilateral Filter



Motivation: noise reduction



- Even multiple images of the **same static scene** will not be identical.



Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



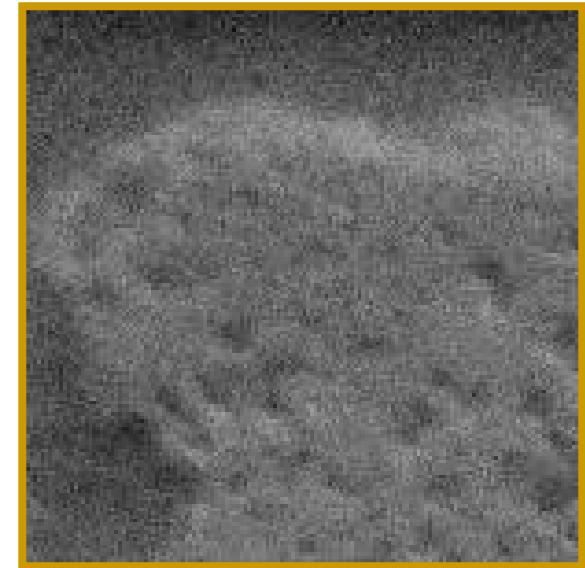
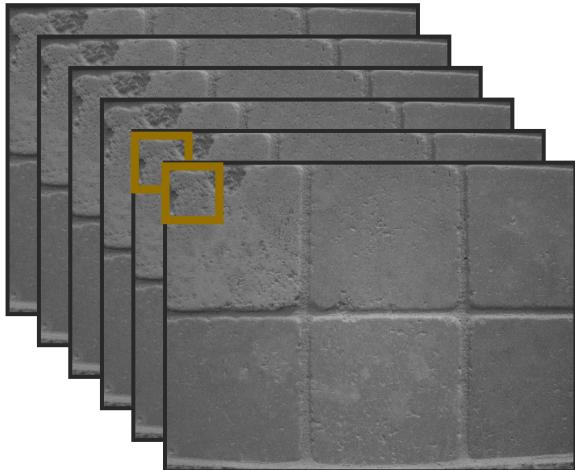
Impulse noise



Gaussian noise



Motivation: noise reduction



- Even multiple images of the same static scene will not be identical.
- How could we reduce the noise, i.e., give an estimate of the true intensities?
- **What if there's only one image?**



First attempt at a solution



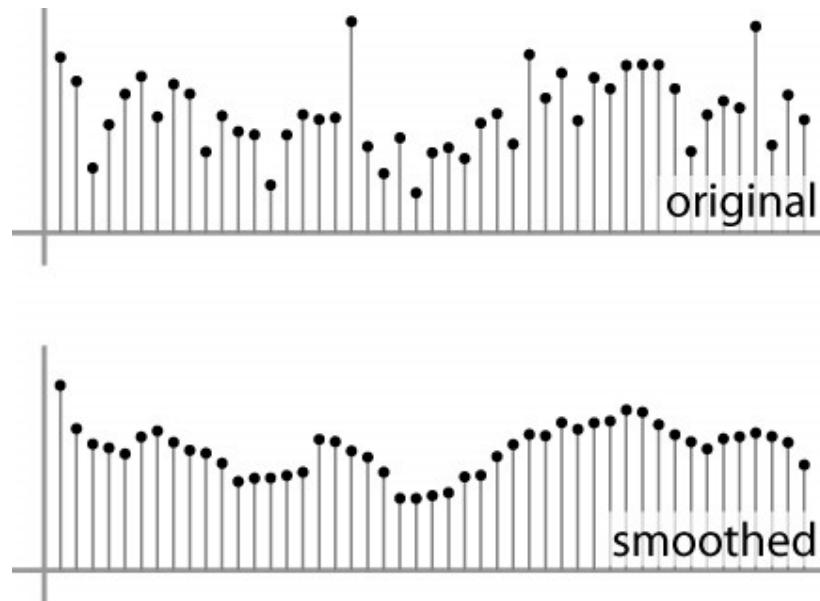
- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel



First attempt at a solution



- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:

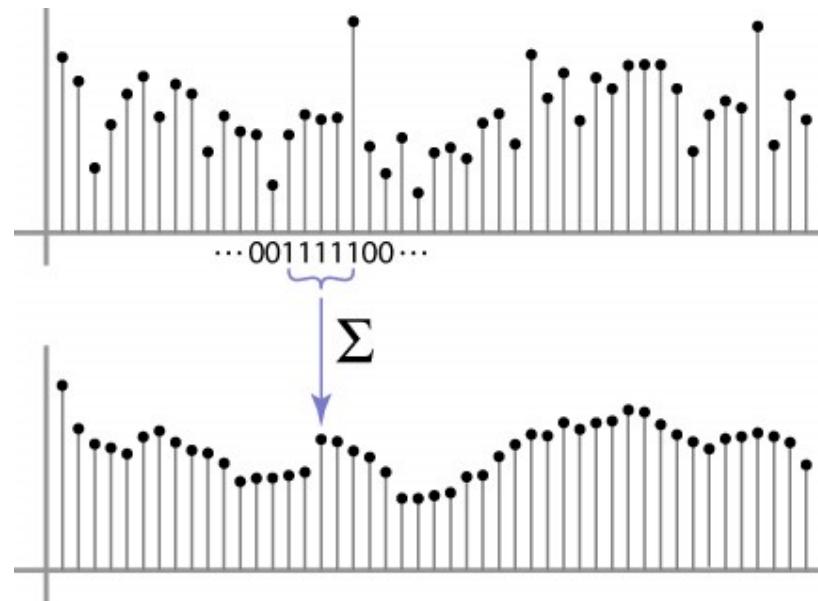




Weighted Moving Average



- Can add weights to our moving average
- *Weights* [1, 1, 1, 1, 1] / 5

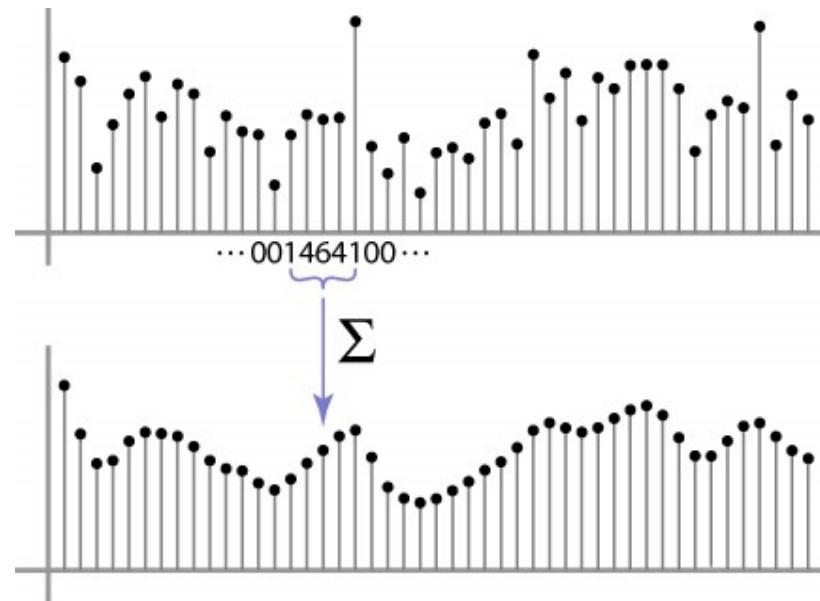




Weighted Moving Average



- Non-uniform weights $[1, 4, 6, 4, 1] / 16$





Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Moving Average In 2D



$$F[x, y]$$

$$G[x, y]$$



Correlation filtering



Say the averaging window size is $2k+1 \times 2k+1$:

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \sum_{u=-k}^k \underbrace{\sum_{v=-k}^k F[i+u, j+v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$



Correlation filtering



$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “kernel” or “mask” $H[u, v]$ is the prescription for the weights in the linear combination.



Averaging filter

- What values belong in the kernel H for the moving average example?

$F[x, y]$									
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$G[x, y]$									
	0	20	40	60	60				

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & ? \\ 1 & 1 & 1 \end{matrix}$$

“box filter”

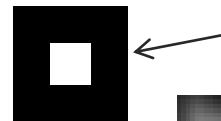
$$G = H \otimes F$$



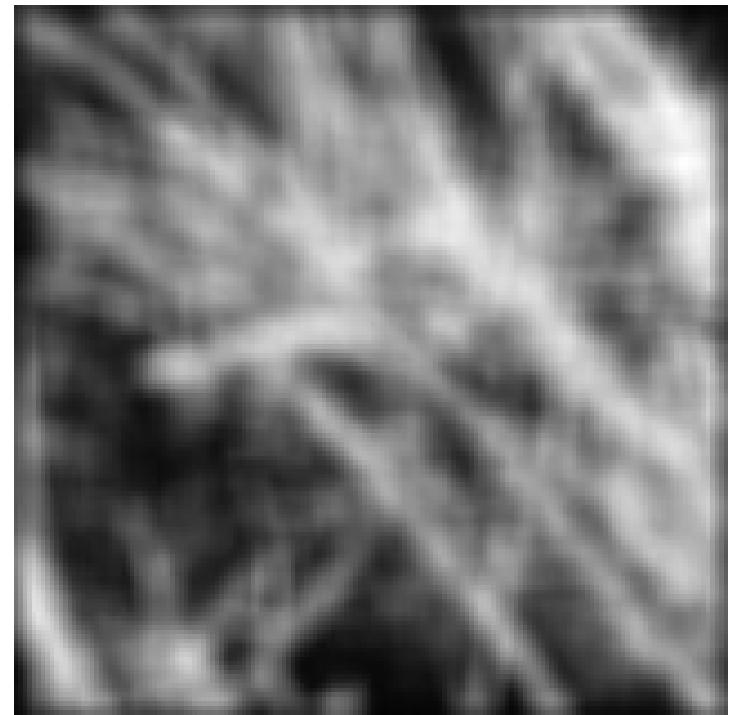
Smoothing by averaging



original



depicts box filter:
white = high value, black = low value



filtered

What if the filter size was 5×5 instead of 3×3 ?

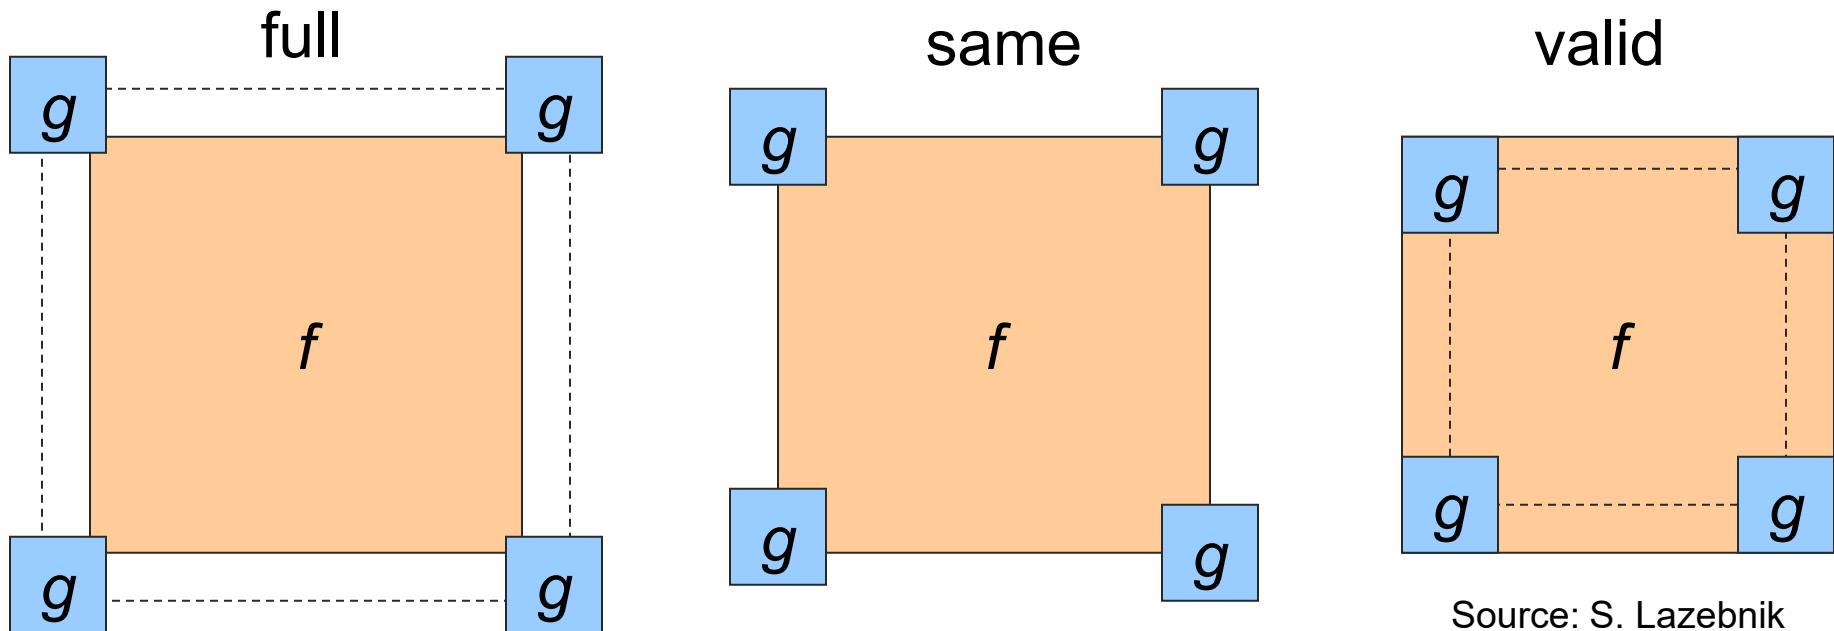


Boundary issues



MATLAB: output size / “shape” options

- *shape = ‘full’*: output size is sum of sizes of f and g
- *shape = ‘same’*: output size is same as f
- *shape = ‘valid’*: output size is difference of sizes of f and g



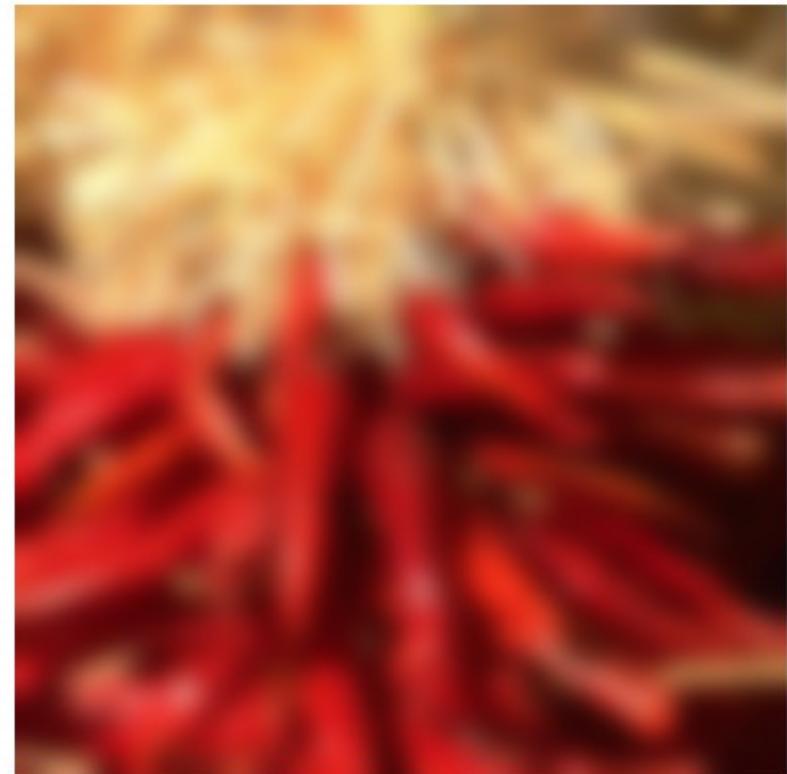


Boundary issues



■ What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge





Boundary issues



■ What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods (MATLAB):
 - clip filter (black): `imfilter(f, g, 0)`
 - wrap around: `imfilter(f, g, 'circular')`
 - copy edge: `imfilter(f, g, 'replicate')`
 - reflect across edge: `imfilter(f, g, 'symmetric')`



Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

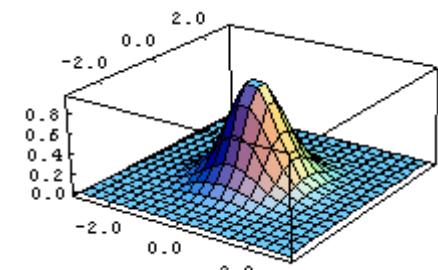
$F[x, y]$

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

This kernel is an approximation of a 2d Gaussian function:

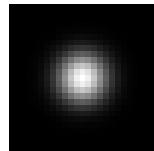
$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



- Removes high-frequency components from the image (“low-pass filter”).



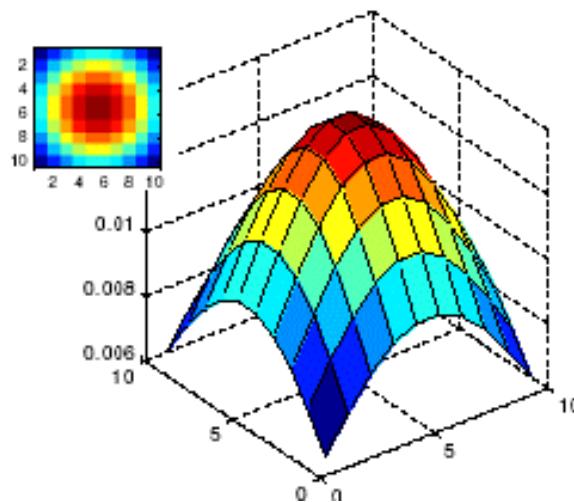
Smoothing with a Gaussian



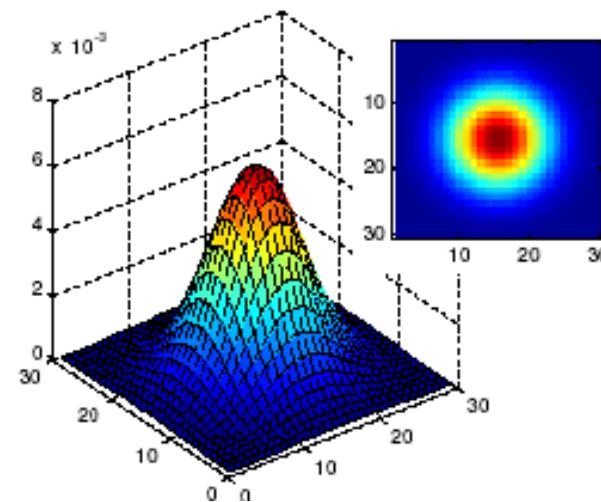


Gaussian filters

- What parameters matter here?
- Size of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels



$\sigma = 5$ with 10×10 kernel



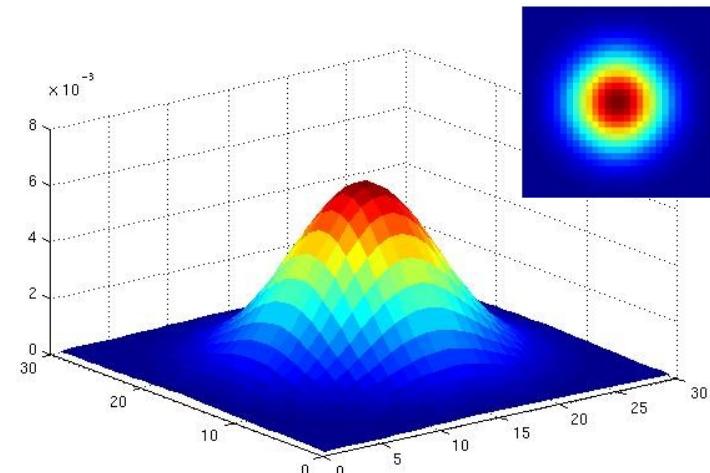
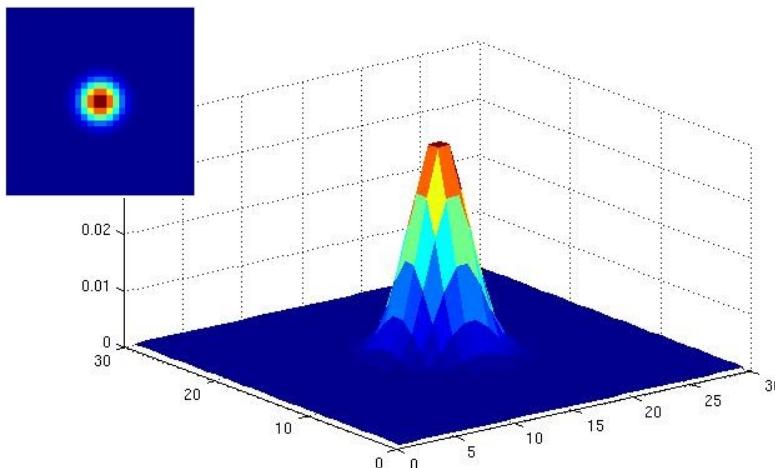
$\sigma = 5$ with 30×30 kernel



Gaussian filters



- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing

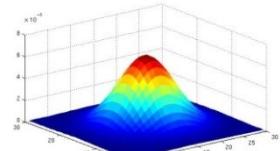




Matlab



```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```



```
>> mesh(h);
```



```
>> imagesc(h);
```

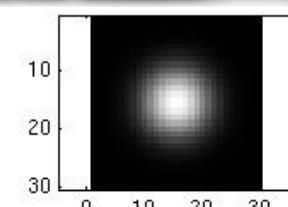
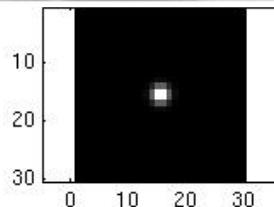


```
>> outim = imfilter(im, h); % correlation outim  
>> imshow(outim);
```

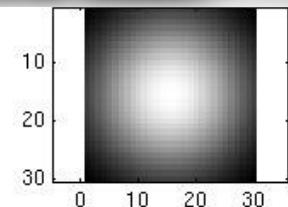


Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize,
    sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```



Properties of smoothing filters



■ Smoothing

- Values positive
- Sum to 1 → constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components; “low-pass” filter



Overview



- What is an image?
- Image formation: light and color
- Image transform
- Image noise and image smoothing
- **Convolution operation**
- Non-linear filter: Media filter, Bilateral Filter



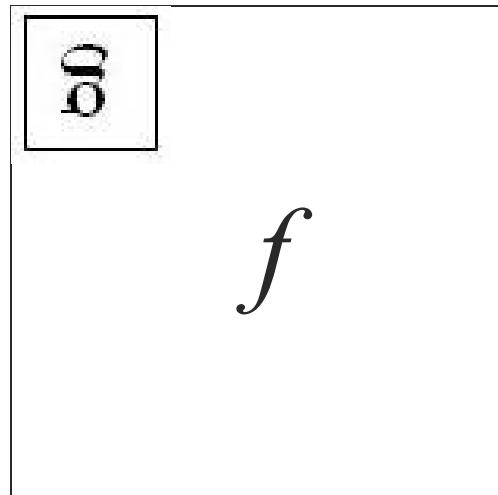
Defining convolution



- Let f be the image and g be the kernel. The output of convolving f with g is denoted $f * g$.

$$(f * g)[m, n] = \sum_{k,l} f[m - k, n - l]g[k, l]$$

Convention:
kernel is “flipped”





Key properties

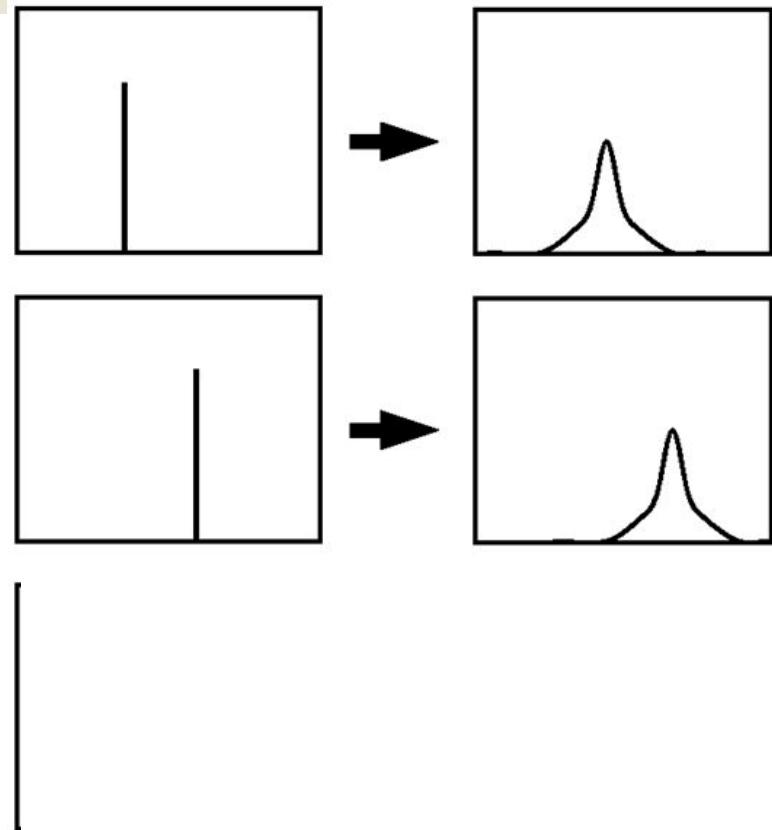


- **Shift invariance:** same behavior regardless of pixel location:

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$

- **Linearity:**

$$\begin{aligned}\text{filter}(f_1 + f_2) &= \\ \text{filter}(f_1) + \text{filter}(f_2)\end{aligned}$$



- Theoretical result: any linear shift-invariant operator can be represented as a convolution



Properties in more detail



- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$,
 $a * e = a$



Convolution vs correlation



Definition of discrete 2D convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

↑ notice the flip

Definition of discrete 2D correlation:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x + i, y + j)$$

↑ notice the lack of a flip

- Most of the time won't matter, because our kernels will be symmetric.



Separability



- In some cases, filter is separable, and we can factor into two steps:
 - Convolve all rows with a 1D filter
 - Convolve all columns with a 1D filter

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix} = \begin{bmatrix} 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \\ 0 & \frac{1}{3} & 0 \end{bmatrix} \circ \begin{bmatrix} 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ 0 & 0 & 0 \end{bmatrix}$$



Overview



- What is an image?
- Image formation: light and color
- Image transform
- Image noise and image smoothing
- Convolution operation
- **Non-linear filter: Media filter, Bilateral Filter**



Effect of smoothing filters



5x5



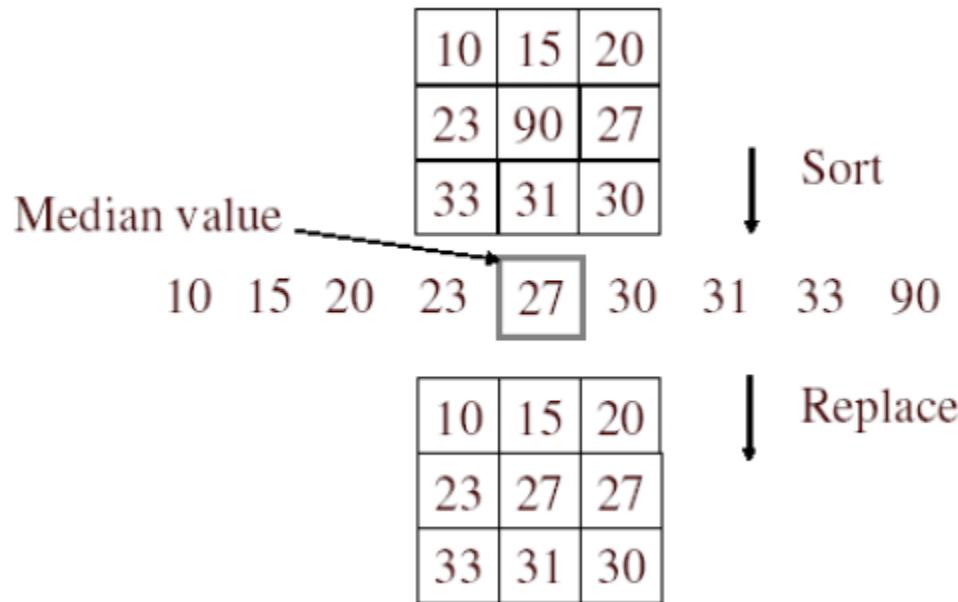
Additive Gaussian noise



Salt and pepper noise



Median filter



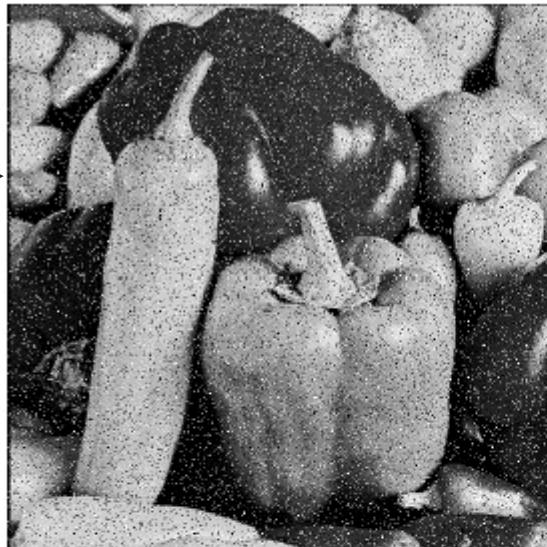
- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise
- Non-linear filter



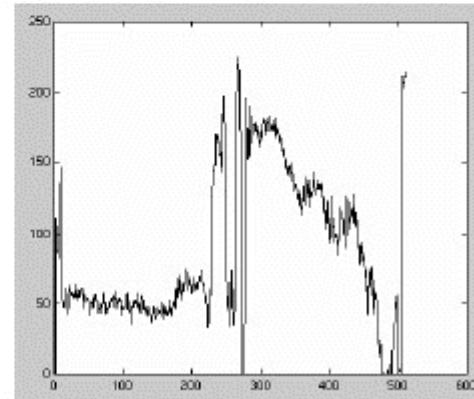
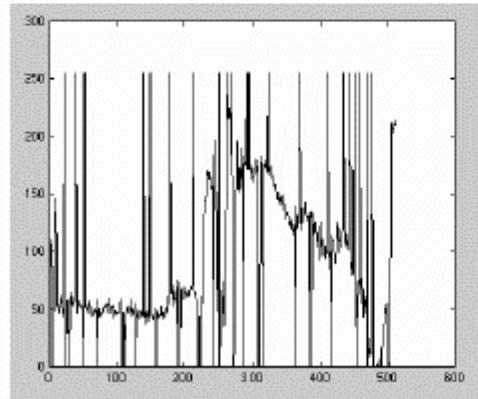
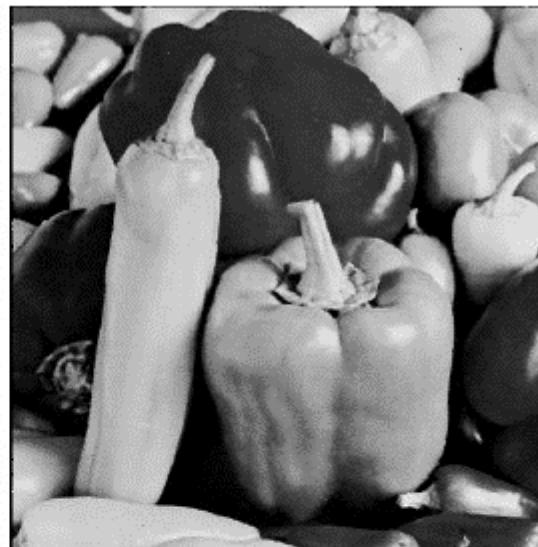
Median filter



Salt and
pepper
noise



Median
filtered



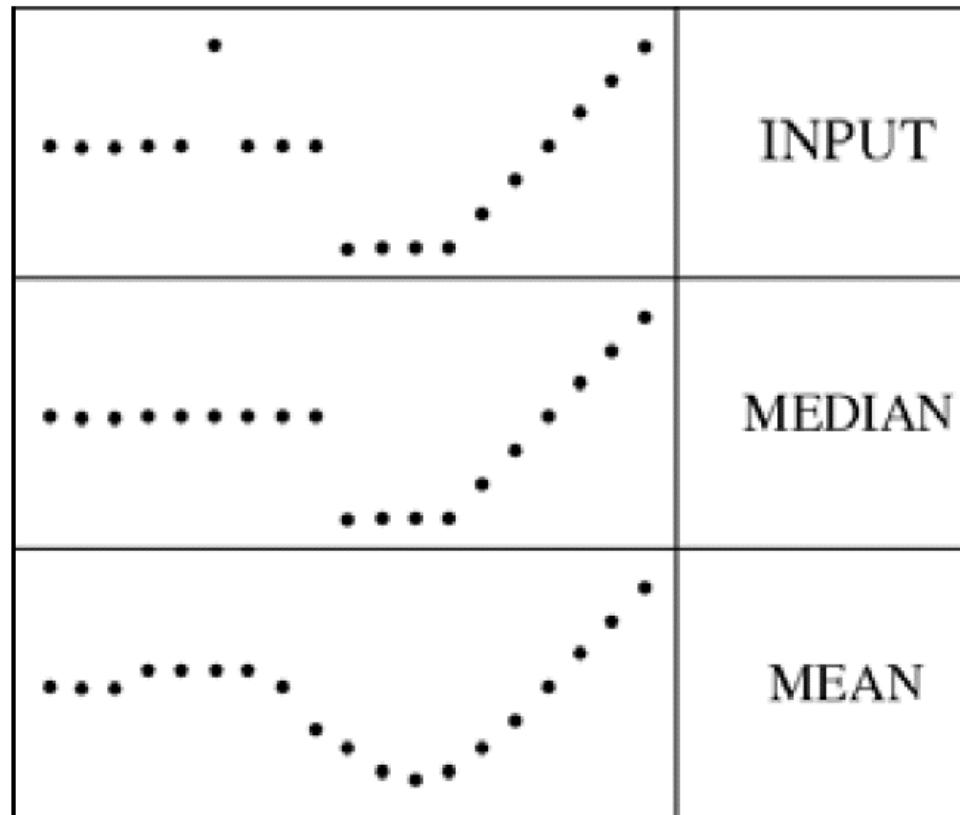
Plots of a row of the image

Matlab: output im = medfilt2(im, [h w]);



Median filter

- Median filter is edge preserving





Objective of bilateral filtering



Smooth texture

Preserve edges



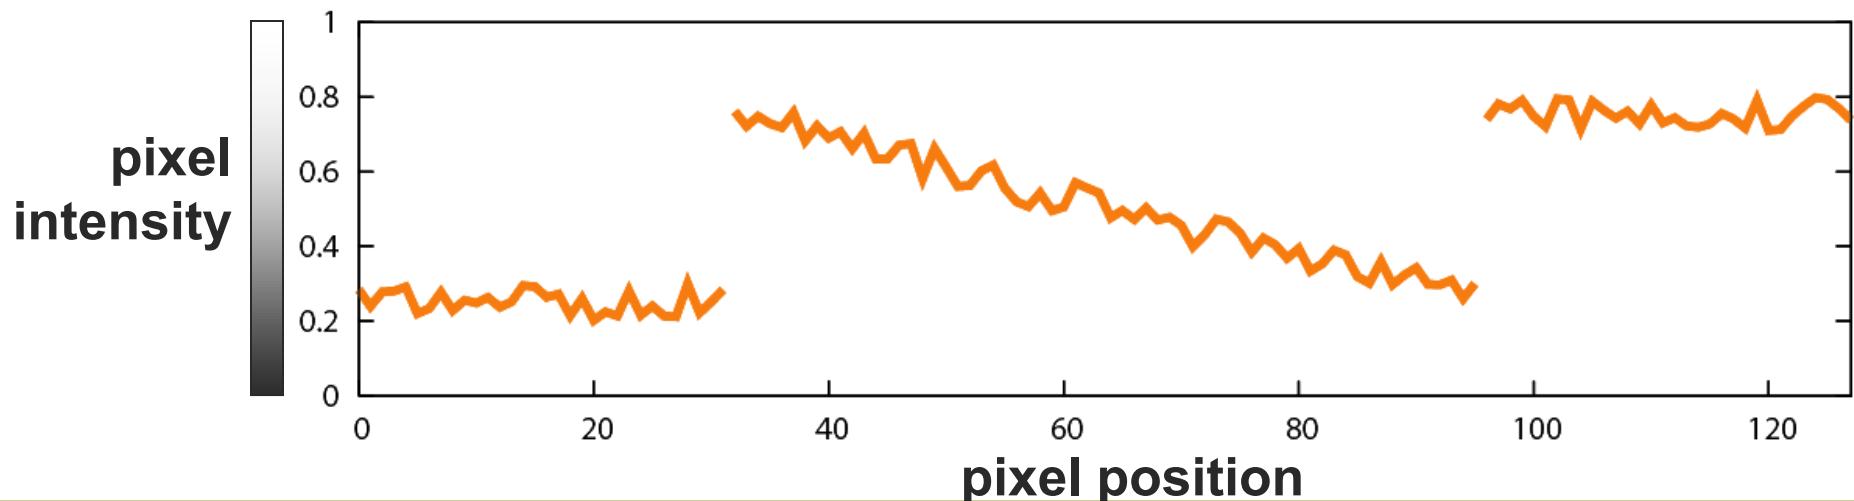
Illustration a 1D Image



- 1D image = line of pixels



- Better visualized as a plot





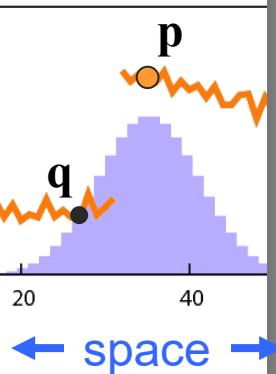
Definition

Gaussian blur

$$I_p^b = \sum_{q \in S} G_{\sigma_s}(\|p - q\|) I_q$$

space

- only spatial distance, intensity ignored

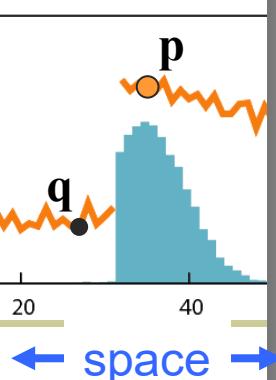


Bilateral filter

[Aurich 95, Smith 97, Tomasi 98]

$$I_p^{bf} = \frac{1}{W_p^{bf}} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

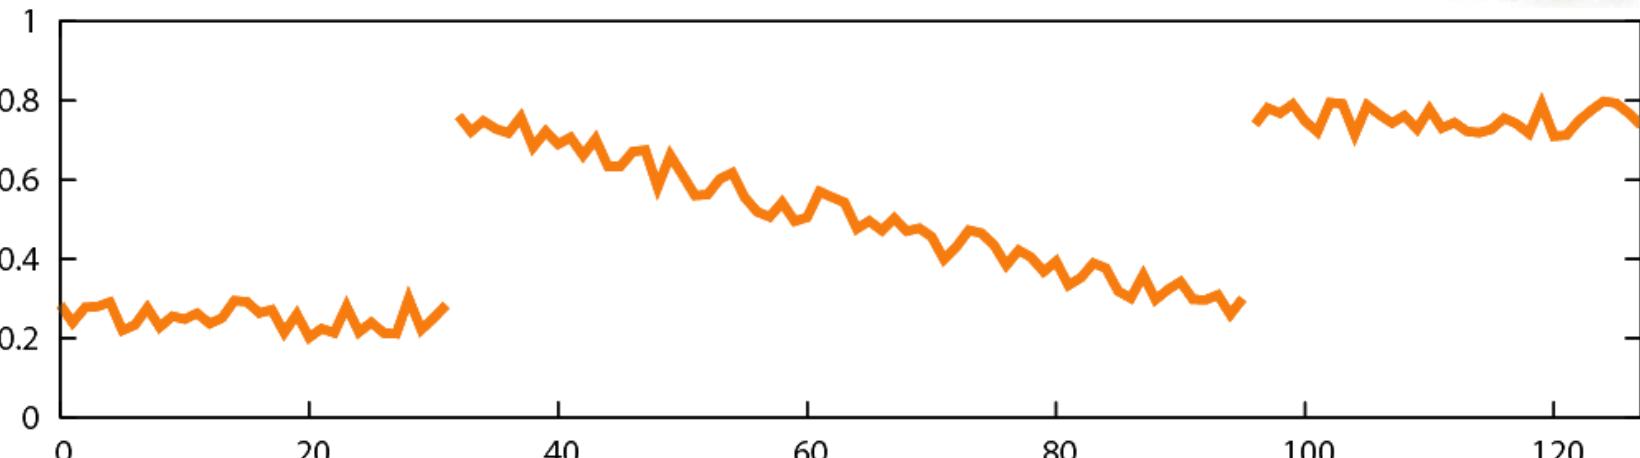
normalization



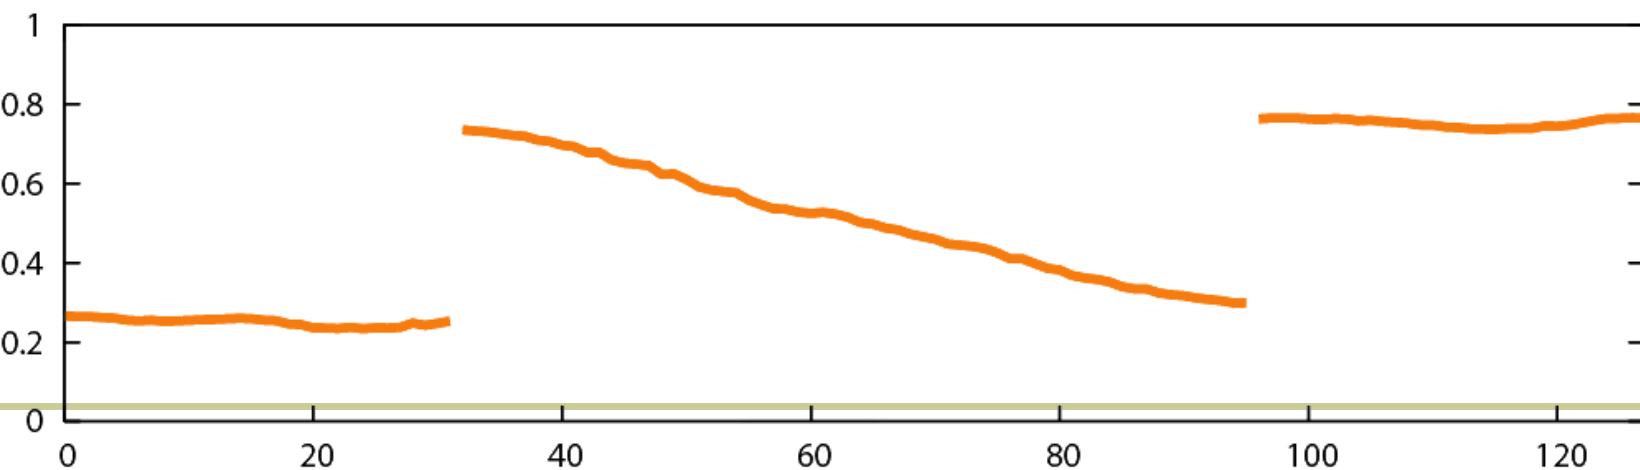
- spatial and range distances
- weights sum to 1



Bilateral Filter on 1D Signal



BF



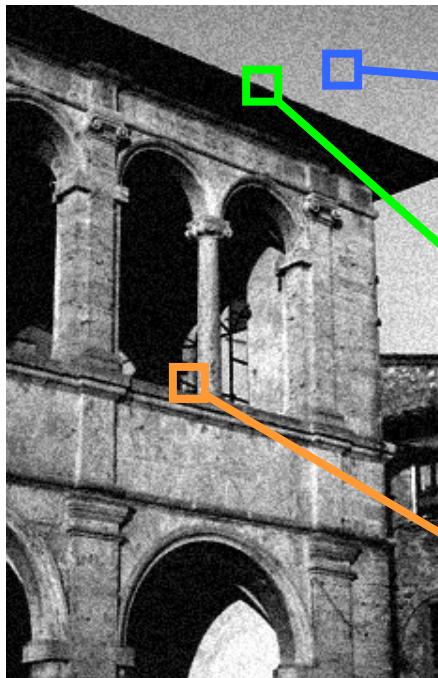


Example on a Real Image

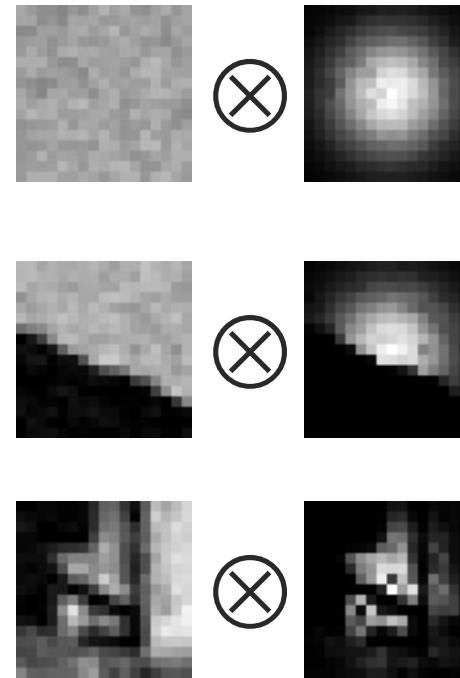


- Kernels can have complex, spatially varying shapes.

input



output





Bilateral Filter is Expensive



- Brute-force computation is slow (several minutes)
 - Two nested for loops:
for each pixel, look at all pixels
 - Non-linear, depends on image content
⇒ no FFT, no pre-computation...
- Fast approximations exist [Durand 02, Weiss 06]
 - Significant **loss of accuracy**
 - **No formal understanding** of accuracy versus speed



(a)



(b)



(c)



(d)



(e)



(f)



(g)



(h)

Figure 3.18 Median and bilateral filtering: (a) original image with Gaussian noise; (b) Gaussian filtered; (c) median filtered; (d) bilaterally filtered; (e) original image with shot noise; (f) Gaussian filtered; (g) median filtered; (h) bilaterally filtered. Note that the bilateral filter fails to remove the shot noise because the noisy pixels are too different from their neighbors.