



计算机视觉表征与识别

Chapter 4: Template, Pyramid, and Filter Banks

王利民

媒体计算课题组

<http://mcg.nju.edu.cn/>



Today's class



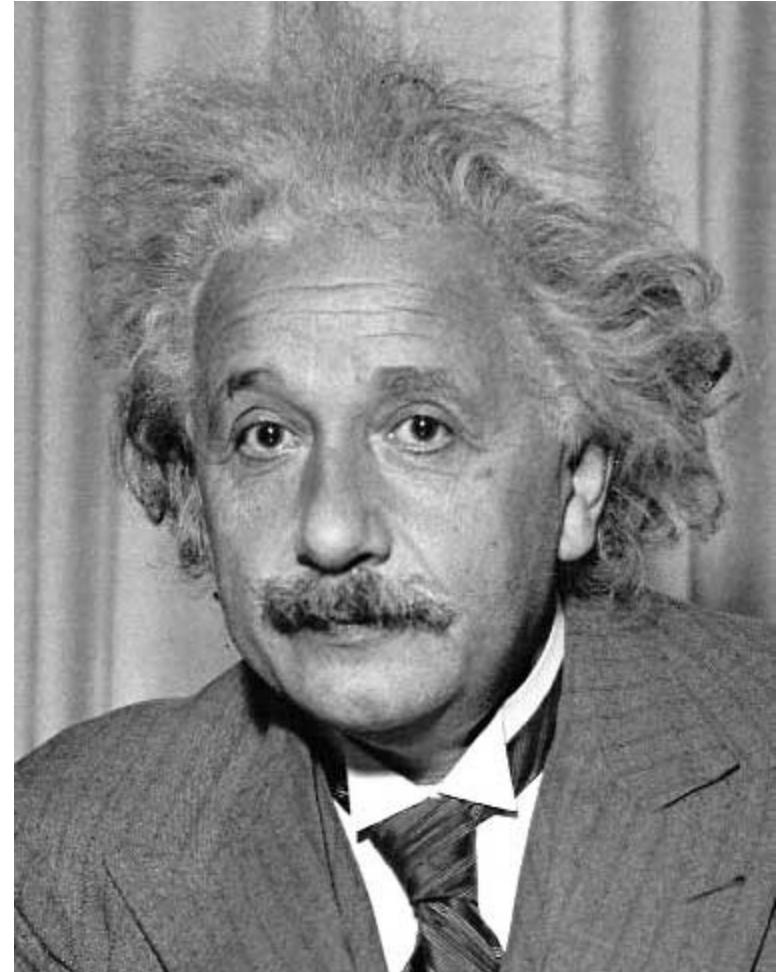
- Template matching
- Gaussian pyramids
 - Application for recognition
 - Pyramids representation in deep learning
- Laplacian pyramids
 - Hybrid images
 - Application for image blending
- Steerable pyramids
 - Steerable filters
 - Orientation analysis
- Human vision system
 - Visual perception
- Filter Banks and Texture Analysis



Template matching



- Goal: find in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation





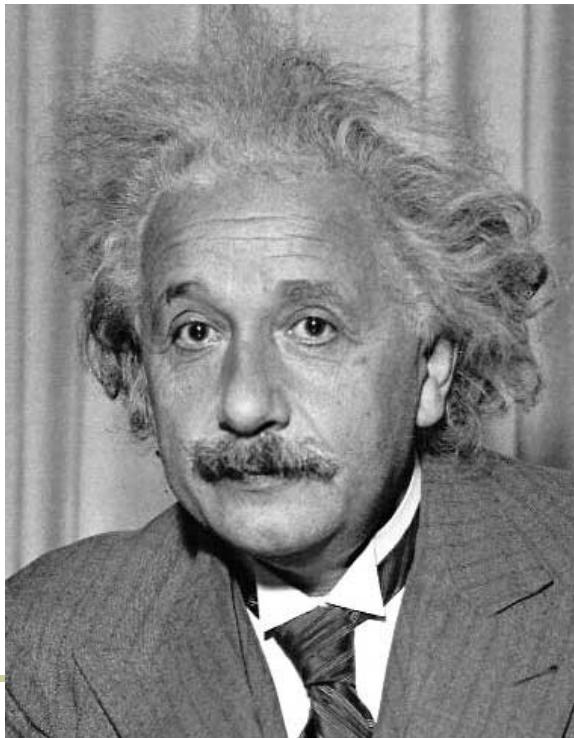
Matching with filters



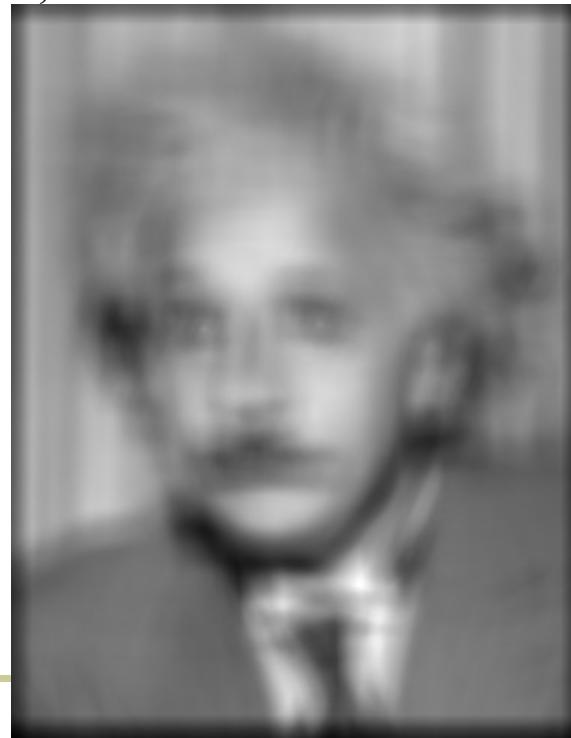
- Goal: find in image
- Method 0: filter the image with eye patch

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

f = image
 g = filter



Input



Filtered Image

What went wrong?

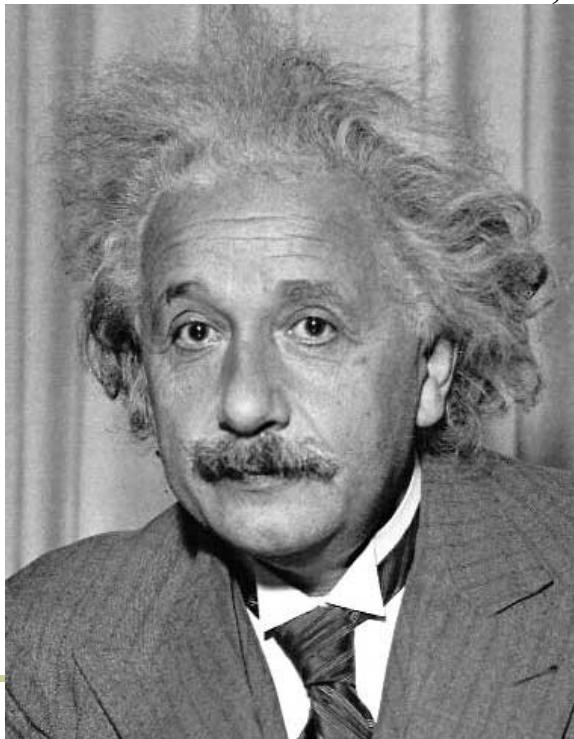


Matching with filters

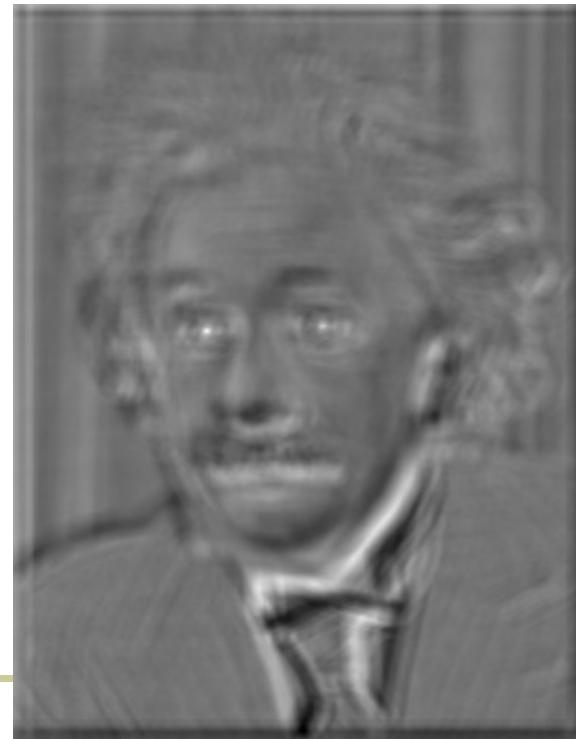


- Goal: find  in image
- Method 1: filter the image with zero-mean eye

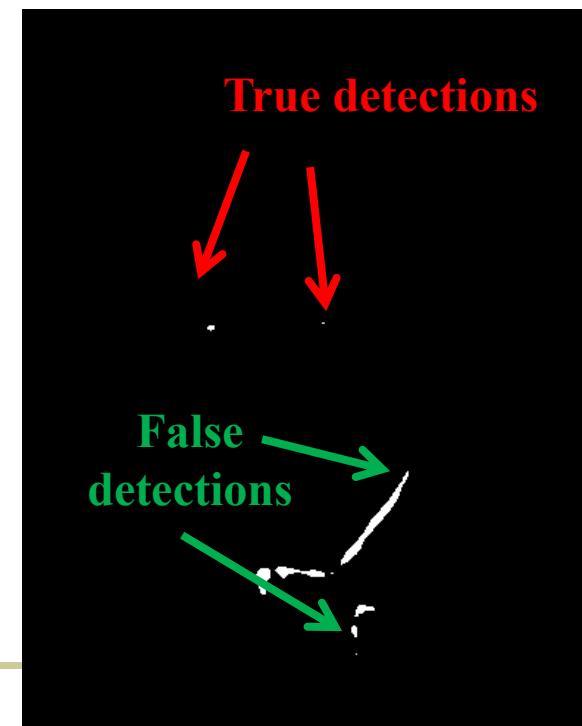
$$h[m, n] = \sum_{k, l} (g[k, l] - \bar{g}) \underbrace{(f[m + k, n + l])}_{\text{mean of template } g}$$



Input



Filtered Image (scaled)



Thresholded Image

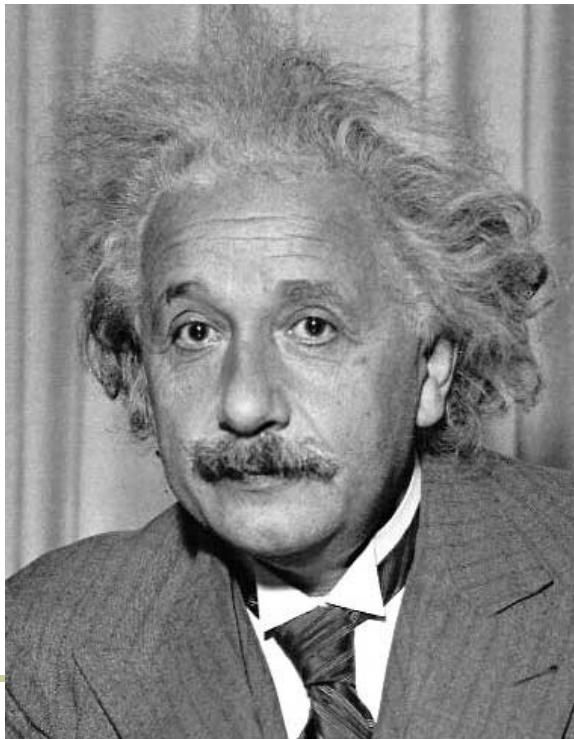


Matching with filters



- Goal: find in image
- Method 2: SSD

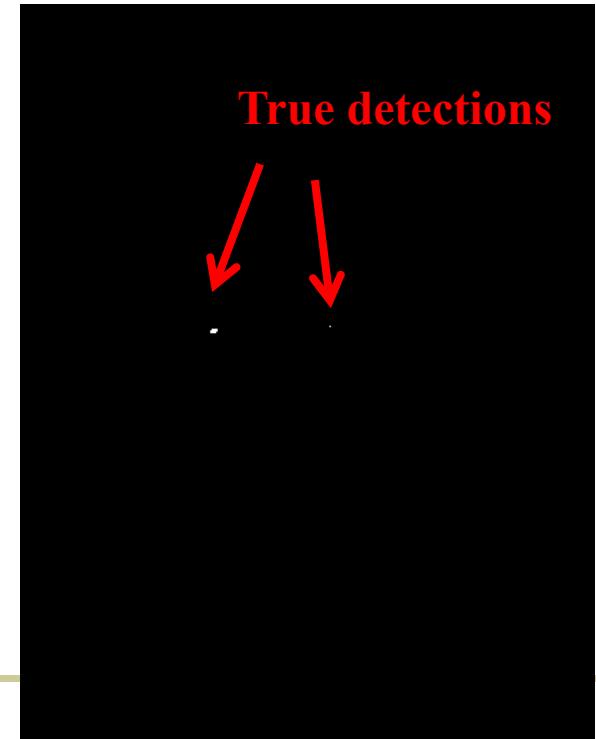
$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



Input



1 - $\text{sqrt}(\text{SSD})$



True detections



Matching with filters



Can SSD be implemented with linear filters?

$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



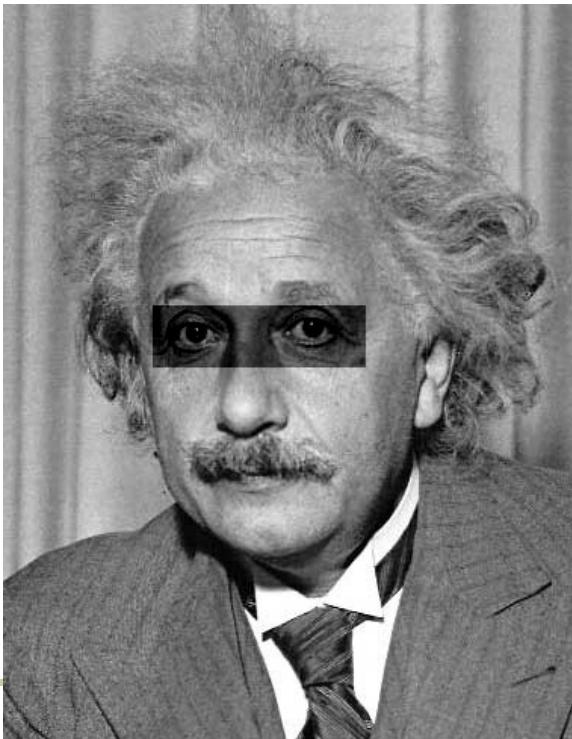
Matching with filters



- Goal: find in image
- Method 2: SSD

What's the potential downside
of SSD?

$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



Input



1 - sqrt(SSD)



Matching with filters



- Goal: find in image
- Method 3: Normalized cross-correlation

$$h[m, n] = \frac{\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m + k, n + l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

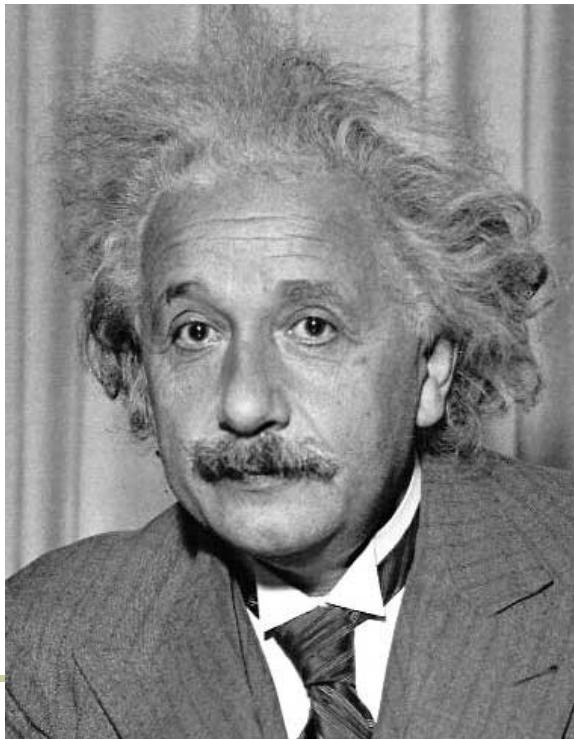
Matlab: `normxcorr2(template, im)`



Matching with filters



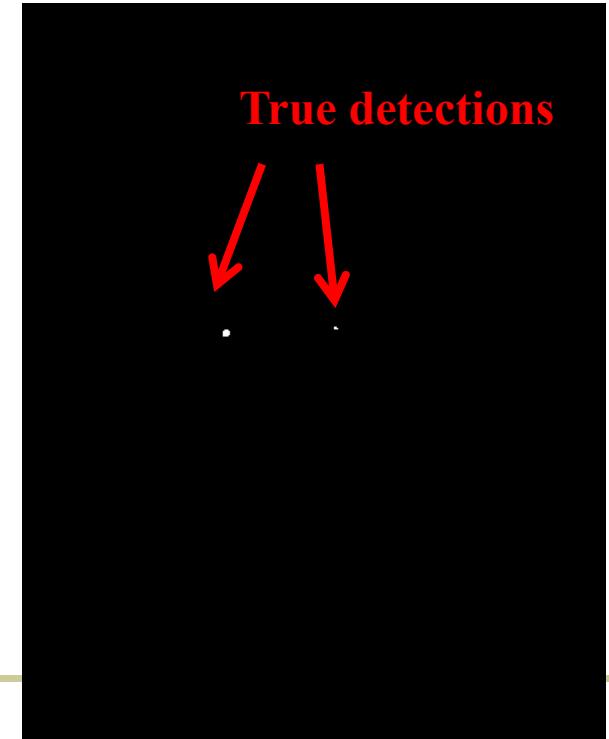
- Goal: find in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

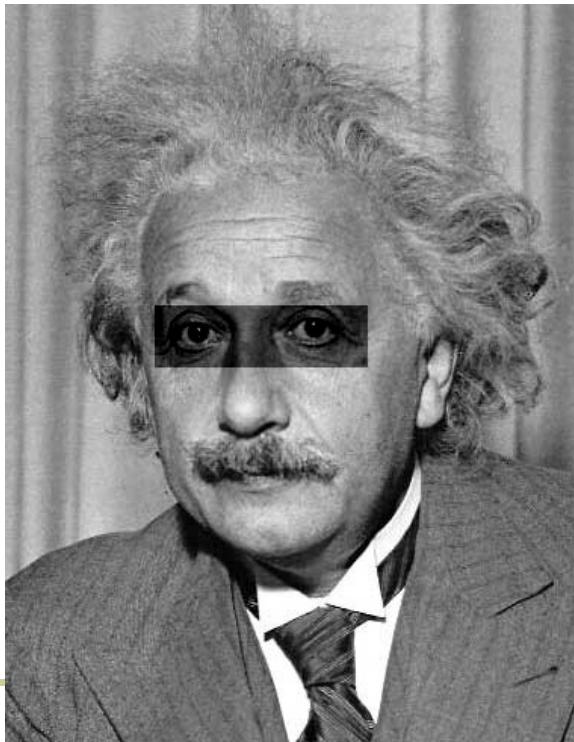
True detections



Matching with filters



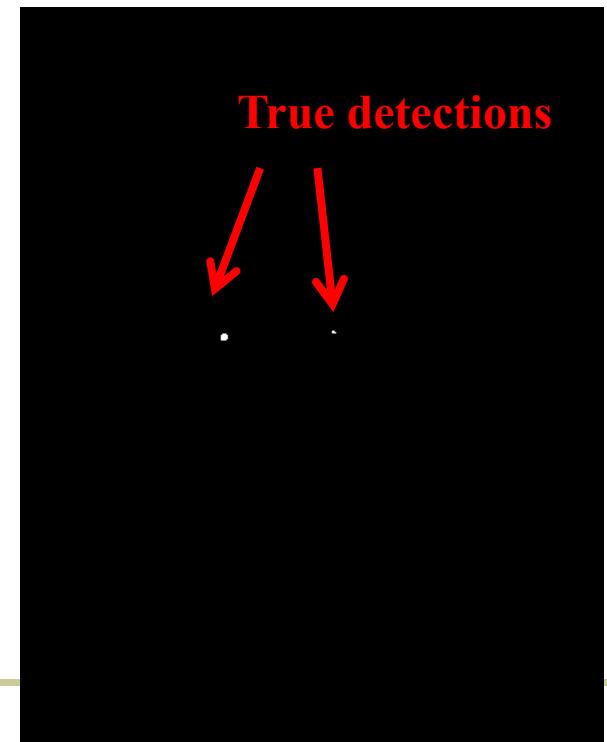
- Goal: find in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



True detections



Q: What is the best method to use?



A: Depends

- Zero-mean filter: fastest but not a great matcher
- SSD: next fastest, sensitive to overall intensity
- Normalized cross-correlation: slowest, invariant to local average intensity and contrast



Today's class



- Template matching
- **Gaussian pyramids**
 - Application for recognition
 - Pyramids representation in deep learning
- Laplacian pyramids
 - Hybrid images
 - Application for image blending
- Steerable pyramids
 - Steerable filters
 - Orientation analysis
- Human vision system
 - Visual perception
- Filter Banks and Texture Analysis



Translation invariance



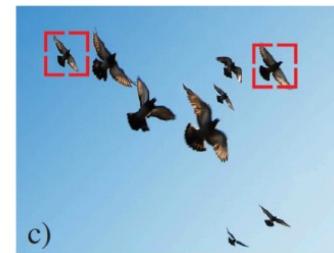
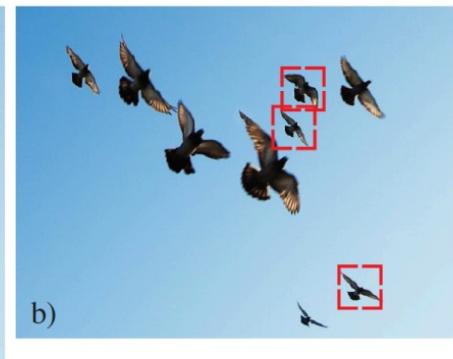
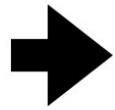


We need translation and scale invariance



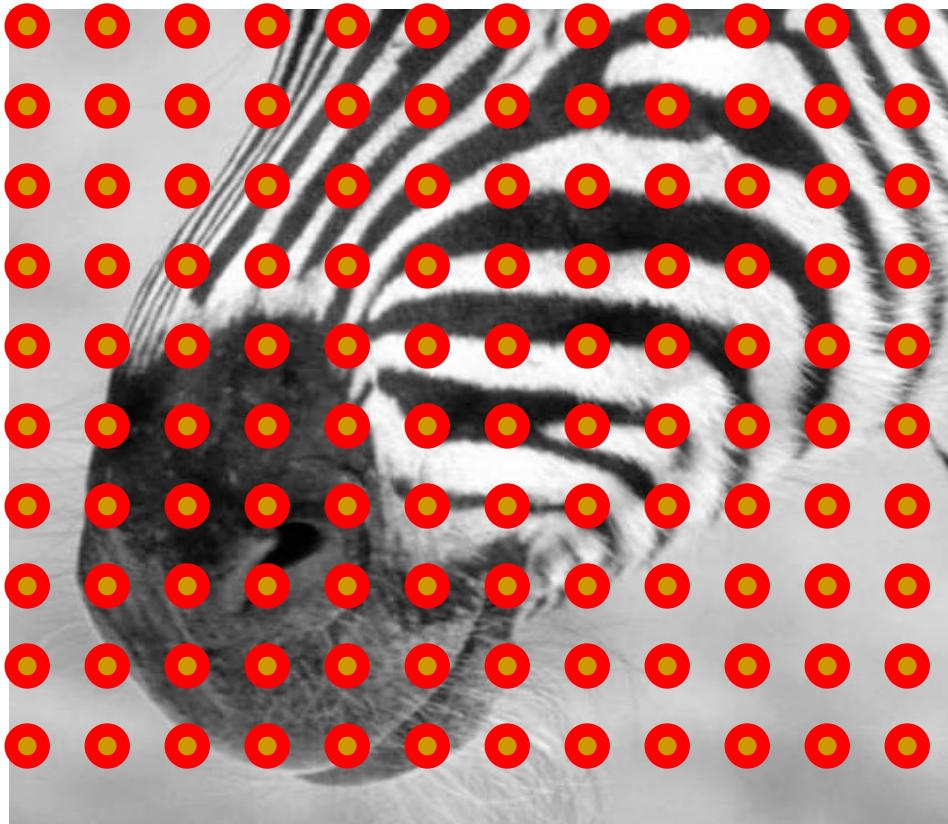


Image pyramids





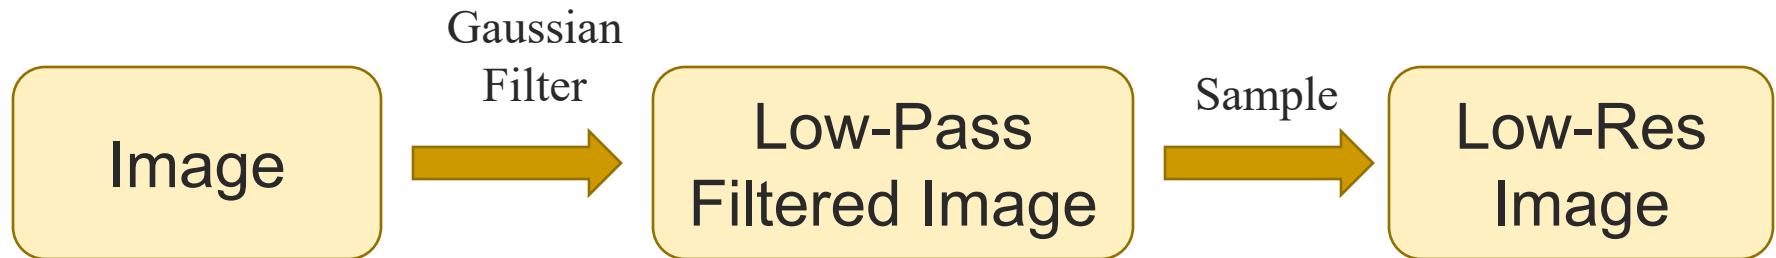
Subsampling by a factor of 2



Throw away every other row and column
to create a 1/2 size image



Recall: sampling



E. H. Adelson | C. H. Anderson | J. R. Bergen | P. J. Burt | J. M. Ogden

Pyramid methods in image processing

http://persci.mit.edu/pub_pdfs/RCA84.pdf

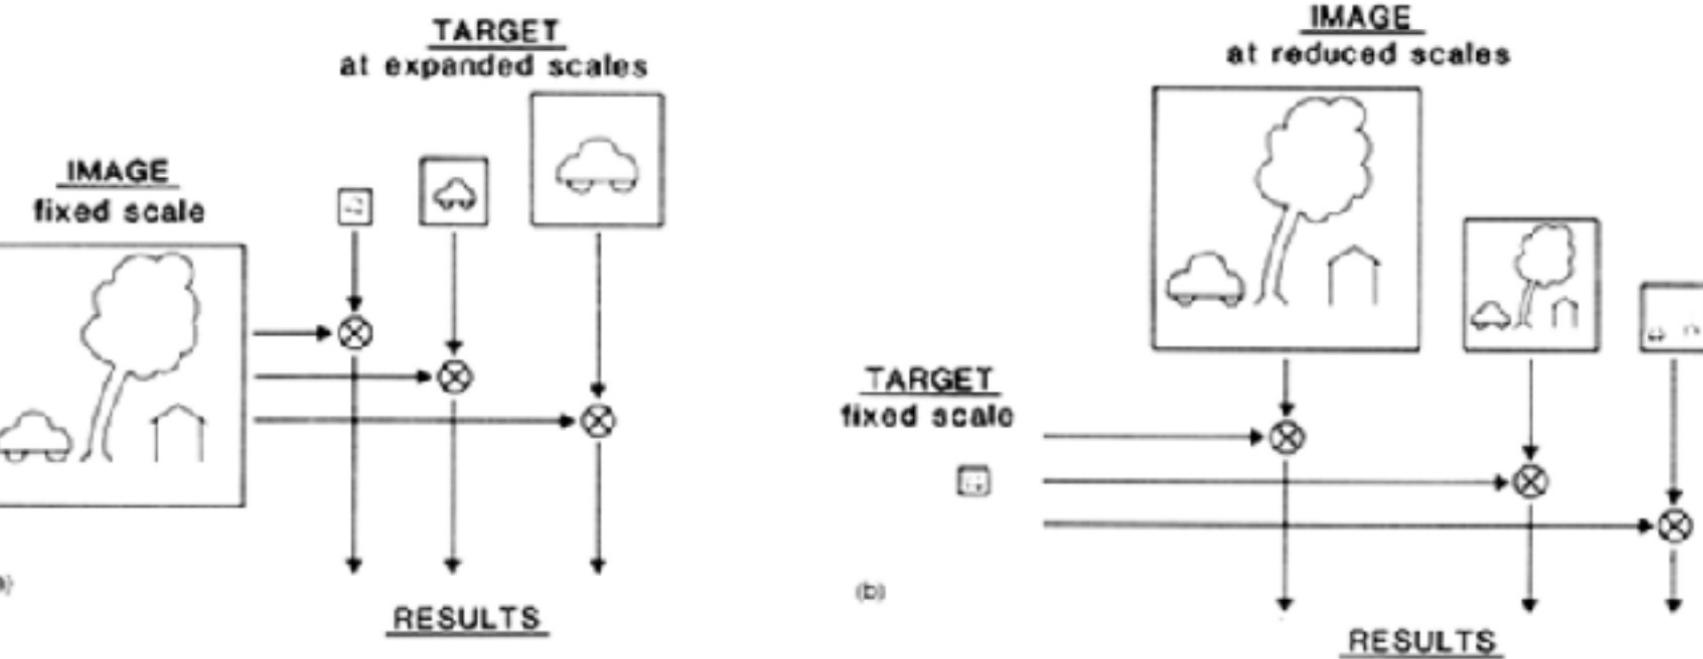


Fig. 1. Two methods of searching for a target pattern over many scales. In the first approach, (a), copies of the target pattern are constructed at several expanded scales, and each is convolved with the original image. In the second approach, (b), a single copy of the target is convolved with

copies of the image reduced in scale. The target should be just large enough to resolve critical details. The two approaches should give equivalent results, but the second is more efficient by the fourth power of the scale factor (image convolutions are represented by 'O').



Fig. 2a. The Gaussian pyramid. The original image, G_0 is repeatedly filtered and subsampled to generate the sequence of reduced resolution image G_1 , G_2 , etc. These comprise a set of lowpass-filtered copies of the original image in which the bandwidth decreases in one-octave steps.



$G_{0,0}$



$G_{1,1}$



$G_{2,2}$

Fig. 2b. Levels of the Gaussian pyramid expanded to the size of the original image.
The effects of lowpass filtering are now clearly apparent.



512 256 128 64 32 16 8



Gaussian Pyramid

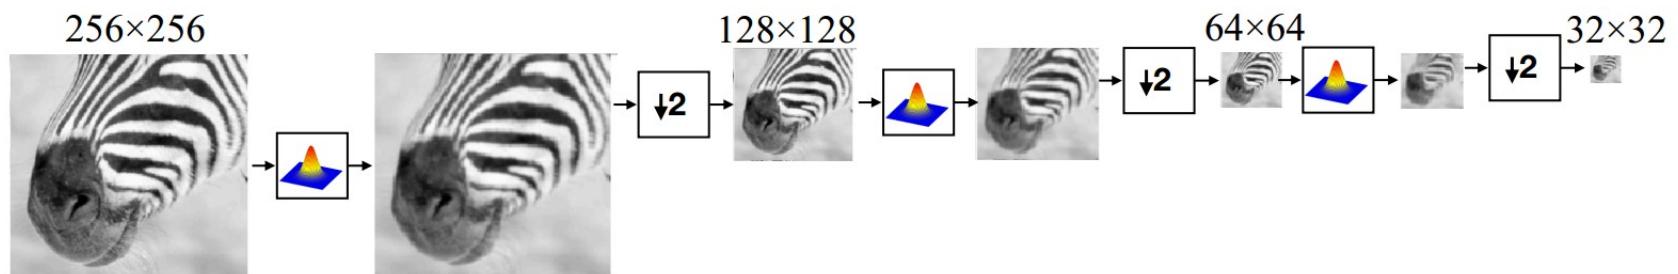
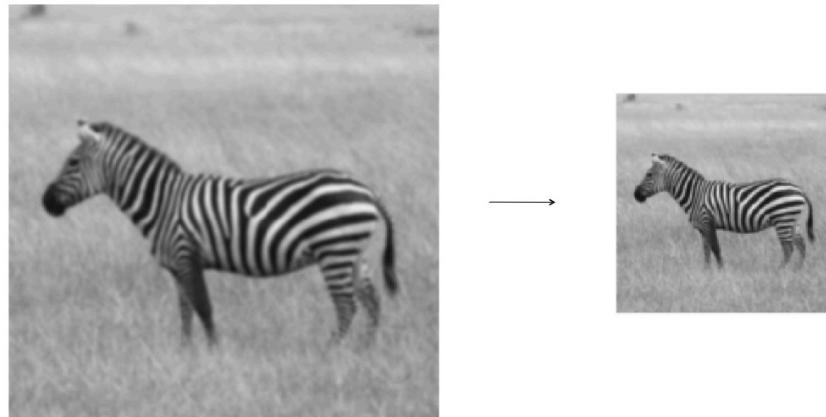


Gaussian Pyramid



For each level

1. Blur input image with a Gaussian filter
2. Downsample image





Downsampling & Upsampling

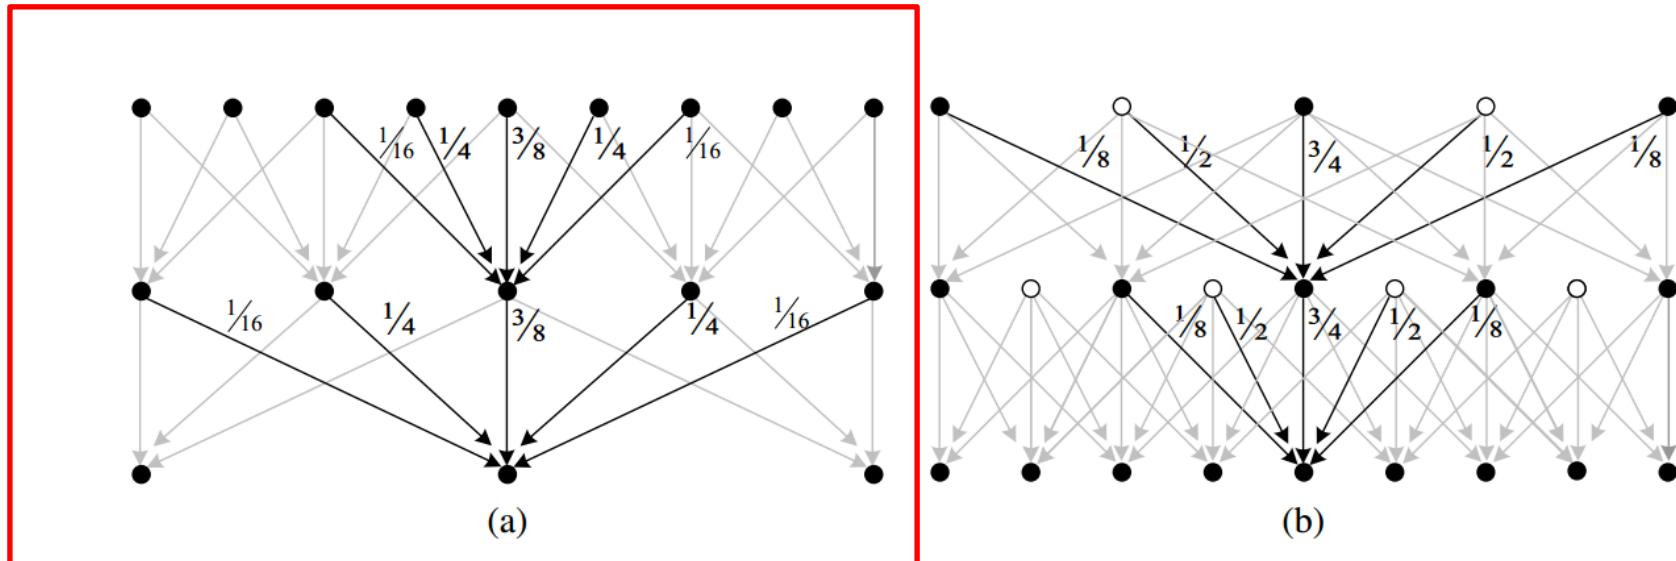
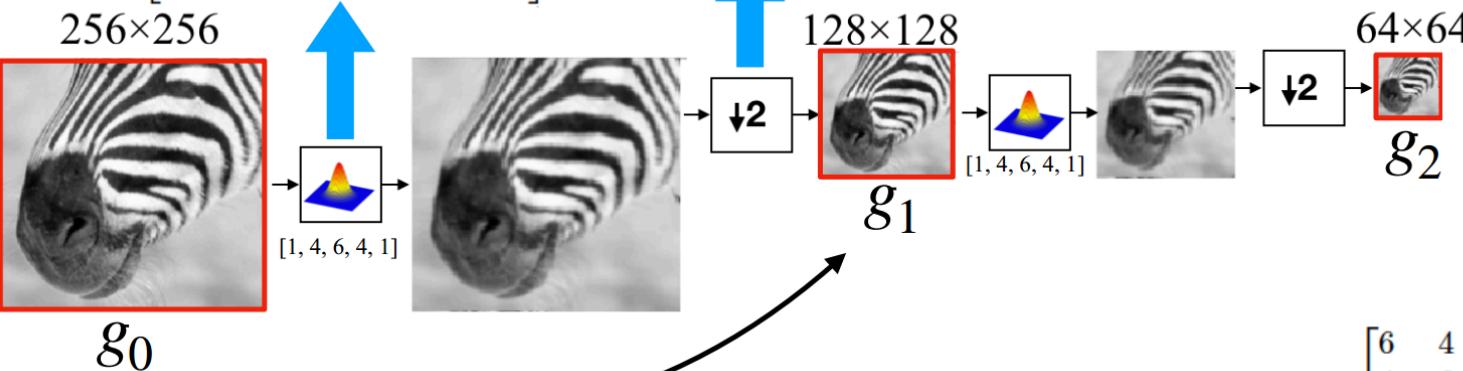


Figure 3.33 The Gaussian pyramid shown as a signal processing diagram: The (a) analysis and (b) re-synthesis stages are shown as using similar computations. The white circles indicate zero values inserted by the $\uparrow 2$ upsampling operation. Notice how the reconstruction filter coefficients are twice the analysis coefficients. The computation is shown as flowing down the page, regardless of whether we are going from coarse to fine or *vice versa*.

The Gaussian pyramid

$$\frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$



g_0

$$g_1 = G_0 g_0$$

$$G_0 = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} \frac{1}{16}$$

$$\begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

The Gaussian pyramid

$$\frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

256×256



g_0

[1, 4, 6, 4, 1]



128×128

g_1



$$G_0 = \frac{1}{16} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \end{bmatrix}$$

$$g_1 = G_0 g_0$$

$$[1, 4, 6, 4, 1]$$

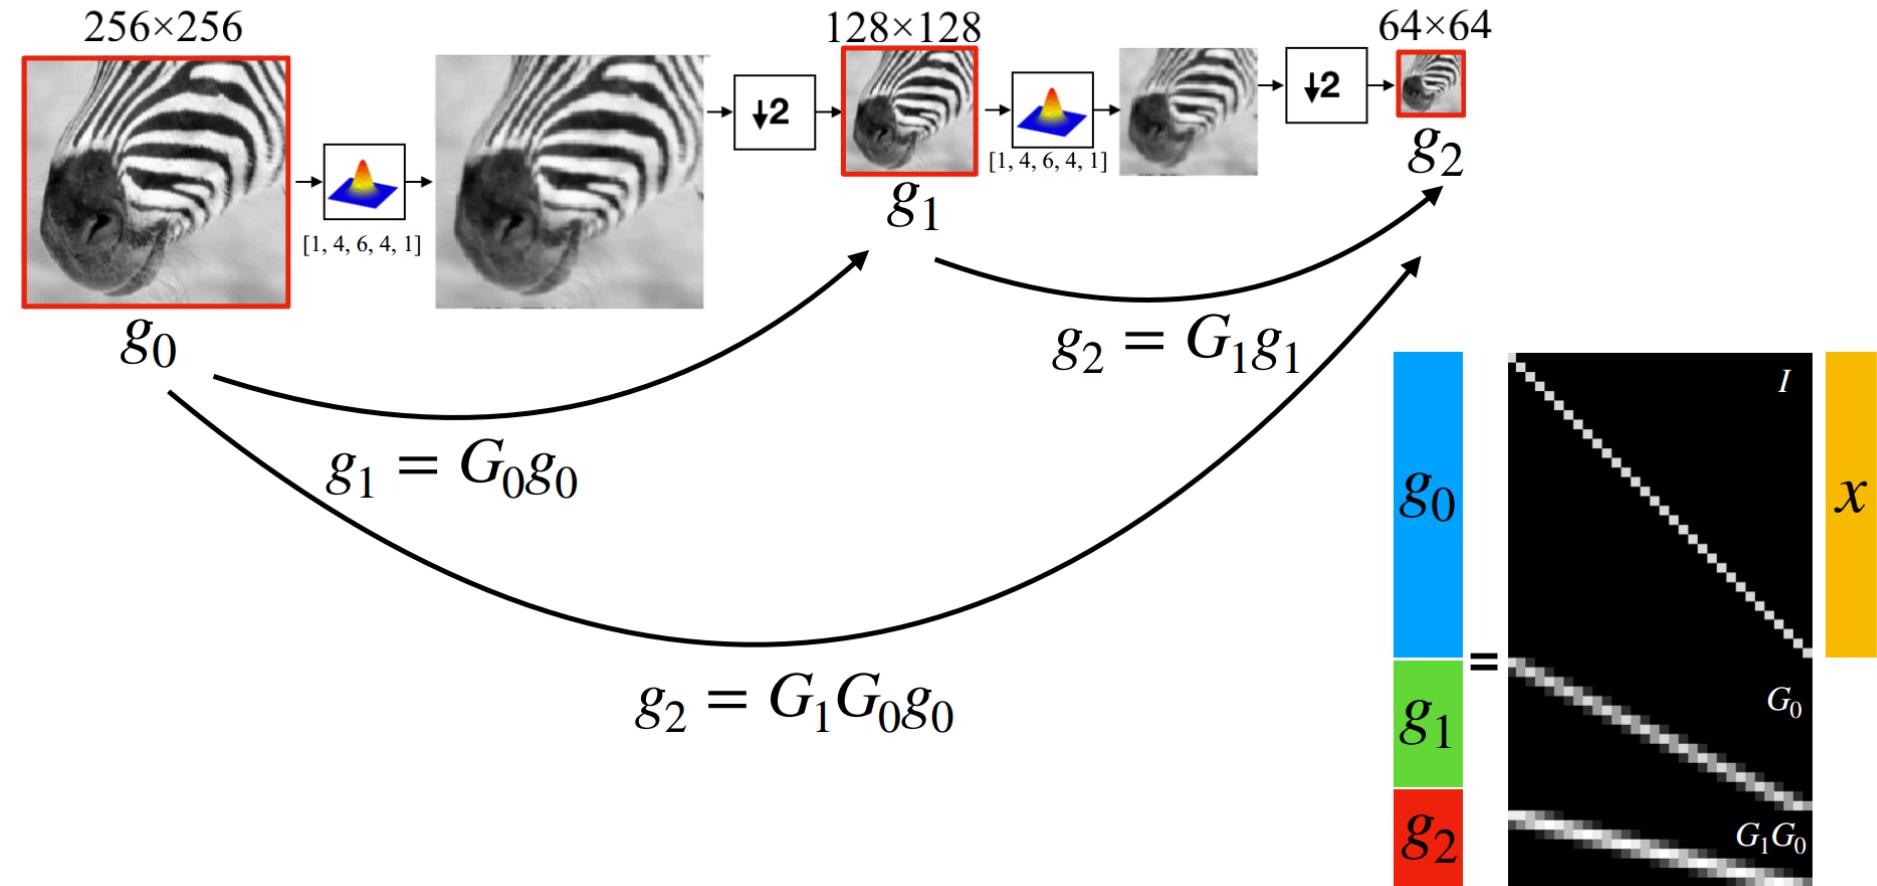
64×64



g_2

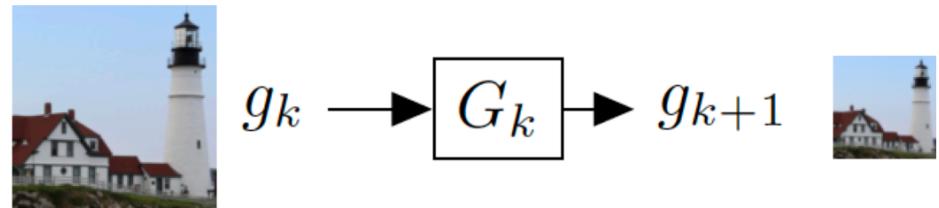


The Gaussian pyramid





Gaussian Pyramid



For each level

1. Blur input image with a Gaussian filter
2. Downsample image



Gaussian pyramids used for



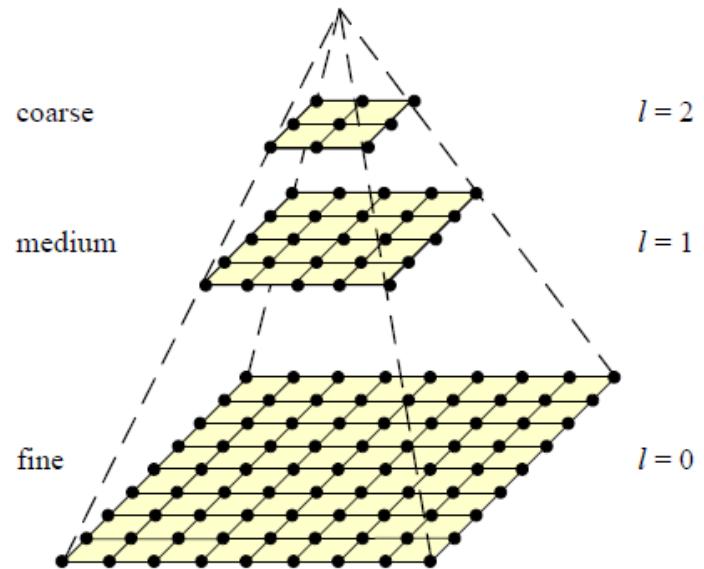
- Up or down sample images
- Multi-resolution image analysis
 - Look for an object over various spatial scales
 - Coarse-to-fine image registration: form blur estimate or the motion analysis on very low-resolution image, upsample and repeat
 - Often a successful strategy for avoiding local minima in complicated estimation tasks



Coarse-to-fine Image Registration



1. Compute Gaussian pyramid
2. Align with coarse pyramid
3. Successively align with finer pyramids
 - Search smaller range



Why is this faster?

Are we guaranteed to get the same result?



Template Matching with Image Pyramids

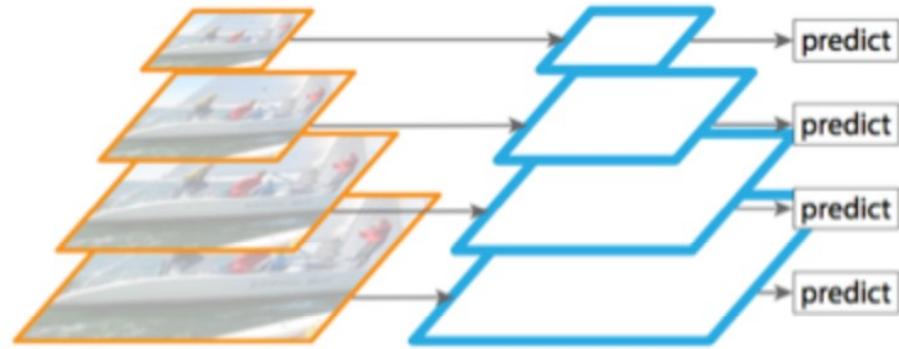
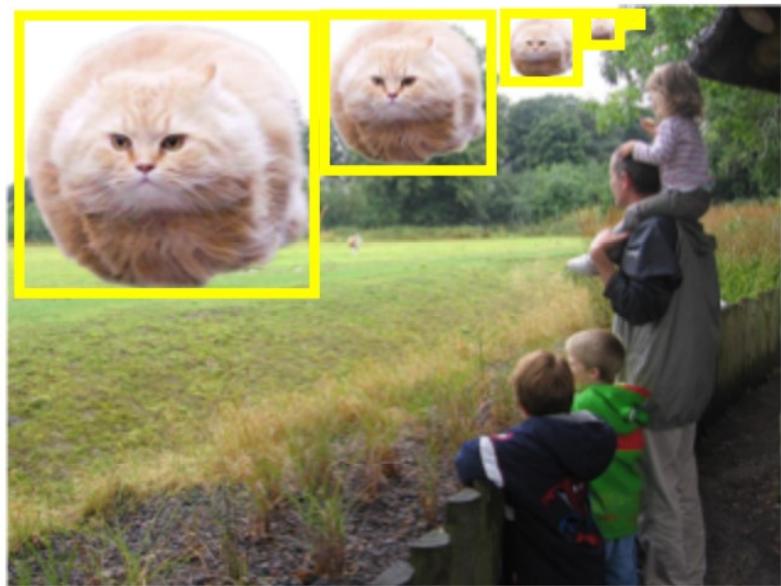


Input: Image, Template

1. Match template at current scale
2. Downsample image
 - In practice, scale step of 1.1 to 1.2
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression



From Image Pyramid to Feature Pyramid



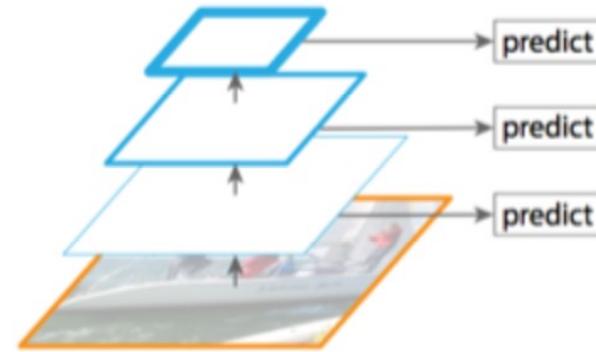
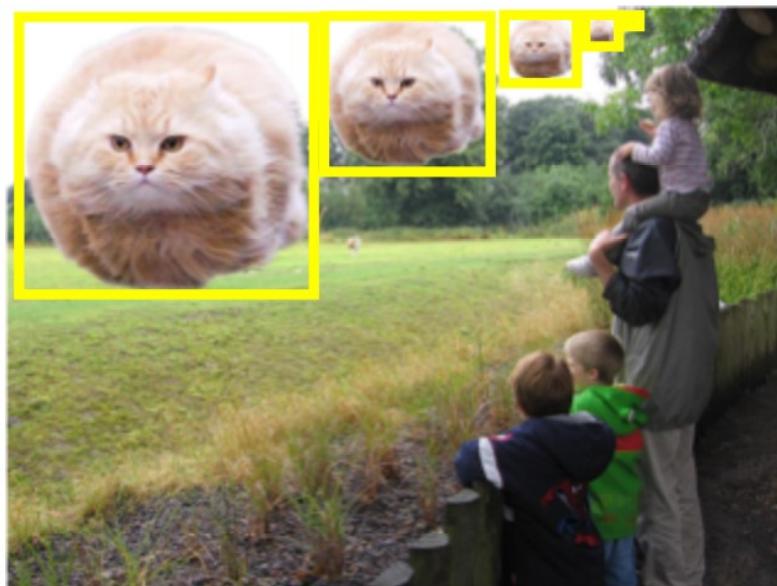
(a) Featurized image pyramid

Standard solution – slow!

(E.g., Viola & Jones, HOG, DPM, SPP-net, multi-scale Fast R-CNN, ...)



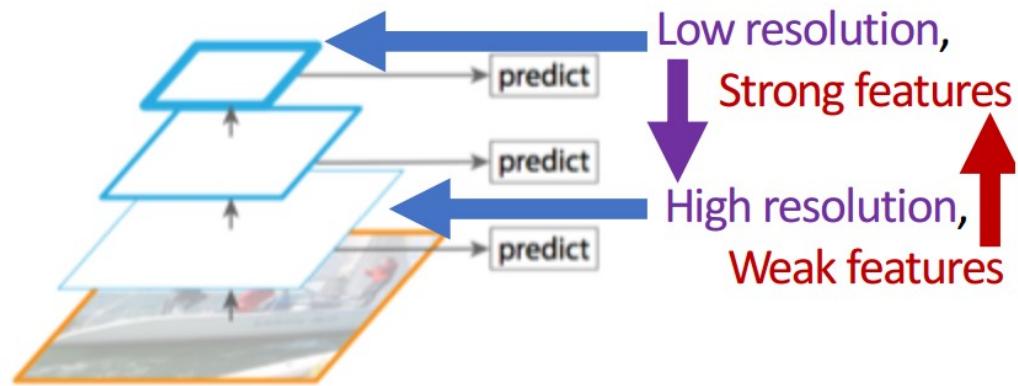
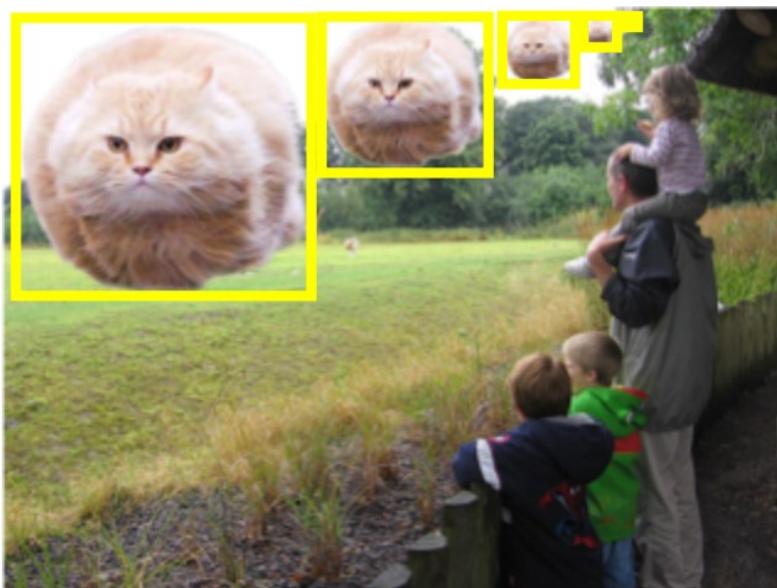
Feature Pyramid Networks



(c) Pyramidal feature hierarchy
Use the internal pyramid – *fast, suboptimal*
(E.g., \approx SSD, ...)



Feature Pyramid Networks



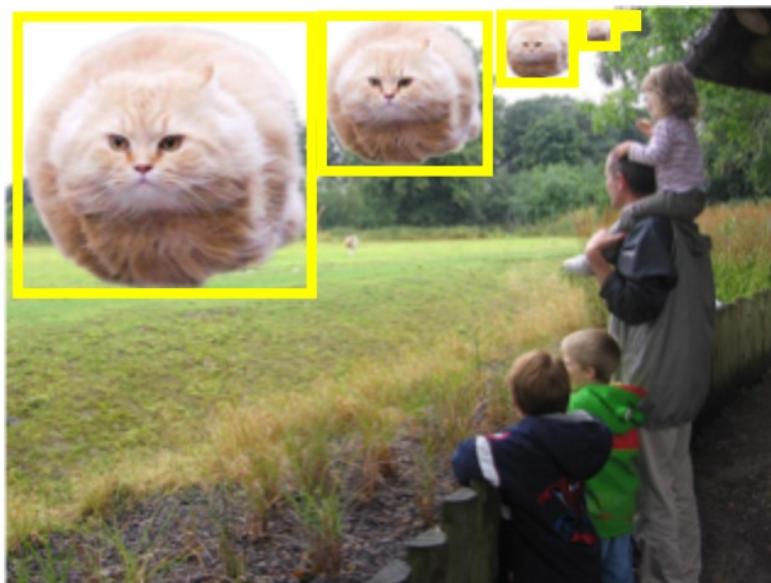
(c) Pyramidal feature hierarchy

Use the internal pyramid – *fast, suboptimal*

(E.g., \approx SSD, ...)



Feature Pyramid Networks

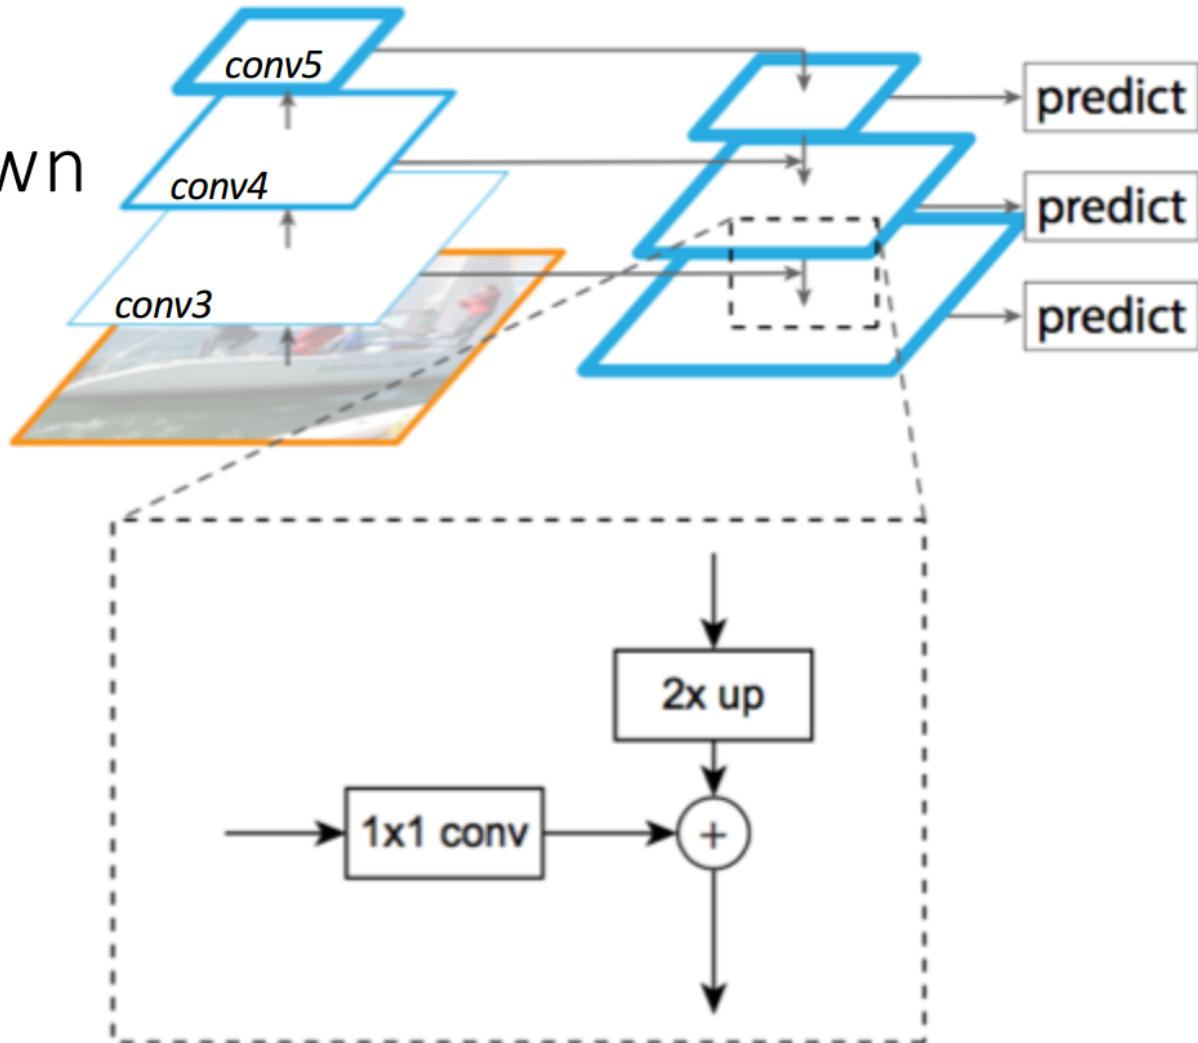


(d) Feature Pyramid Network

Top-down enrichment of high-res features –
fast, less suboptimal

Lin et al. Feature Pyramid Networks for Object Detection. CVPR 2017.

FPN Top-down Refinement Module



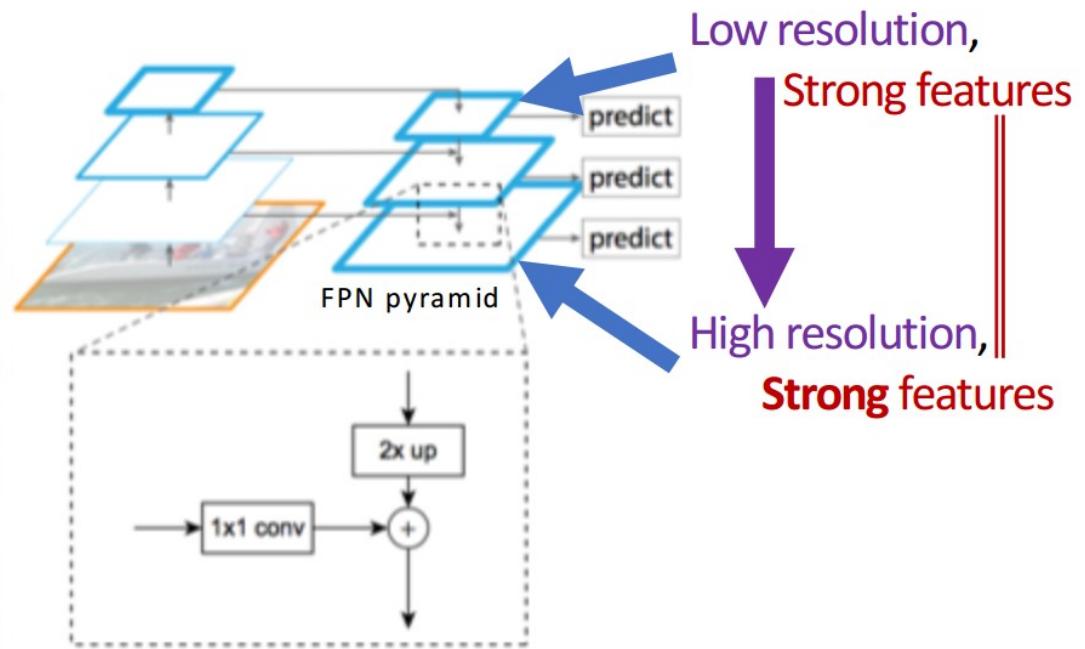
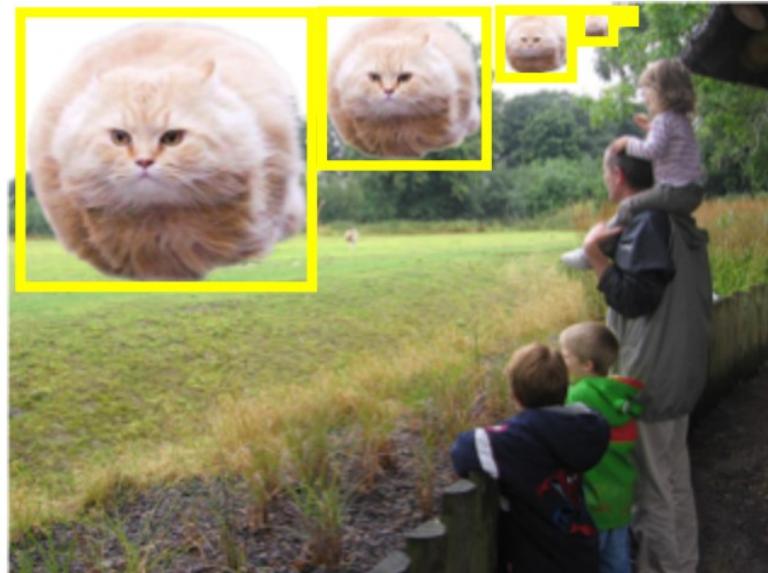
Lin et al. Feature Pyramid Networks for Object Detection. CVPR 2017.



Feature Pyramid Networks



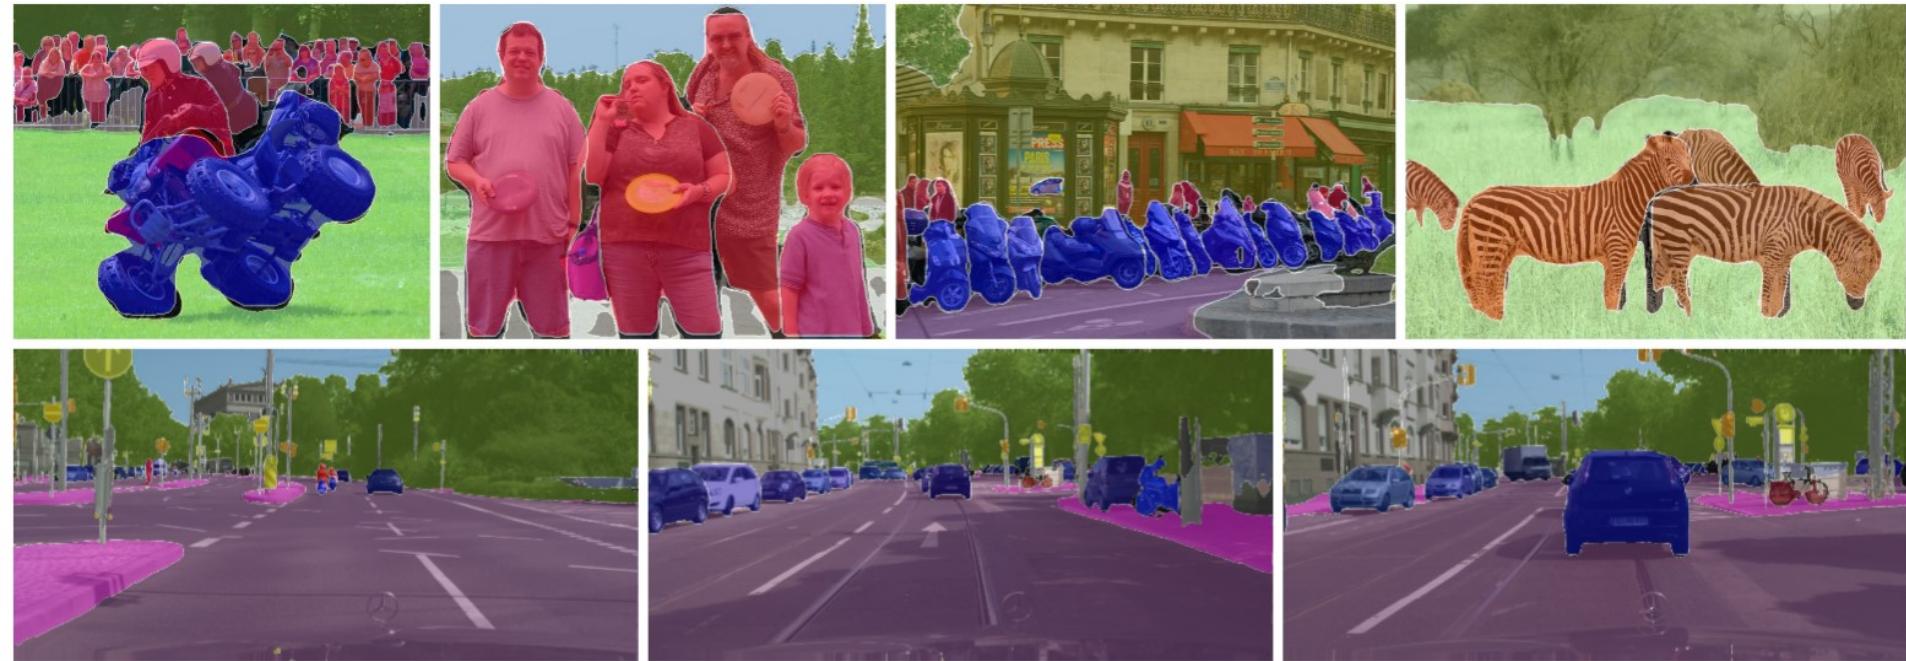
No Compromise on Feature Quality, still Fast



Lin et al. Feature Pyramid Networks for Object Detection. CVPR 2017. See also: Shrivastava's TDM.

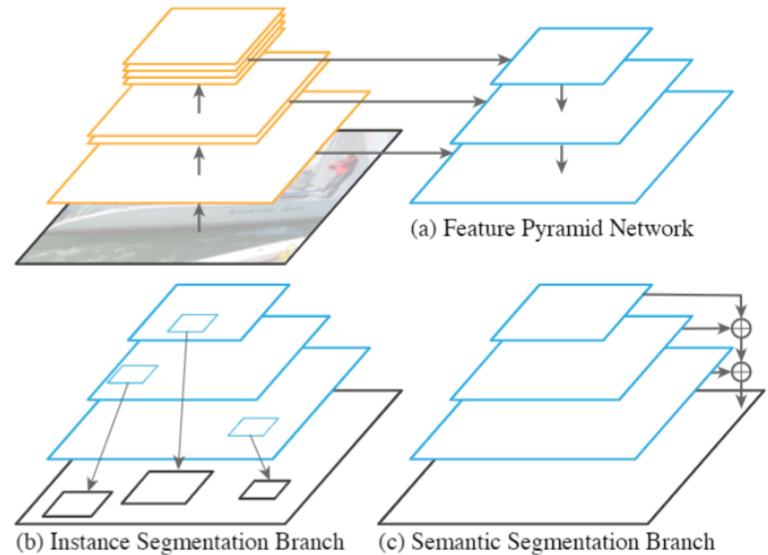
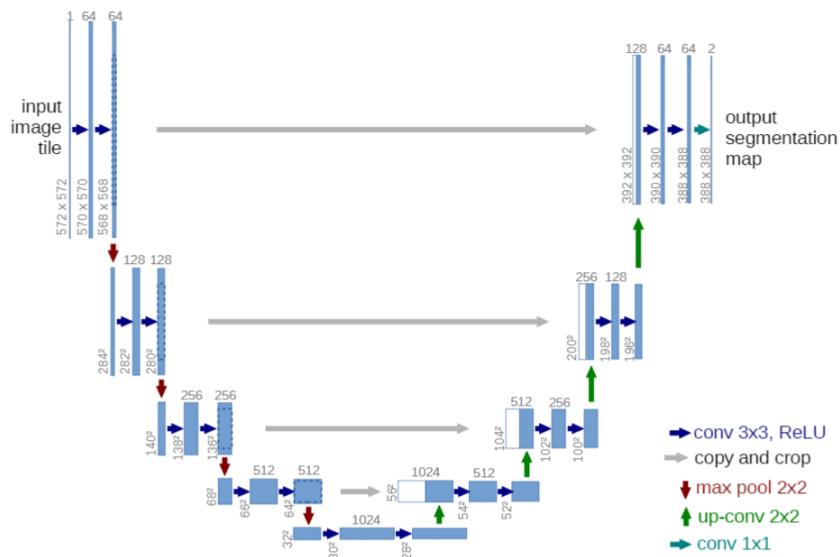


Segmentation





Pyramid in segmentation

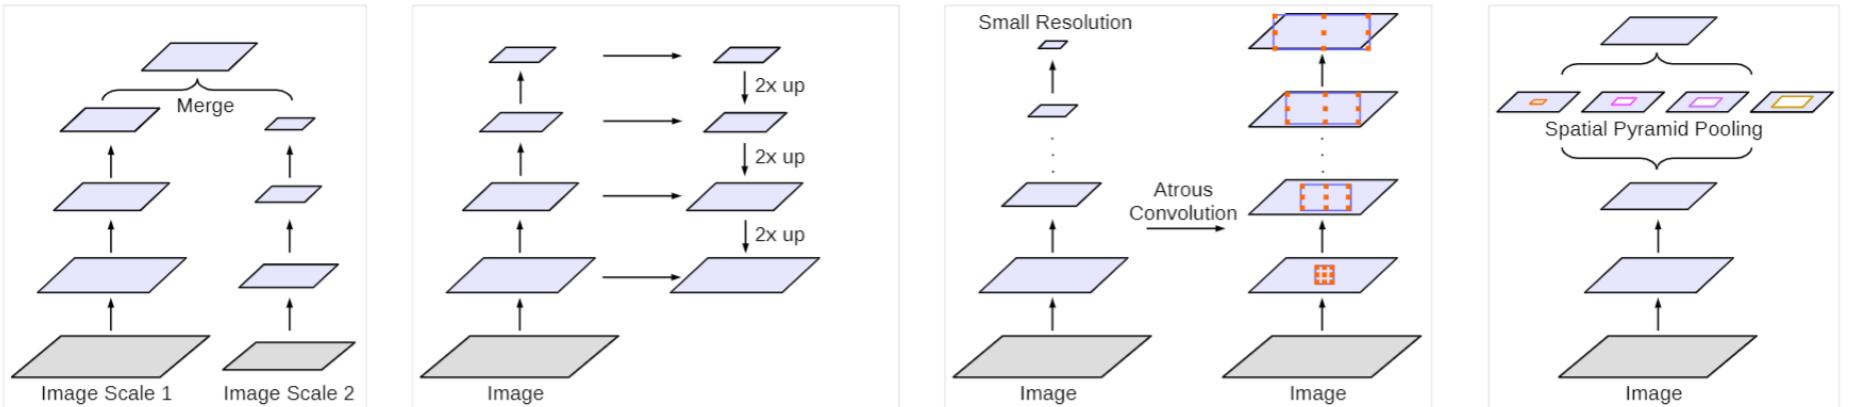


U-Net: Convolutional Networks for Biomedical Image Segmentation

Panoptic Feature Pyramid Networks



Atrous Convolution



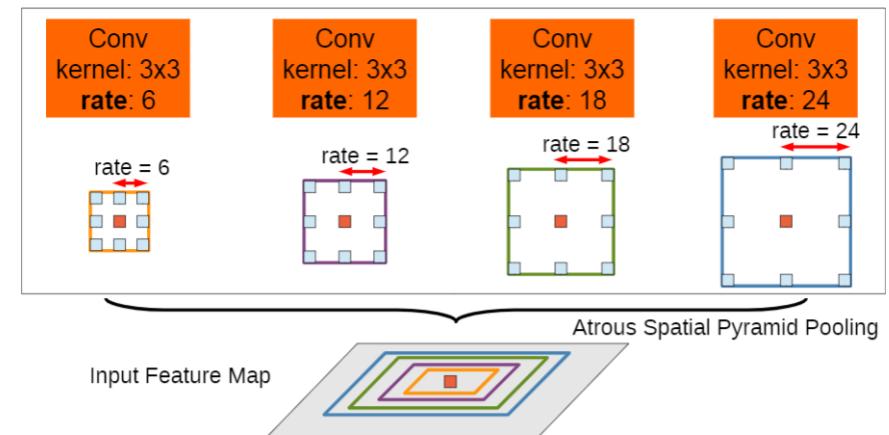
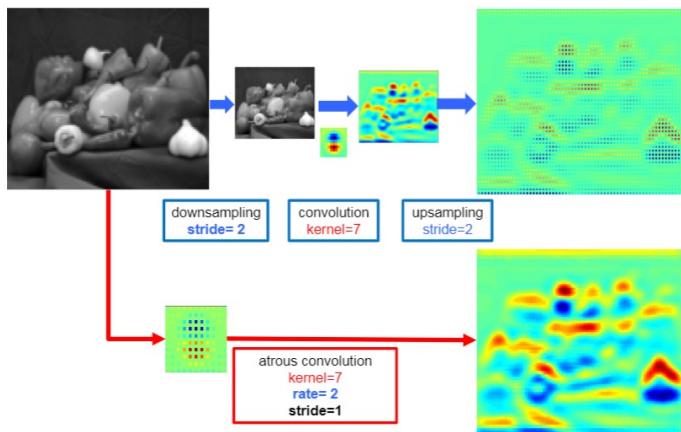
(a) Image Pyramid

(b) Encoder-Decoder

(c) Deeper w. Atrous Convolution

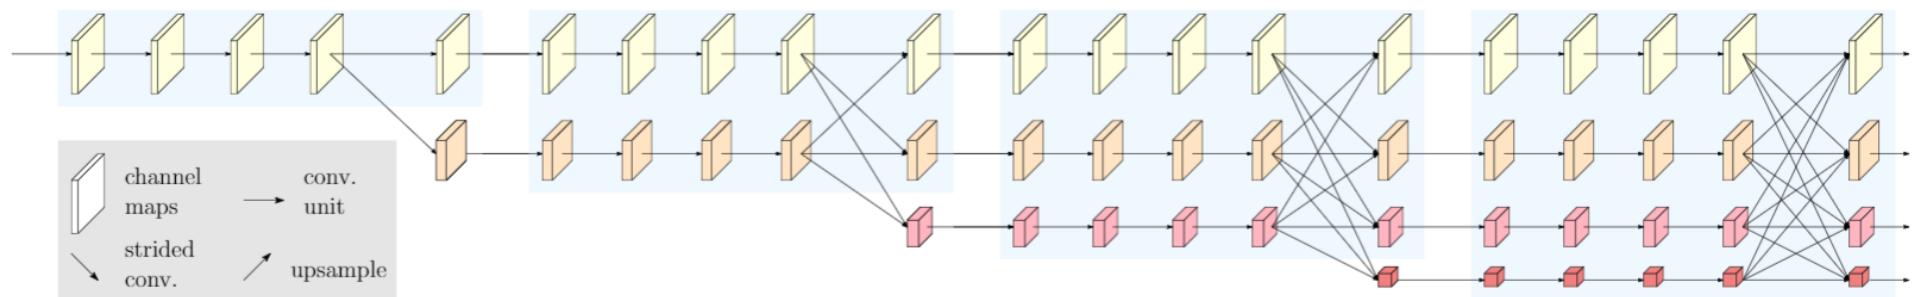
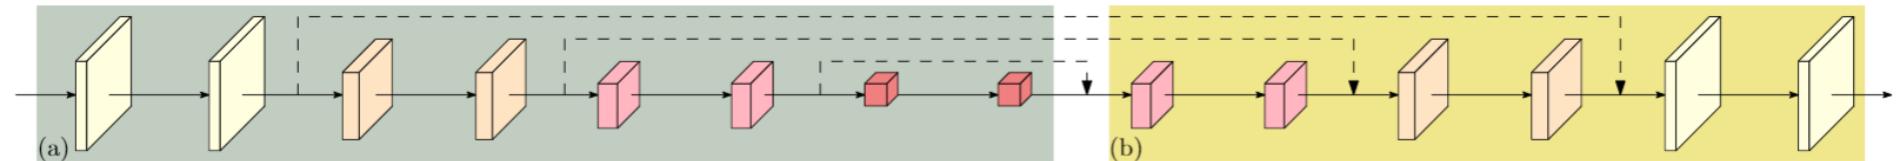
(d) Spatial Pyramid Pooling

Figure 2. Alternative architectures to capture multi-scale context.

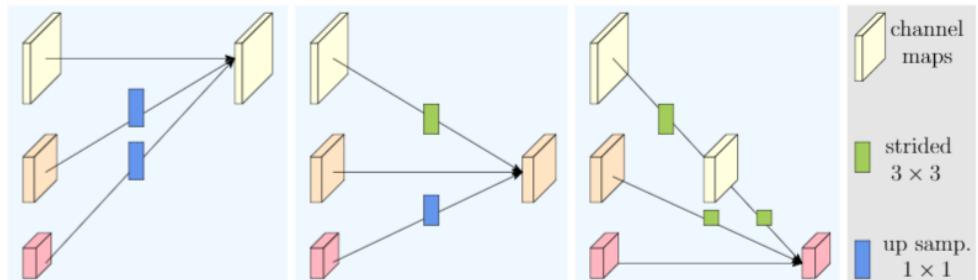




HRNet



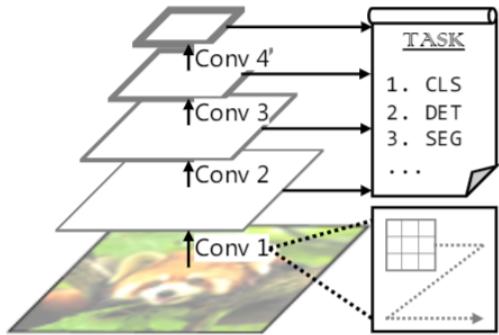
$$\begin{array}{ccccccccc} \mathcal{N}_{11} & \rightarrow & \mathcal{N}_{21} & \rightarrow & \mathcal{N}_{31} & \rightarrow & \mathcal{N}_{41} \\ & \searrow & & & \searrow & & \\ & & \mathcal{N}_{22} & \rightarrow & \mathcal{N}_{32} & \rightarrow & \mathcal{N}_{42} \\ & & & & \searrow & & \\ & & & & \mathcal{N}_{33} & \rightarrow & \mathcal{N}_{43} \\ & & & & & & \searrow \\ & & & & & & \mathcal{N}_{44}, \end{array}$$



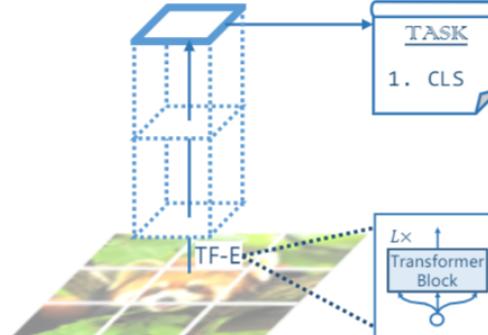
Deep High-Resolution Representation Learning for Visual Recognition



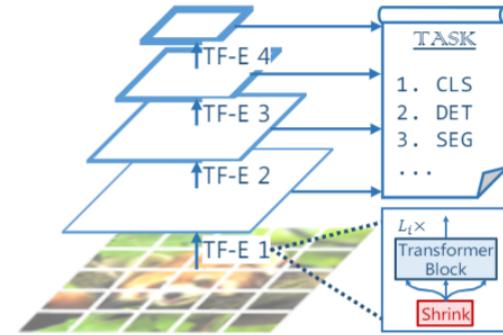
Pyramid in Transformer



(a) CNNs: VGG [41], ResNet [15], etc.

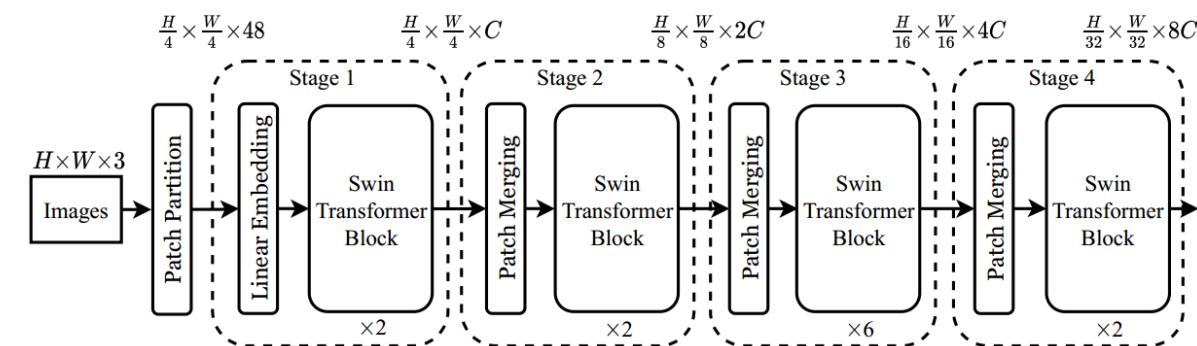
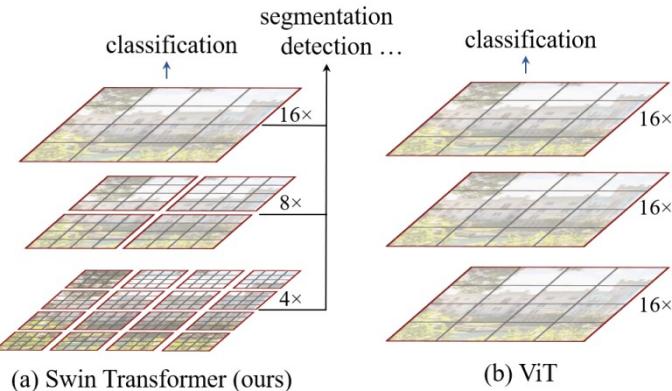


(b) Vision Transformer [10]



(c) Pyramid Vision Transformer (ours)

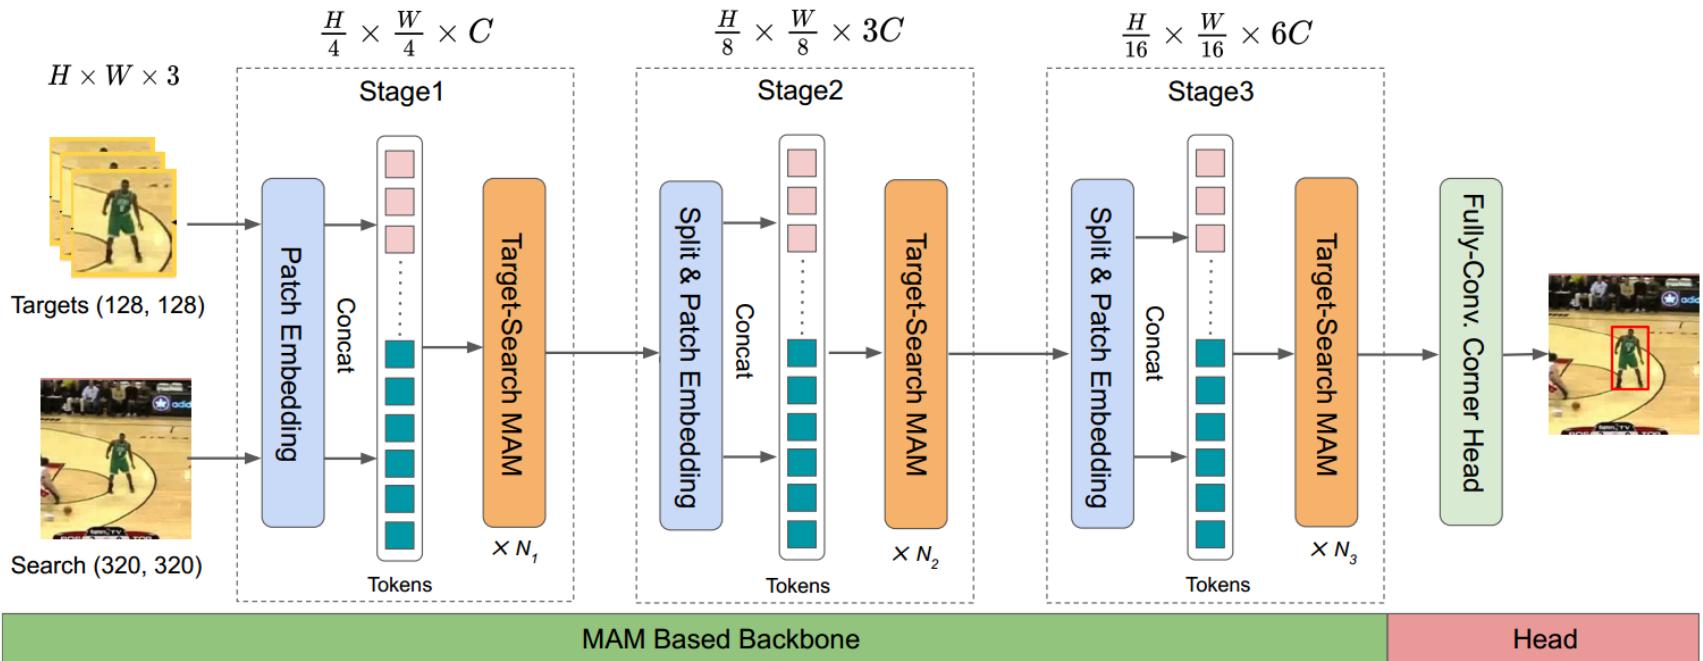
Pyramid Vision Transformer: A Versatile Backbone for Dense Prediction without Convolutions



Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

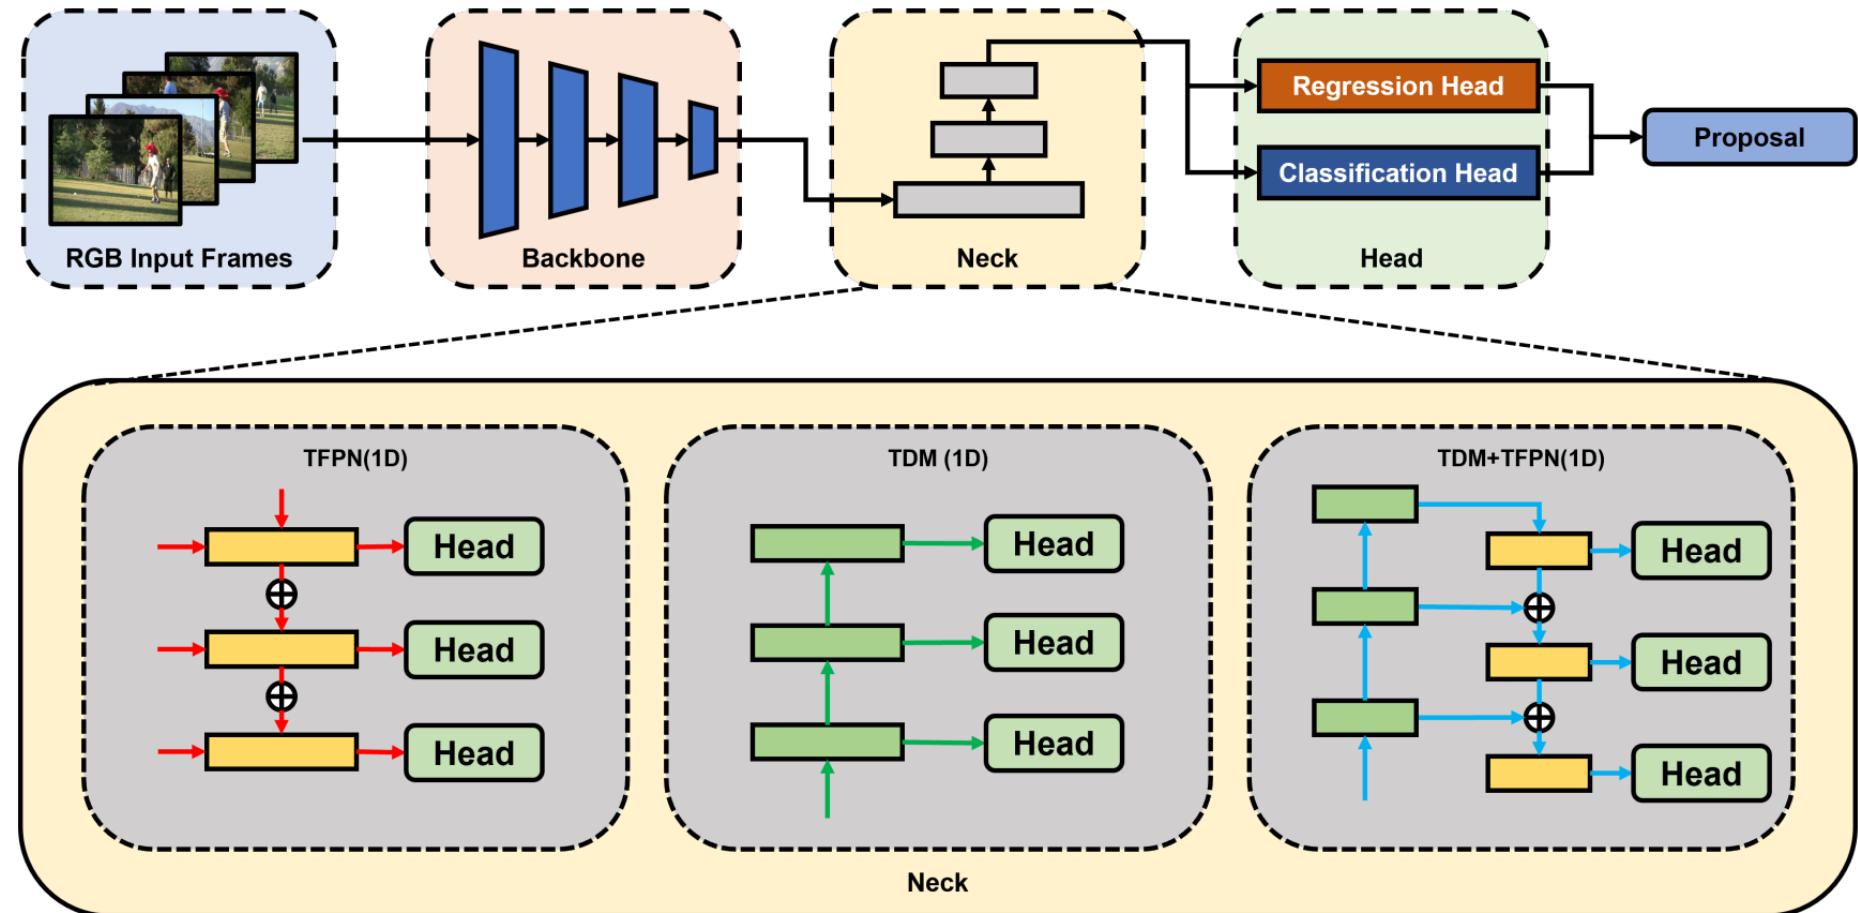


MixFormer in Tracking



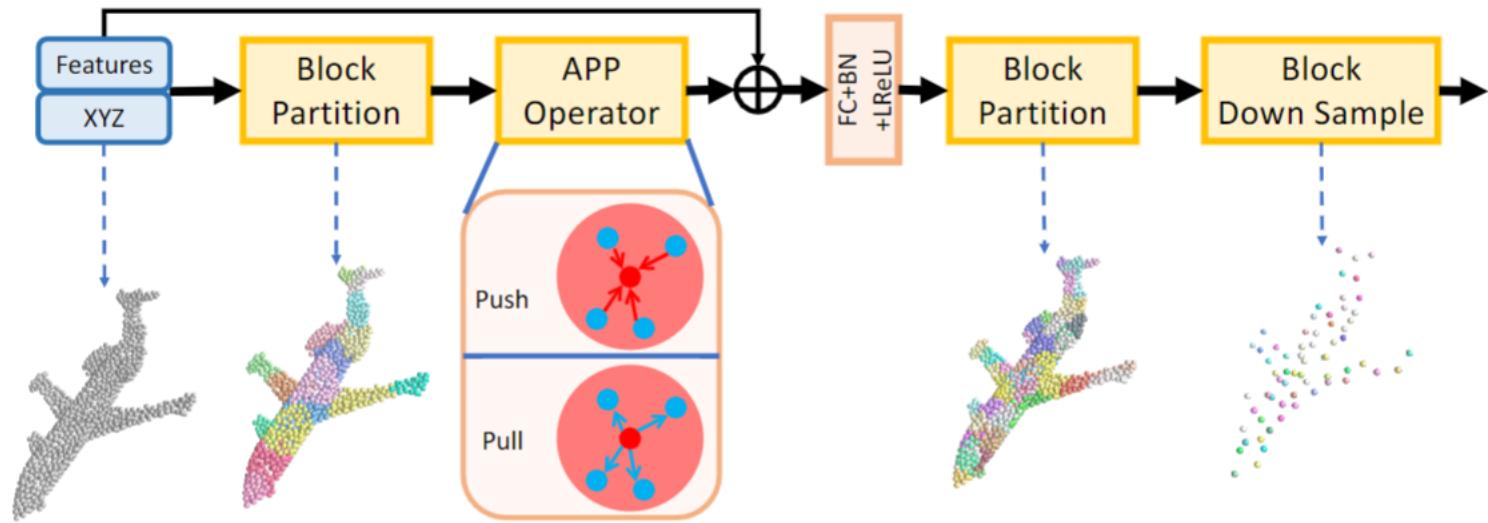


Temporal Pyramid: Image to Video





Pyramid in Point Cloud





Today's class



- Template matching
- Gaussian pyramids
 - Application for recognition
 - Pyramids representation in deep learning
- **Laplacian pyramids**
 - Hybrid images
 - Application for image blending
- Steerable pyramids
 - Steerable filters
 - Orientation analysis
- Human vision system
 - Visual perception
- Filter Banks and Texture Analysis



L_0



L_1



L_2



L_3



L_4

Fig. 4a. The Laplacian pyramid. Each level of this band-pass pyramid represents the difference between successive levels of the Gaussian pyramid.



$L_{0.0}$



$L_{1.0}$



$L_{2.2}$

Fig. 4b. Levels of the Laplacian pyramid expanded to the size of the original image.
Note that edge and bar features are enhanced and segregated by size.



512

256

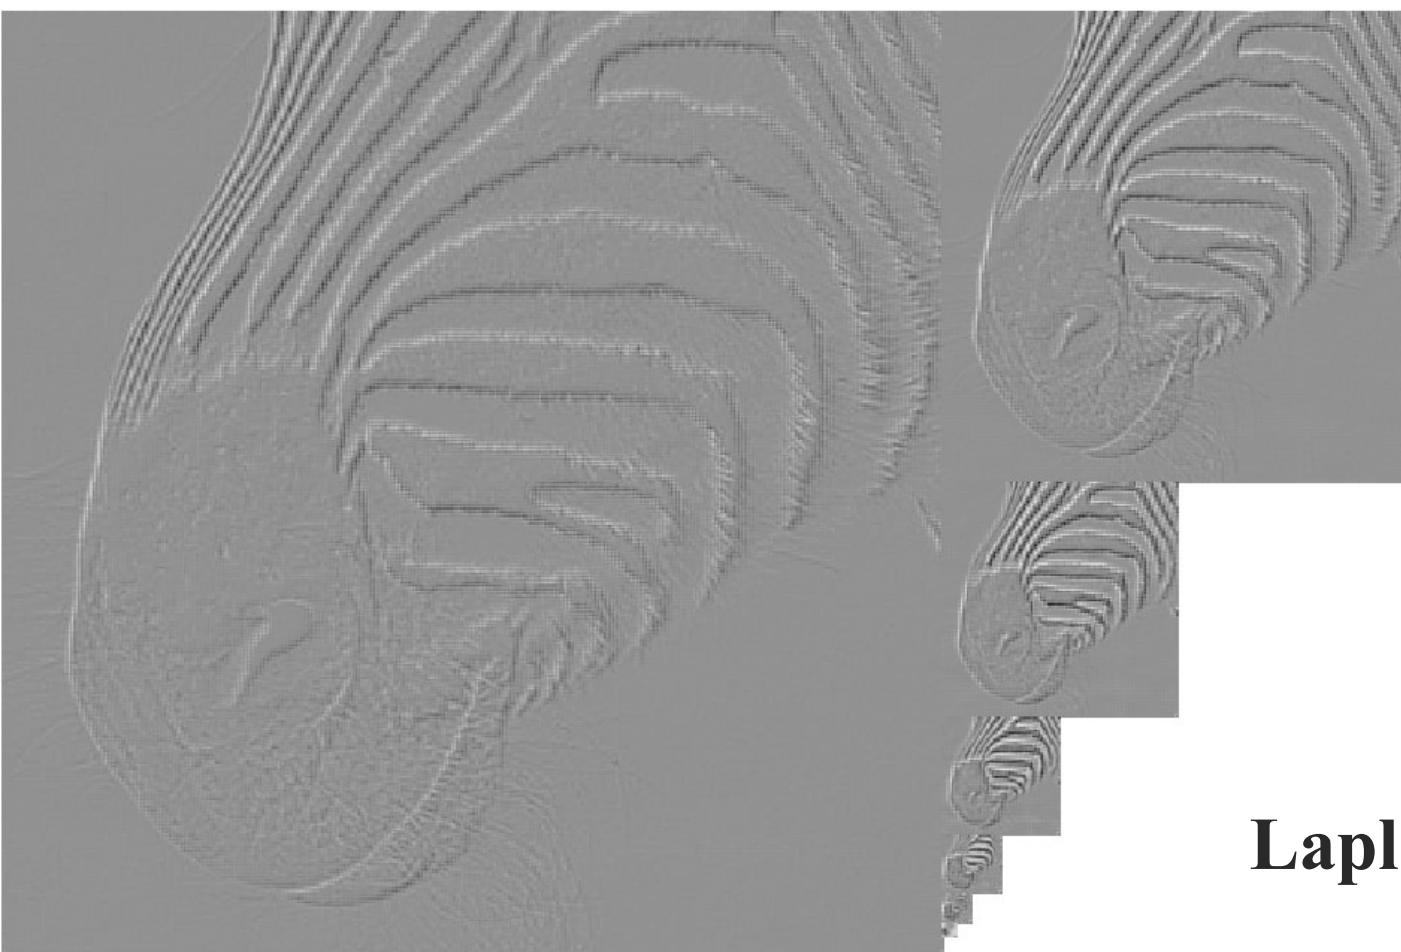
128

64

32

16

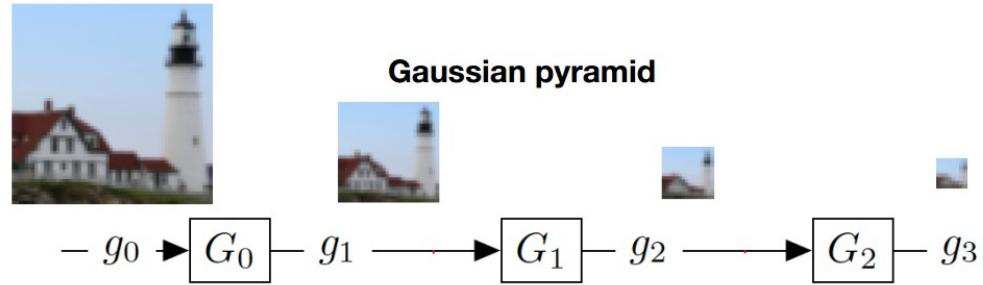
8



Laplacian Pyramid

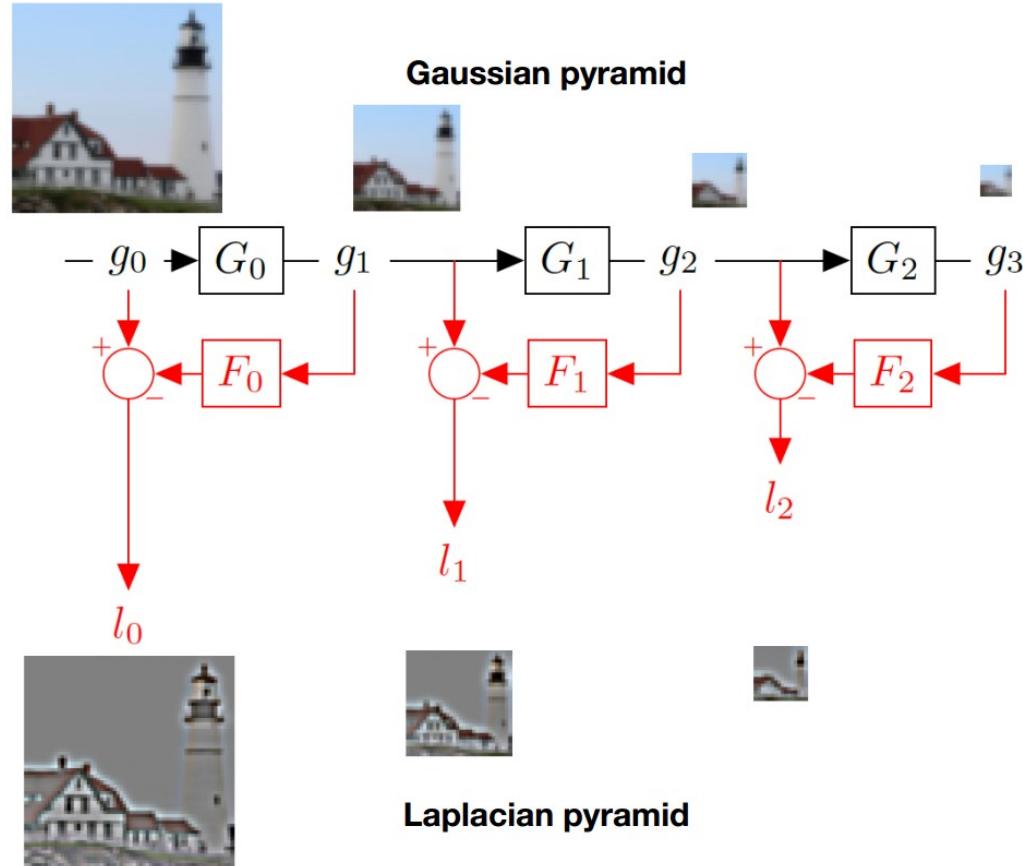


The Laplacian Pyramid





The Laplacian Pyramid

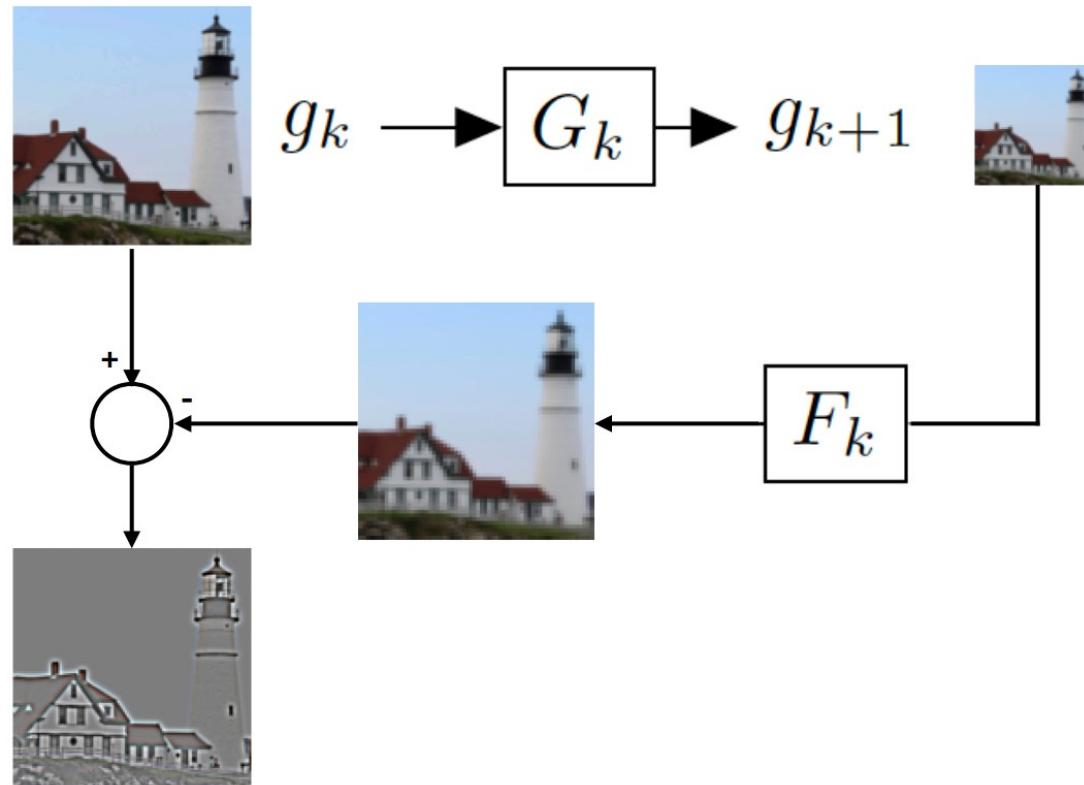




The Laplacian Pyramid

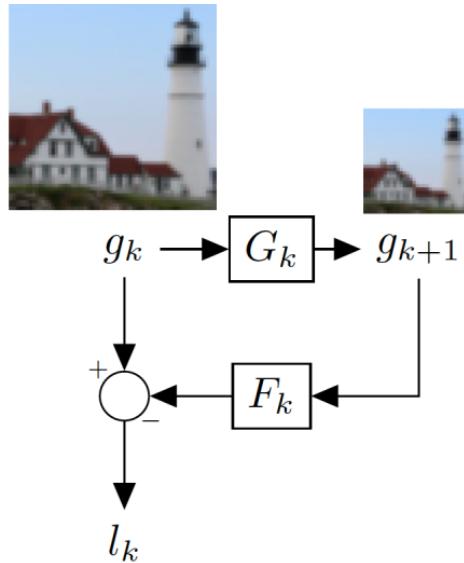


Compute the difference between upsampled Gaussian pyramid level $k+1$ and Gaussian pyramid level k .





The Laplacian Pyramid



Blurring and downsampling:

$$G_0 = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} \frac{1}{16}$$

(Downsampling by 2)

Upsampling and blurring:

$$\begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

(blur)

$$F_0 =$$





Downsampling & Upsampling

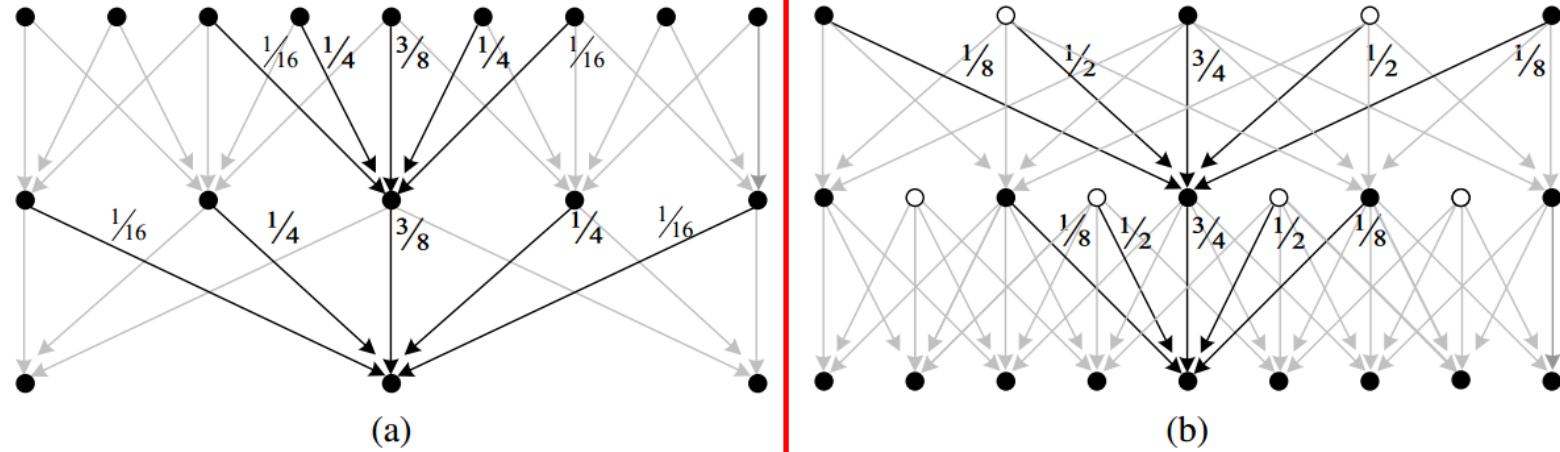


Figure 3.33 The Gaussian pyramid shown as a signal processing diagram: The (a) analysis and (b) re-synthesis stages are shown as using similar computations. The white circles indicate zero values inserted by the $\uparrow 2$ upsampling operation. Notice how the reconstruction filter coefficients are twice the analysis coefficients. The computation is shown as flowing down the page, regardless of whether we are going from coarse to fine or *vice versa*.



Upsampling



Insert zeros
→



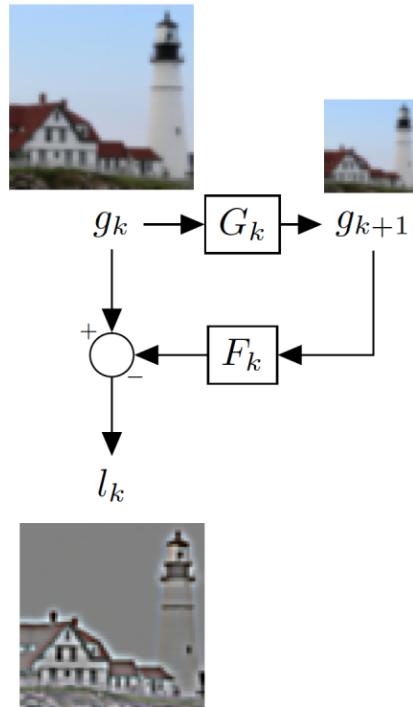
128x128

$$\odot \begin{bmatrix} 0.25 & 0.5 & 0.25 \\ 0.5 & 1 & 0.5 \\ 0.25 & 0.5 & 0.25 \end{bmatrix} =$$





The Laplacian Pyramid



Blurring and downsampling:

$$G_0 = \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

(blur)

Upsampling and blurring:

$$F_0 = \frac{1}{8} \begin{bmatrix} 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 \\ 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 \end{bmatrix}$$

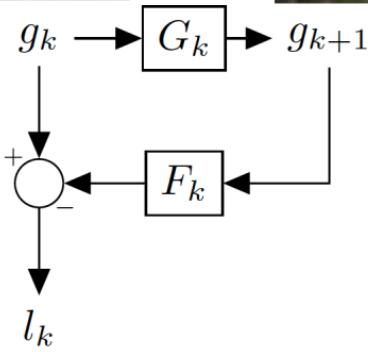
(blur)

$$l_0 = (I_0 - F_0 G_0) g_0$$

22



The Laplacian Pyramid



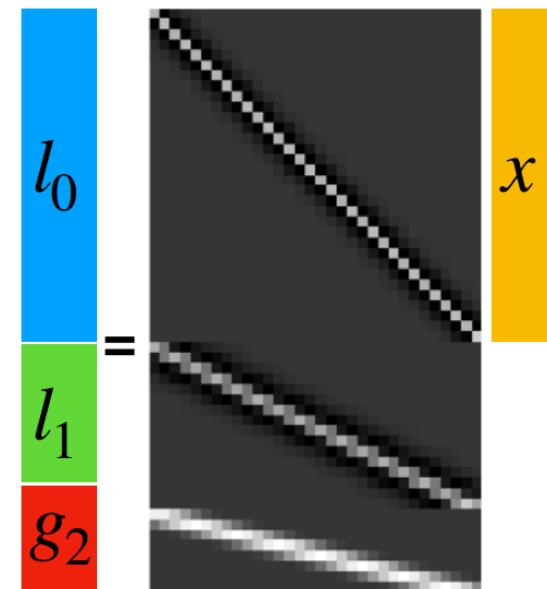
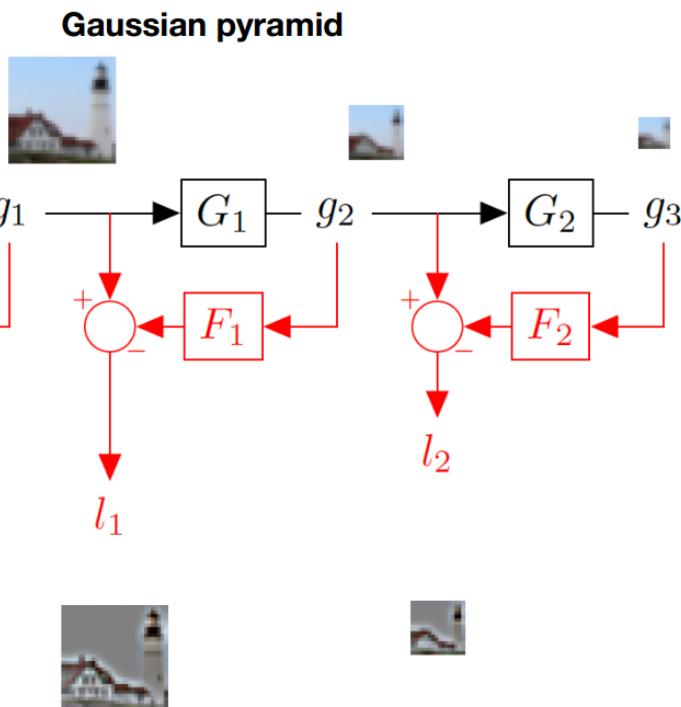
$$l_0 = (I_0 - F_0 G_0) g_0$$

$$= \frac{1}{256} \begin{bmatrix} 182 & -56 & -24 & -8 & -2 & 0 & 0 & 0 \\ -56 & 192 & -56 & -32 & -8 & 0 & 0 & 0 \\ -24 & -56 & 180 & -56 & -24 & -8 & -2 & 0 \\ -8 & -32 & -56 & 192 & -56 & -32 & -8 & 0 \\ -2 & -8 & -24 & -56 & 180 & -56 & -24 & -8 \\ 0 & 0 & -8 & -32 & -56 & 192 & -56 & -32 \\ 0 & 0 & -2 & -8 & -24 & -56 & 182 & -48 \\ 0 & 0 & 0 & 0 & -8 & -32 & -48 & 224 \end{bmatrix} x$$



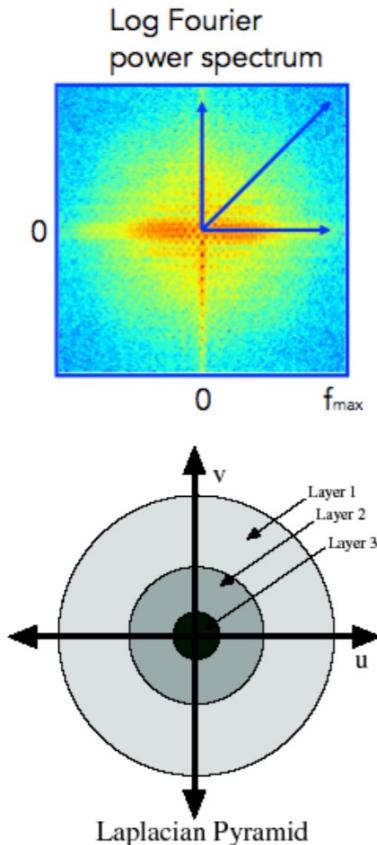


The Laplacian Pyramid





The Laplacian Pyramid



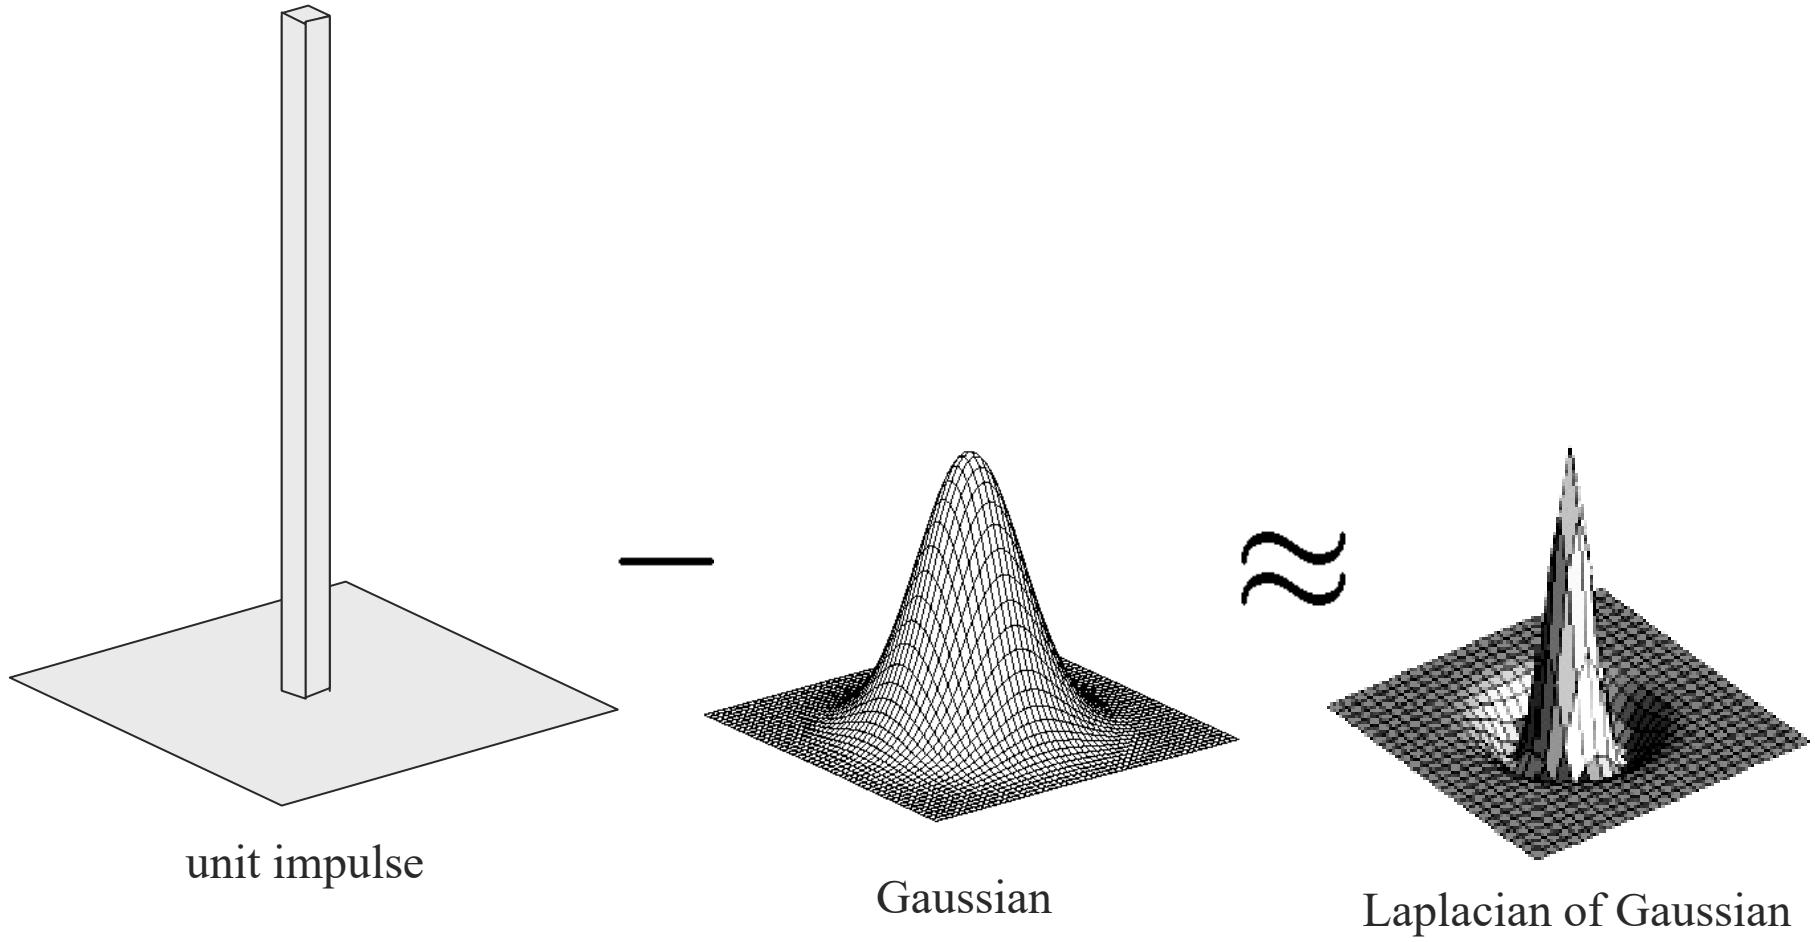
- Each layer of a laplacian pyramid consists in the elements of a smoothed and resampled image.
- **Band pass filters** - each level represents spatial frequencies (largely) unrepresented at other levels
- The fourier transform of each layer is an annulus
- Laplacian pyramid is orientation independent



SubbandDoGLaplacian/subbanddecomposition.m



Laplacian filter





LoG vs. DoG

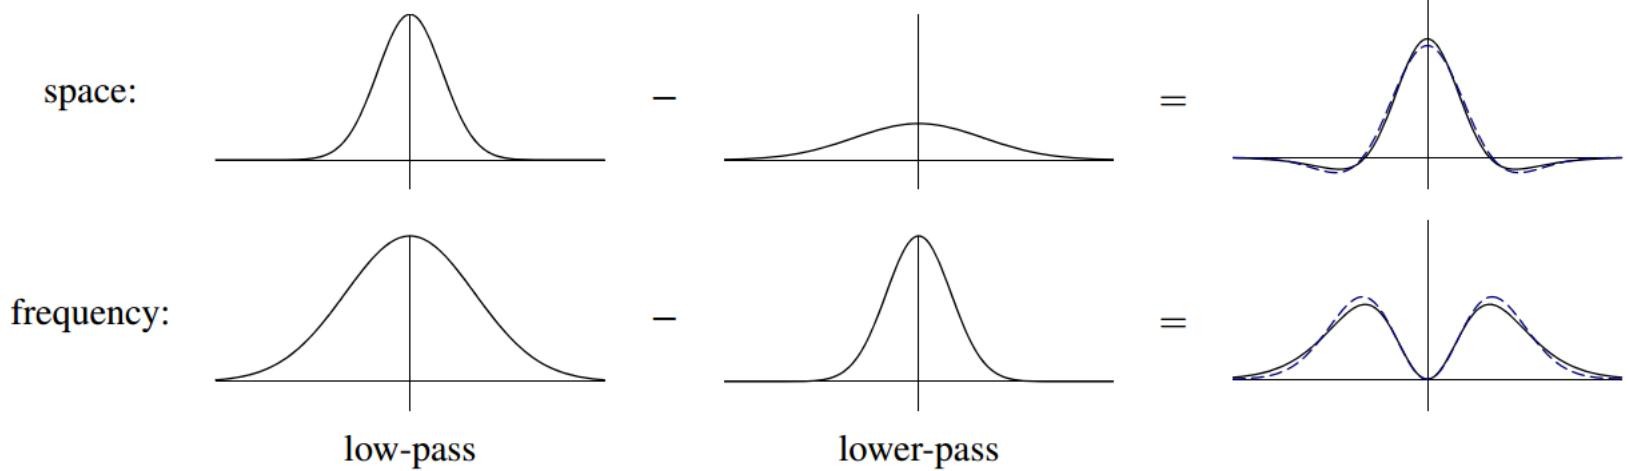


Figure 3.35 The difference of two low-pass filters results in a band-pass filter. The dashed blue lines show the close fit to a half-octave Laplacian of Gaussian.

$$\text{DoG}\{I; \sigma_1, \sigma_2\} = G_{\sigma_1} * I - G_{\sigma_2} * I = (G_{\sigma_1} - G_{\sigma_2}) * I.$$

$$\text{LoG}\{I; \sigma\} = \nabla^2(G_{\sigma} * I) = (\nabla^2 G_{\sigma}) * I,$$

$$\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$



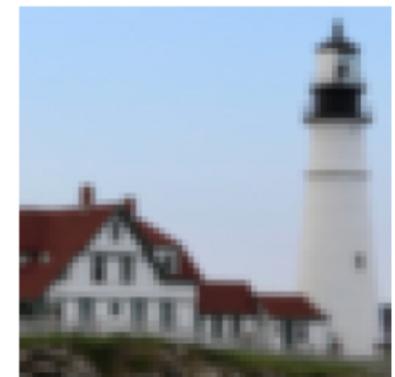
The Laplacian Pyramid



Laplacian pyramid

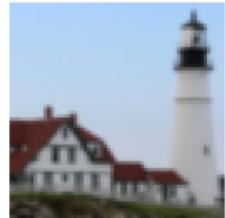


Gaussian
residual

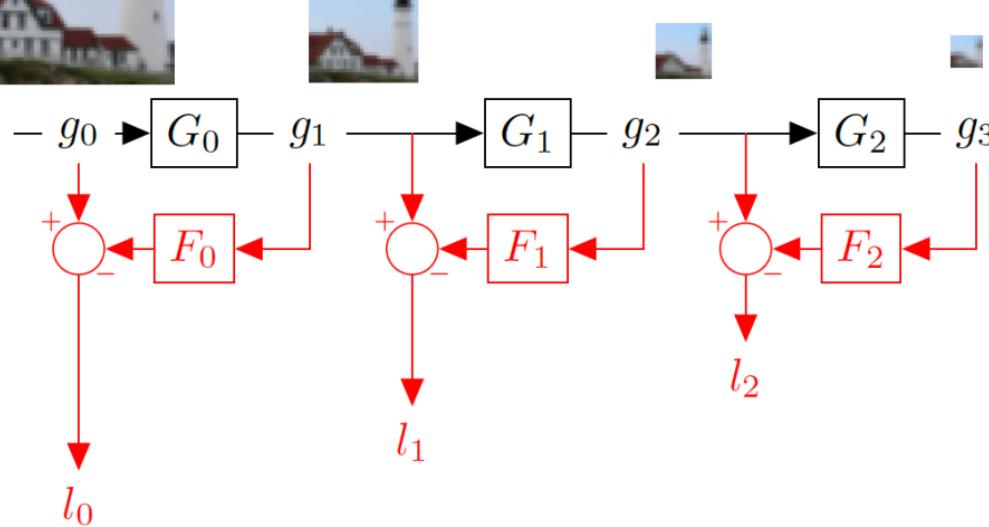


Can we invert the
Laplacian Pyramid?

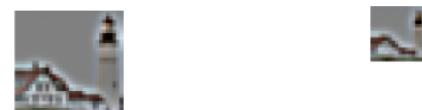
The Laplacian Pyramid



Gaussian pyramid



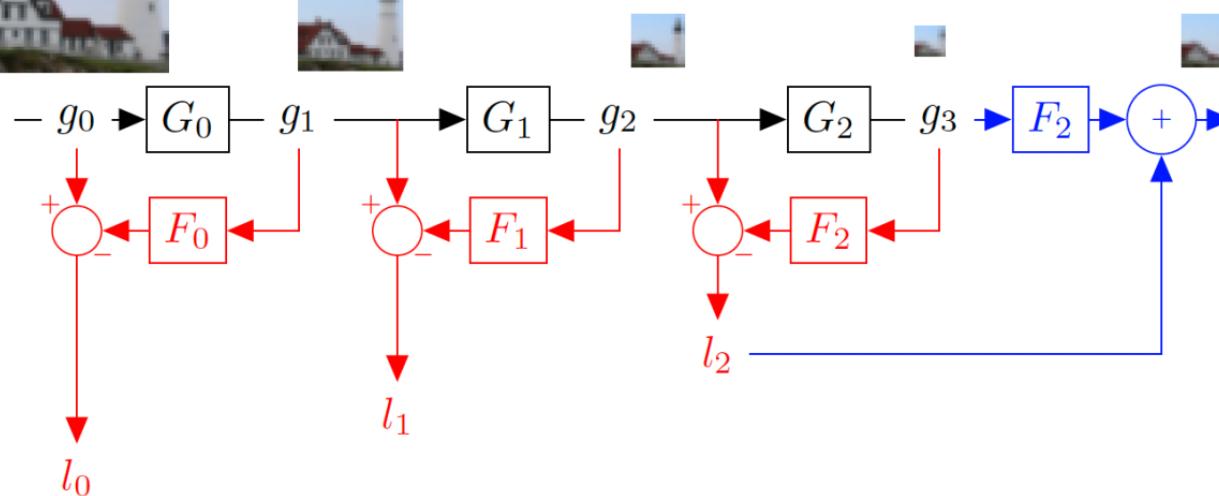
Laplacian pyramid



The Laplacian Pyramid

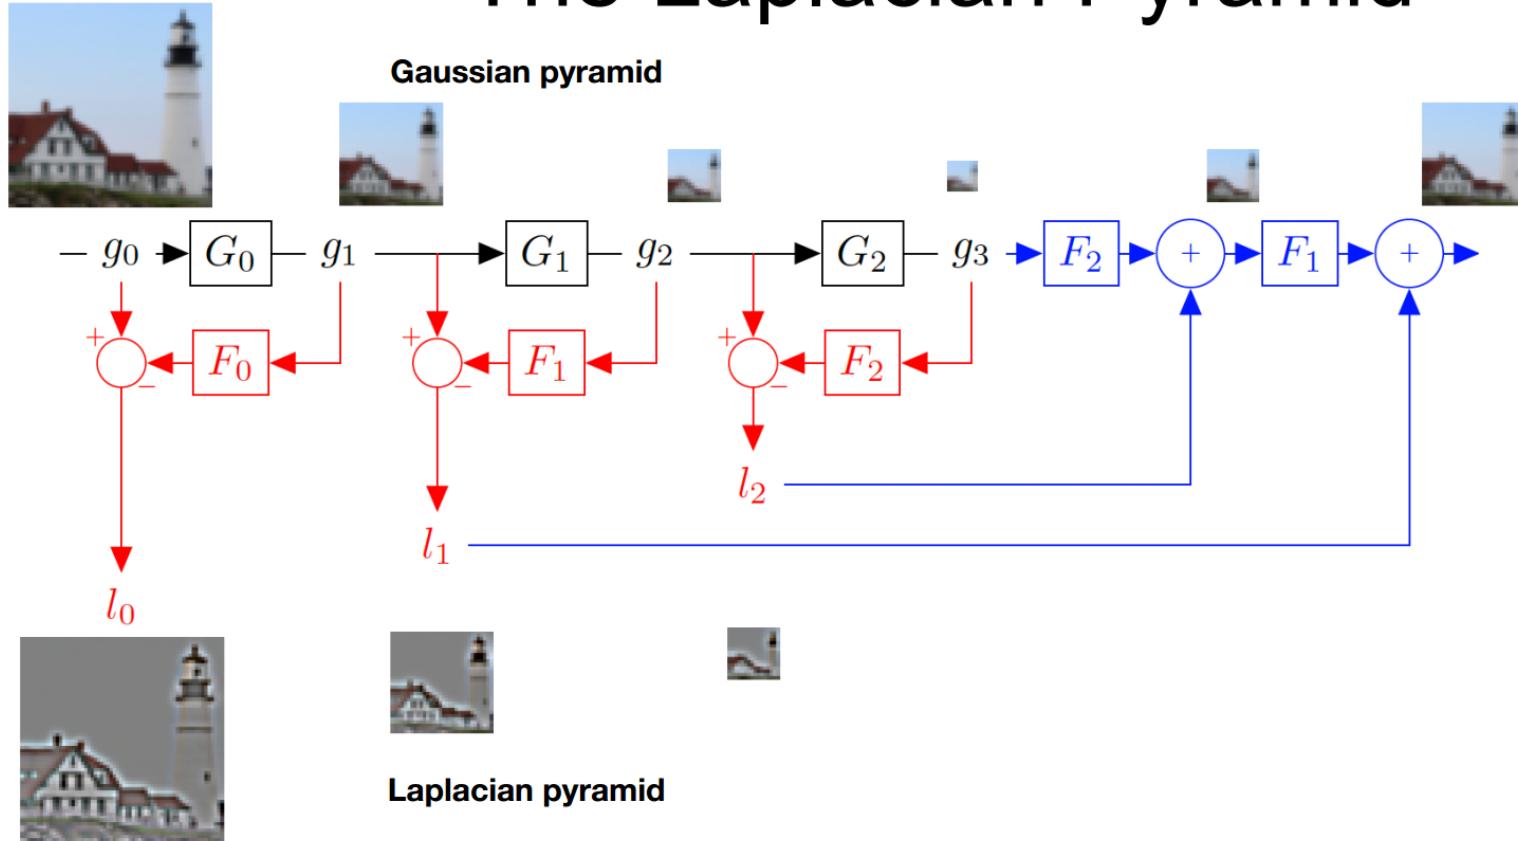


Gaussian pyramid

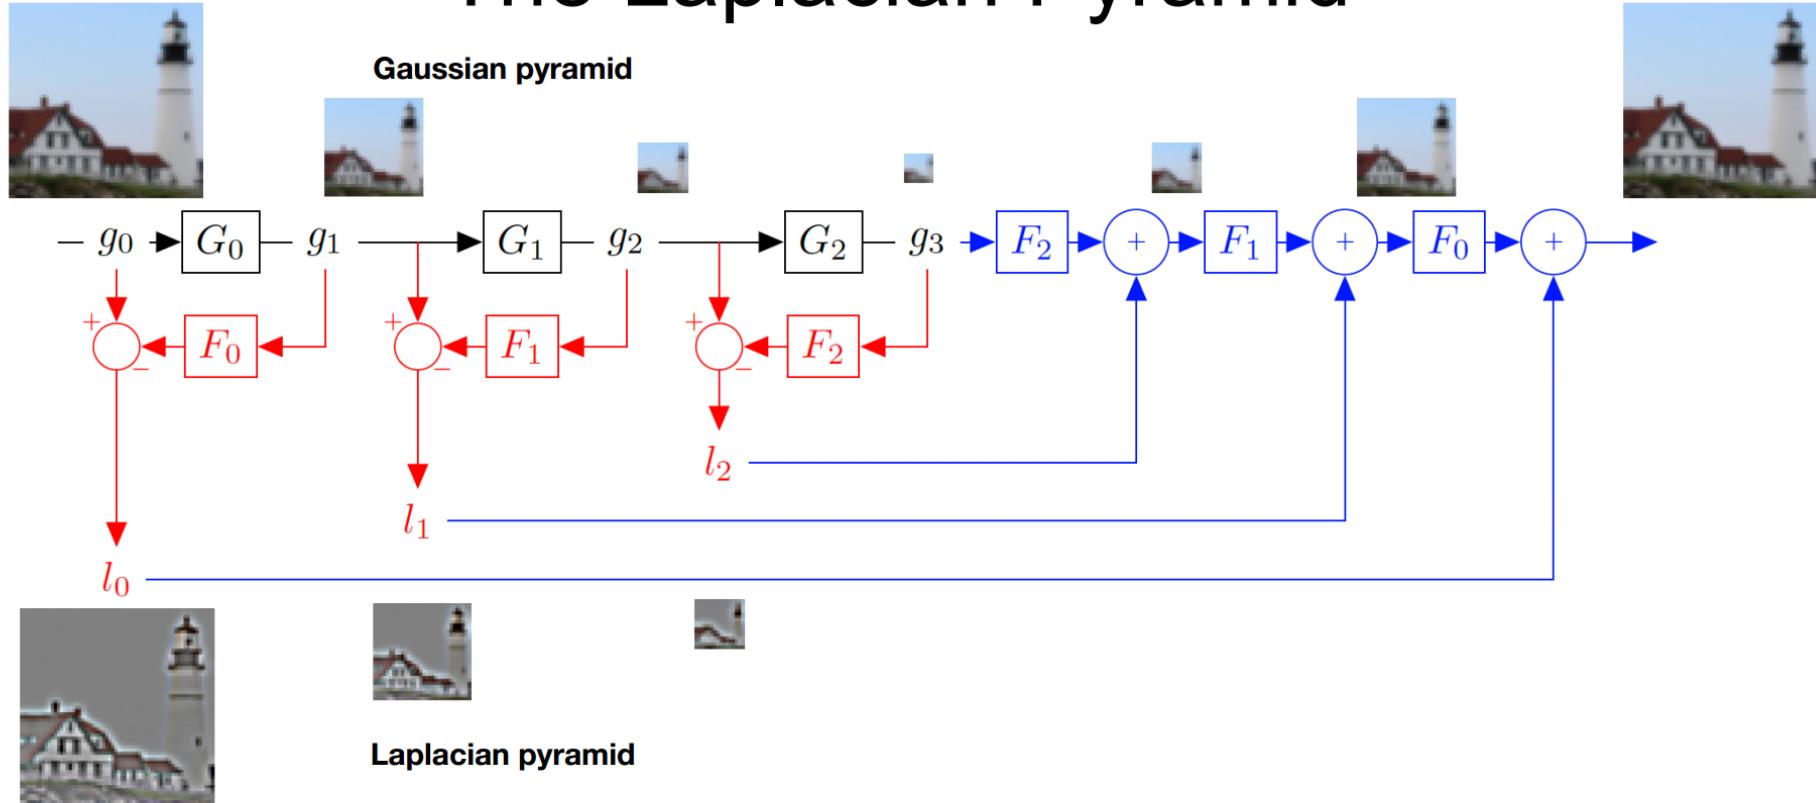


Laplacian pyramid

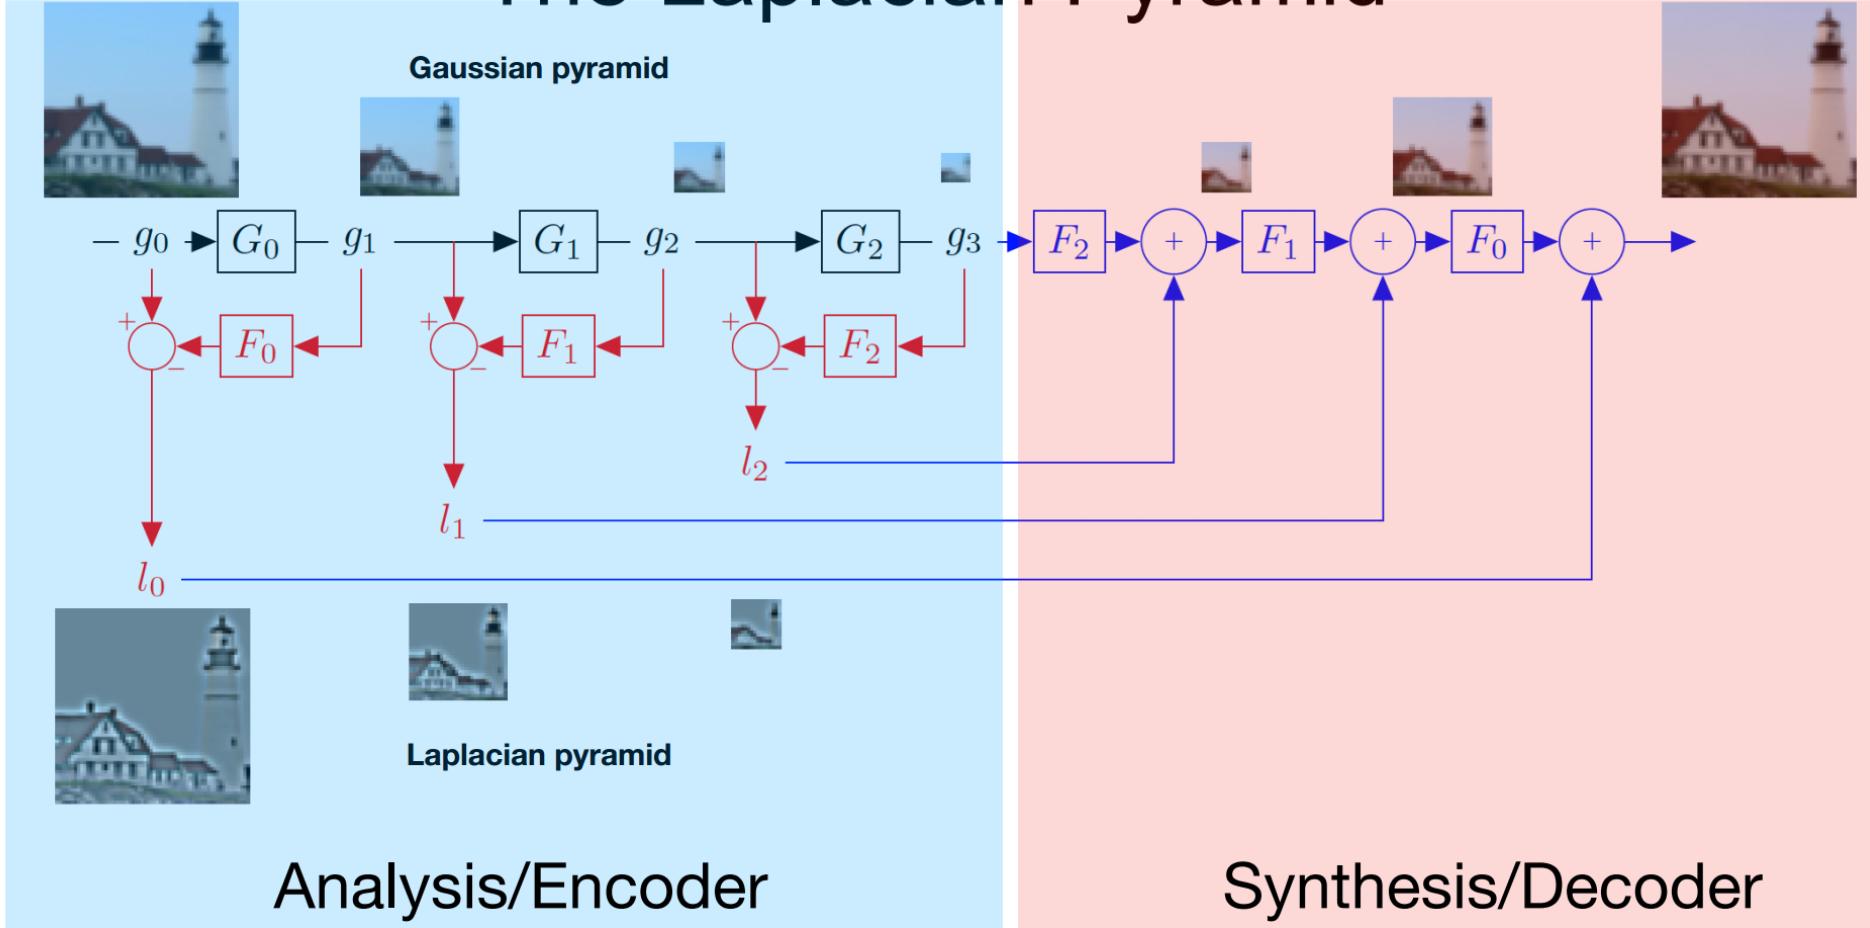
The Laplacian Pyramid



The Laplacian Pyramid

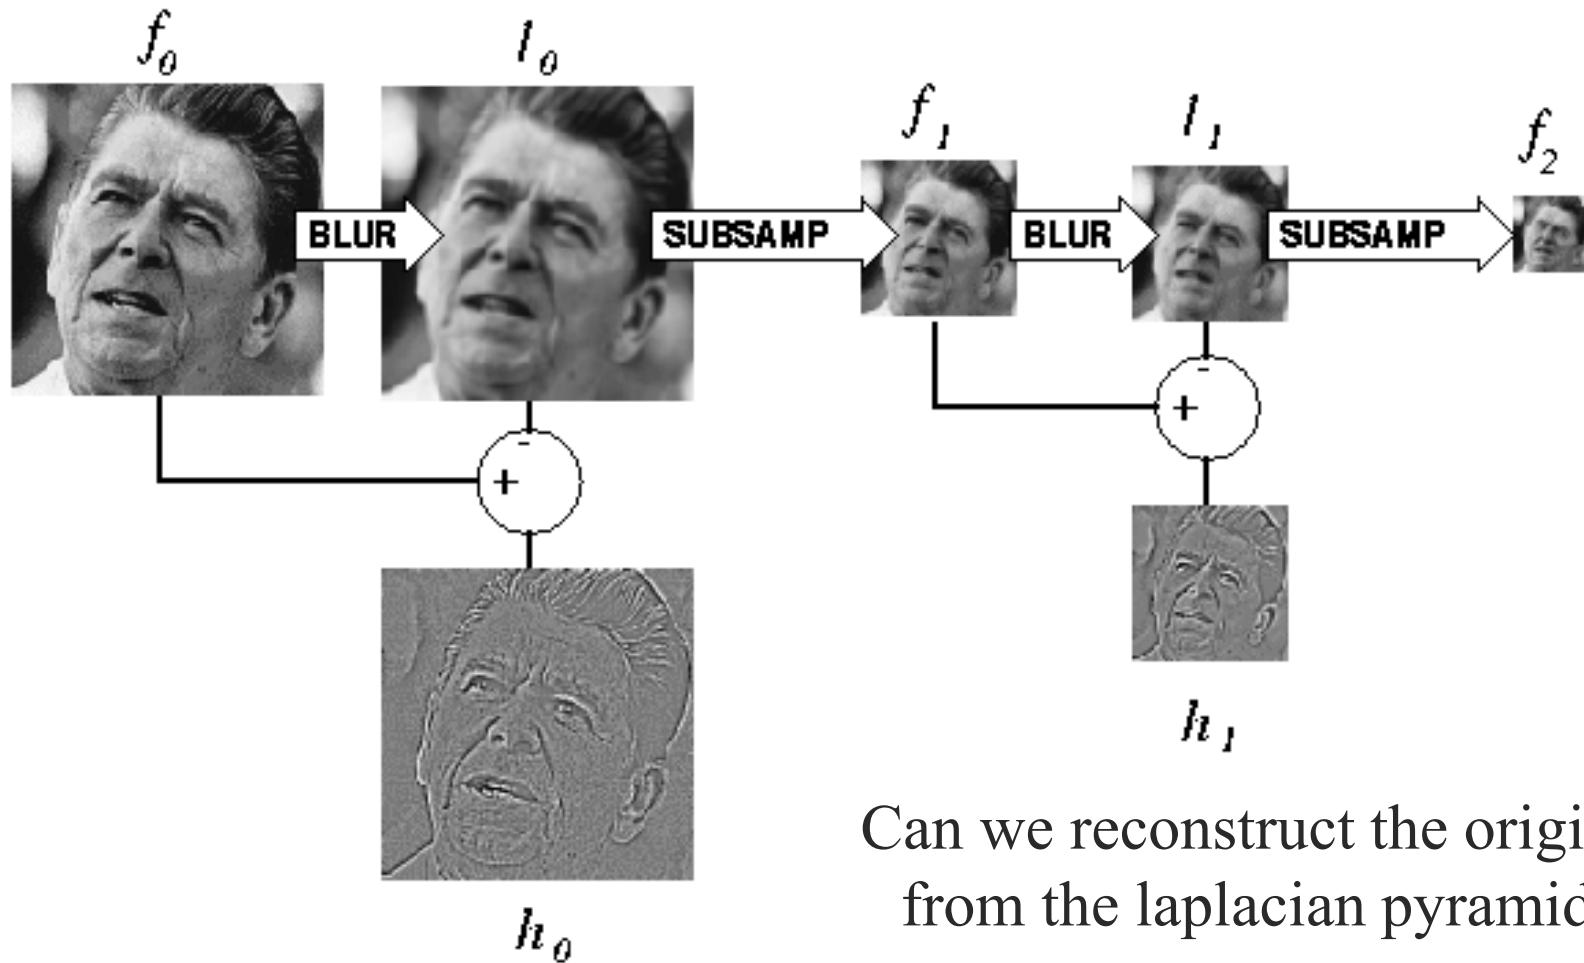


The Laplacian Pyramid





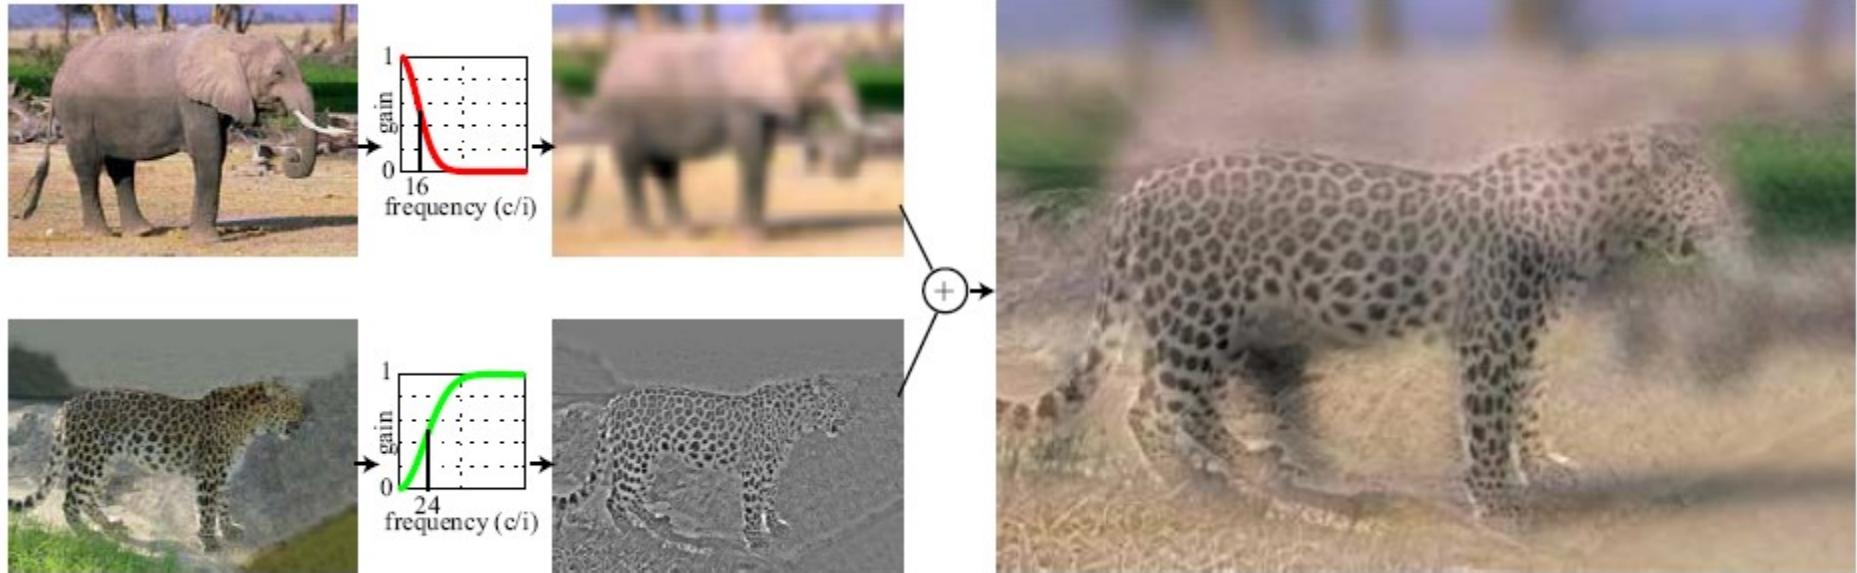
Computing Gaussian/Laplacian Pyramid



Can we reconstruct the original
from the laplacian pyramid?



Hybrid Images

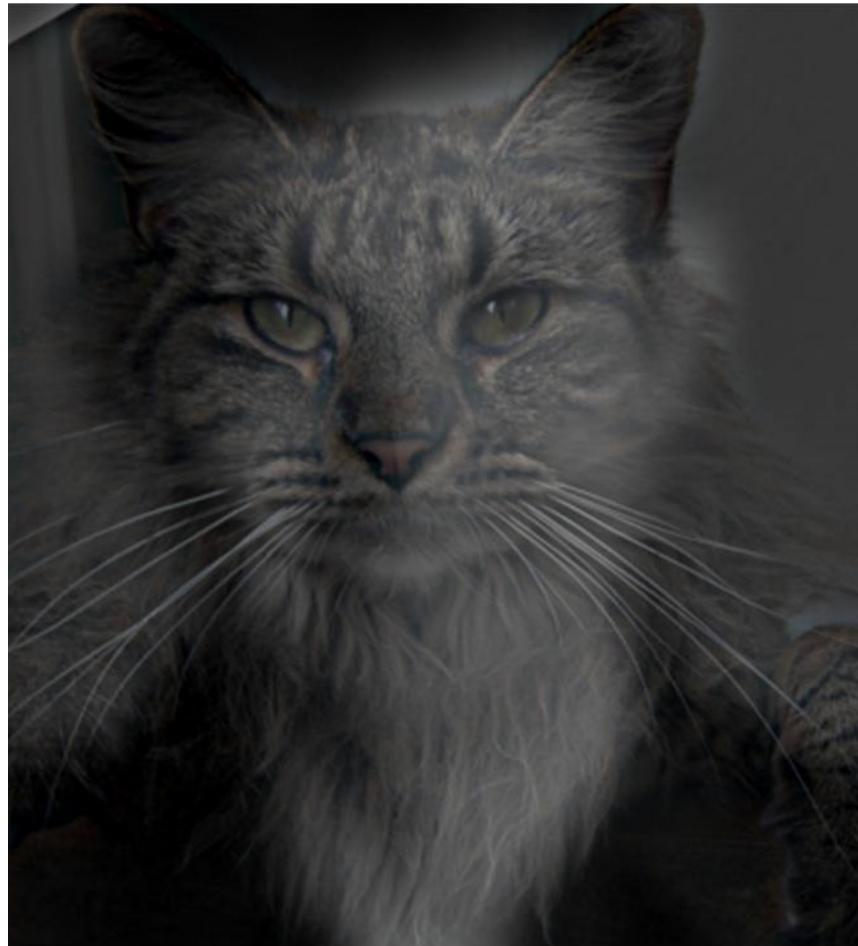


- A. Oliva, A. Torralba, P.G. Schyns,
“Hybrid Images,” SIGGRAPH 2006





Hybrid Image





Creating the Gaussian/Laplacian Pyramid



Image = G_1

Smooth, then downsample



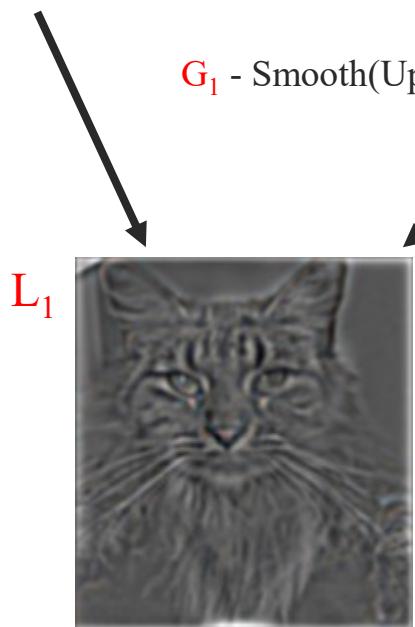
Downsample
(Smooth(G_1))



Downsample
(Smooth(G_2))



$\dots G_N = L_N$



$G_1 - \text{Smooth}(\text{Upsample}(G_2))$



$G_2 - \text{Smooth}(\text{Upsample}(G_3))$



$G_3 - \text{Smooth}(\text{Upsample}(G_4))$

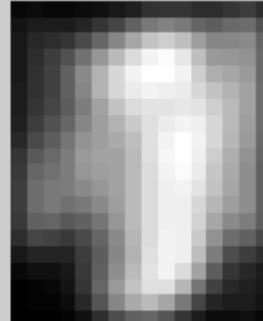
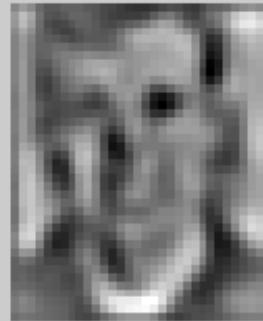
- Use same filter for smoothing in each step (e.g., Gaussian with $\sigma = 2$)
 - Downsample/upsample with “nearest” interpolation



Hybrid Image in Laplacian Pyramid



High frequency → Low frequency





Reconstructing image from Laplacian pyramid



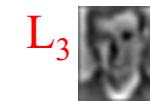
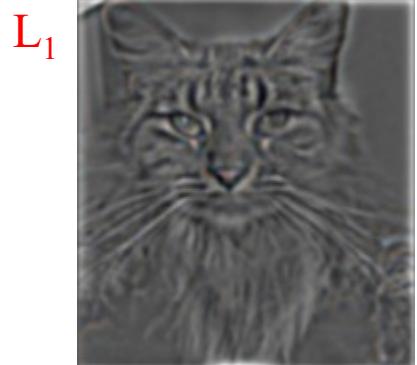
$$\text{Image} = L_1 + \text{Smooth}(\text{Upsample}(G_2))$$



$$G_2 = L_2 + \text{Smooth}(\text{Upsample}(G_3))$$



$$G_3 = L_3 + \text{Smooth}(\text{Upsample}(L_4))$$



- Use same filter for smoothing as in deconstruction
 - Upsample with “nearest” interpolation
 - Reconstruction will be nearly lossless



Laplacian pyramid applications



- Texture synthesis
- Image compression
- Noise removal
- Computing image features (e.g., SIFT)
- Image Blending...



Image Blending



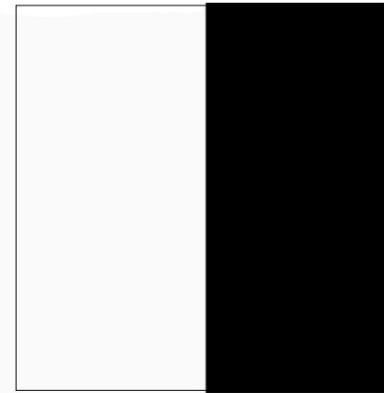


Image Blending





Image Blending

 I^A  I^B  m  I

$$I = m * I^A + (1 - m) * I^B$$



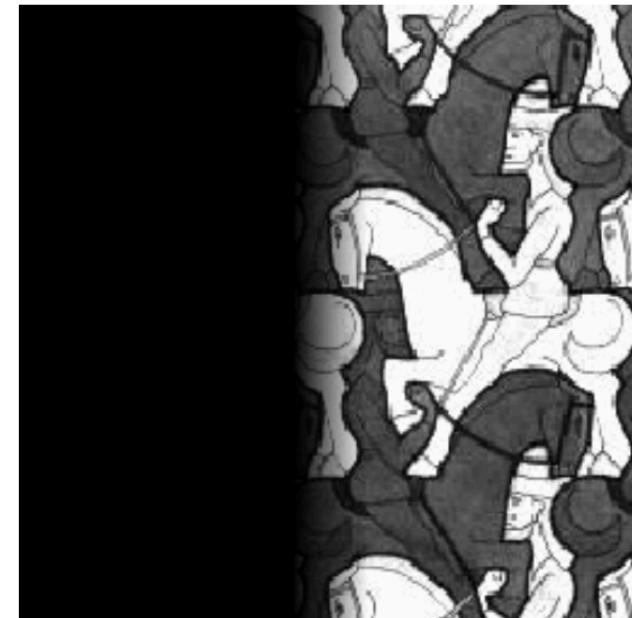
Image Blending



Feathering

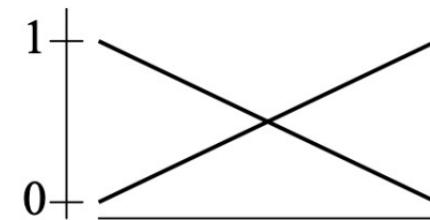
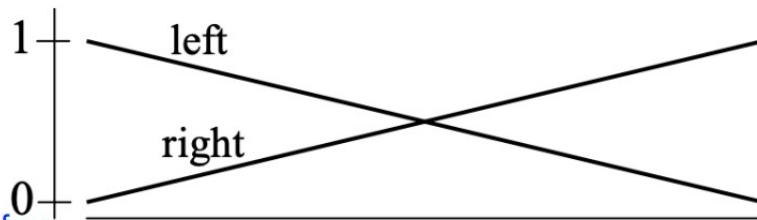
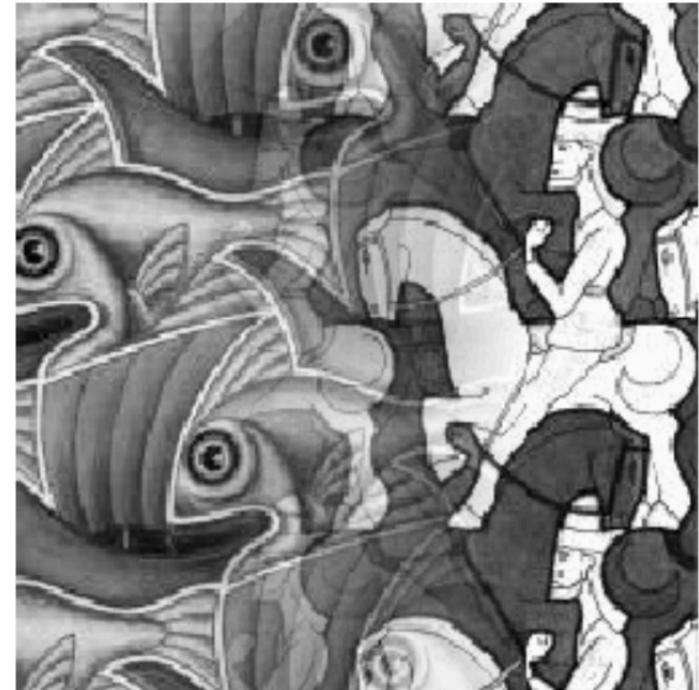


+





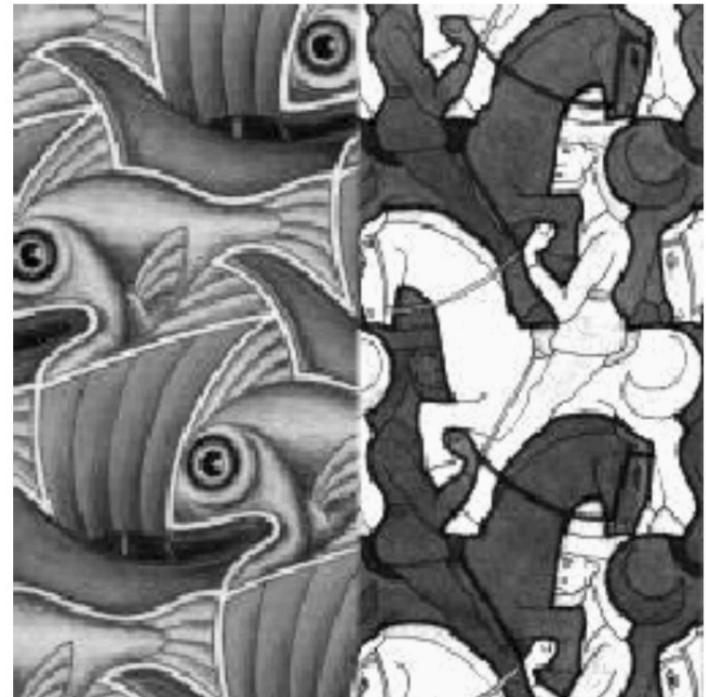
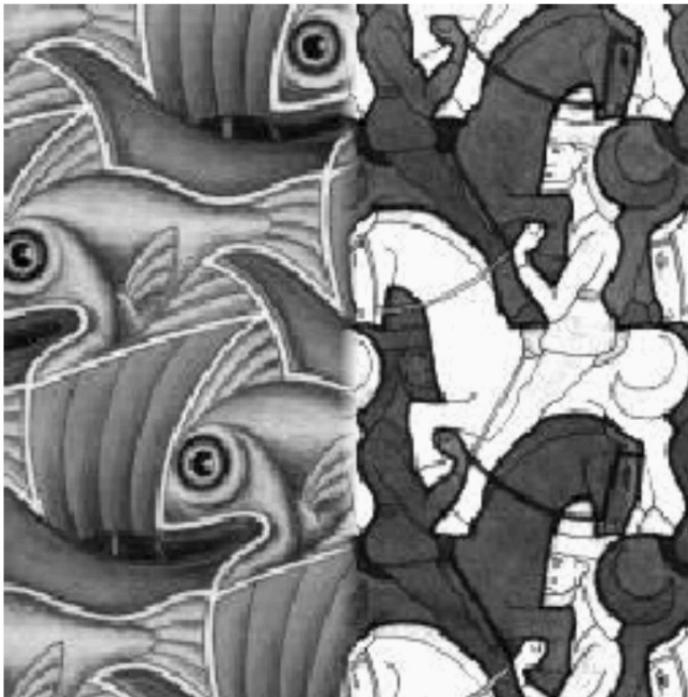
Affect of window size



Slide by A. Efros



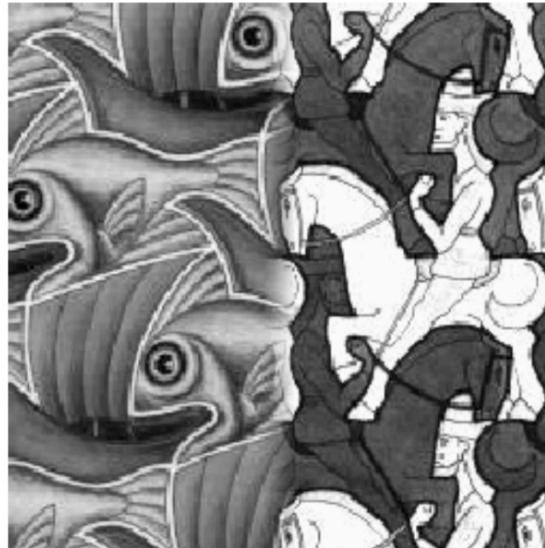
Affect of window size



Slide by A. Efros



Good Window Size



“Optimal” Window: smooth but not ghosted



What is the Optimal Window



To avoid seams

- $\text{window} \geq \text{size of largest prominent feature}$

To avoid ghosting

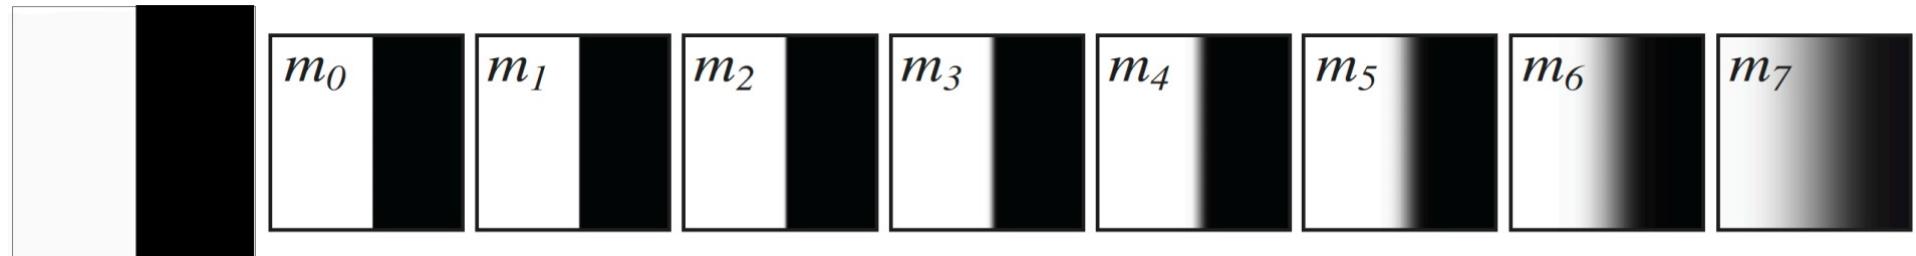
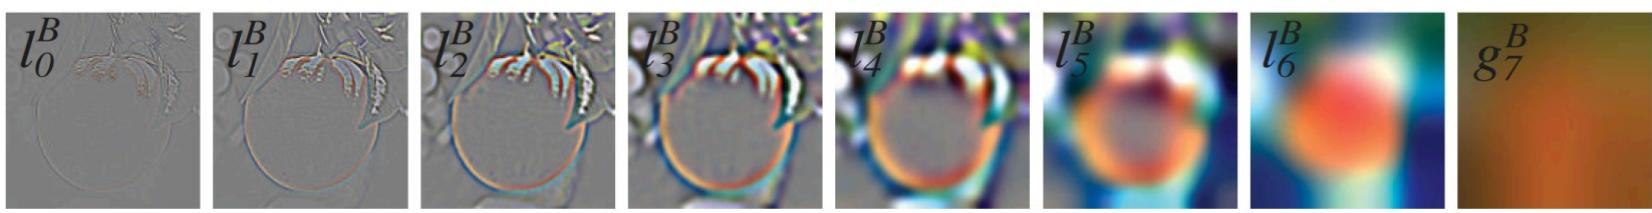
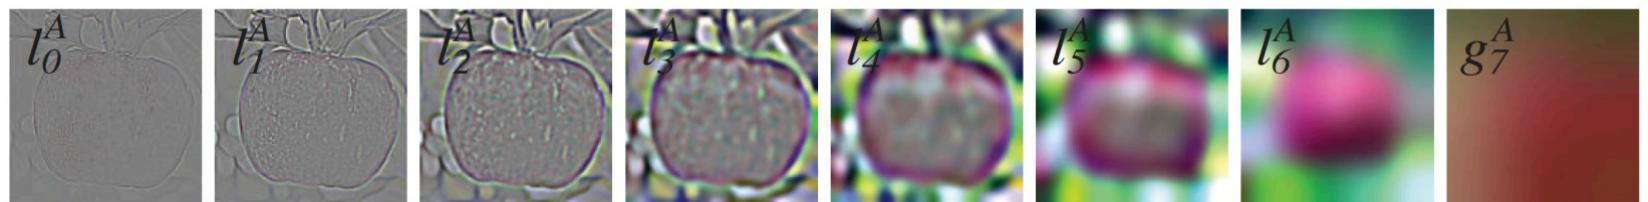
- $\text{window} \leq 2 * \text{size of smallest prominent feature}$

Natural to cast this in the *Fourier domain*

- largest frequency $\leq 2 * \text{size of smallest frequency}$
- image frequency content should occupy one “octave” (power of two)



Image Blending with the Laplacian Pyramid



$$l_k = l_k^A * m_k + l_k^B * (1 - m_k)$$



Image Blending with the Laplacian Pyramid



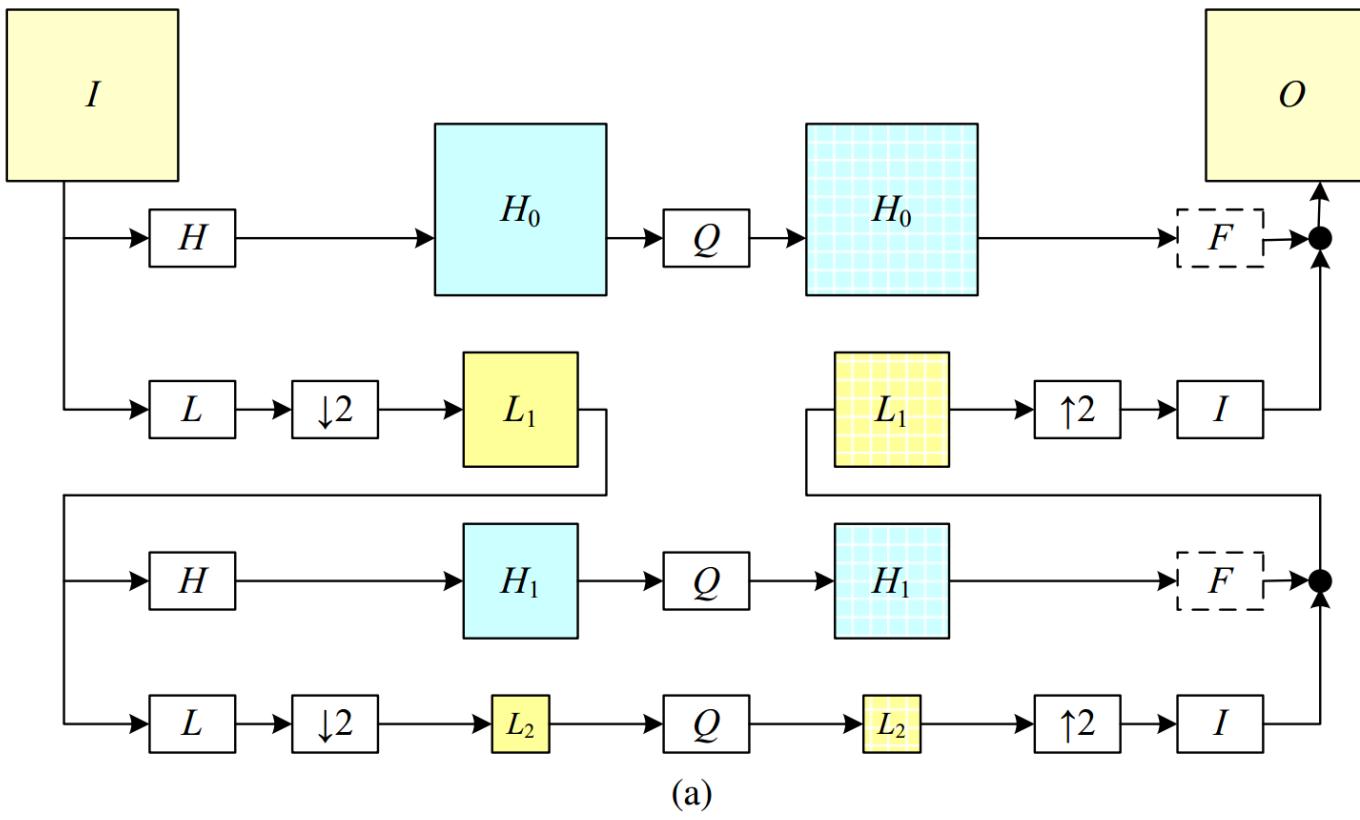


Image Blending with the Laplacian Pyramid

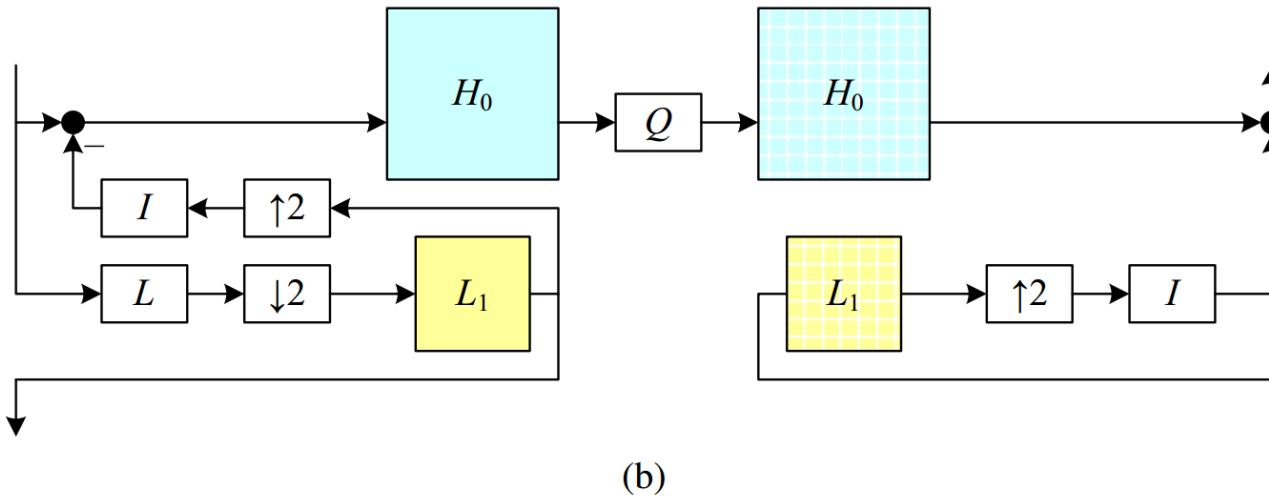


- Build Laplacian pyramid for both images: LA, LB
- Build Gaussian pyramid for mask: G
- Build a combined Laplacian pyramid:
- Collapse L to obtain the blended image





(a)



(b)

Image pyramids

- Gaussian
- Laplacian



Progressively blurred and subsampled versions of the image. Adds scale invariance to fixed-size algorithms.



Shows the information added in Gaussian pyramid at each spatial scale. Useful for noise reduction & coding.



Today's class



- Template matching
- Gaussian pyramids
 - Application for recognition
 - Pyramids representation in deep learning
- Laplacian pyramids
 - Hybrid images
 - Application for image blending
- **Steerable pyramids**
 - Steerable filters
 - Orientation analysis
- Human vision system
 - Visual perception
- Filter Banks and Texture Analysis

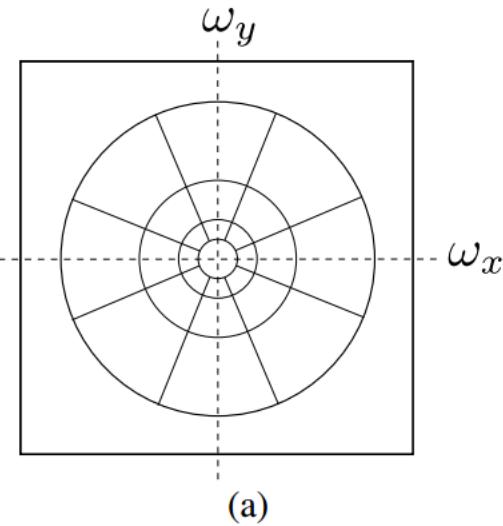


Orientations

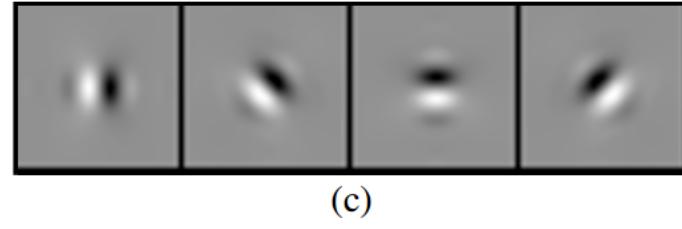




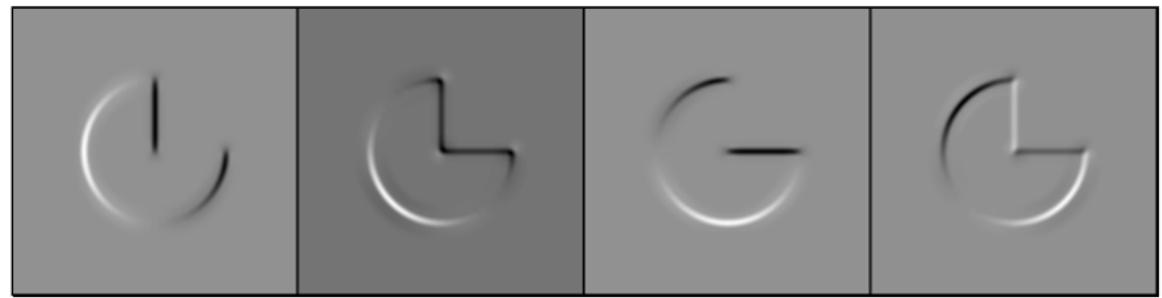
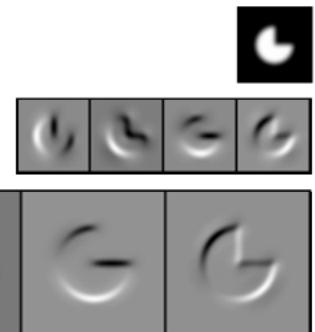
Steerable Pyramid



(b)



(c)



(d)

Figure 3.40 Steerable shiftable multiscale transforms (Simoncelli, Freeman, Adelson *et al.* 1992) © 1992 IEEE: (a) radial multi-scale frequency domain decomposition; (b) original image; (c) a set of four steerable filters; (d) the radial multi-scale wavelet decomposition.

Steerable Pyramid



$$x \rightarrow B_0 \rightarrow b_{0,0}$$



$$x \rightarrow B_1 \rightarrow b_{0,1}$$



$$x \rightarrow B_n \rightarrow b_{0,n}$$



$$x \rightarrow L \rightarrow D \rightarrow B_0 \rightarrow b_{1,0}$$



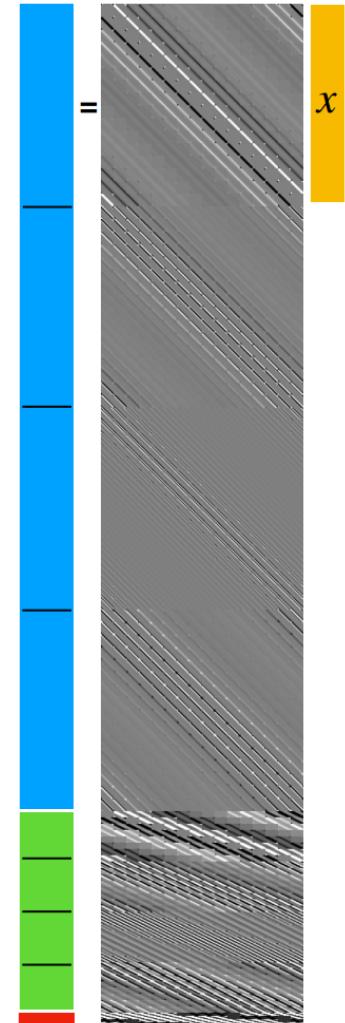
$$x \rightarrow B_1 \rightarrow b_{1,1}$$



$$x \rightarrow B_n \rightarrow b_{1,n}$$

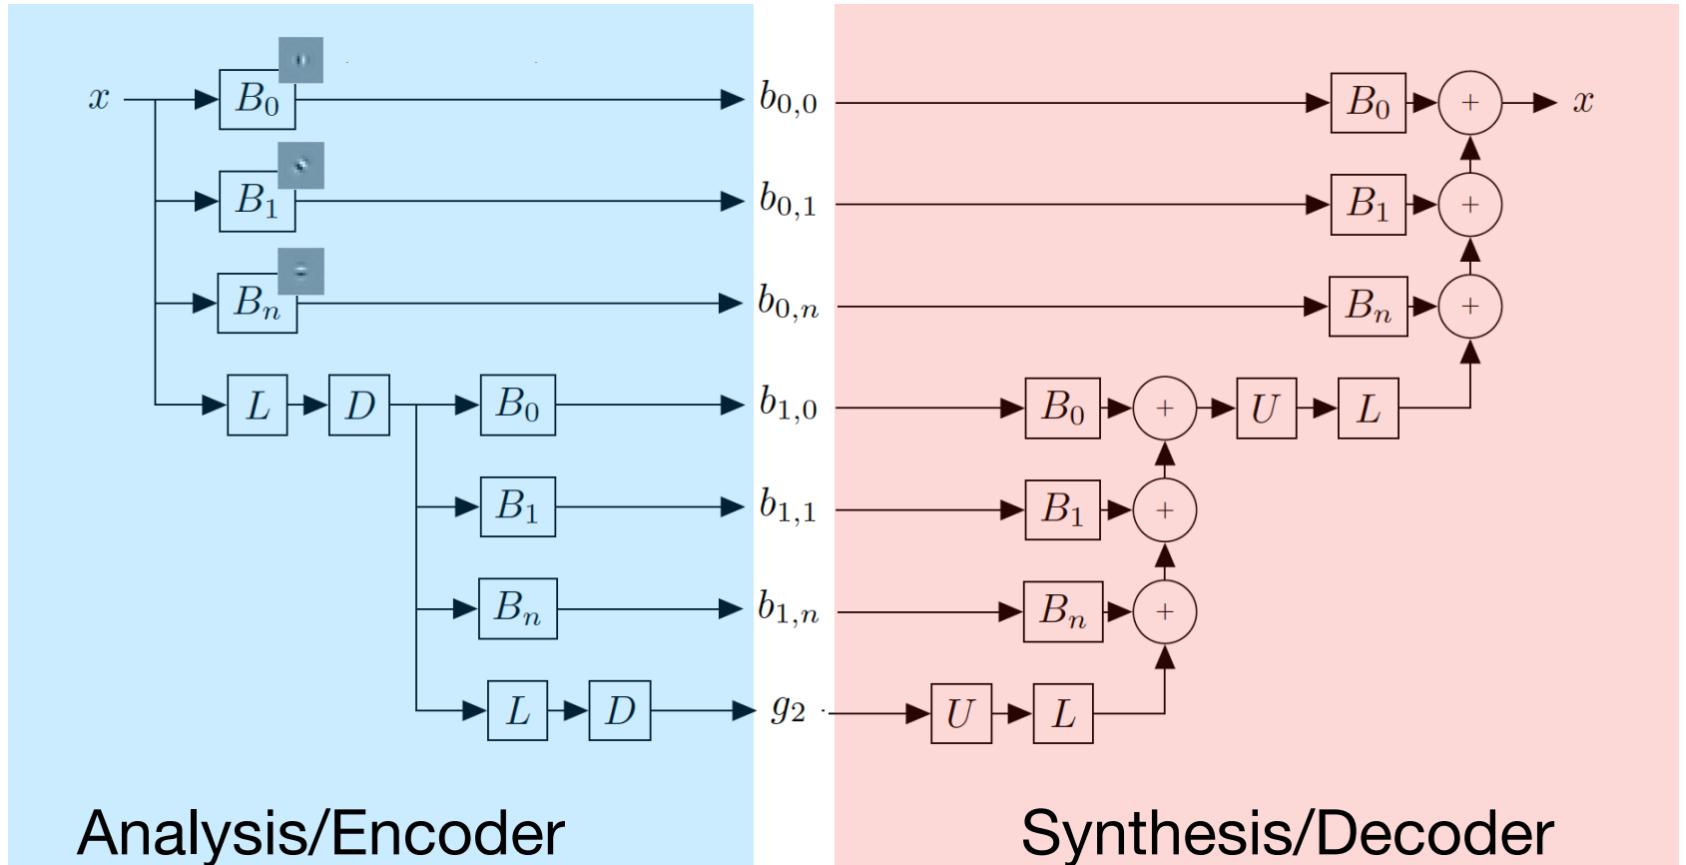


$$x \rightarrow L \rightarrow D \rightarrow g_2$$





Steerable Pyramid



Analysis/Encoder

Synthesis/Decoder

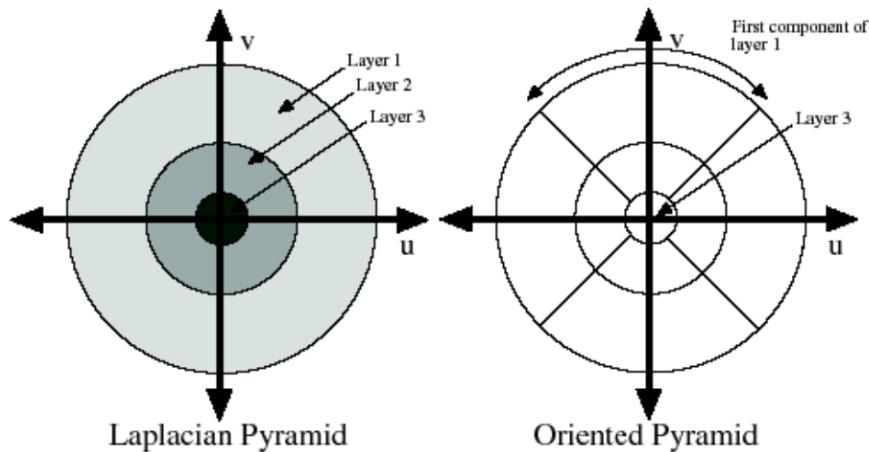
The steerable pyramid



- Preserves all image information (we can go back to the image)
- Provides more independent channels of information than pixel values (we can mess with each band independently)

Oriented Pyramids in spectral space

- Apply an oriented filter to determine orientations at each layer
 - Important for texture synthesis
 - this represents image information at a particular scale and orientation, similar to V1 cells in the human brain



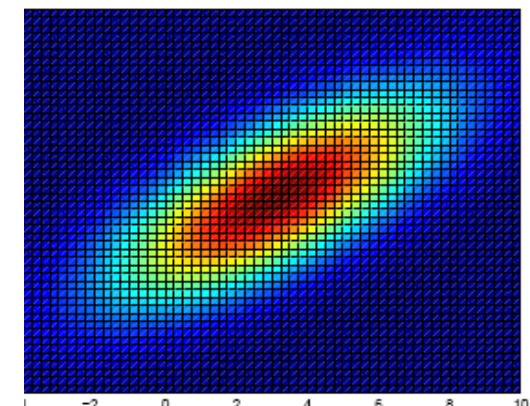
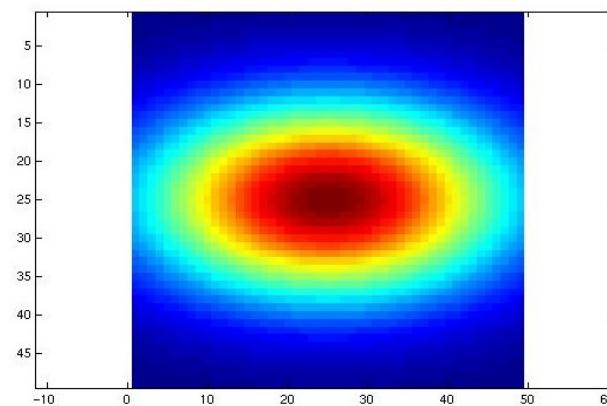
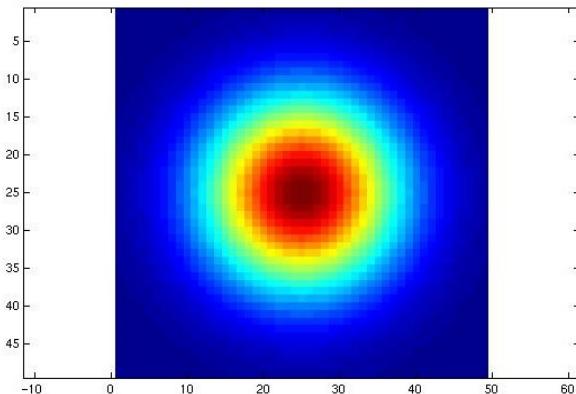
An oriented pyramid cuts each annulus into a set of wedges. If (u,v) Fourier space is represented in polar coordinates, each wedge corresponds to an interval of radius values and an interval of angle values.



Multivariate Gaussian



$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$



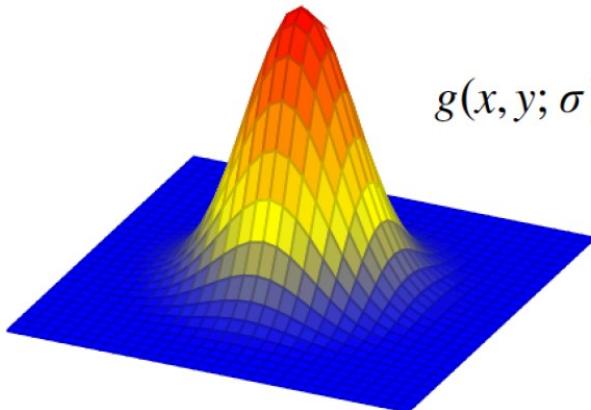
$$\Sigma = \begin{bmatrix} 9 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 16 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 10 & 5 \\ 5 & 5 \end{bmatrix}$$



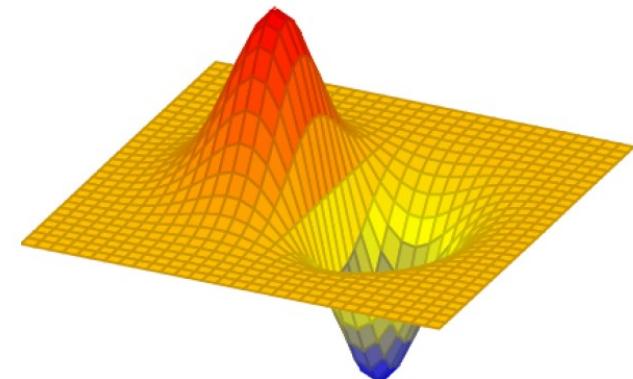
Gaussian Derivative



$$g(x, y; \sigma) = \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2}$$

The continuous derivative is:

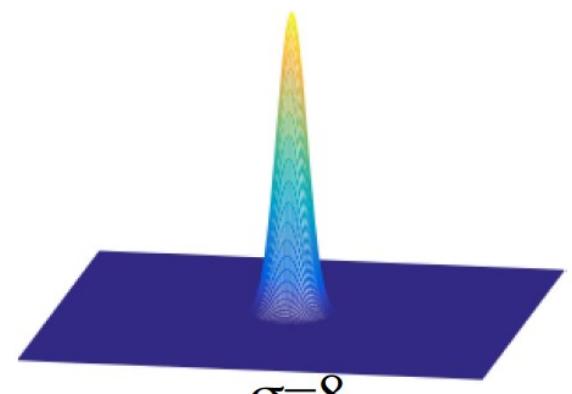
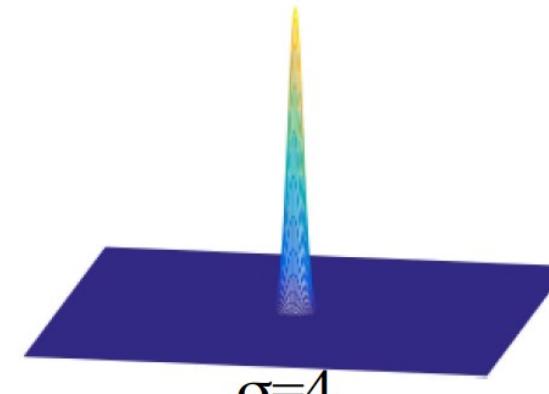
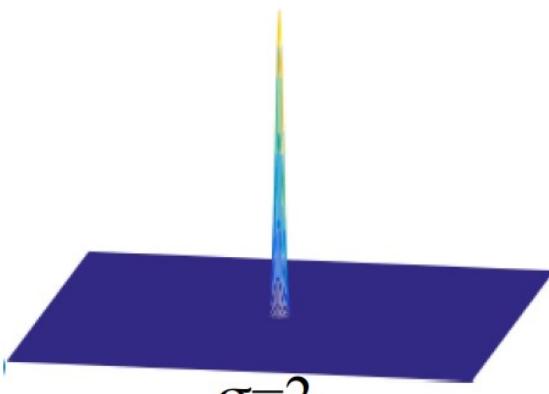
$$\begin{aligned} g_x(x, y; \sigma) &= \frac{\partial g(x, y; \sigma)}{\partial x} = \\ &= \frac{-x}{2\pi\sigma^4} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \frac{-x}{\sigma^2} g(x, y; \sigma) \end{aligned}$$



104

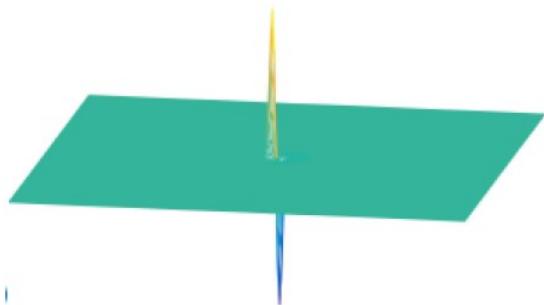
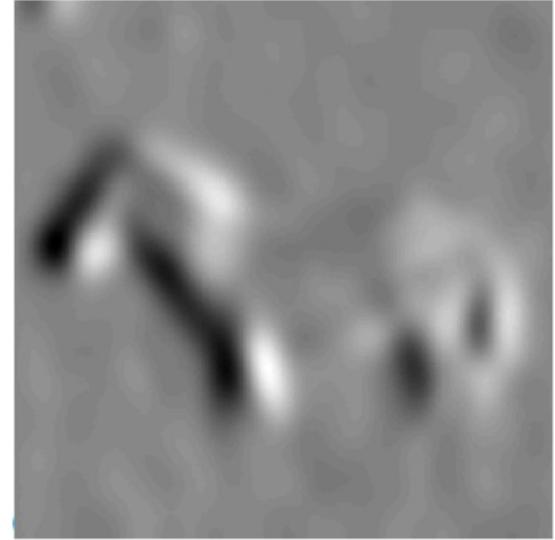
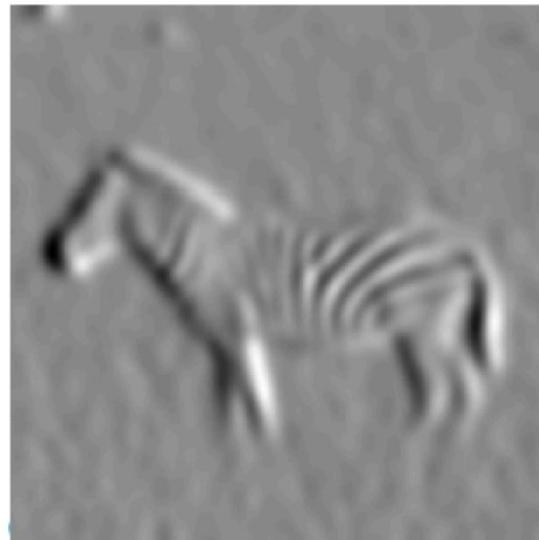
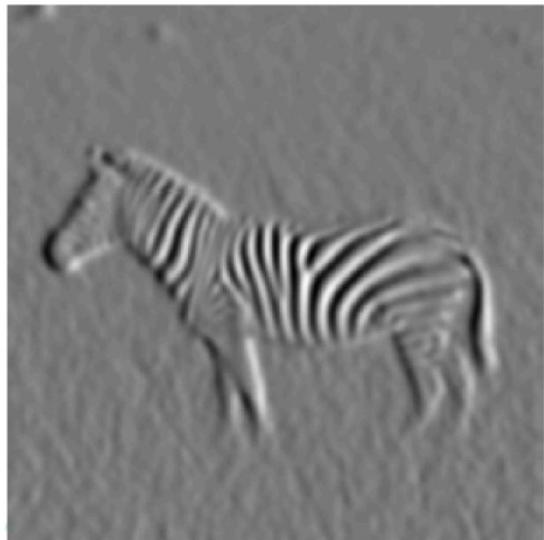


Gaussian Scale

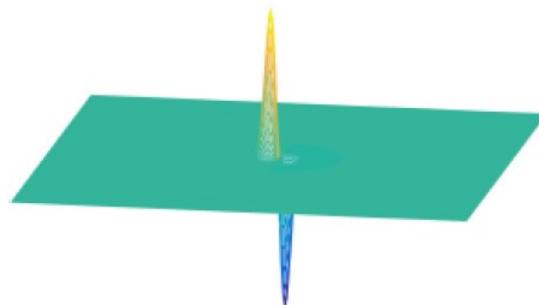




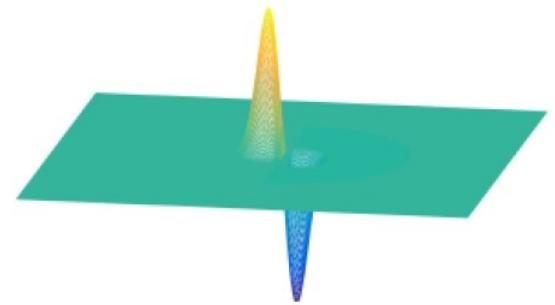
Derivatives of Gaussian: Scale



$\sigma=2$



$\sigma=4$

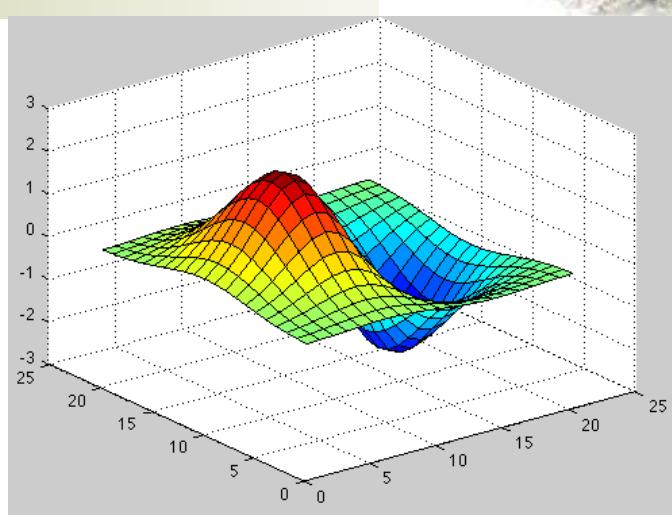


$\sigma=8$

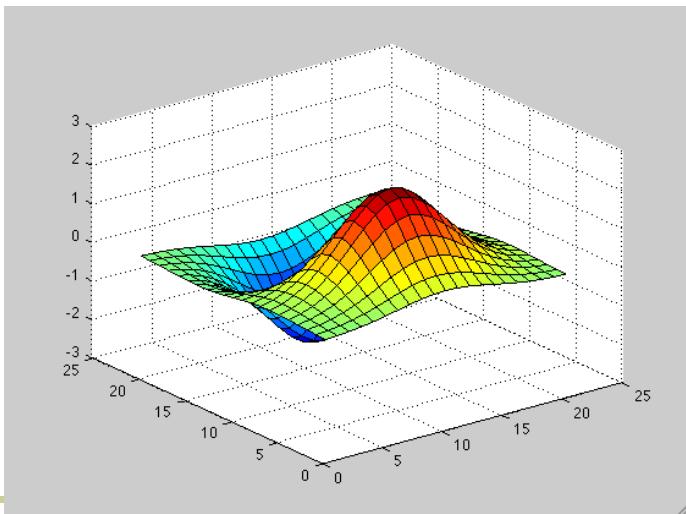


Orientation

$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

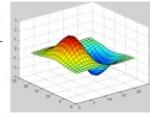




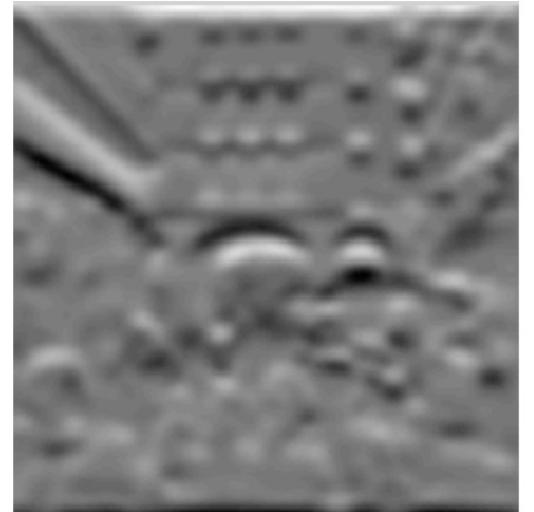
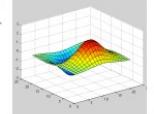
Orientation



$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



What about other orientations not axis aligned?



Why are they steerable filter

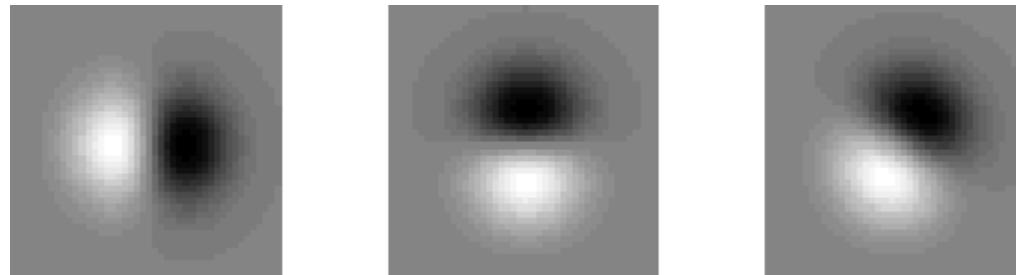


“Steerability”-- the ability to synthesize a filter of any orientation from a linear combination of filters at fixed orientations.

$$G_{\theta}^1 = \cos(\theta)G_0^1 + \sin(\theta)G_{90}^1$$

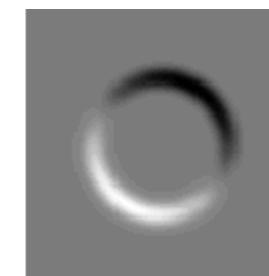
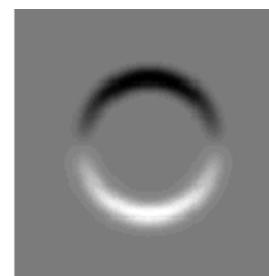
0° 90° Synthesized 30°

Filter Set:



Response:

Raw Image



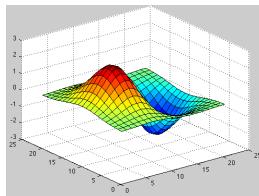
Taken from:
W. Freeman, T.
Adelson, “The Design
and Use of Steerable
Filters”, IEEE
Trans. Patt. Anal. and
Machine Intell.,
vol 13, #9, pp 891-900,
Sept 1991



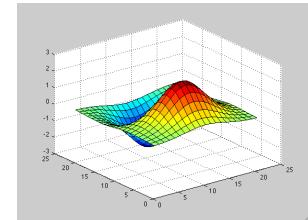
Steerable filters



$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



The smoothed directional gradient is a linear combination of two kernels

$$u^T \nabla g \otimes I = (\cos(\alpha)g_x(x,y) + \sin(\alpha)g_y(x,y)) \otimes I(x,y) =$$

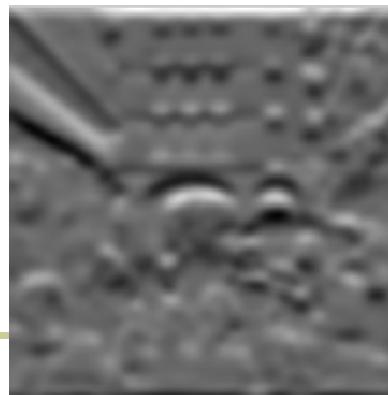
Any orientation can be computed as a linear combination of two filtered images

$$= \cos(\alpha)g_x(x,y) \otimes I(x,y) + \sin(\alpha)g_y(x,y) \otimes I(x,y) =$$

$$= \cos(\alpha)$$



$$+ \sin(\alpha)$$



$$=$$





Orientation analysis

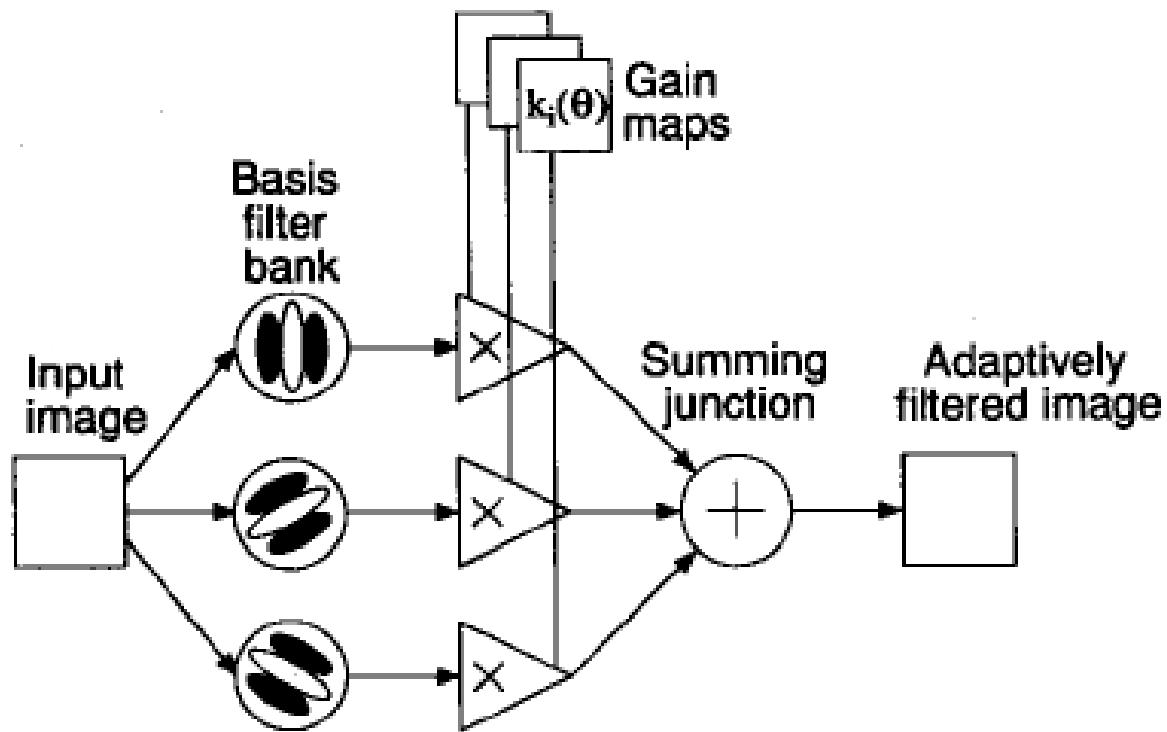
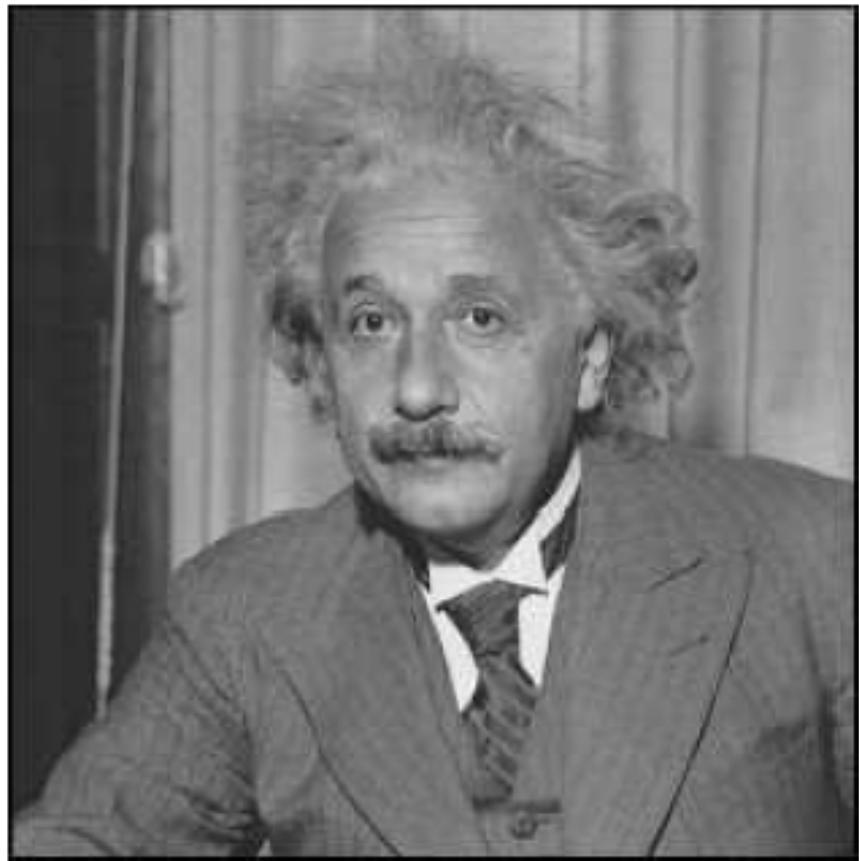


Fig. 3. Steerable filter system block diagram. A bank of dedicated filters process the image. Their outputs are multiplied by a set of gain maps that adaptively control the orientation of the synthesized filter.



Orientation analysis



(a)



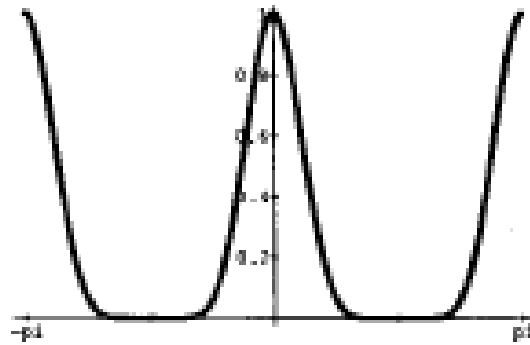
(b)



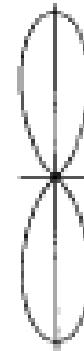
Orientation analysis



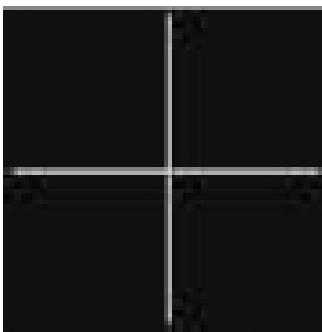
(a)



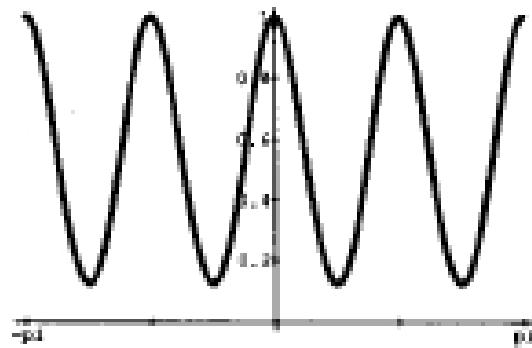
(c)



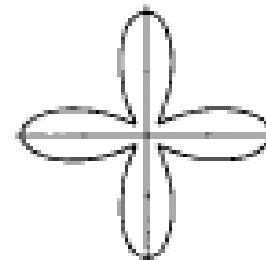
(e)



(b)



(d)



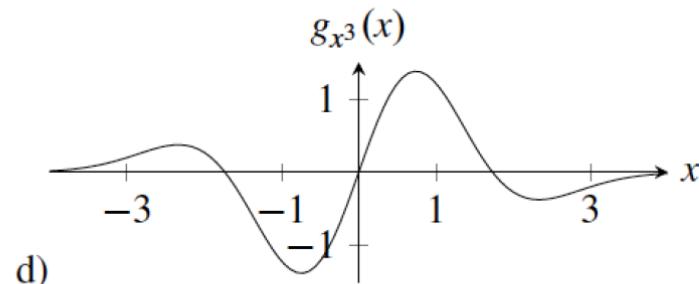
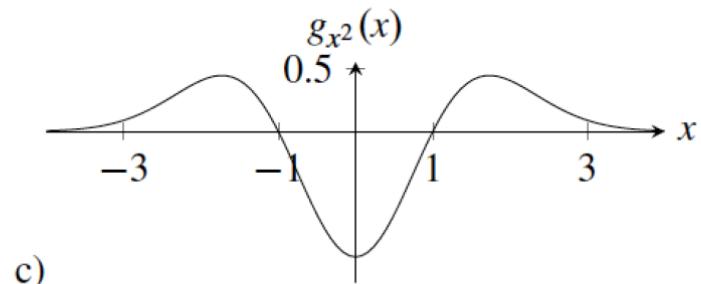
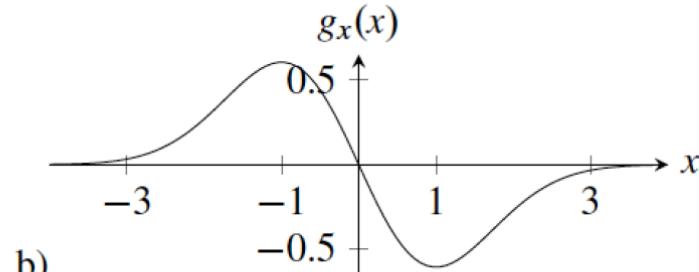
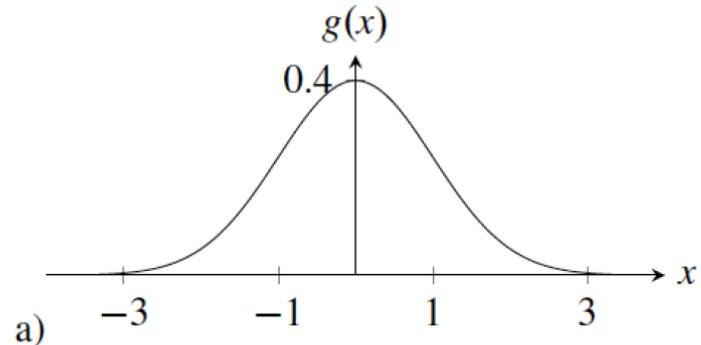
(f)

High resolution in orientation requires many oriented filters as basis (high order gaussian derivatives or fine-tuned Gabor wavelets).





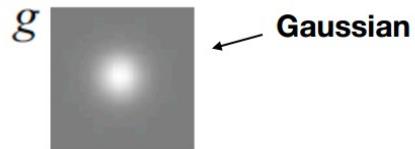
N-th order Gaussian derivatives



$$g_{x^n, y^m}(x, y; \sigma) = \frac{\partial^{n+m} g(x, y)}{\partial x^n \partial y^m} = \left(\frac{-1}{\sigma \sqrt{2}} \right)^{n+m} H_n \left(\frac{x}{\sigma \sqrt{2}} \right) H_m \left(\frac{y}{\sigma \sqrt{2}} \right) g(x, y; \sigma)$$



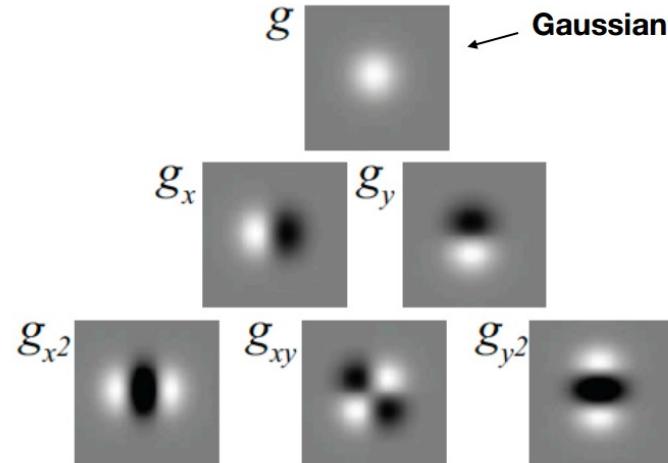
N-th order Gaussian derivatives



$$g_{x^n, y^m}(x, y; \sigma) = \frac{\partial^{n+m} g(x, y)}{\partial x^n \partial y^m} = \left(\frac{-1}{\sigma \sqrt{2}} \right)^{n+m} H_n \left(\frac{x}{\sigma \sqrt{2}} \right) H_m \left(\frac{y}{\sigma \sqrt{2}} \right) g(x, y; \sigma)$$



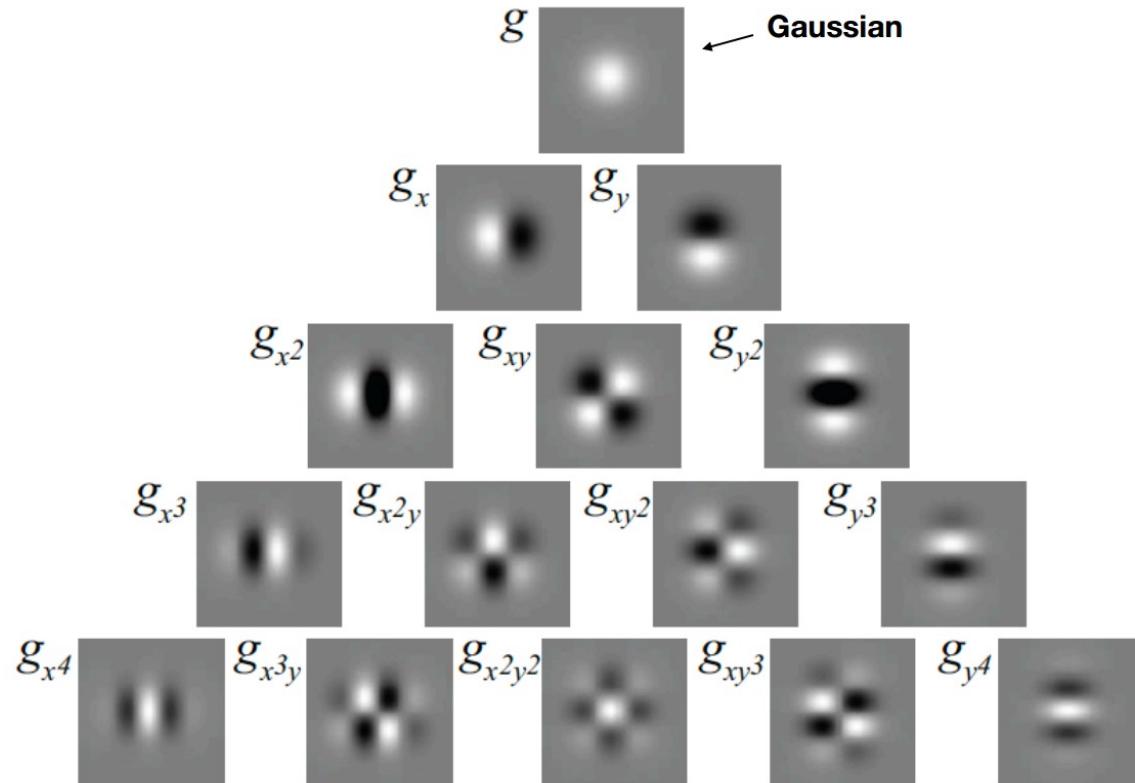
N-th order Gaussian derivatives



$$g_{x^n, y^m}(x, y; \sigma) = \frac{\partial^{n+m} g(x, y)}{\partial x^n \partial y^m} = \left(\frac{-1}{\sigma \sqrt{2}} \right)^{n+m} H_n \left(\frac{x}{\sigma \sqrt{2}} \right) H_m \left(\frac{y}{\sigma \sqrt{2}} \right) g(x, y; \sigma)$$



N-th order Gaussian derivatives



$$g_{x^n, y^m}(x, y; \sigma) = \frac{\partial^{n+m} g(x, y)}{\partial x^n \partial y^m} = \left(\frac{-1}{\sigma \sqrt{2}} \right)^{n+m} H_n \left(\frac{x}{\sigma \sqrt{2}} \right) H_m \left(\frac{y}{\sigma \sqrt{2}} \right) g(x, y; \sigma)$$



Steerable pyramid applications



- Texture synthesis
- Noise removal
- Motion analysis
- Motion synthesis, motion magnification



Summary: Image pyramids



Gaussian



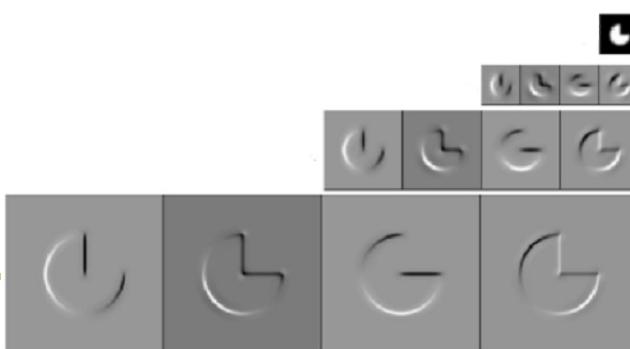
Progressively blurred and subsampled versions of the image. Adds scale invariance to fixed-size algorithms.

Laplacian



Shows the information added in Gaussian pyramid at each spatial scale. Useful for noise reduction & coding.

Steerable pyramid



Shows components at each scale and orientation separately. Non-aliased subbands. Good for texture and feature analysis.

Linear Image Transforms

1297x256

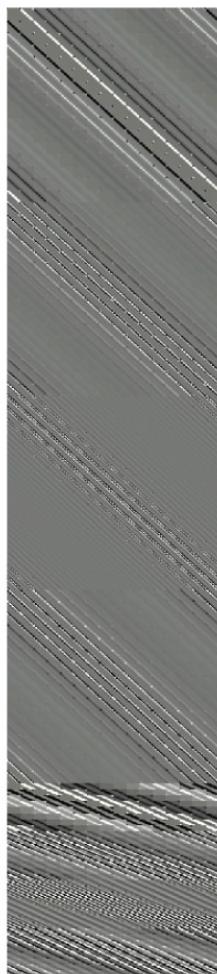
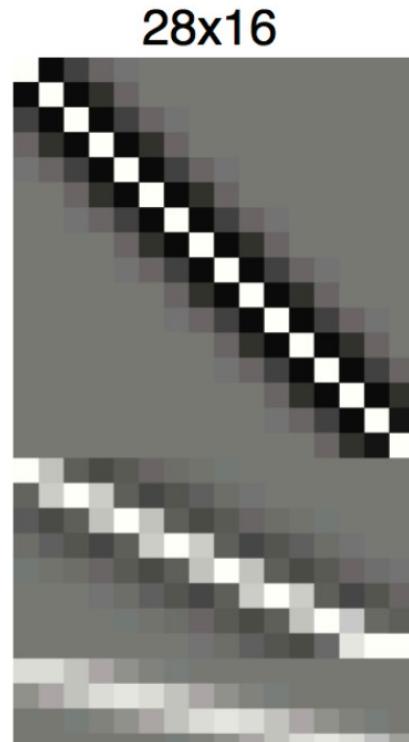
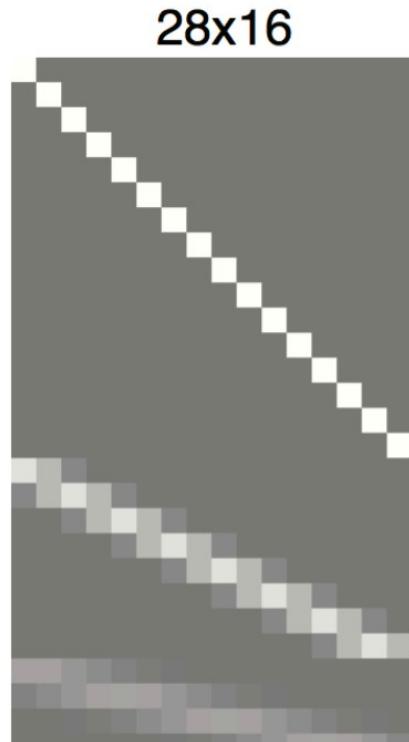
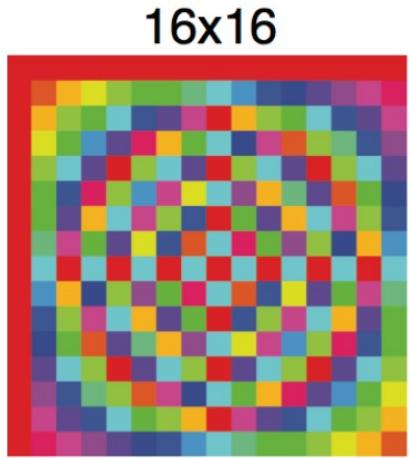




Image representation



- Pixels: great for spatial resolution, poor access to frequency
- Fourier transform: great for frequency, not for spatial information
- Pyramids/filter banks: balance between spatial and frequency information

Why use these representations?

- Handle real-world size variations with a constant-size vision algorithm.
- Remove noise
- Analyze texture
- Recognize objects, scenes
- Label image features

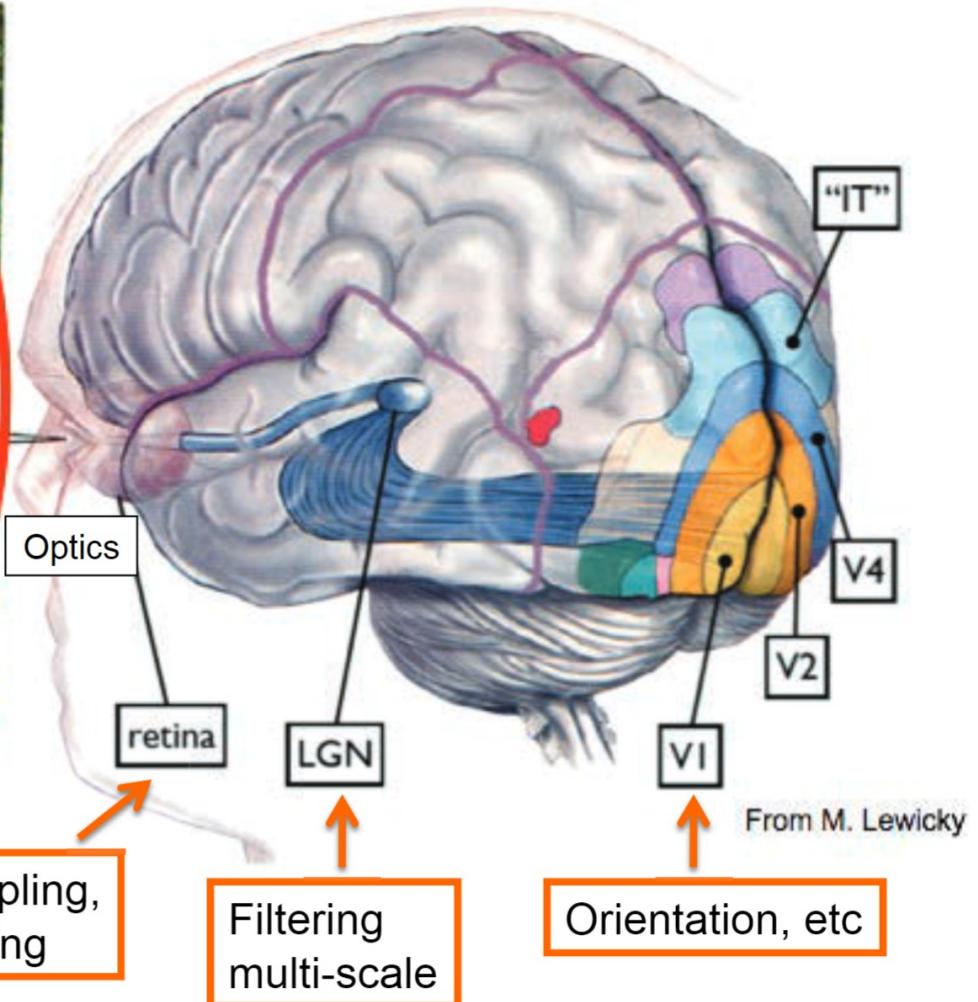


Today's class

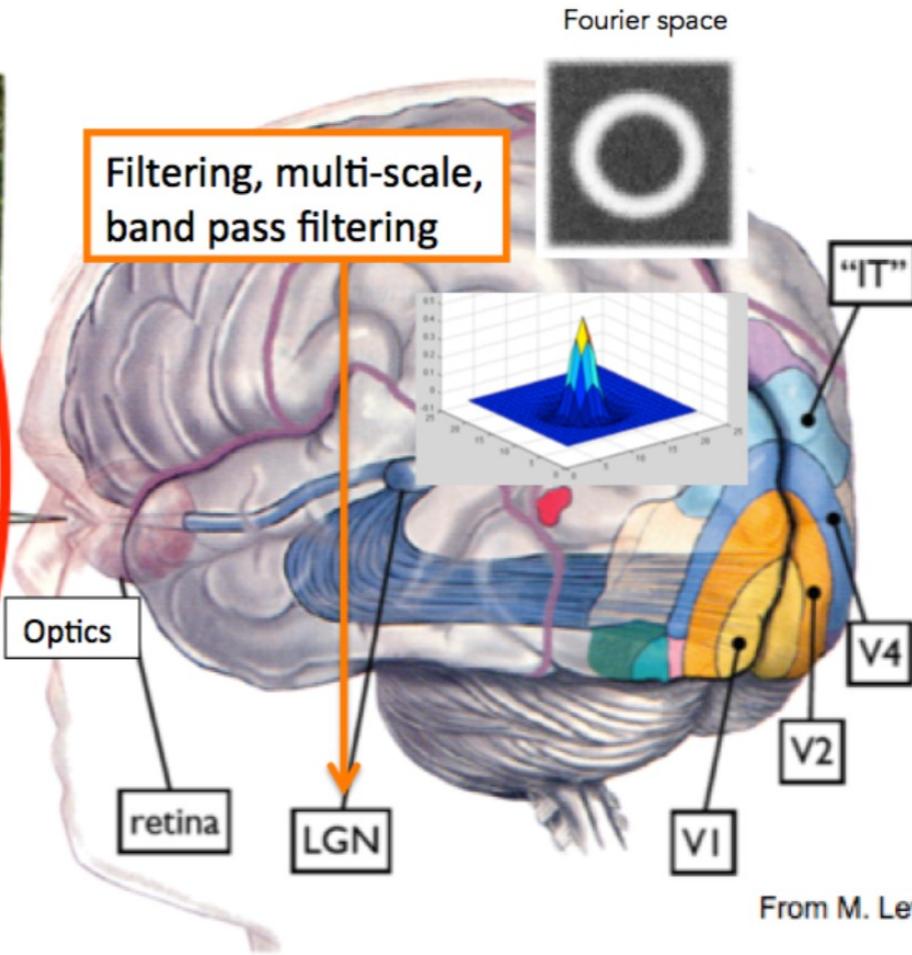


- Template matching
- Gaussian pyramids
 - Application for recognition
 - Pyramids representation in deep learning
- Laplacian pyramids
 - Hybrid images
 - Application for image blending
- Steerable pyramids
 - Steerable filters
 - Orientation analysis
- **Human vision system**
 - Visual perception
- Filter Banks and Texture Analysis

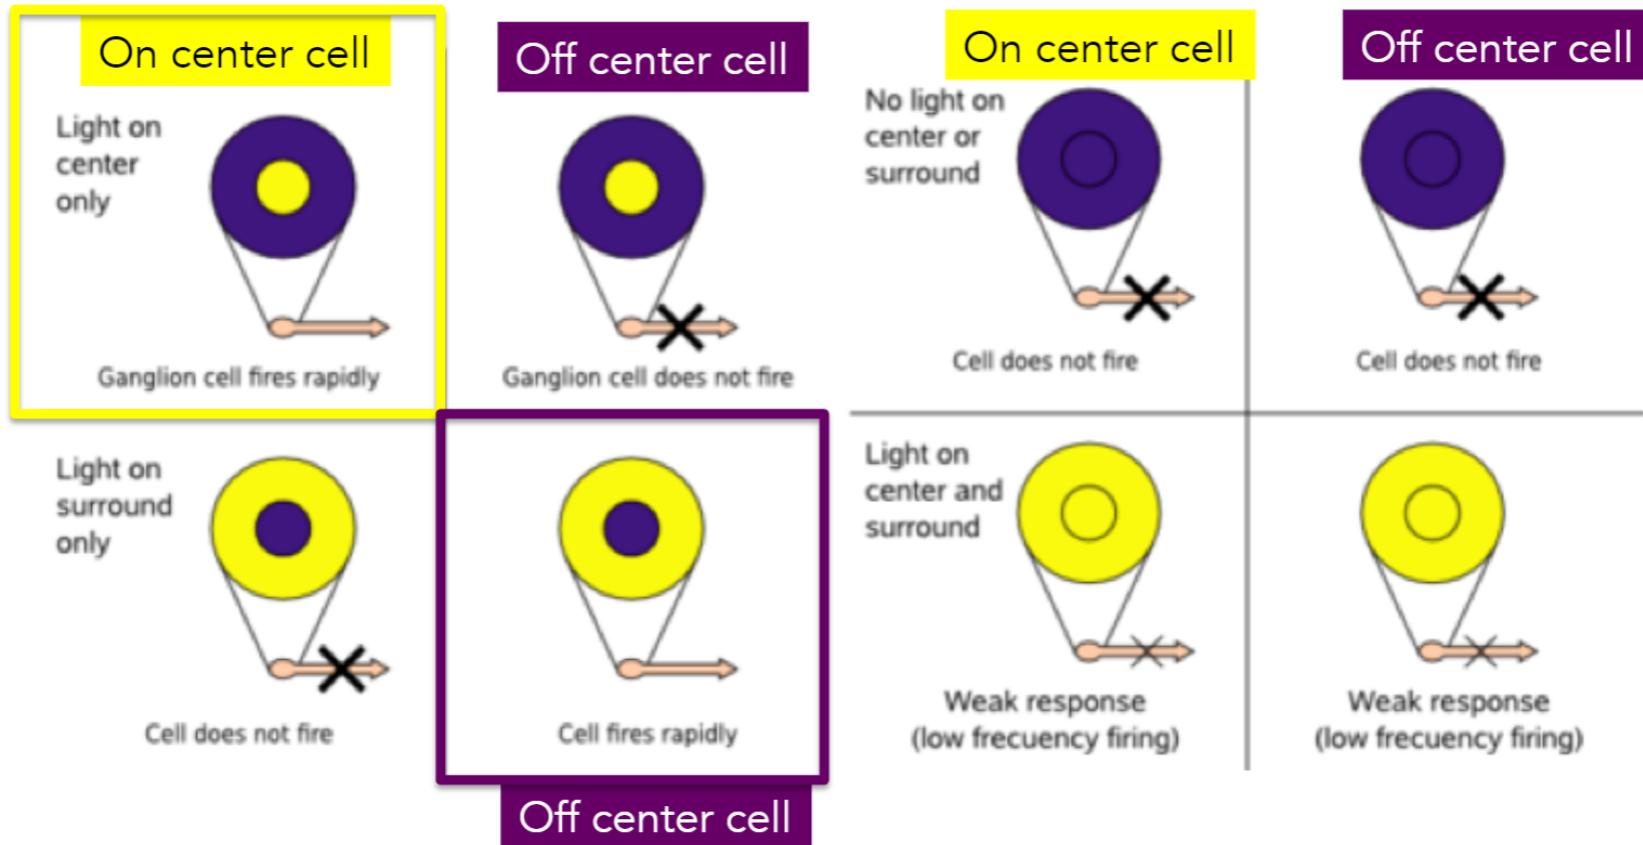
Vision: a multi-stage network



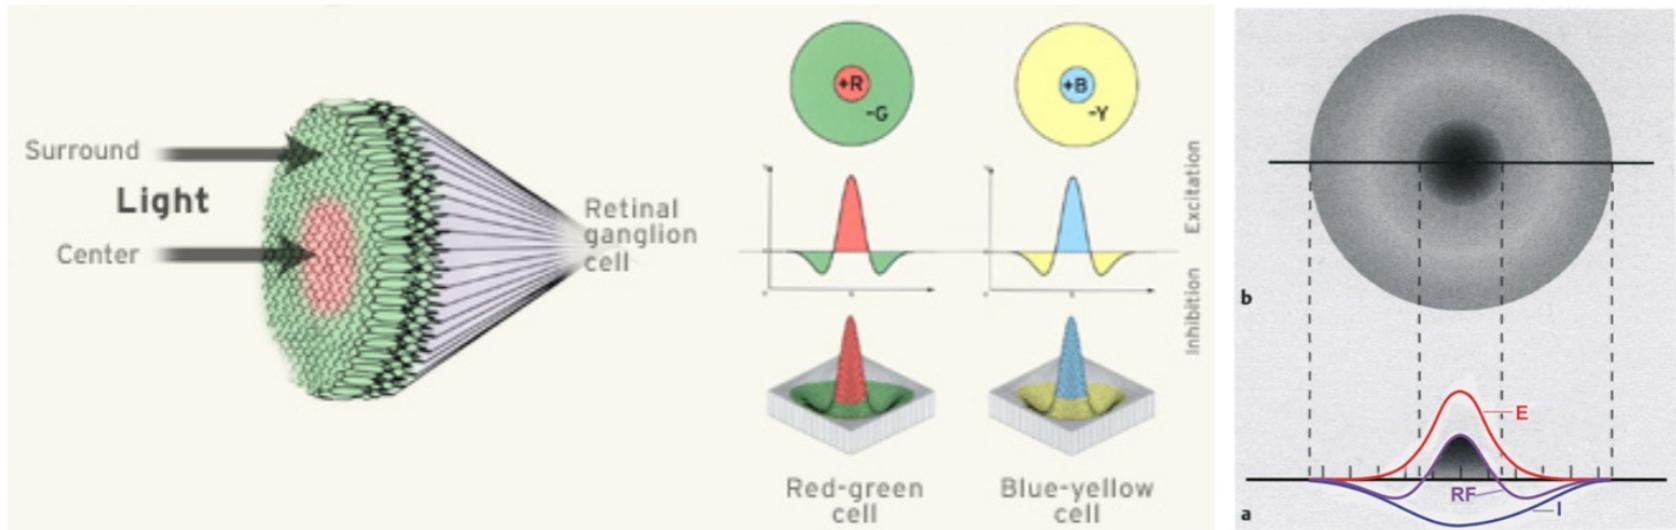
Vision: a multi-stage network



Model of Retinal and LGN cells

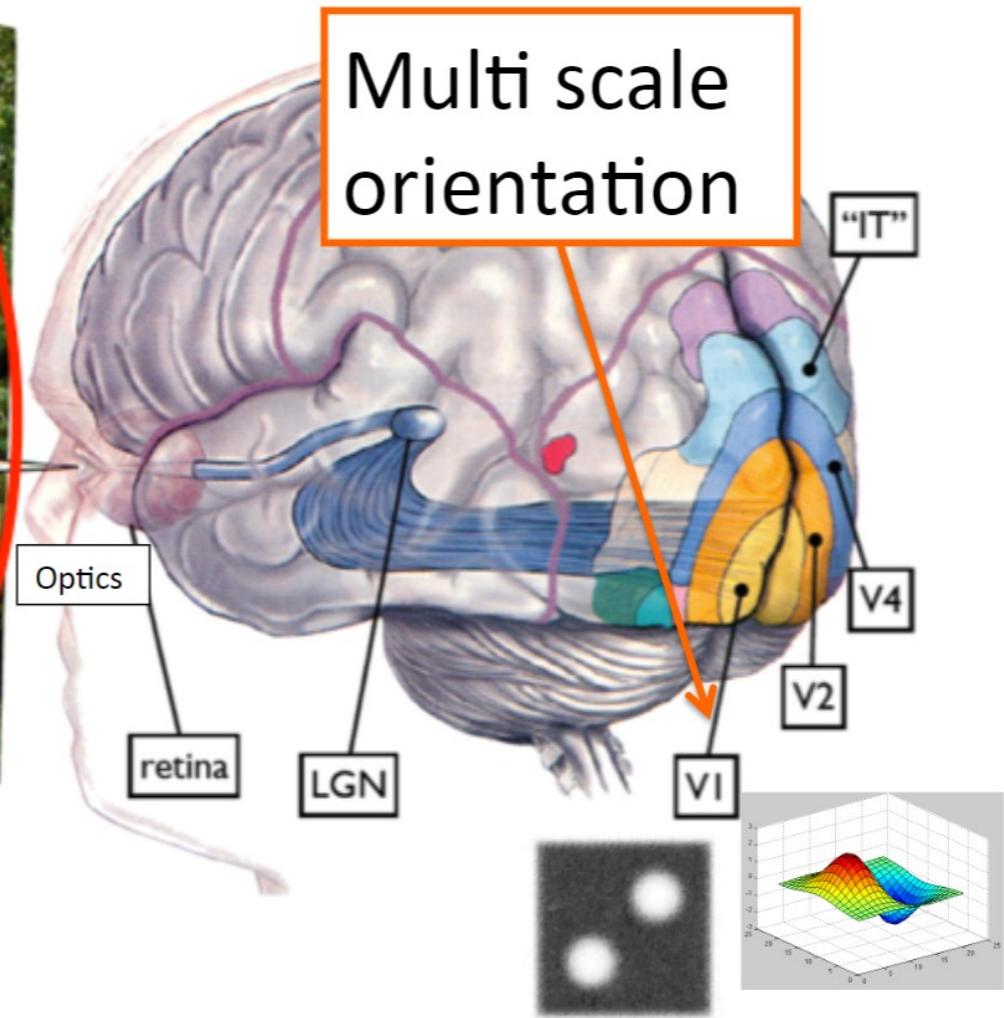
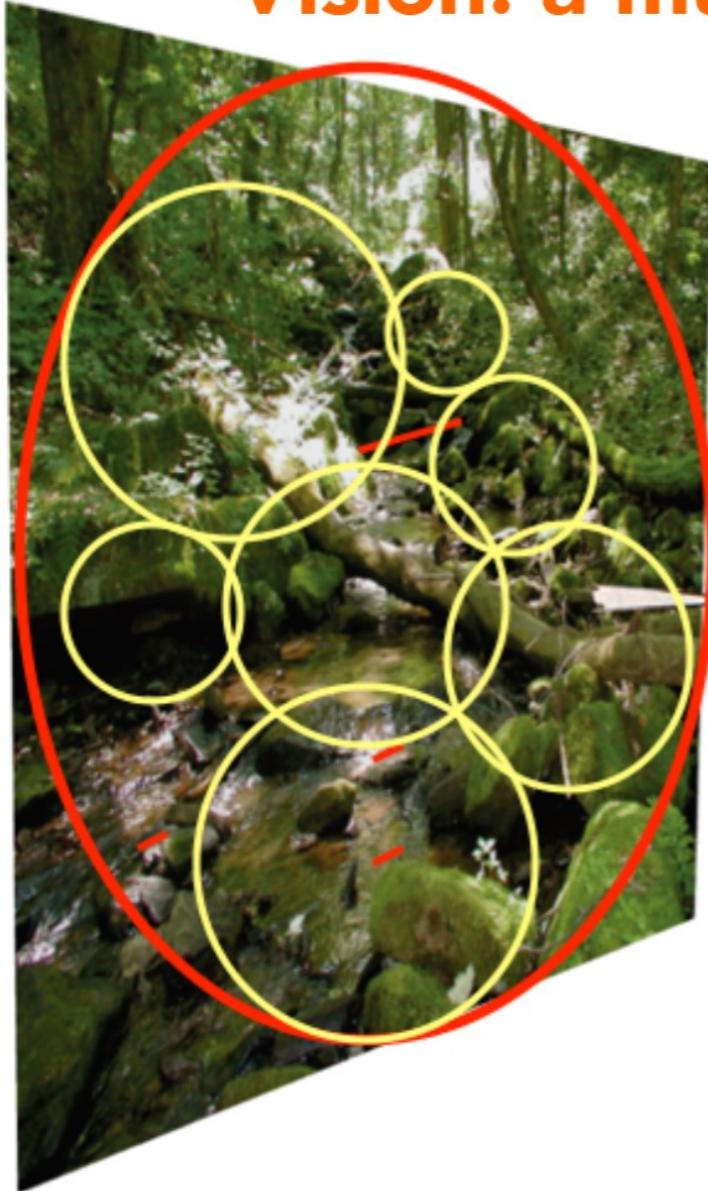


Difference of Gaussian Model



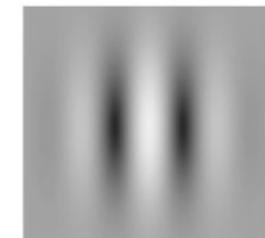
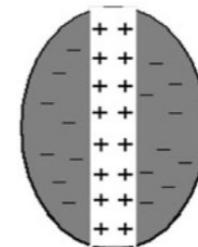
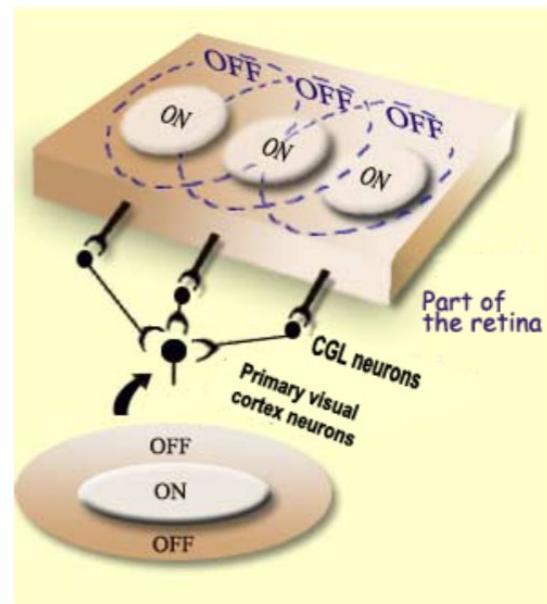
- The output of a retinal ganglion cell is a weighted sum of its inputs. You can treat the cell as a shift-invariant linear system
- **Difference of Gaussian (DoG):** A model of the operation performed by retinal and ganglion cells (LGN). The DoG model supposes that the neural response results from the combined signal of two separate mechanisms
- LGN and retinal neurons have circular receptive fields: **they respond equally well to all stimulus orientations, and at different spatial frequency (scale)**
- Multiple DoG is called a Laplacian Pyramid

Vision: a multi-stage network



Oriented representation in V1

- Simple-cell receptive fields (RFs) are constructed from the output of LGN cells
- They are selective to oriented contours and edges



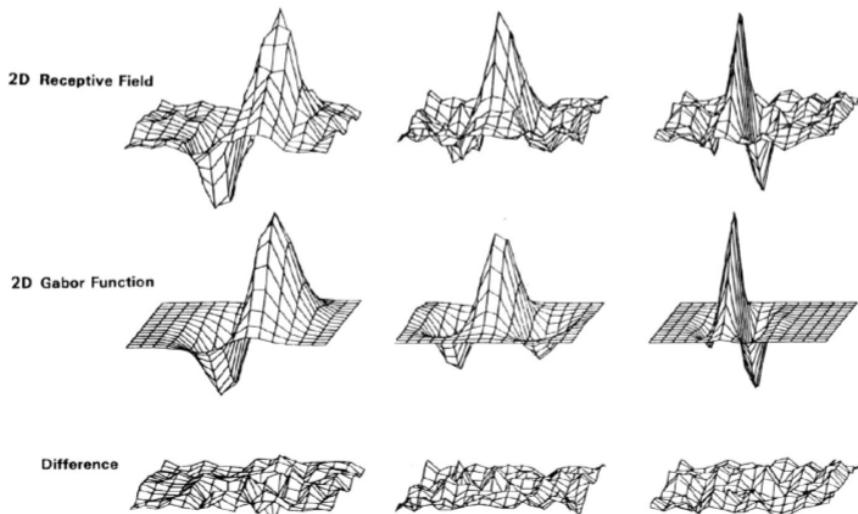
RECEPTIVE FIELD STRUCTURE

GABOR PATCH

V1 cells as Gabor Filters

The receptive fields of simple cells in V1 reflect the orientation and spatial frequency preferences of neurons. One way to model this, is to use a Gabor function, which is basically a two-dimensional gaussian modulated by a sinusoid.

2D Receptive fields in primary visual cortex

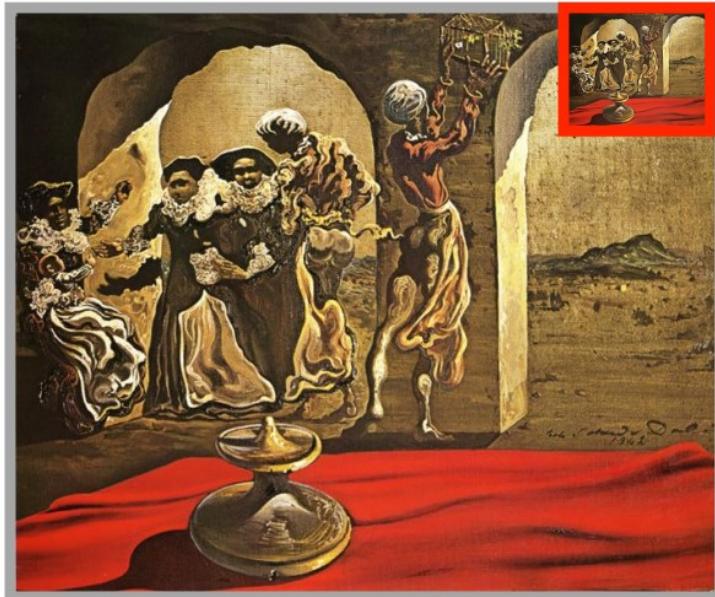


Fit of 2D Gabor wavelet is indistinguishable from noise.

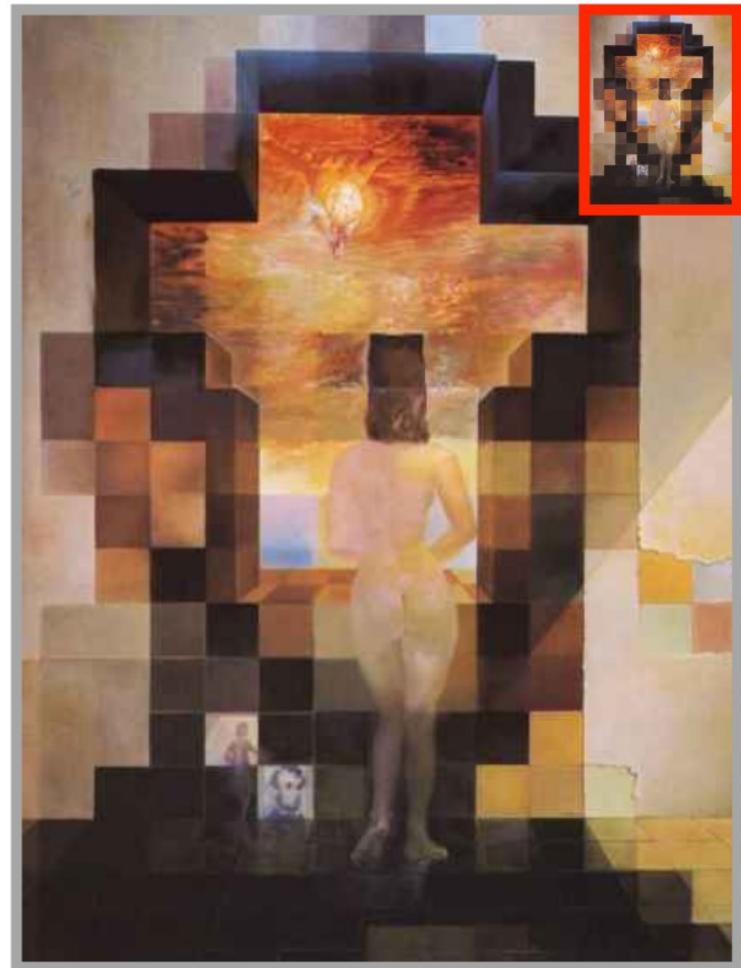
figure from Daugman, 1990
data from Jones and Palmer, 1987

Visual Perception

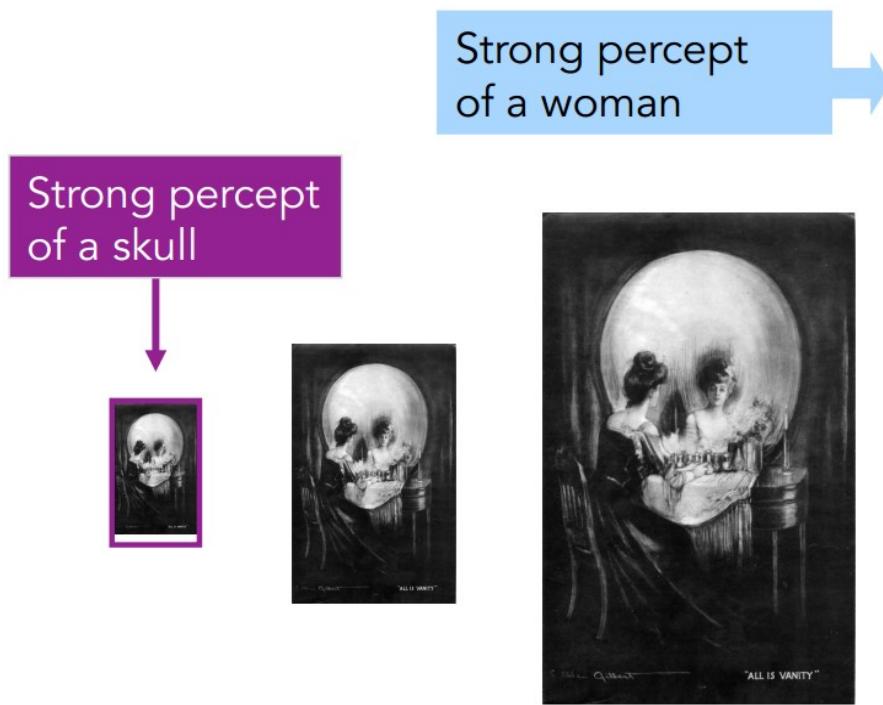
For computer vision systems to “understand art” other properties of human visual perception must be taken into account



Salvador Dalí



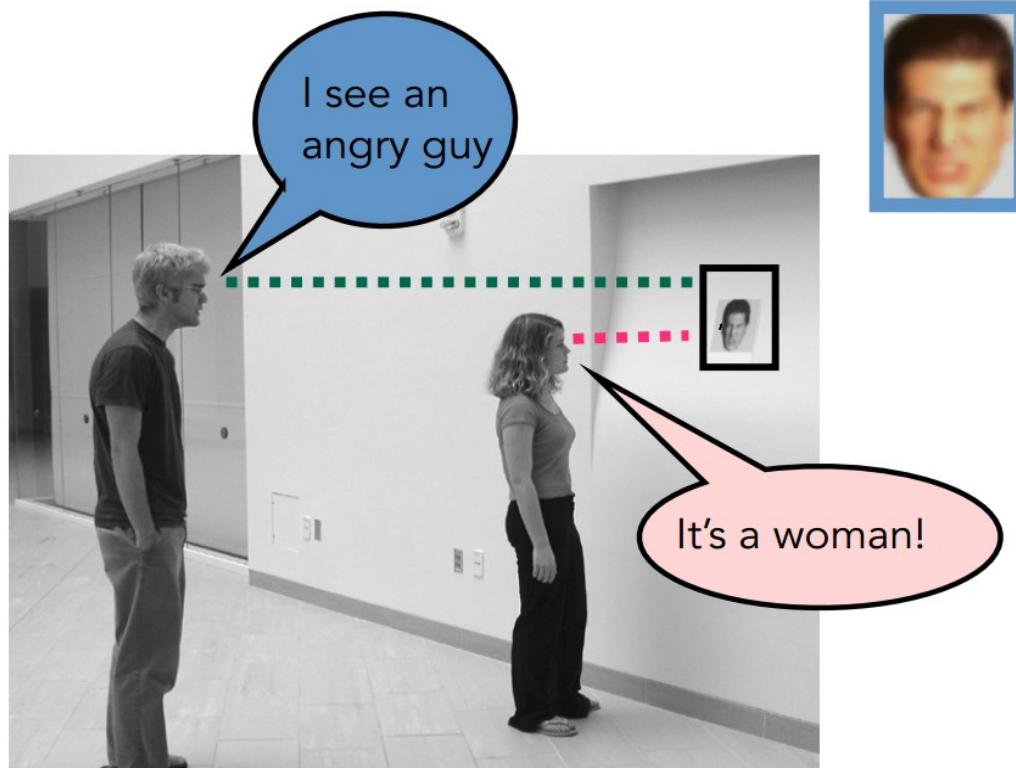
Changes of Meaning with Scale



Looking at the image one way, a skull appears, comprised of the woman's form and her reflection in a mirror (the outline of the skull itself).



Hybrid Images



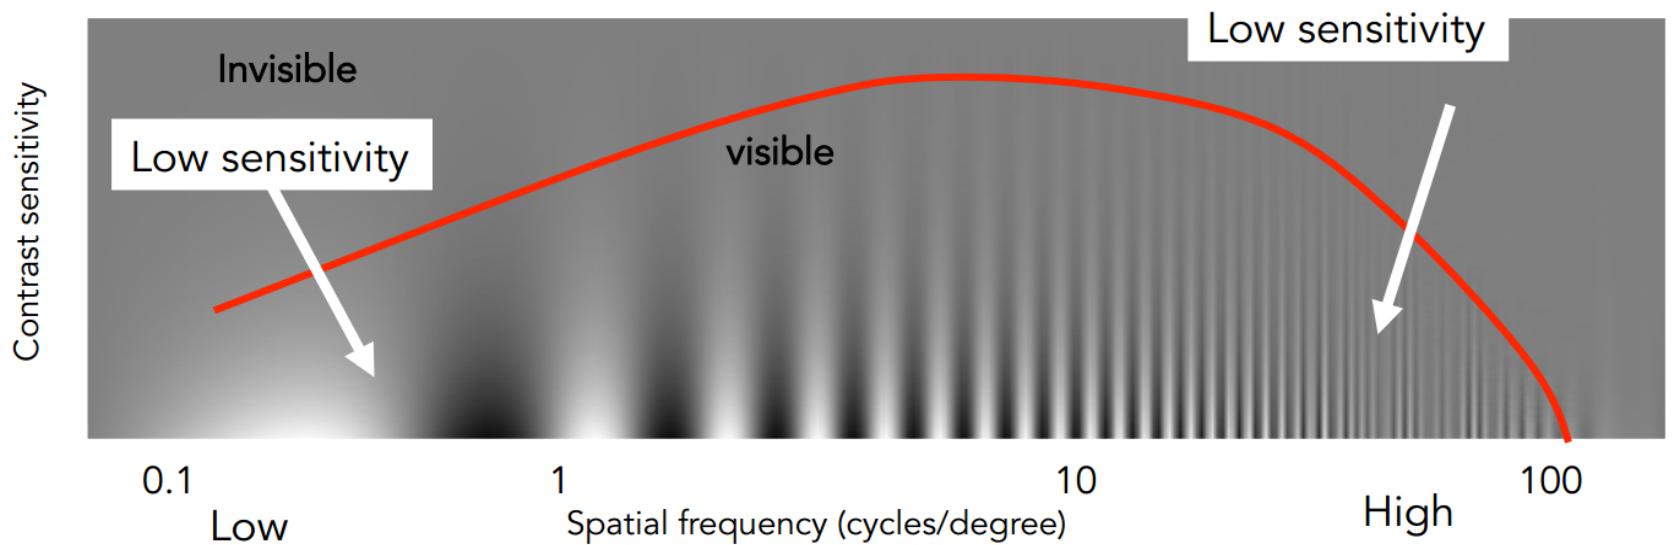
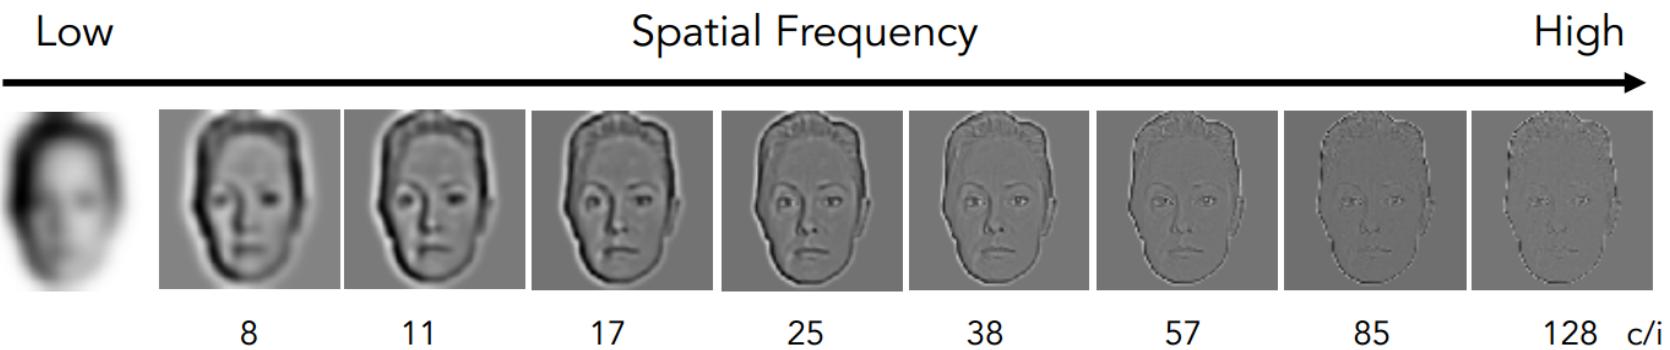
From Far Away



Up Close



Multiscale subband decomposition + Human Contrast Sensitivity Function





Today's class



- Template matching
- Gaussian pyramids
 - Application for recognition
 - Pyramids representation in deep learning
- Laplacian pyramids
 - Hybrid images
 - Application for image blending
- Steerable pyramids
 - Steerable filters
 - Orientation analysis
- Human vision system
 - Visual perception
- **Filter Banks and Texture Analysis**

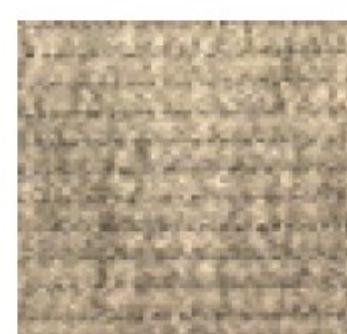
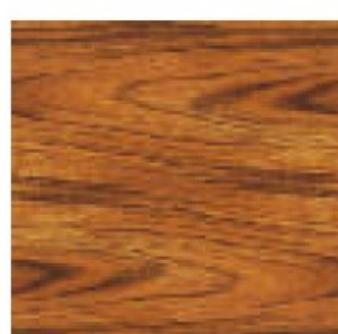
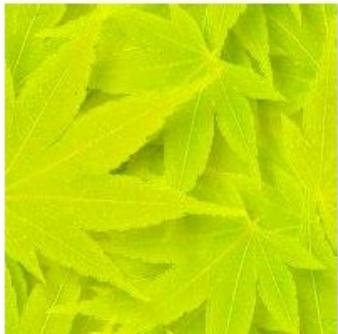
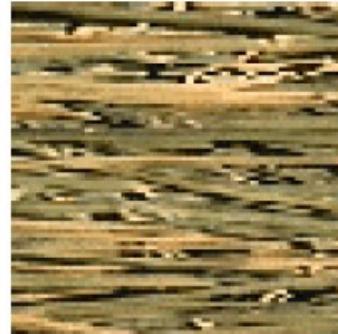


Application: Representing Texture





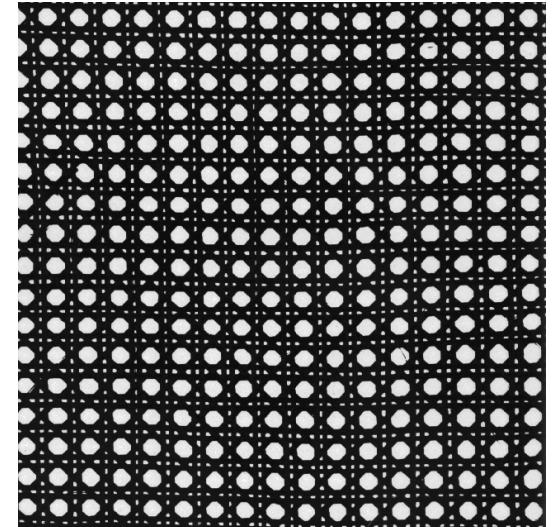
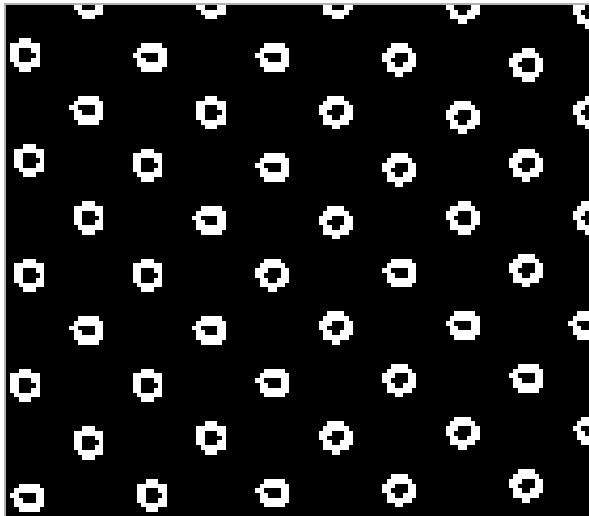
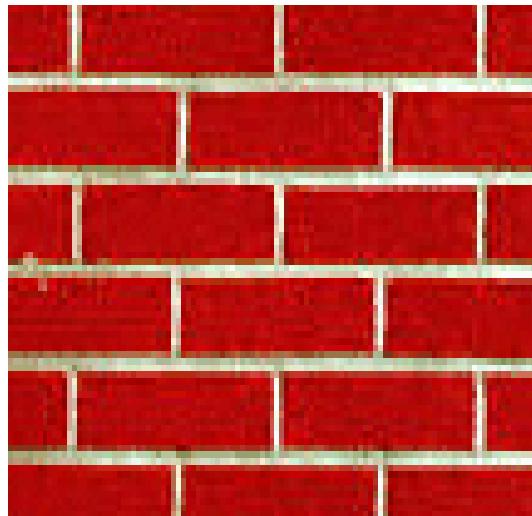
Texture



What defines a texture?

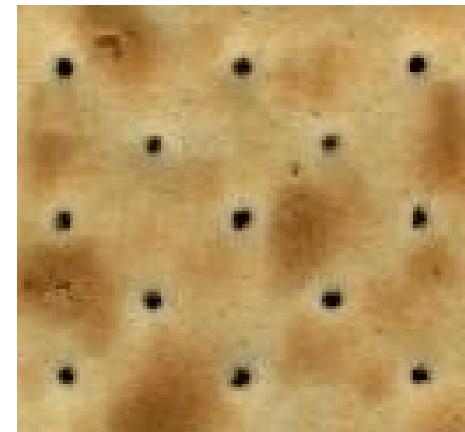
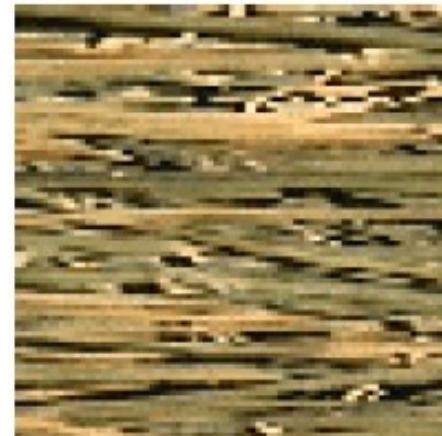


Includes: more regular patterns



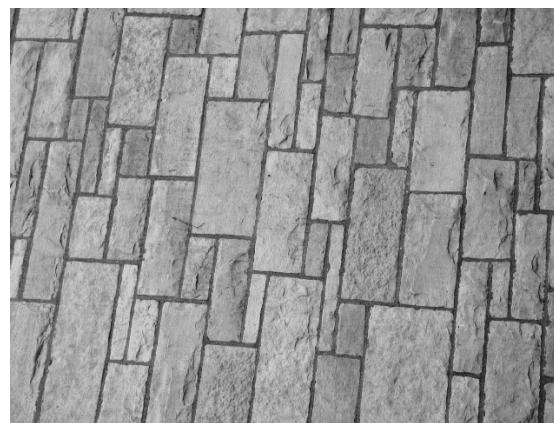
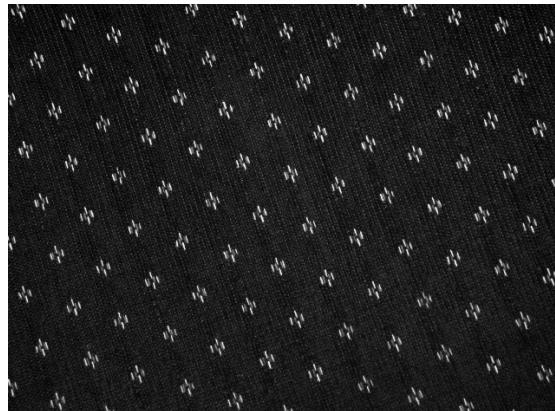
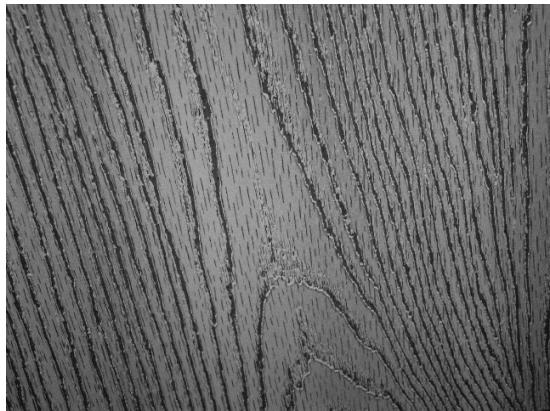


Includes: more random patterns





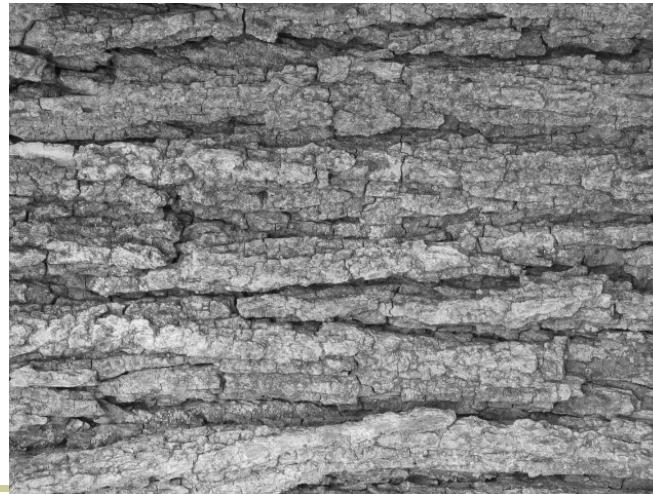
Texture and Material



http://www-cvr.ai.uiuc.edu/ponce_grp/data/texture_database/samples/



Texture and Orientation





Texture and Scale

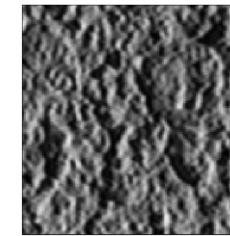
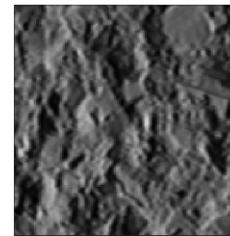
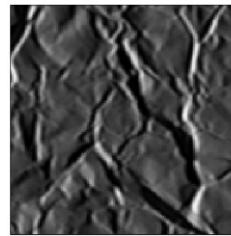
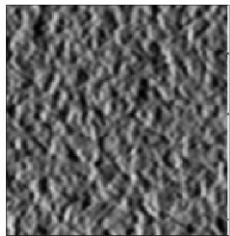
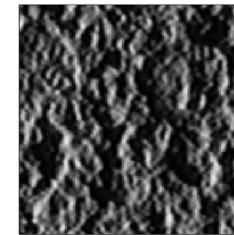
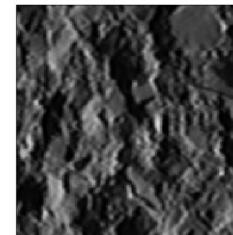
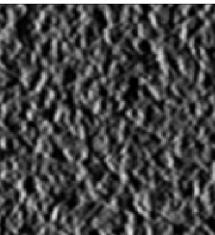
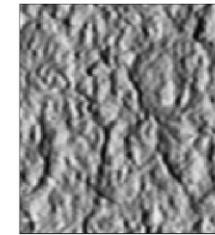
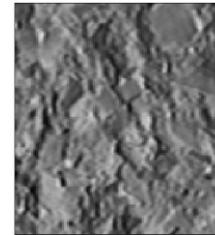
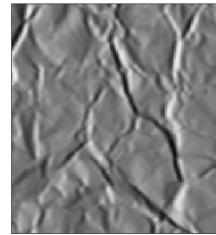
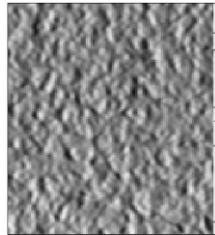




Materials under different illumination and viewing directions



Different
illumination
and viewing
directions



Plaster-a

Crumpled
Paper

Concrete

Plaster-b
(zoomed)



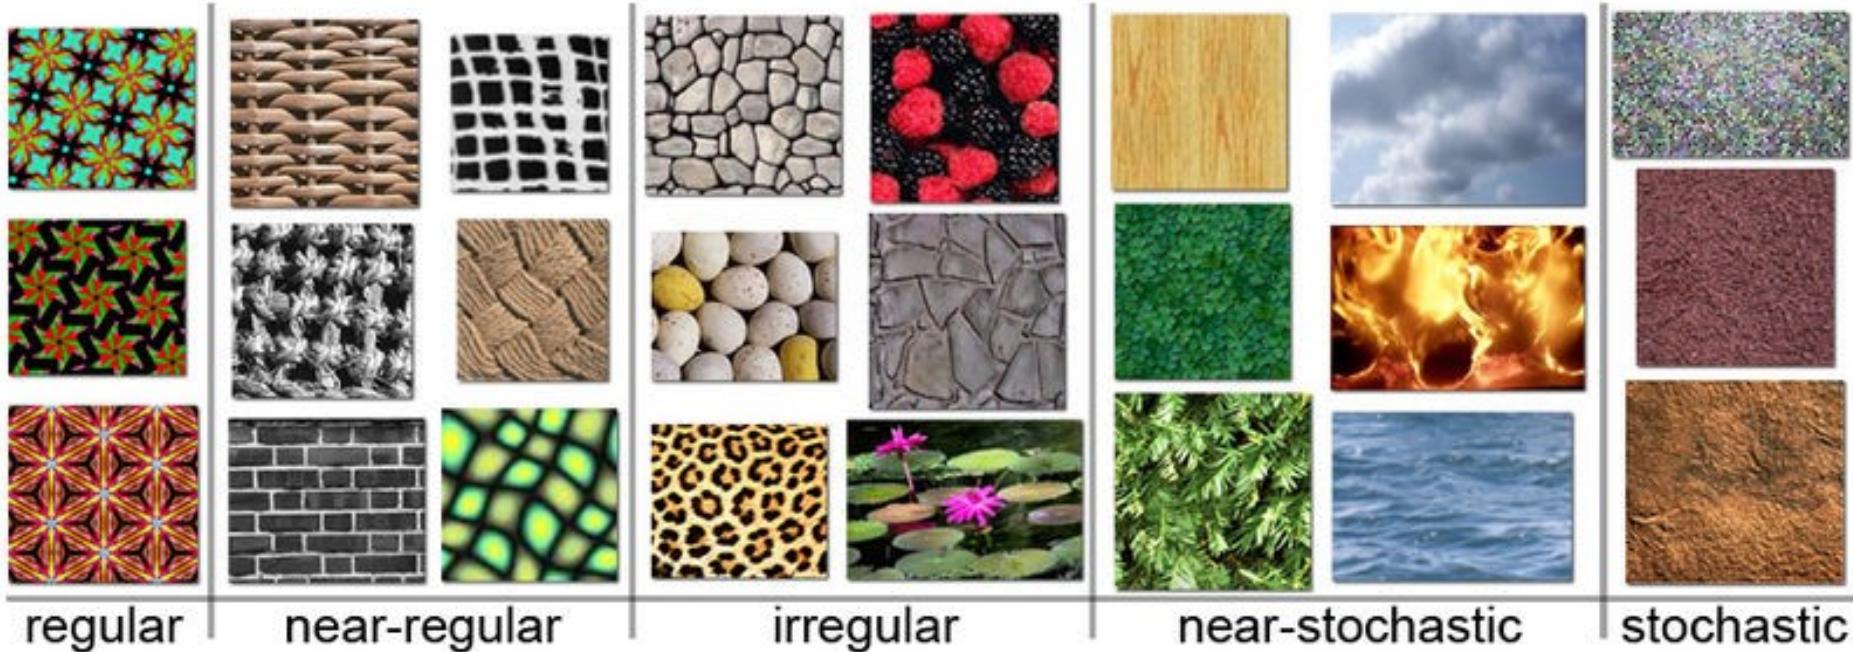
What is texture?



- **Texture is a phenomenon that is widespread, easy to recognize, and hard to define.**
- Views of large numbers of small objects
- Regular or stochastic patterns caused by bumps, grooves, and/or markings
- Textures tend to show **repetition**: the same local patch appears again and again.



Texture overview





Texture-related tasks



■ Shape from texture

- Estimate surface orientation or shape from image texture

■ Segmentation/classification from texture cues

- Analyze, represent texture
- Group image regions with consistent texture

■ Synthesis

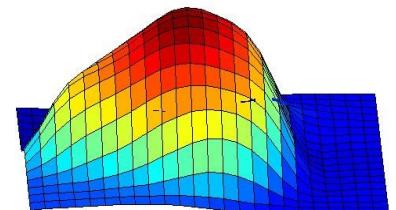
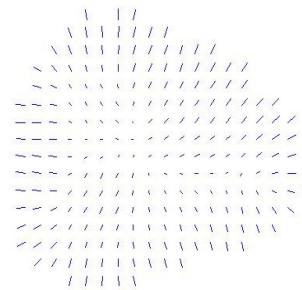
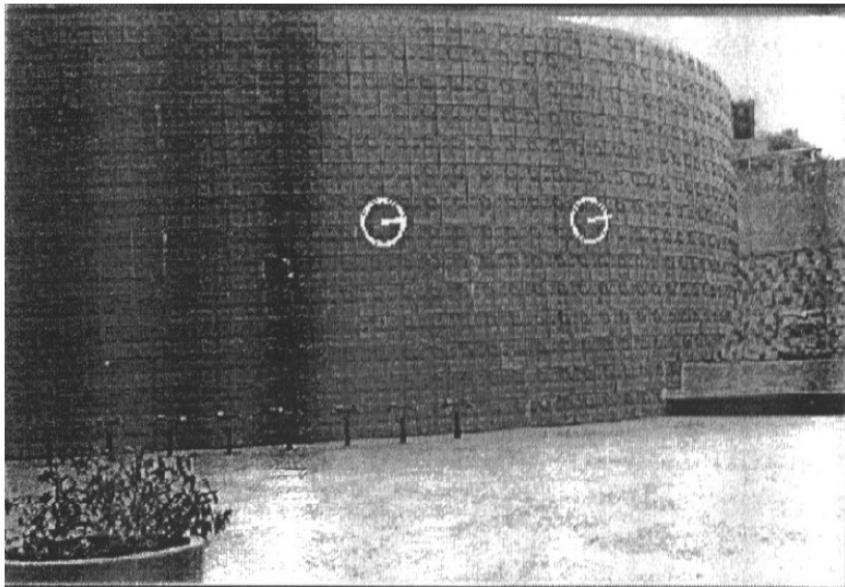
- Generate new texture patches/images given some examples



Shape from texture

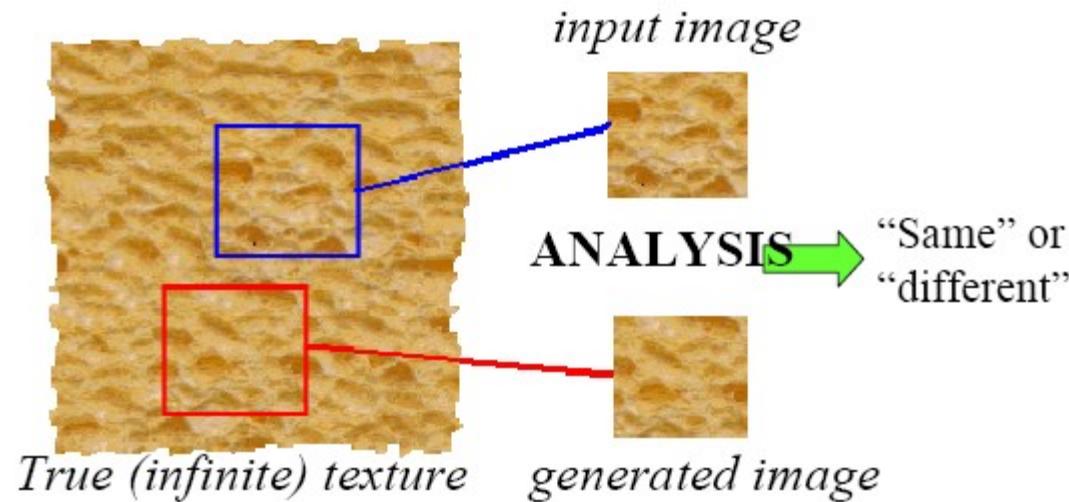


- Use deformation of texture from point to point to estimate surface shape

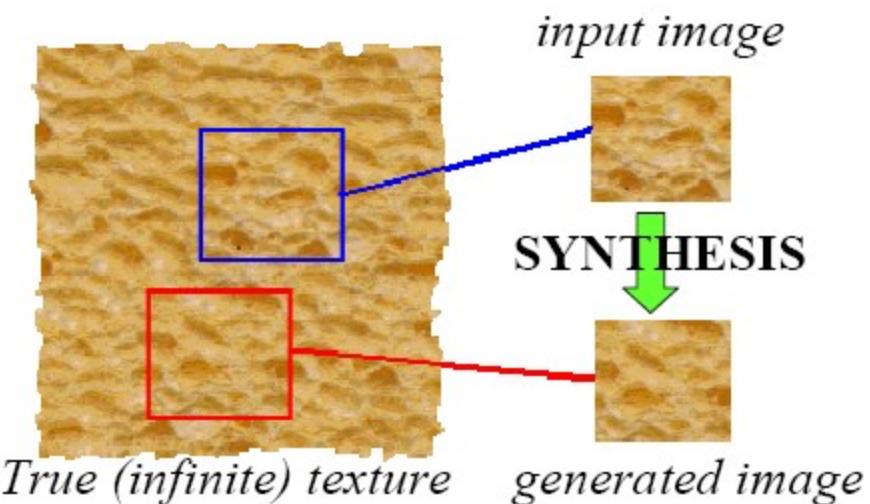




Analysis vs. Synthesis

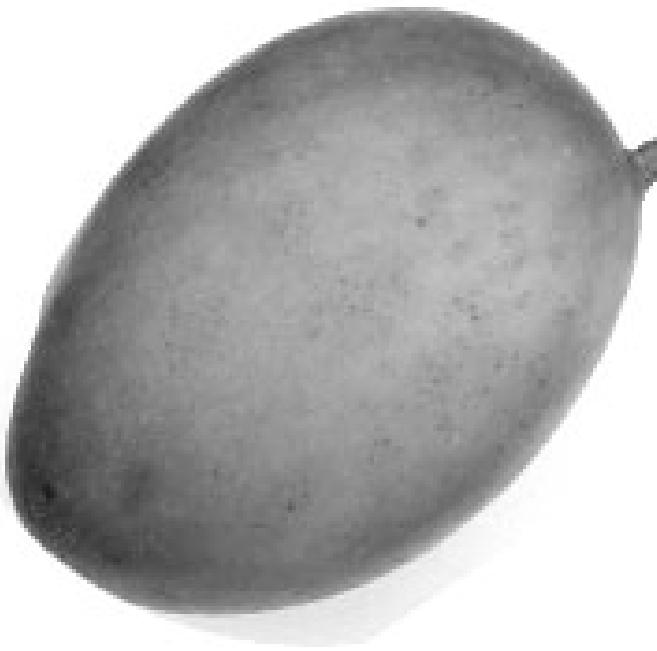
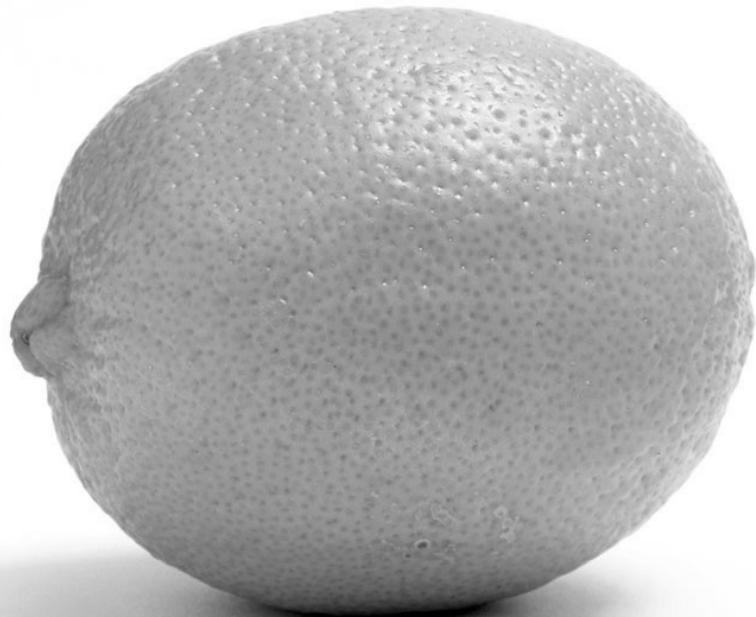


Why analyze
texture?





Slide credit:
Kristen Grauman



Slide credit:
Kristen Grauman



Slide credit:
Kristen Grauman



Why analyze texture?



Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.



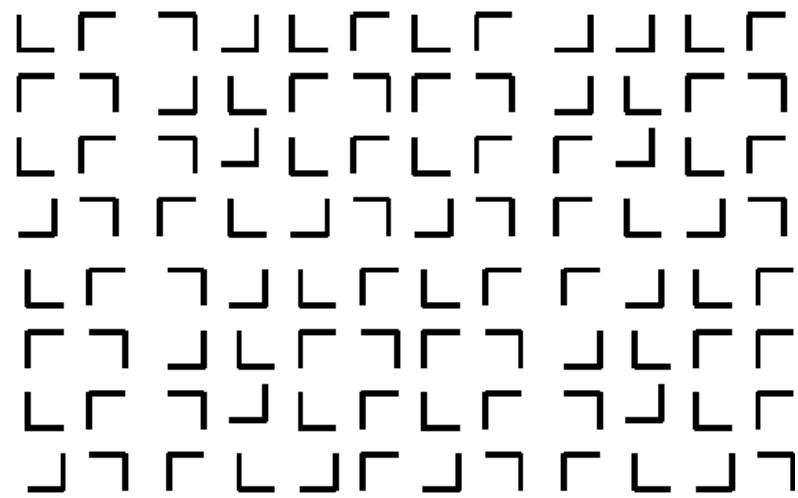
Psychophysics of texture

- Some textures distinguishable with *preattentive* perception— without scrutiny, eye movements [Julesz 1975]

Same or different?



丁 戊 戣 丙 戌 戍 戲 戌 戌
戊 戌 戔 丙 戌 戌 戌 戌 戌
丁 戊 戔 丙 戌 戌 戌 戌 戌
戊 戌 戔 丙 戌 戌 戌 戌 戌
丁 戊 戔 丙 戌 戌 戌 戌 戌
戊 戌 戔 丙 戌 戌 戌 戌 戌
丁 戊 戔 丙 戌 戌 戌 戌 戌
戊 戌 戔 丙 戌 戌 戌 戌 戌





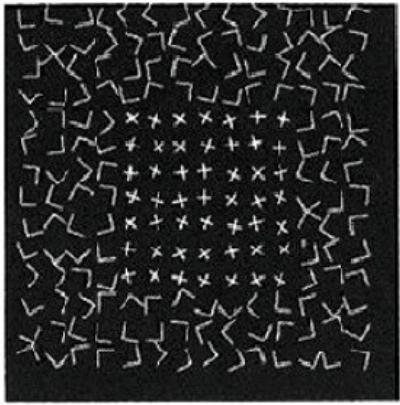
A 10x10 grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, where each block is positioned such that its vertical stem is aligned with the horizontal bar of the block directly above it. This creates a continuous, flowing pattern across the entire grid.

A large, dense grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, creating a complex, woven texture. The grid spans most of the page, with some blocks at the top and bottom edges appearing slightly smaller due to the perspective.





Capturing the local patterns with image measurements



[Bergen &
Adelson,
Nature 1988]

Scale of
patterns
influences
discriminability

Size-tuned
linear filters



Texture representation



- Textures are made up of repeated local patterns, so:
 - Find the patterns
 - Use filters that look like patterns (spots, bars, raw patches...)
 - Consider magnitude of response
 - Describe their statistics within each local window, e.g.,
 - Mean, standard deviation
 - Histogram
 - Histogram of “prototypical” feature occurrences



Texture representation: example



original image



derivative filter
responses, squared

	<u>mean d/dx value</u>	<u>mean d/dy value</u>
Win. #1	4	10

⋮

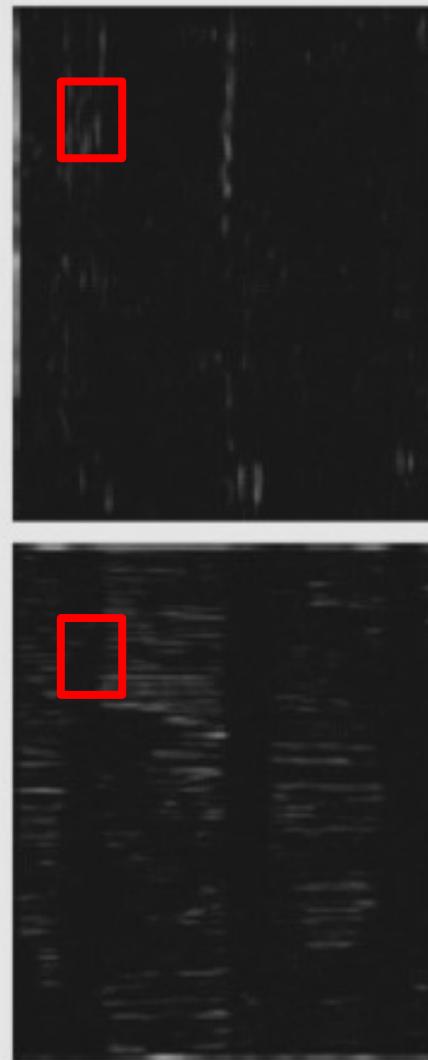
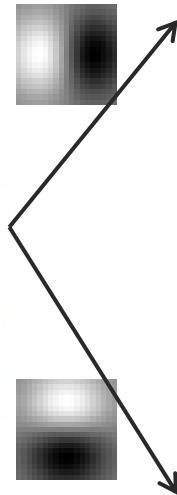
statistics to summarize
patterns in small
windows



Texture representation: example



original image



derivative filter
responses, squared

	<u>mean d/dx value</u>	<u>mean d/dy value</u>
Win. #1	4	10
Win.#2	18	7

⋮

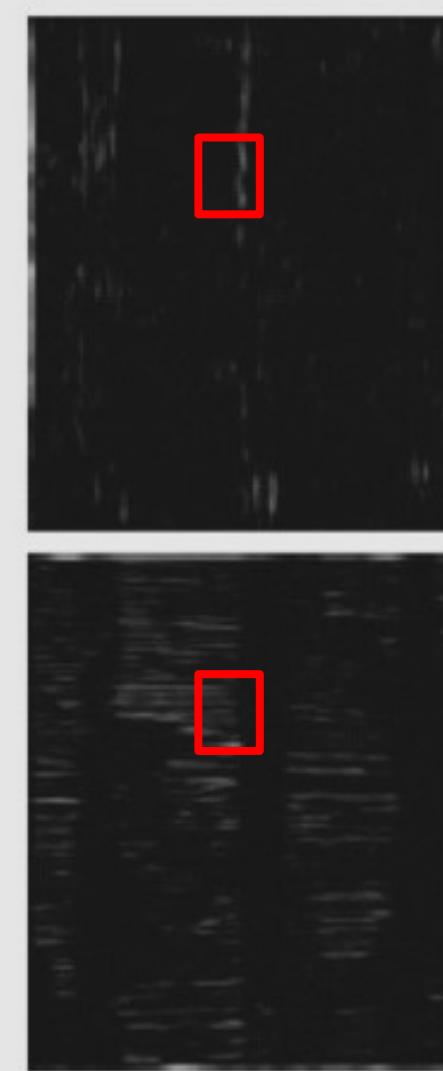
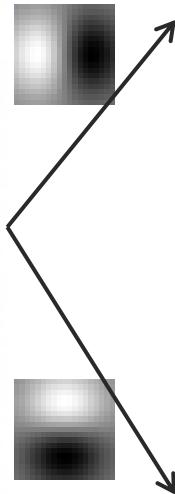
statistics to summarize
patterns in small
windows



Texture representation: example



original image



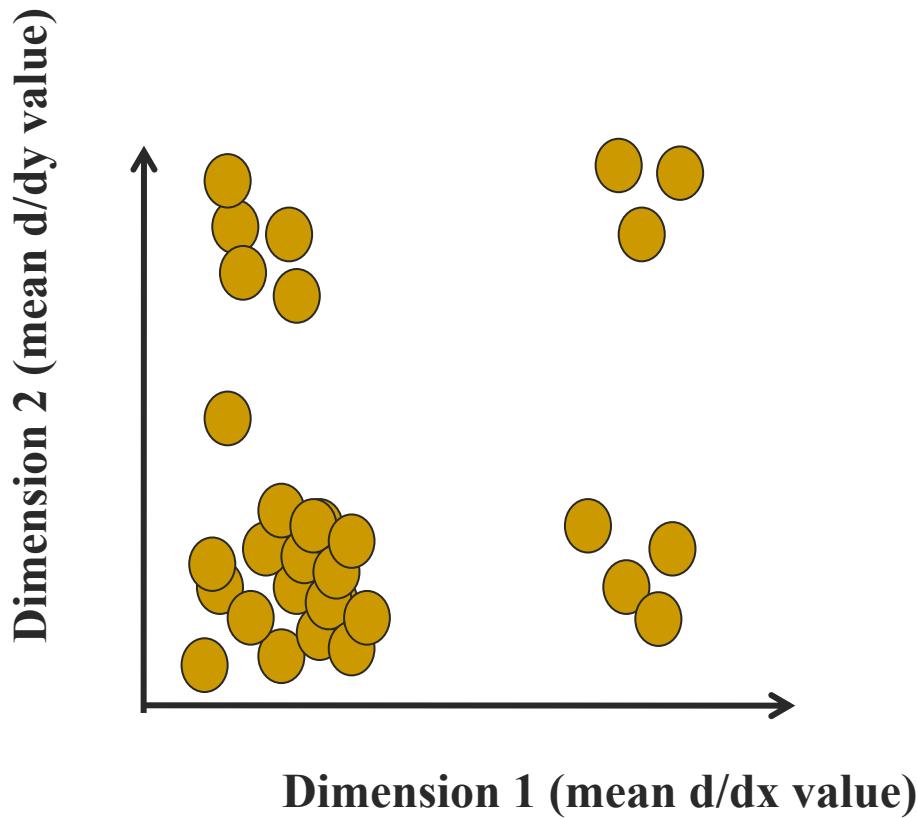
derivative filter
responses, squared

	<u>mean d/dx value</u>	<u>mean d/dy value</u>
Win. #1	4	10
Win.#2	18	7
Win.#9	20	20
⋮	⋮	⋮

statistics to summarize
patterns in small
windows



Texture representation: example



	<u>mean d/dx value</u>	<u>mean d/dy value</u>
Win. #1	4	10
Win.#2	18	7
Win.#9	20	20
⋮	⋮	⋮

statistics to summarize
patterns in small
windows



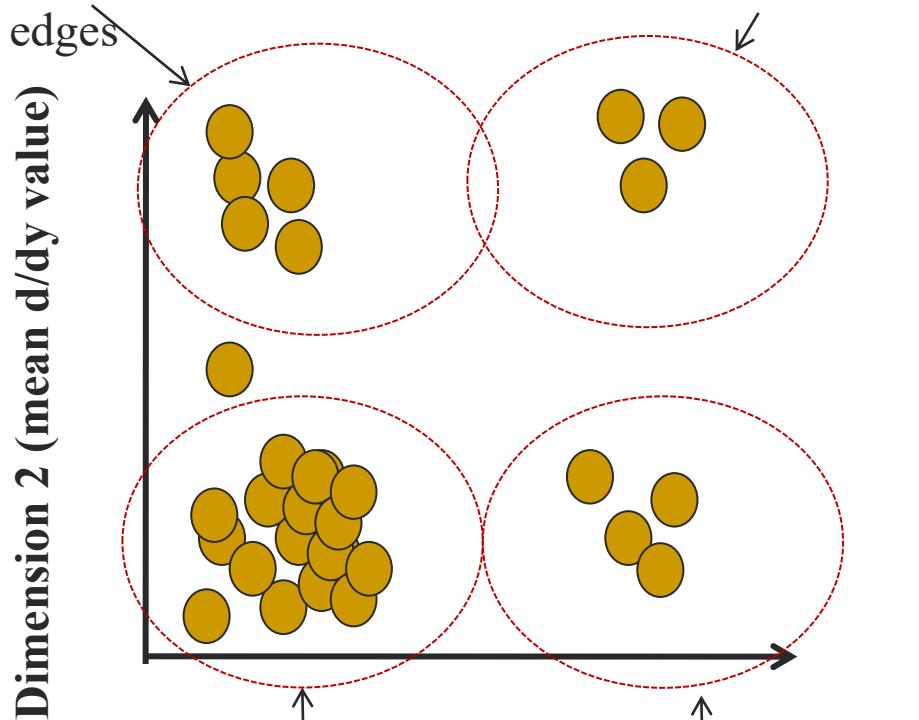
Texture representation: example



Windows with

primarily horizontal
edges

Both



Windows with
small gradient in
both directions

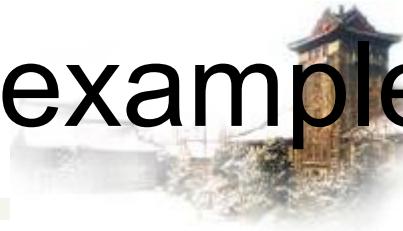
Windows with
primarily vertical
edges

statistics to summarize
patterns in small
windows

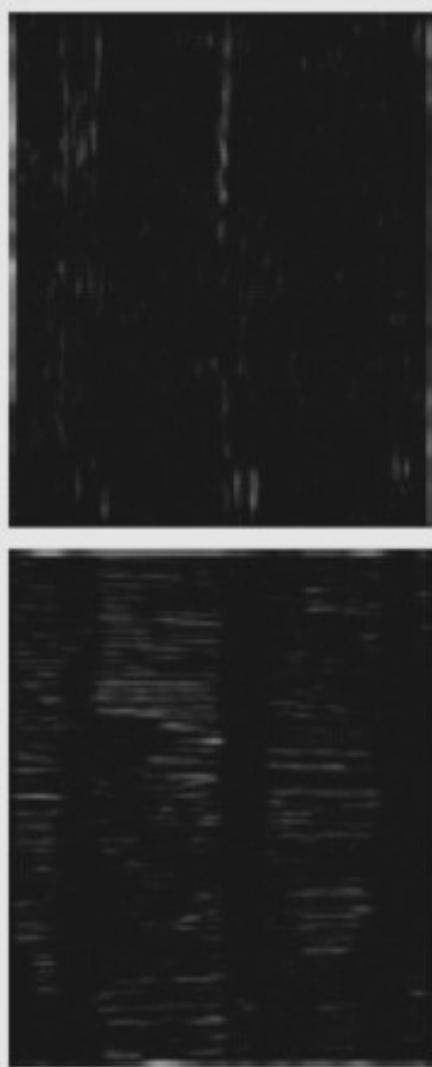
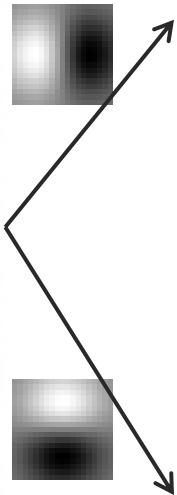
	<u>mean</u> <u>d/dx</u> <u>value</u>	<u>mean</u> <u>d/dy</u> <u>value</u>
Win. #1	4	10
Win.#2	18	7
Win.#9	20	20
⋮	⋮	⋮



Texture representation: example



original image



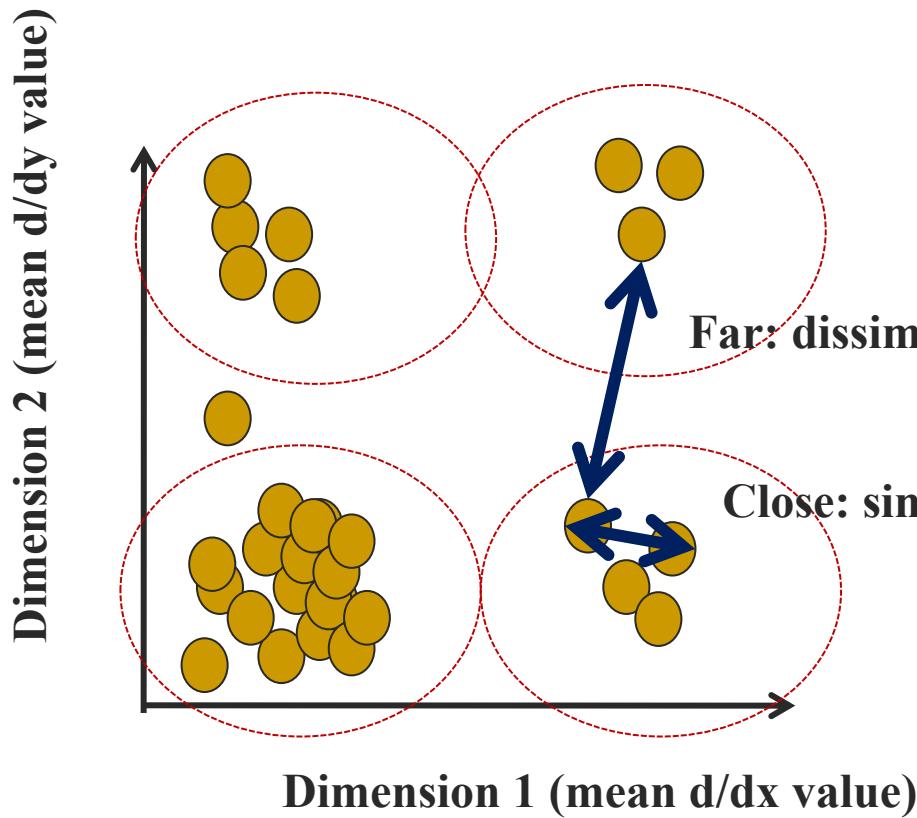
derivative filter
responses, squared



visualization of the
assignment to texture
“types”



Texture representation: example

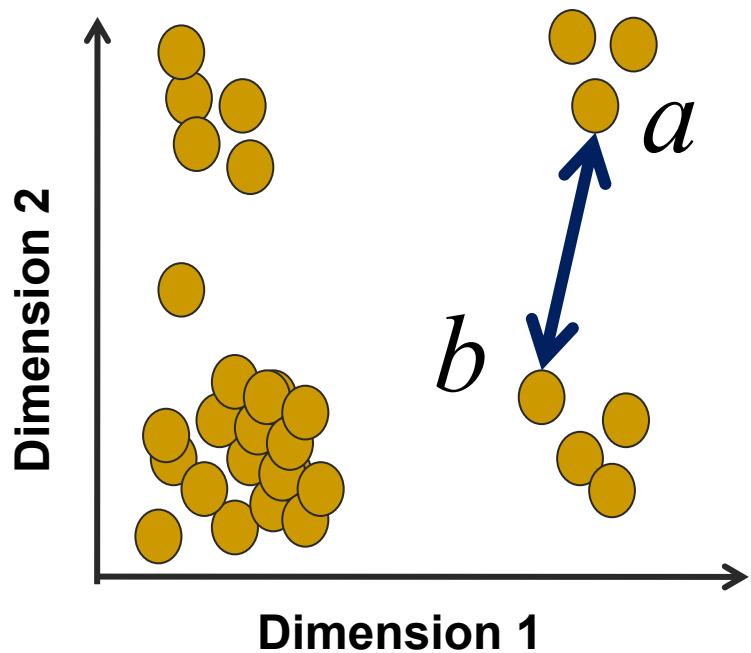


	mean d/dx value	mean d/dy value
Win. #1	4	10
Win. #2	18	7
Win. #9	20	20
⋮	⋮	⋮

statistics to summarize
patterns in small
windows



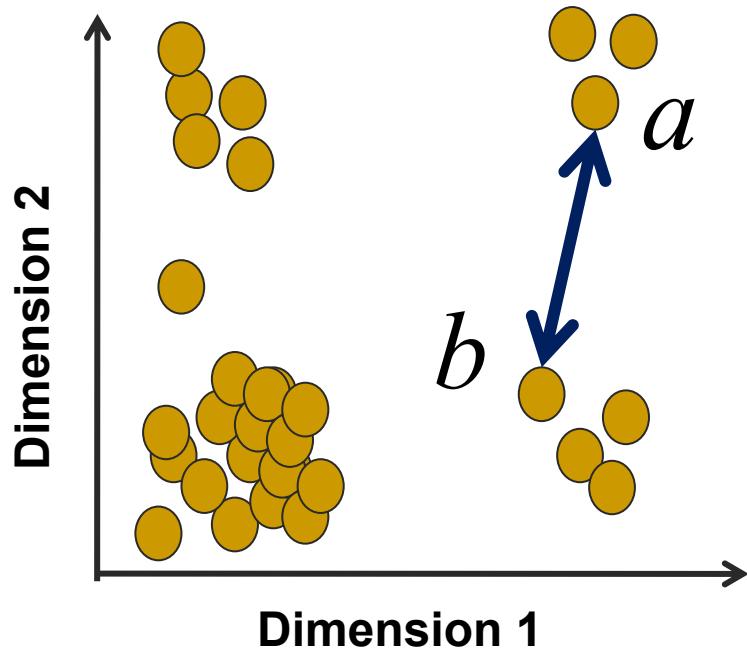
Texture representation: example



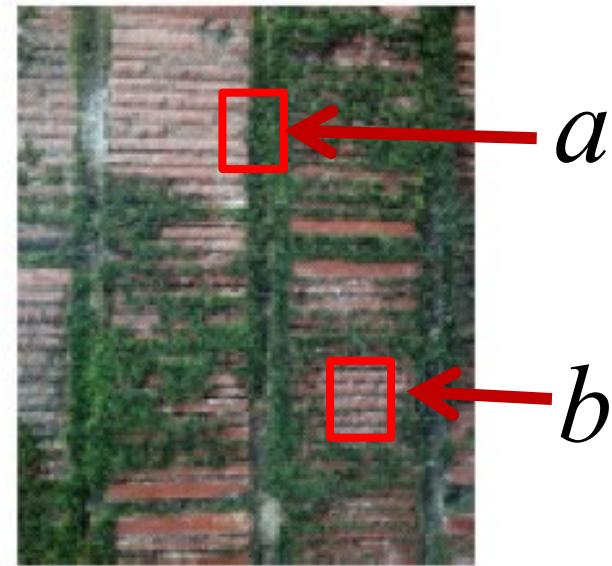
$$D(a,b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2}$$



Texture representation: example



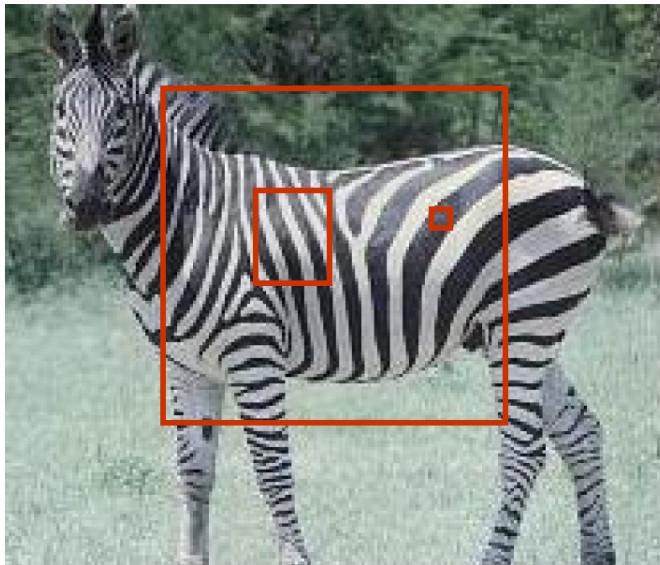
Distance reveals how dissimilar texture from window a is from texture in window b.





Texture representation: window scale

- We're assuming we know the relevant window size for which we collect these statistics.



Possible to perform **scale selection** by looking for window scale where texture description not changing.



Filter banks

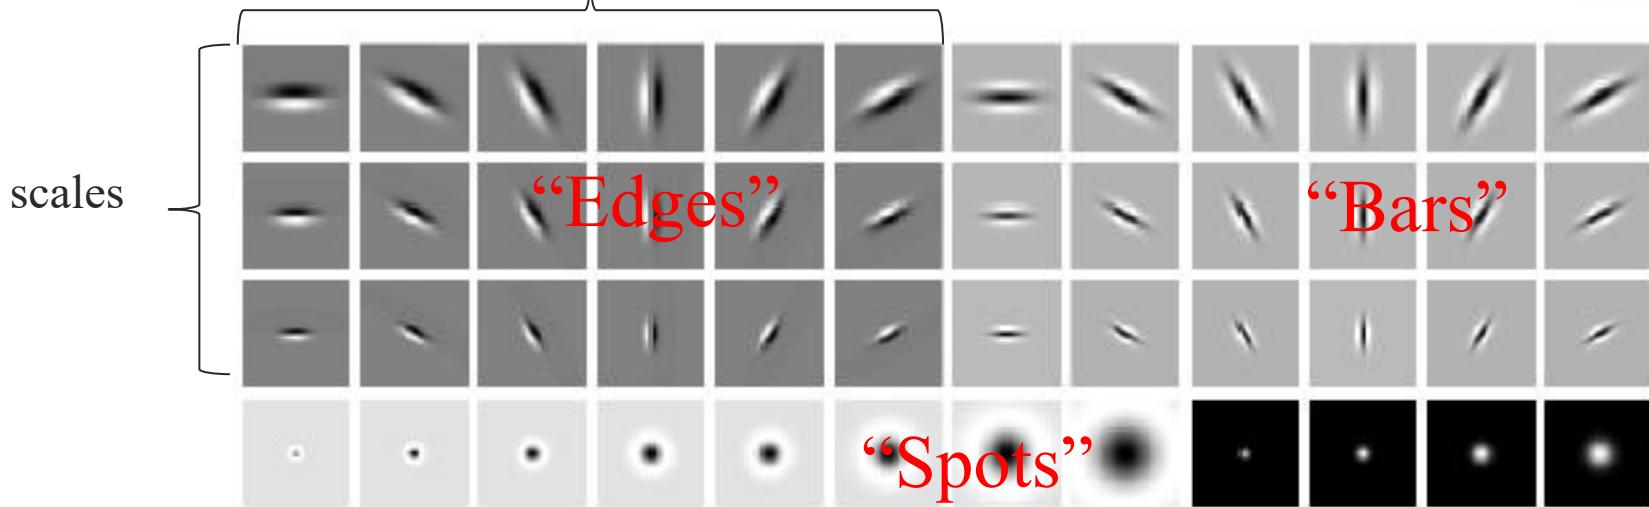


- Our previous example used two filters, and resulted in a 2-dimensional feature vector to describe texture in a window.
 - x and y derivatives revealed something about local structure.
- We can generalize to apply a collection of multiple (d) filters: a “filter bank”
- Then our feature vectors will be d -dimensional.
 - still can think of nearness, farness in feature space



Filter banks

orientations



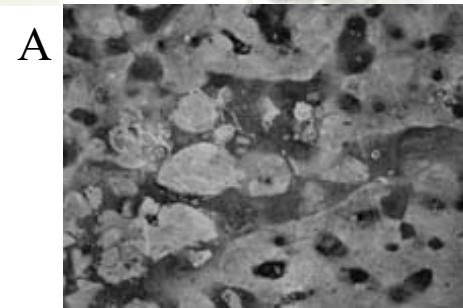
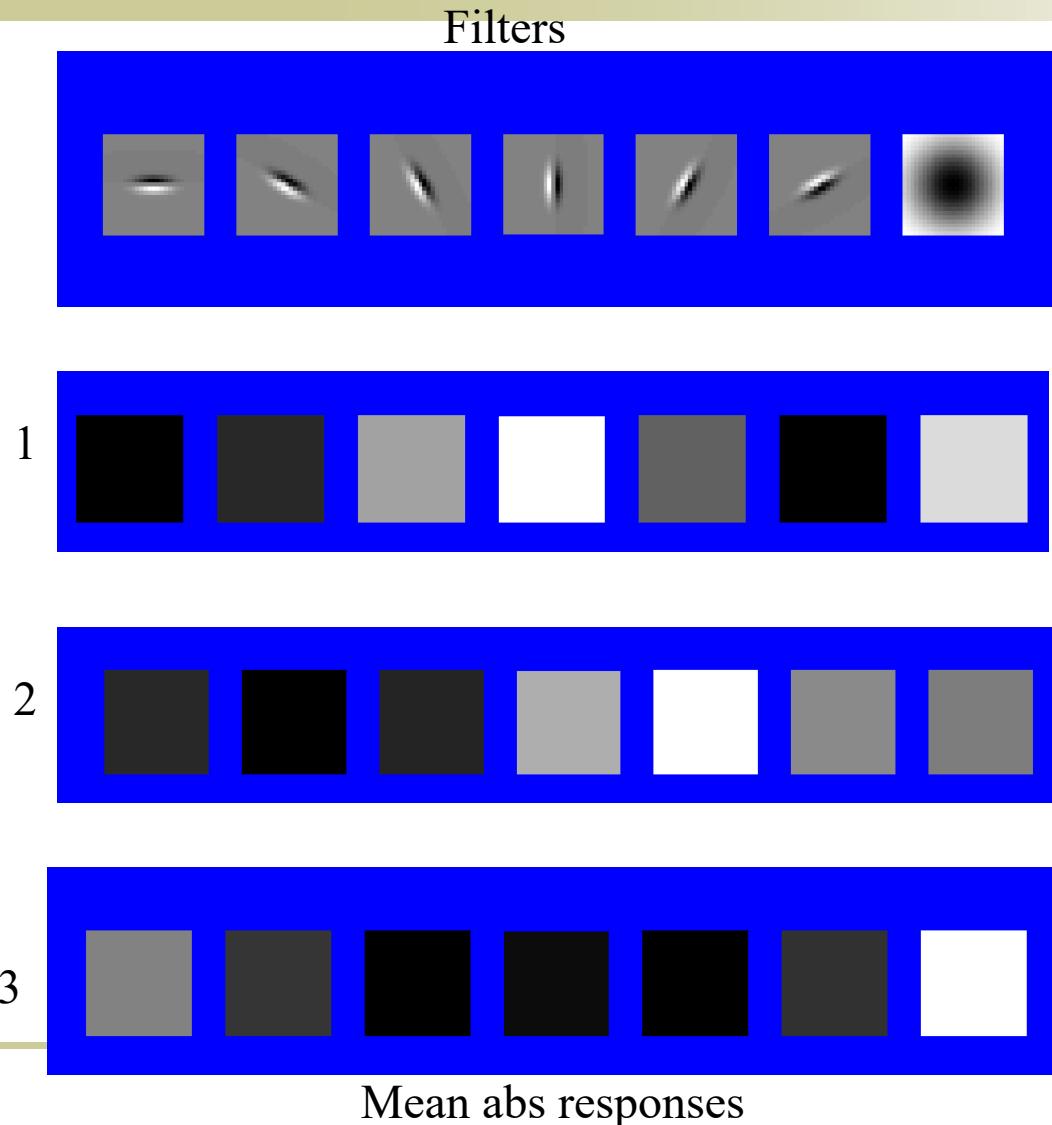
- What filters to put in the bank?
 - Typically we want a combination of scales and orientations, different types of patterns.

Matlab code available for these examples:

<http://www.robots.ox.ac.uk/~vgg/research/texclass/filters.html>

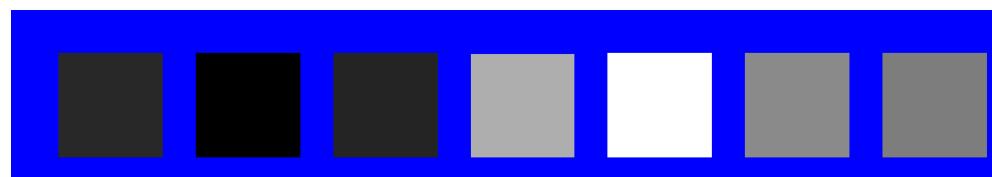
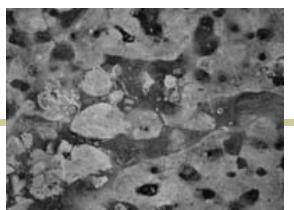
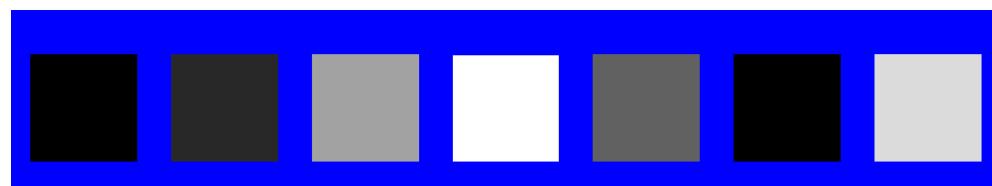
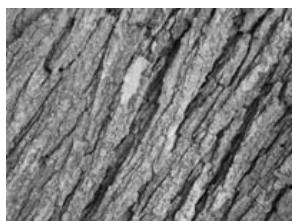
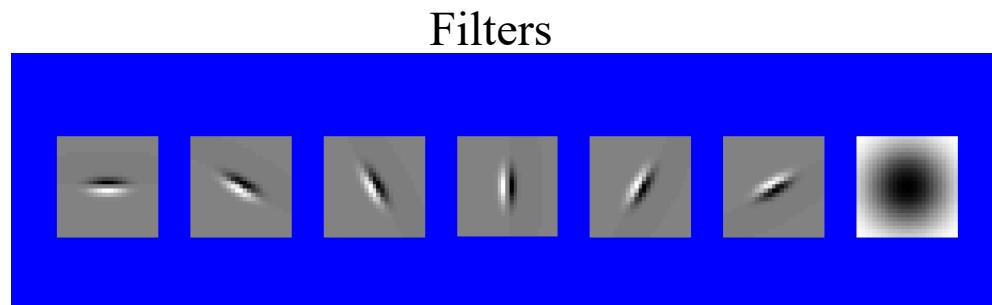


You try: Can you match the texture to the response?



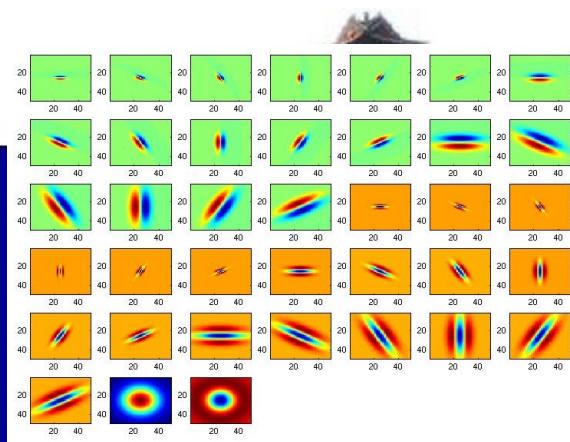
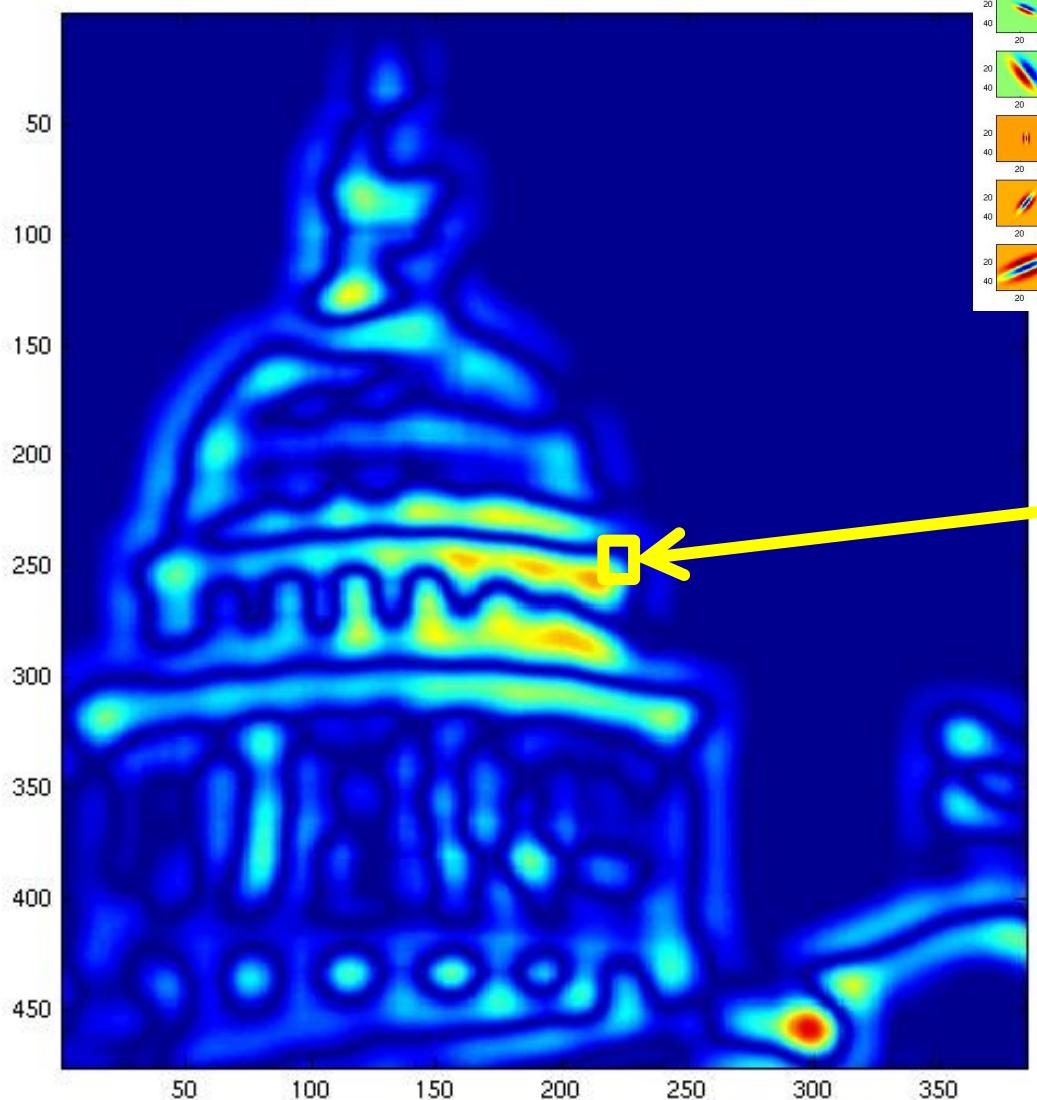


Representing texture by mean abs response



Mean abs responses

Derek Hoiem



$[r_1, r_2, \dots, r_{38}]$

We can form a feature vector from the list of responses at each pixel.

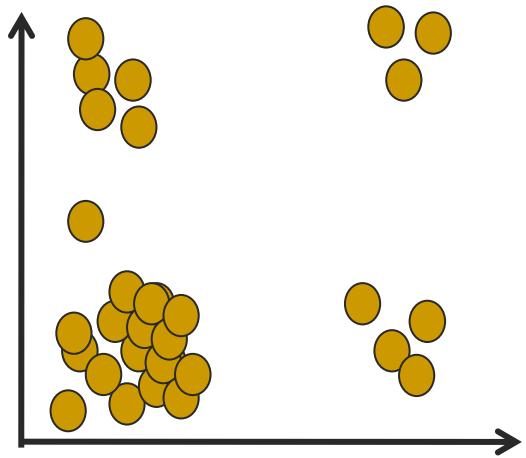


d -dimensional features



$$D(a, b) = \sqrt{\sum_{i=1}^d (a_i - b_i)^2}$$

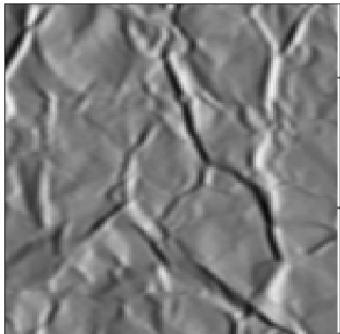
Euclidean distance (L_2)



2d



Texture Recognition: Local to Global



?



Felt?

Polyester?

Terrycloth?

Rough Plaster?

Leather?

Plaster?

Concrete?

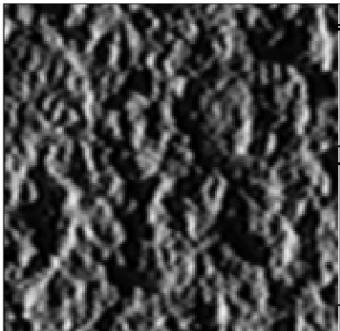
Crumpled Paper?

Sponge?

Limestone?

Brick?

⋮



?



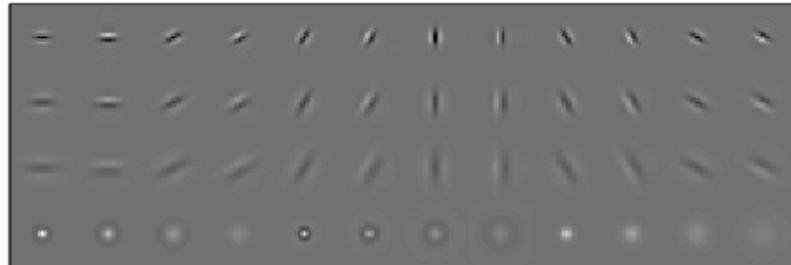


2D Textons



- Goal: find canonical local features in a texture;

- 1) Filter image with linear filters:

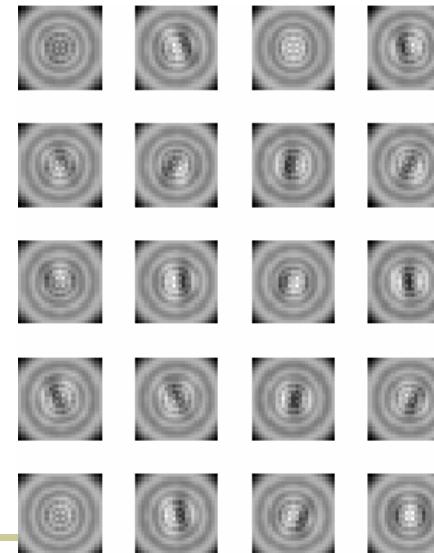
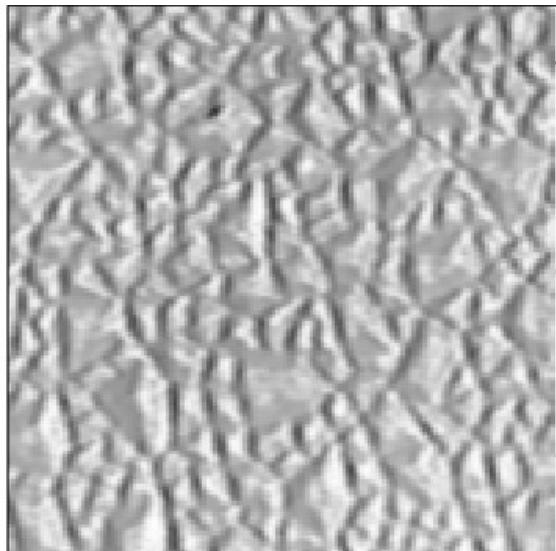
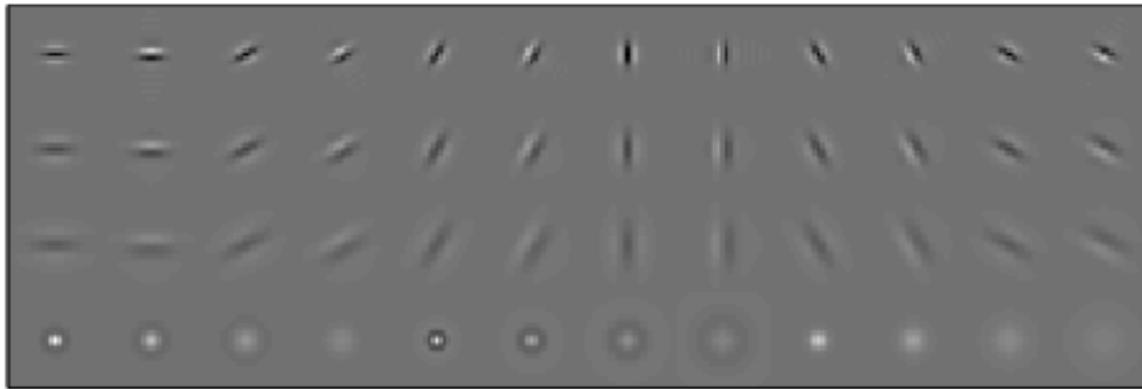


- 2) Vector quantization (k-means) on filter outputs;
- 3) Quantization centers are the textons.

- Spatial distribution of textons defines the texture;



2D Textons (cont'd)

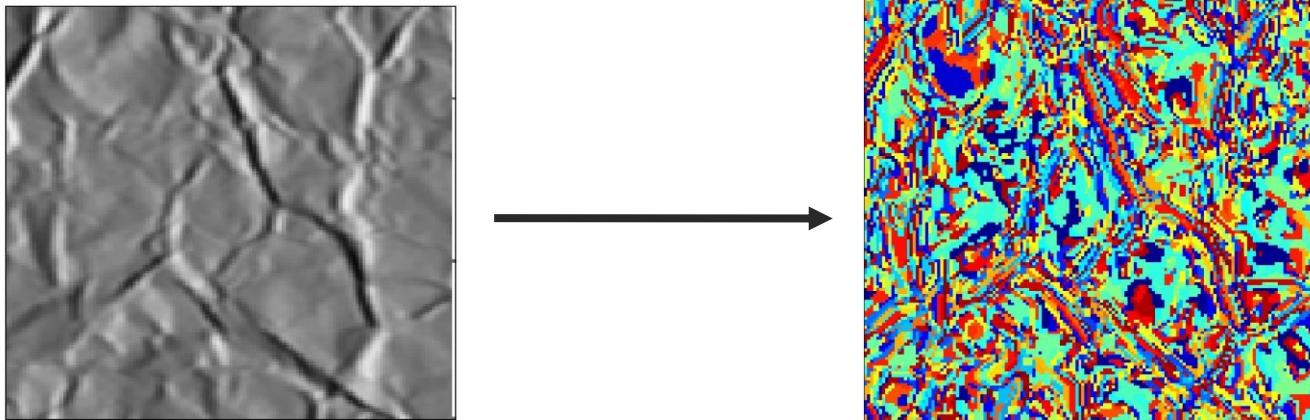




Texton Labeling



- Each pixel labeled to texton i (1 to K) which is *most similar in appearance*;
- Similarity measured by the Euclidean distance between the filter responses;





Material Representation



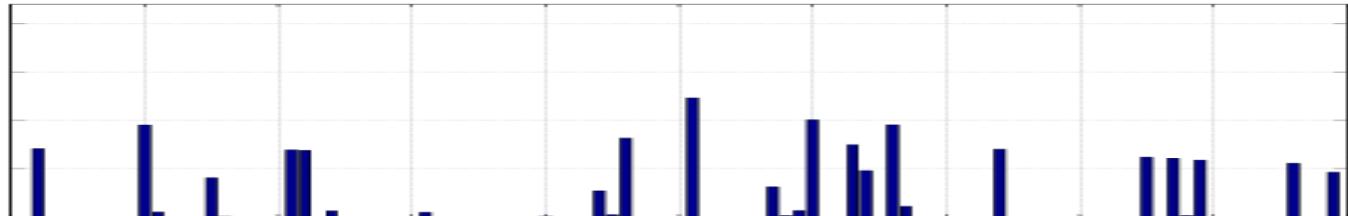
- Each material is now represented as a spatial arrangement of symbols from the texton vocabulary;
- Recognition: ignore spatial arrangement, use histogram ($K=100$);



Histogram Models for Recognition (Leung & Malik, 1999)



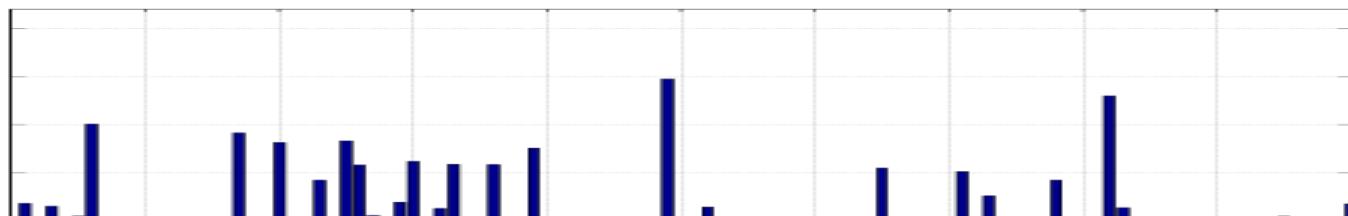
Rough Plastic



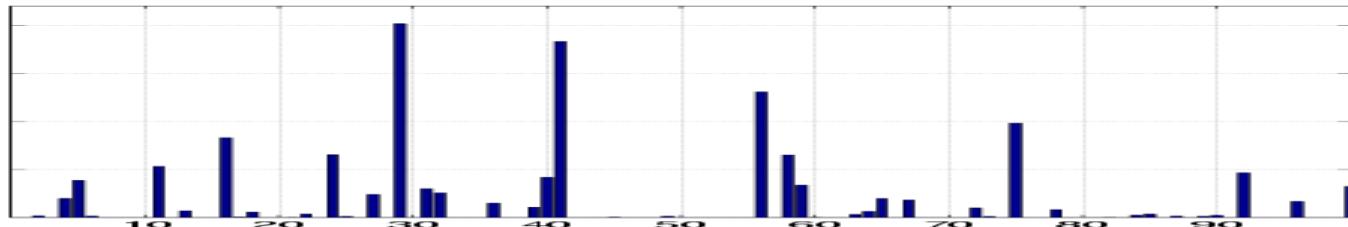
Pebbles



Plaster-b



Terrycloth





Similarity of materials



- Similarity between histograms measured using chi-square difference:

$$\chi^2(h_1, h_2) = \sum_{n=1}^N \frac{(h_1(n) - h_2(n))^2}{h_1(n) + h_2(n)}$$

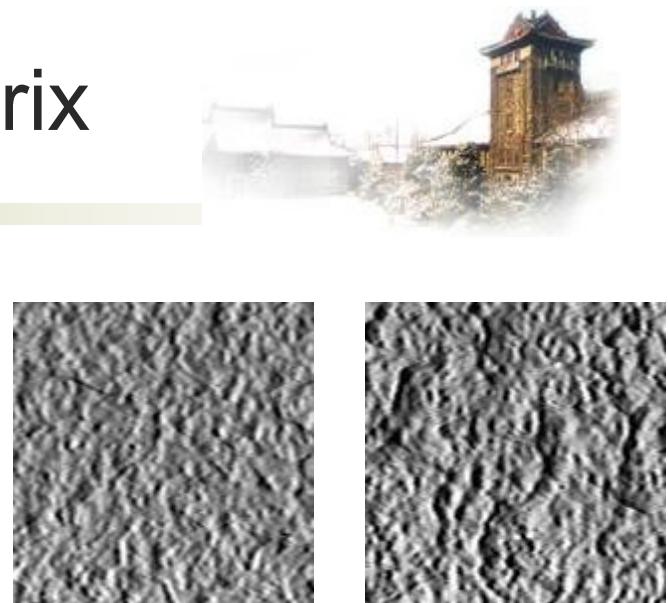


Similarity Matrix

Felt	0.6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Terrycloth	0.0	1.0	0.0	0.0	0.3	0.0	0.1	0.2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Rough Plastic	0.0	0.0	0.9	0.0	0.0	0.0	0.2	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Leather	0.2	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Sandpaper	0.0	0.1	0.0	0.0	1.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Pebbles	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1	0.0	0.0
Plastera	0.0	0.1	0.2	0.0	0.1	0.0	1.0	0.7	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Plasterb	0.0	0.2	0.1	0.0	0.0	0.0	0.8	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Rough Paper	0.0	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Artificial Grass	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.1	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Roof Shingle	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.2	1.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Aluminum Foil	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Cork	0.0	0.0	0.0	0.0	0.0	0.3	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.2	0.0	0.0	0.0	0.0	0.0
Rough Tile	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.9	0.0	0.0	0.0	0.0
	Felt	Terrycloth	Rough Plastic	Leather	Sandpaper	Pebbles	Plastera	Plasterb	Rough Paper	Artificial Grass	Rough Shingle	Aluminum Foil	Cork	Rough Tile					

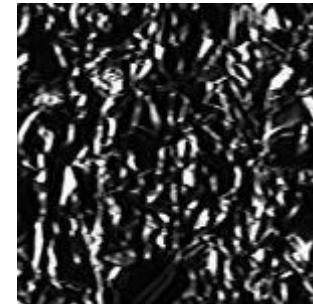
$$e_{ij} = \text{Similarity}(\text{material} = i, \text{sample} = j)$$

198

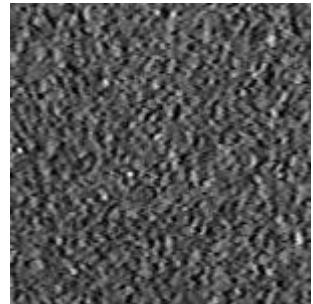


Plaster-a

Plaster-b



Aluminum
Foil



Cork



Example uses of texture in vision: analysis



Classifying materials, “stuff”



Figure by Varma &
Zisserman



(a)



1) 130066



2) 130070



3) 130068



4) 130051



5) 130038

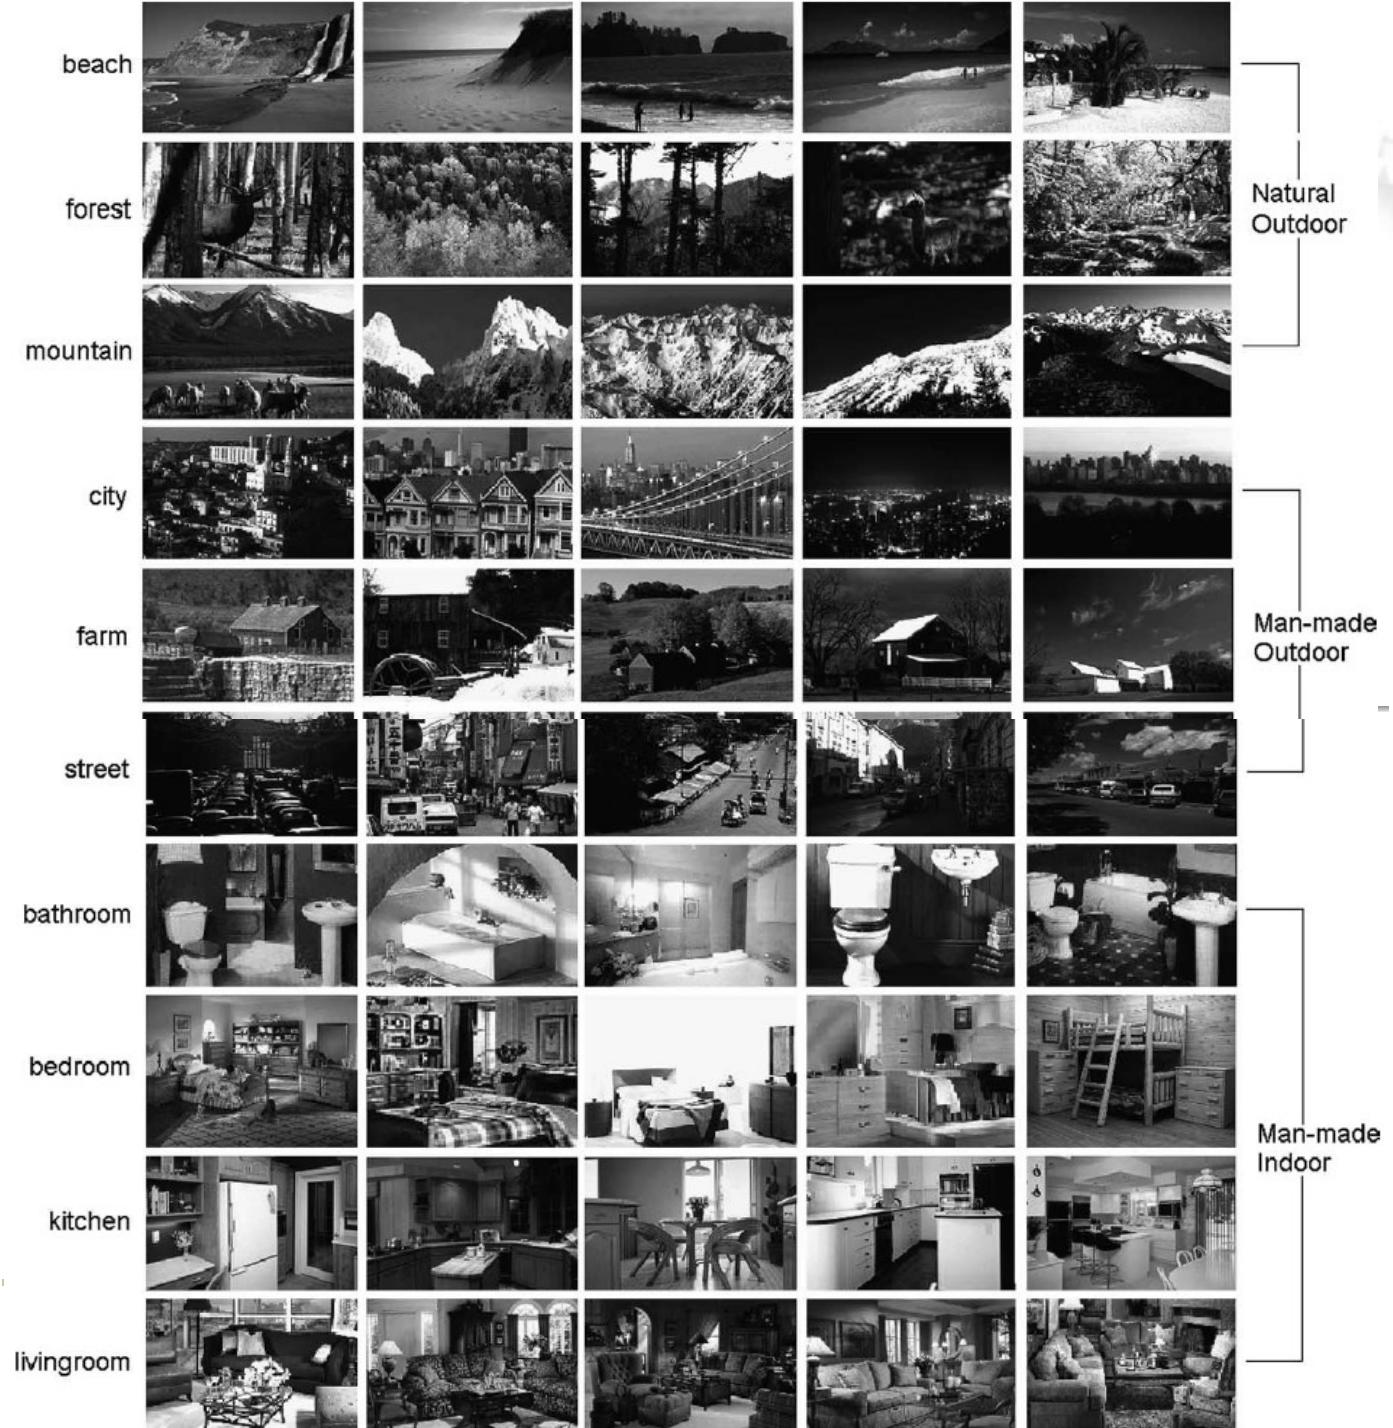


6) 130039

Texture features for image retrieval



Y. Rubner, C. Tomasi, and L. J. Guibas. The earth mover's distance as a metric for image retrieval. *International Journal of Computer Vision*, 40(2):99-121, November 2000,



Characterizing scene categories by texture

L. W. Renninger and J. Malik. When is scene identification just texture recognition? Vision Research 44 (2004) 2301–2311



Segmenting aerial imagery by textures



Texture-related tasks

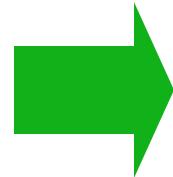


- **Shape from texture**
 - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
 - Analyze, represent texture
 - Group image regions with consistent texture
- **Synthesis**
 - Generate new texture patches/images given some examples



Texture synthesis

- Goal: create new samples of a given texture
- Many applications: virtual environments, hole-filling, texturing surfaces

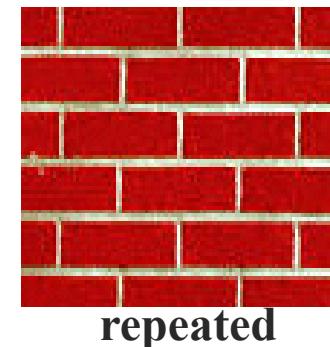




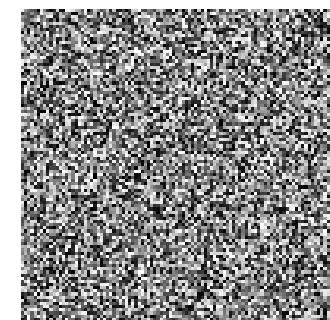
The Challenge

- Need to model the whole spectrum: from repeated to stochastic texture

Alexei A. Efros and Thomas K. Leung, “Texture Synthesis by Non-parametric Sampling,” Proc. International Conference on Computer Vision (ICCV), 1999.



repeated



stochastic



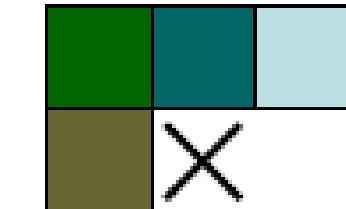
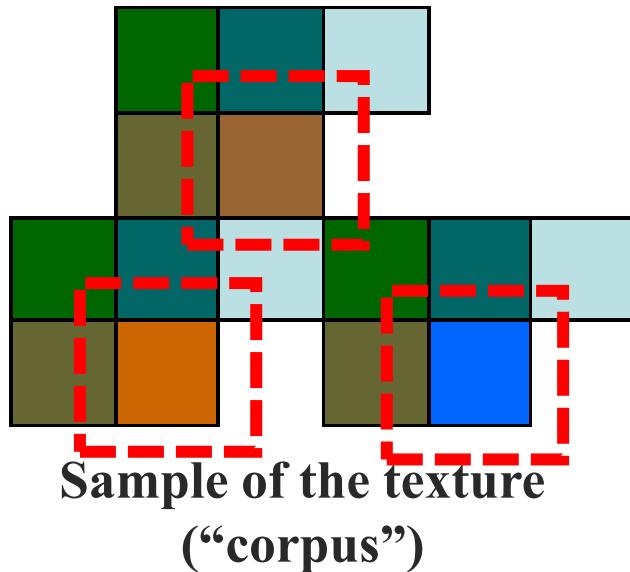
Both?



Texture synthesis: intuition

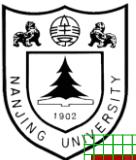


Now we want to insert **pixel intensities** based on existing nearby pixel values.

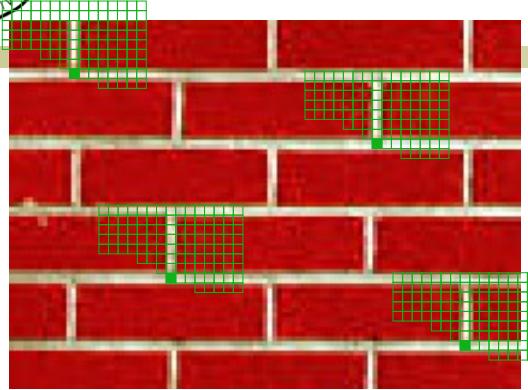


Place we want to
insert next

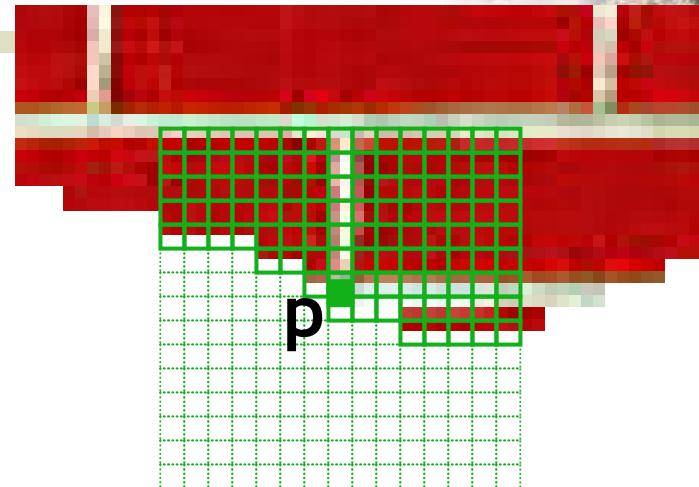
Distribution of a value of a pixel is conditioned on its neighbors alone.



Synthesizing One Pixel



input image

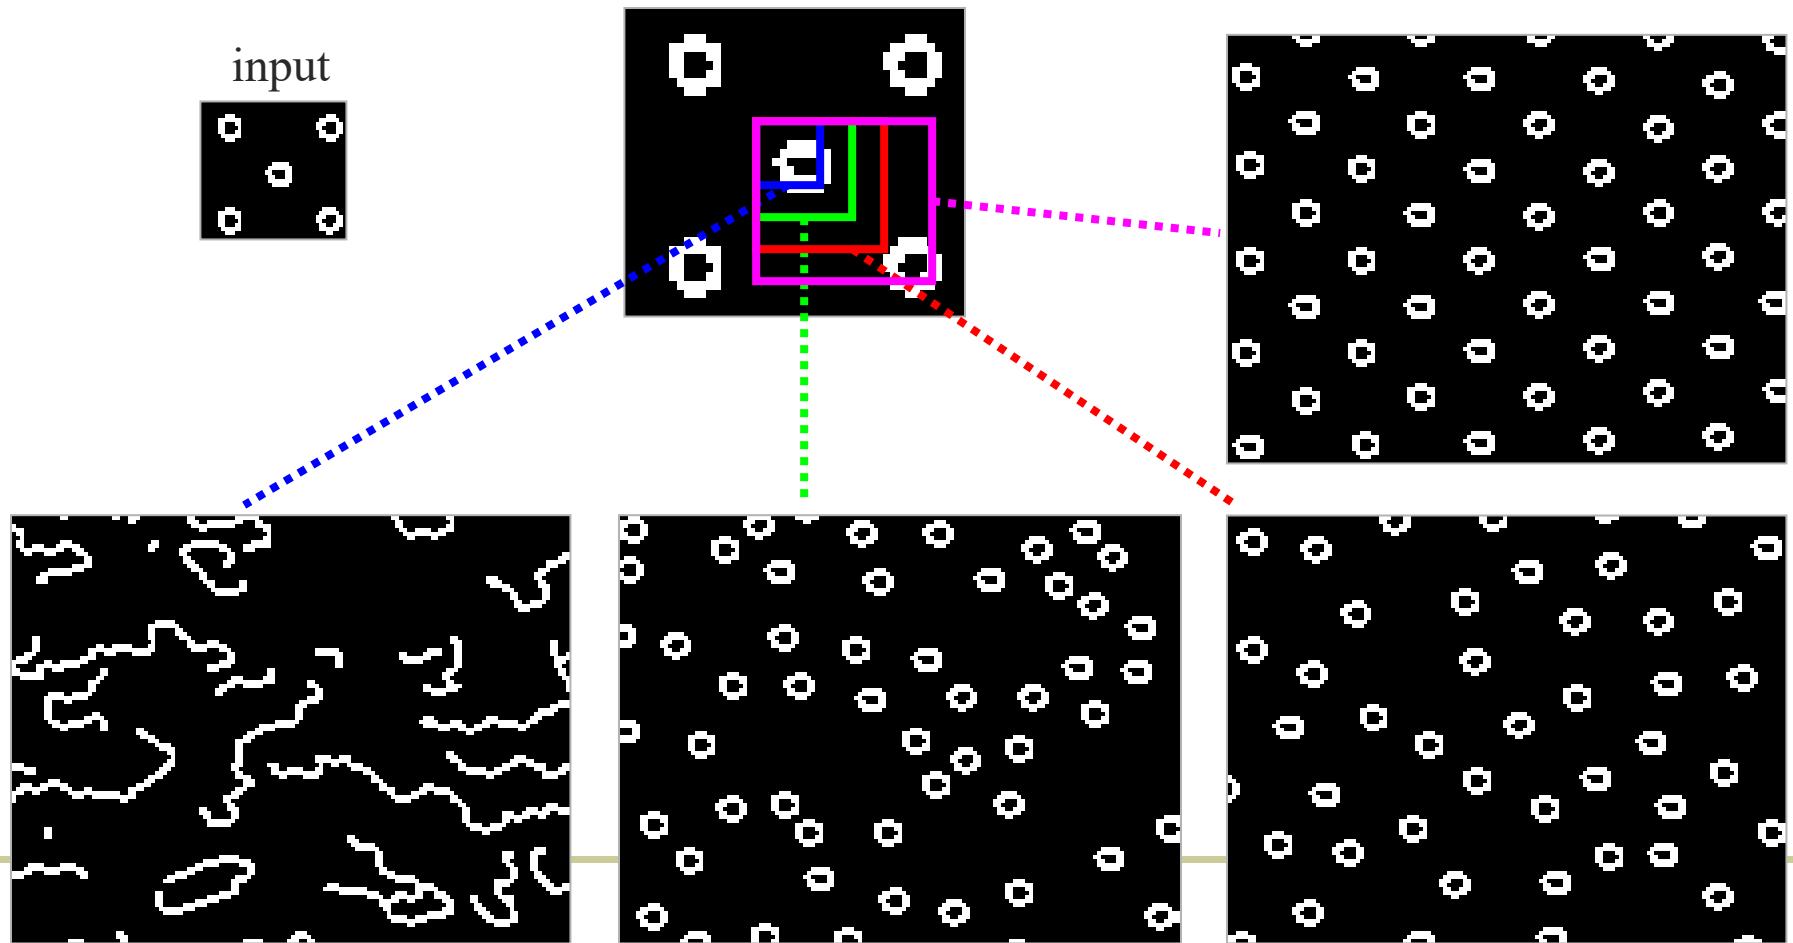


synthesized image

- What is $P(x|\text{neighborhood of pixels around } x)$?
- Find all the windows in the image that match the neighborhood
- To synthesize x
 - pick one matching window at random
 - assign x to be the center pixel of that window

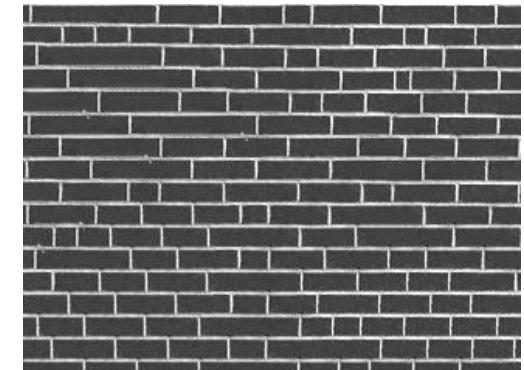
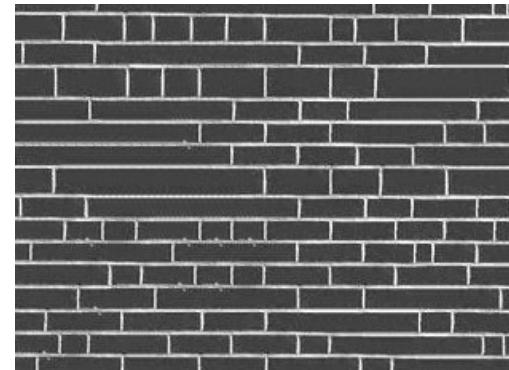
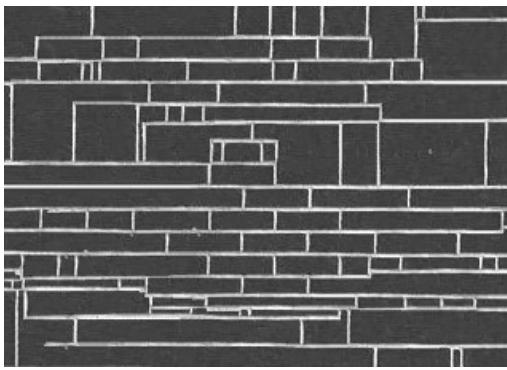
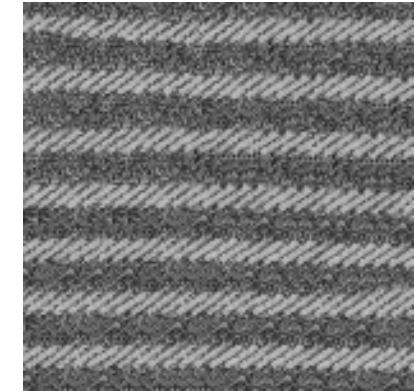
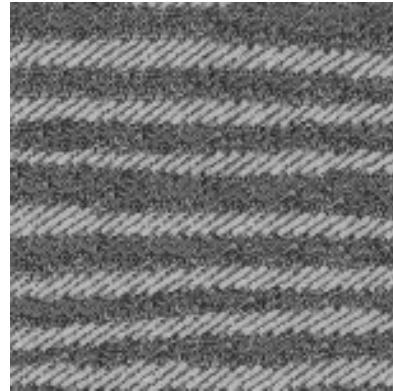
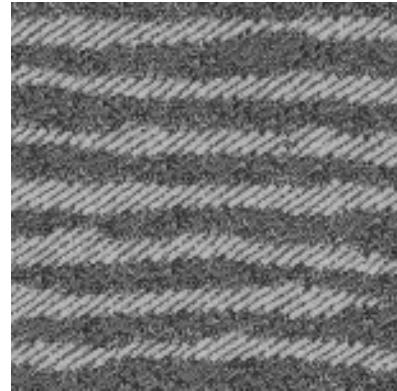
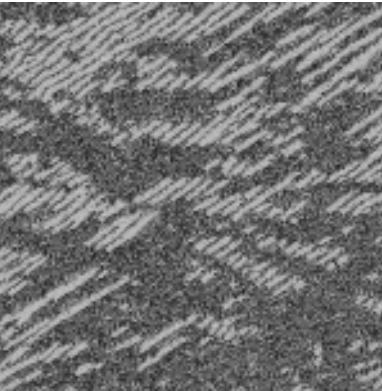


Neighborhood Window





Varying Window Size



Increasing window size

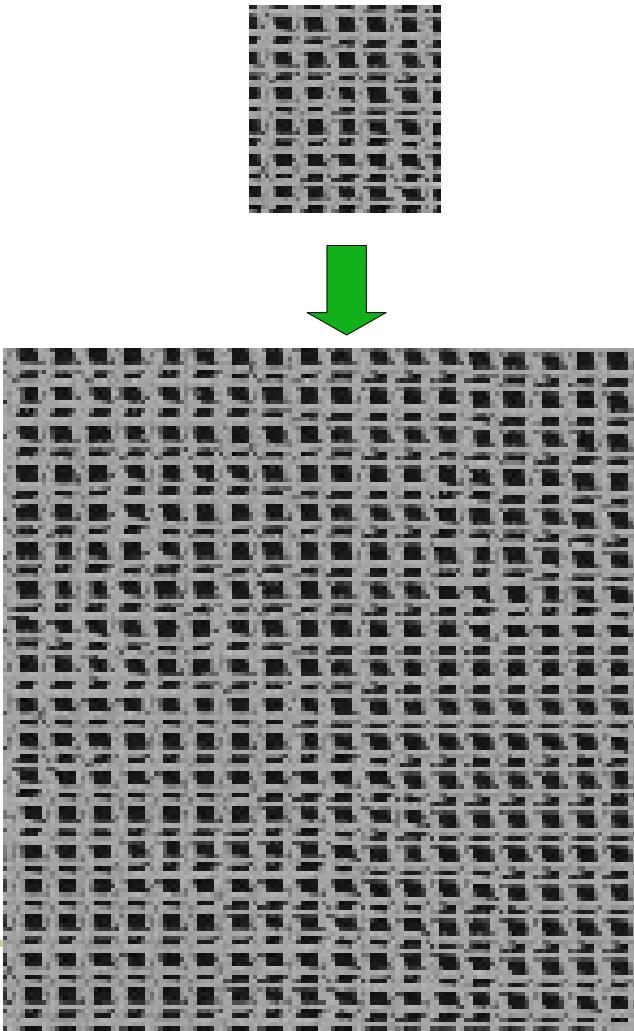




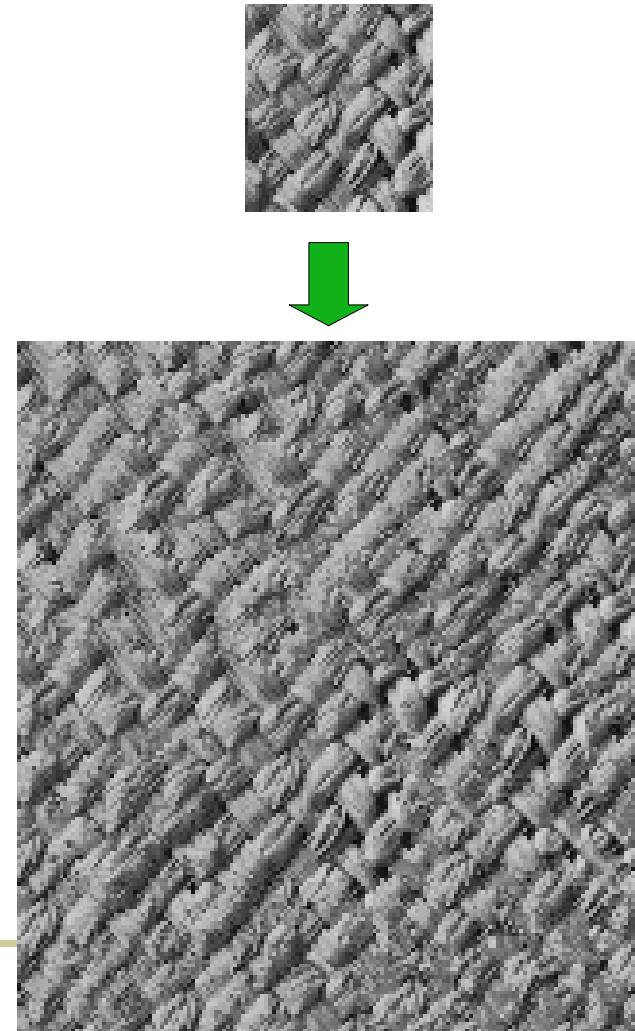
Synthesis results



french canvas



rafia weave



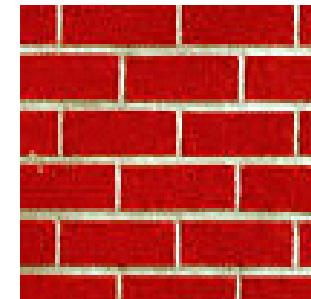
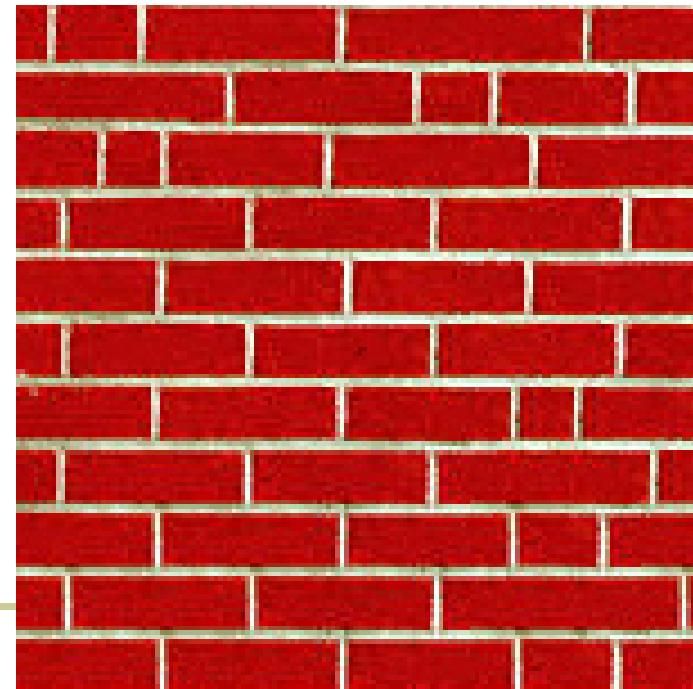


Synthesis results

white bread



brick wall





Synthesis results

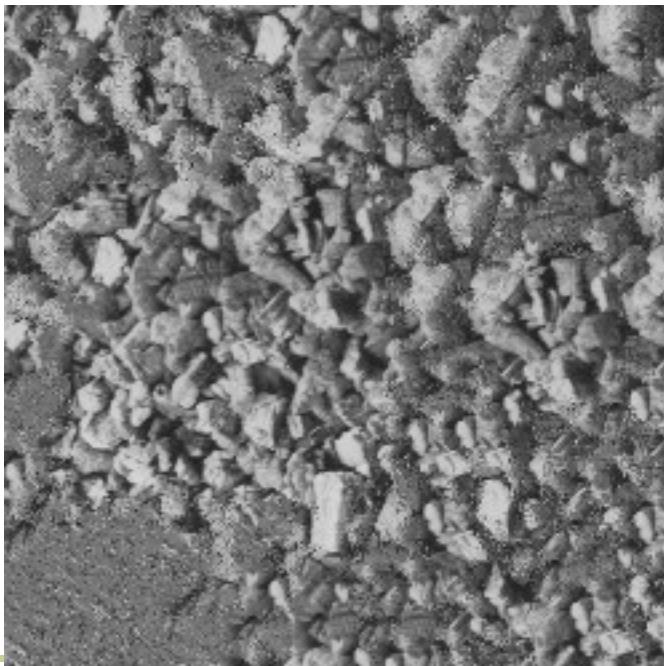


uring in the unsensato
r Dick Gephardt was fai
rful riff on the looming
nly asked, "What's your
tions?" A heartfelt sigh
story about the emergenc
es against Clinton. "Boy
g people about continuin
ardt began, patiently obs
, that the legal system h
e with this latest tanger

ithairm, them . "Whnephartfe lartifelintomimen
el ck Clirtioout omaim thatfelins. f ait s anetc
the ry onst wartfe lck Gephntoomimeationl sigab
Chiooufit Clinut Cll riff on, hat's yordn, parut tly
ons ycontonsteht wasked, painf sahe loo' riff on l
nskoneploourtfeas leil A nst Clit, "Wleontongal s
k Cirtioouirtfepe óng pme abegal fartfenstemem
tiensteneltorydt telemeaphinsverdt was agemer
ff ons artientont Cling peme asurtfe atish, "Boui s
hal s fartfelt sig pedr tl'dt ske abounutie aboutioo
itfaonewwas yous aboonthardt thatins fain, ped,
ains, them, pabout wasy arfiut courtly d, ln A h
ole emthrdingbooreme agas fa bontinsyst Clinut
ory about continst Clipeouinst Cloke agatiff out C
stome zinemen tly ardt beoraboul n, thenly as t C
cons faimeme Diontont wat coutlyohgans as fan
ien, phrtfaul, "Wbaut cout congagal còmingan
mifmst Cliiy abon al coountha ermungairt tf our
The looorystan loontieph. Intly on, theoplegatick
aul tatteeontly atie Diontiomt wal s f thegæ ener
mthahsat's enenhhmas fan. "intchthorw ahons v



Failure Cases



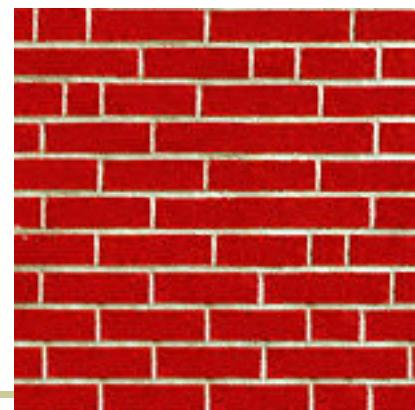
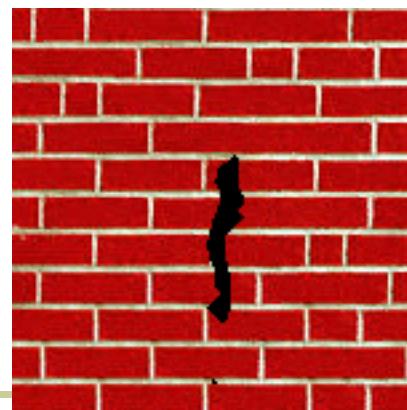
Growing garbage



Verbatim copying

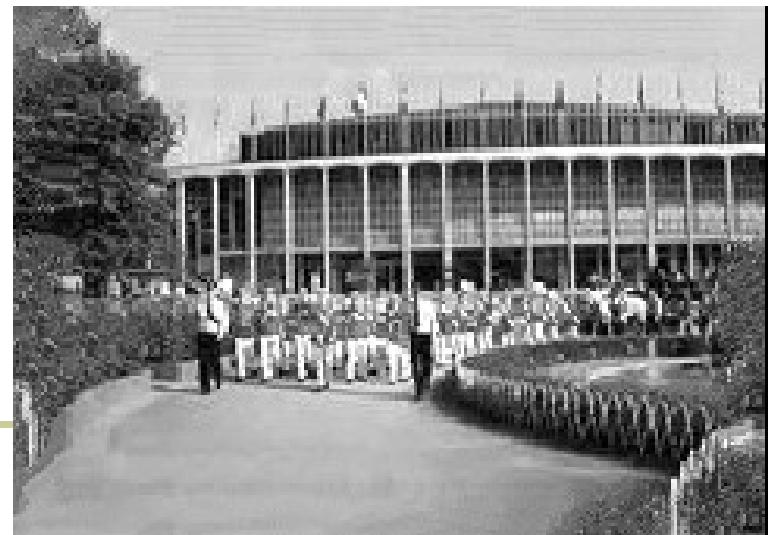
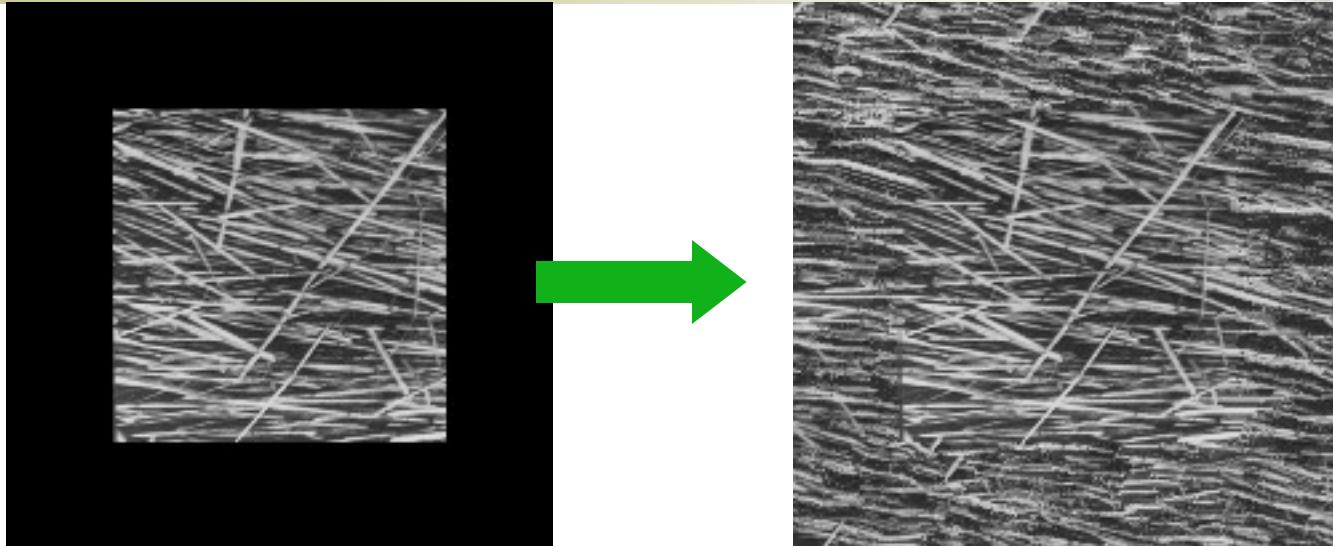


Hole Filling





Extrapolation

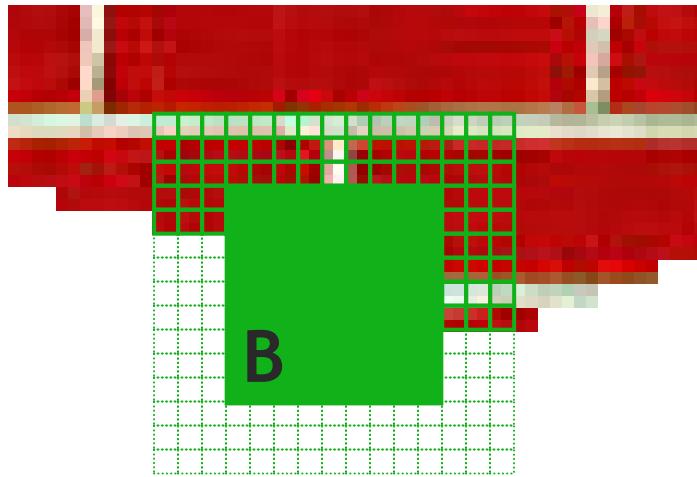




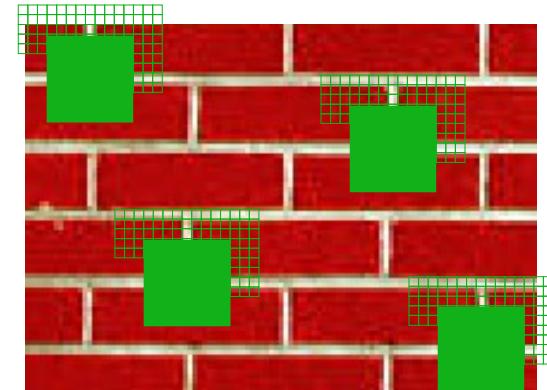
- The Efros & Leung algorithm
 - Simple
 - Surprisingly good results
 - Synthesis is easier than analysis!
 - ...but very slow



Image Quilting [Efros & Freeman 2001]



non-parametric
sampling



Input image

Synthesizing a block

- Observation: neighbor pixels are highly correlated
- Idea: unit of synthesis = block

- Exactly the same but now we want $P(B|N(B))$
- Much faster: synthesize all pixels in a block at once



Summary



- Template, Pyramid, and Texture
 - Template matching (SSD or Normxcorr2)
 - SSD can be done with linear filters, is sensitive to overall intensity
 - Gaussian pyramid
 - Coarse-to-fine search, multi-scale detection
 - Laplacian pyramid
 - More compact image representation
 - Can be used for compositing in graphics
 - Steerable pyramid
 - Filter banks for representing texture