



计算机视觉表征与识别

Chapter 9: Alignment and Transformation

王利民

媒体计算课题组

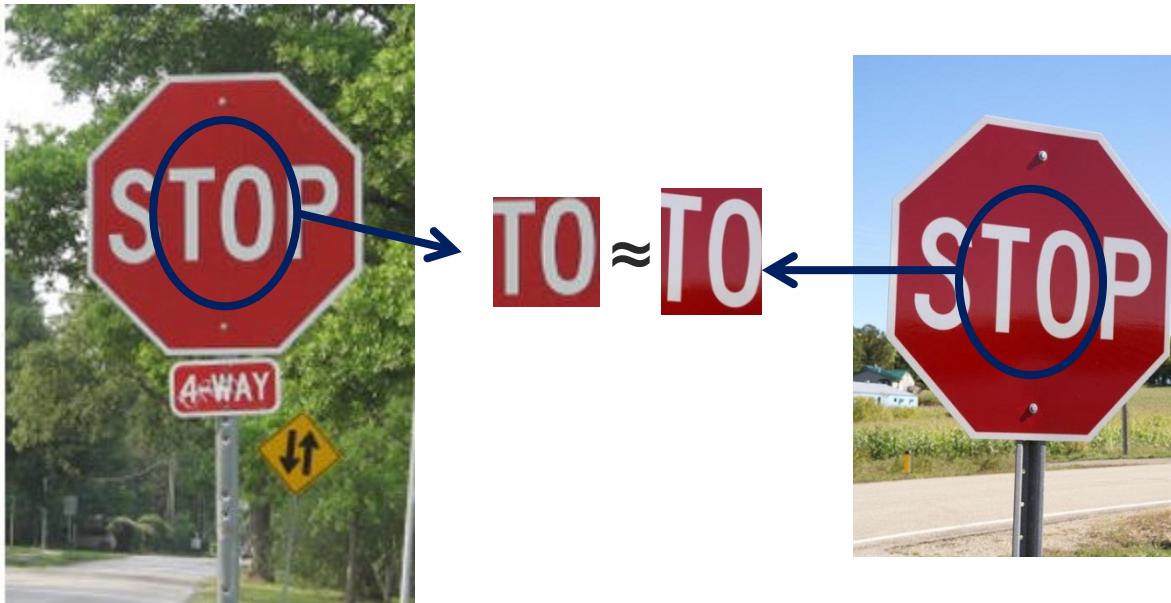
<http://mcg.nju.edu.cn/>



Correspondence and alignment

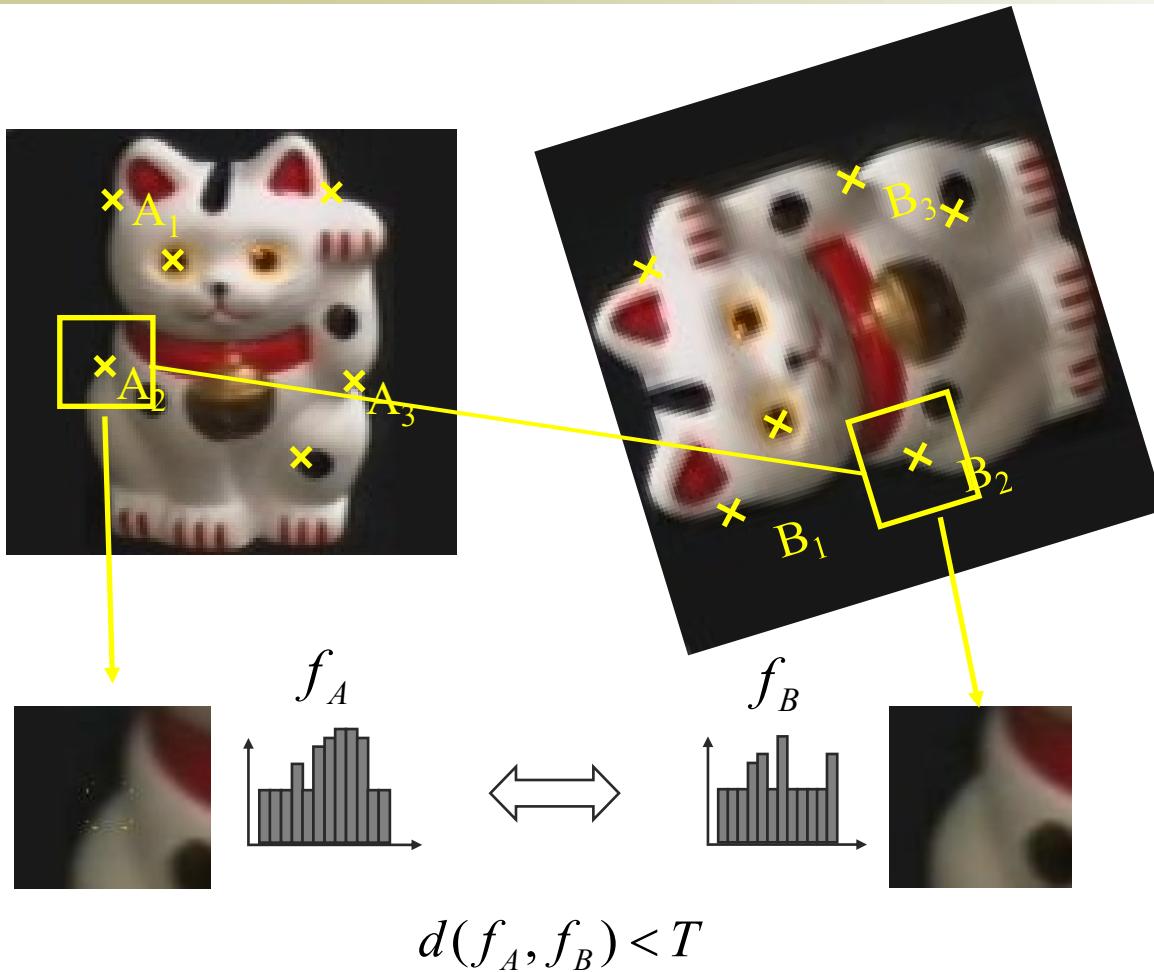


Correspondence: matching points, patches, edges, or regions across images





Recap: Keypoint Matching

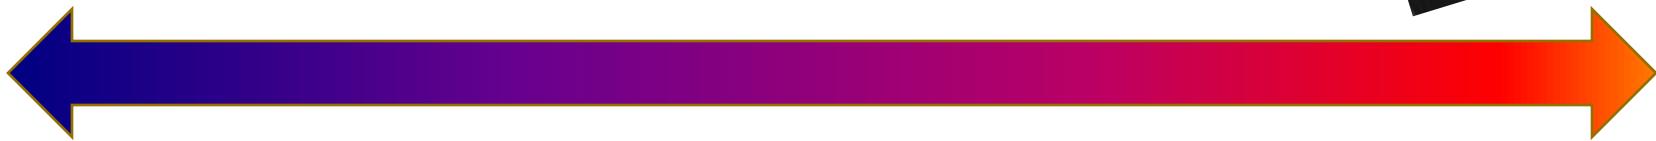


1. Find a set of distinctive keypoints
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor from the normalized region
5. Match local descriptors



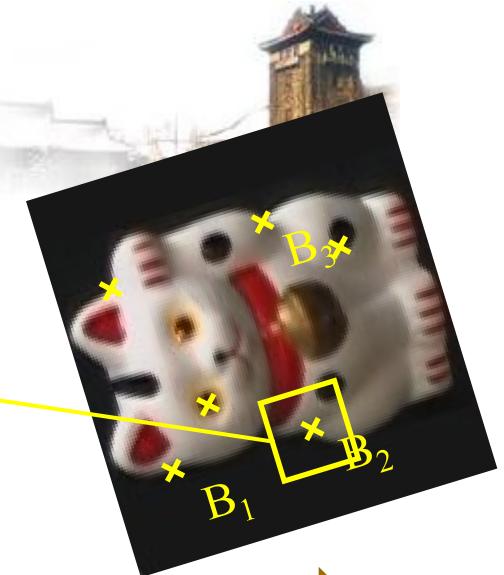
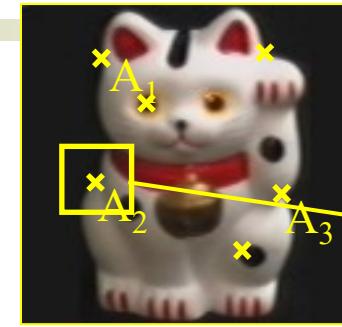
Recap: Key trade-offs

Detection



More Repeatable

Robust detection
Precise localization



More Points

Robust to occlusion
Works with less texture

Description



More Distinctive

Minimize wrong matches

More Flexible

Robust to expected variations
Maximize correct matches



Summary: Scale Invariant Detection

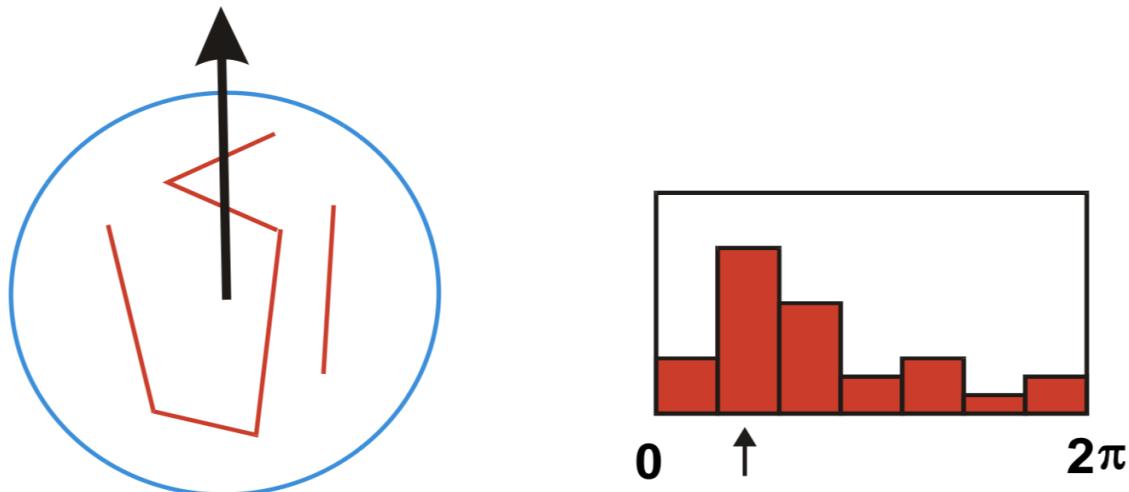
- **Given:** Two images of the same scene with a large *scale difference* between them.
- **Goal:** Find *the same* interest points *independently* in each image.
- **Solution:** Search for *maxima* of suitable functions in *scale* and in *space* (over the image).
- Two strategies
 - Laplacian-of-Gaussian (LoG)
 - Difference-of-Gaussian (DoG) as a fast approximation
 - *These can be used either on their own, or in combinations with single-scale keypoint detectors (Harris, Hessian).*



Orientation Normalization



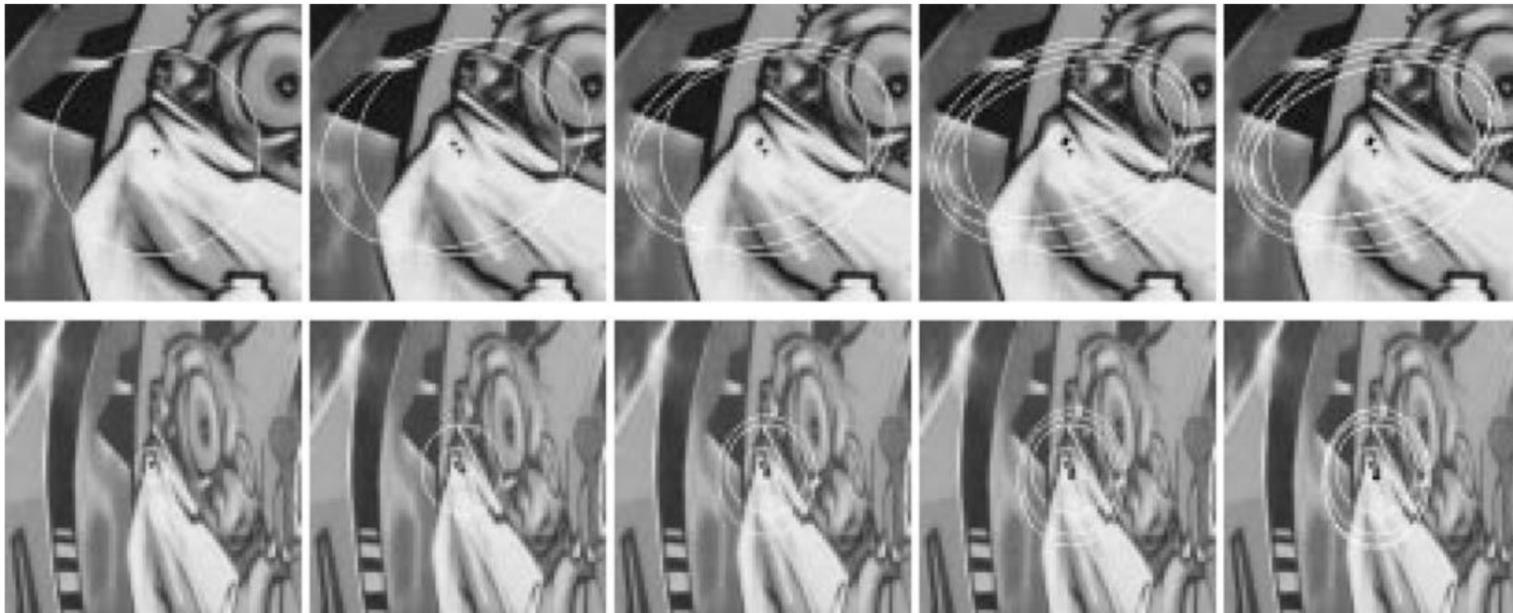
- Compute orientation histogram [Lowe, SIFT, 1999]
- Select dominant orientation
- Normalize: rotate to fixed orientation



Slide adapted from David Lowe



Iterative Adaption



1. Detect keypoints, e.g. multi-scale Harris
2. Automatically select the scales
3. Adapt affine shape based on second order moment matrix
4. Refine point location

K. Mikolajczyk and C. Schmid, [Scale and affine invariant interest point detectors](#), 56
IJCV 60(1):63-86, 2004.
Slide credit: Tinne Tuytelaars



Summary: Affine Invariance



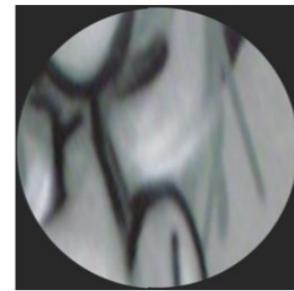
Extract affine regions



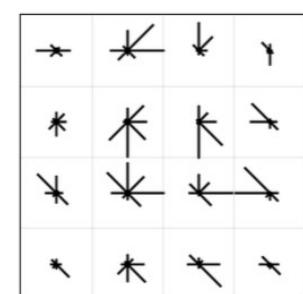
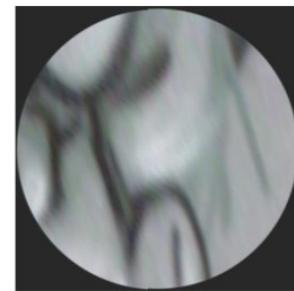
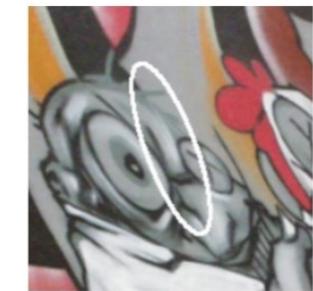
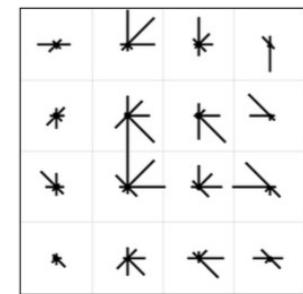
Normalize regions



Eliminate rotational ambiguity



Compare descriptors

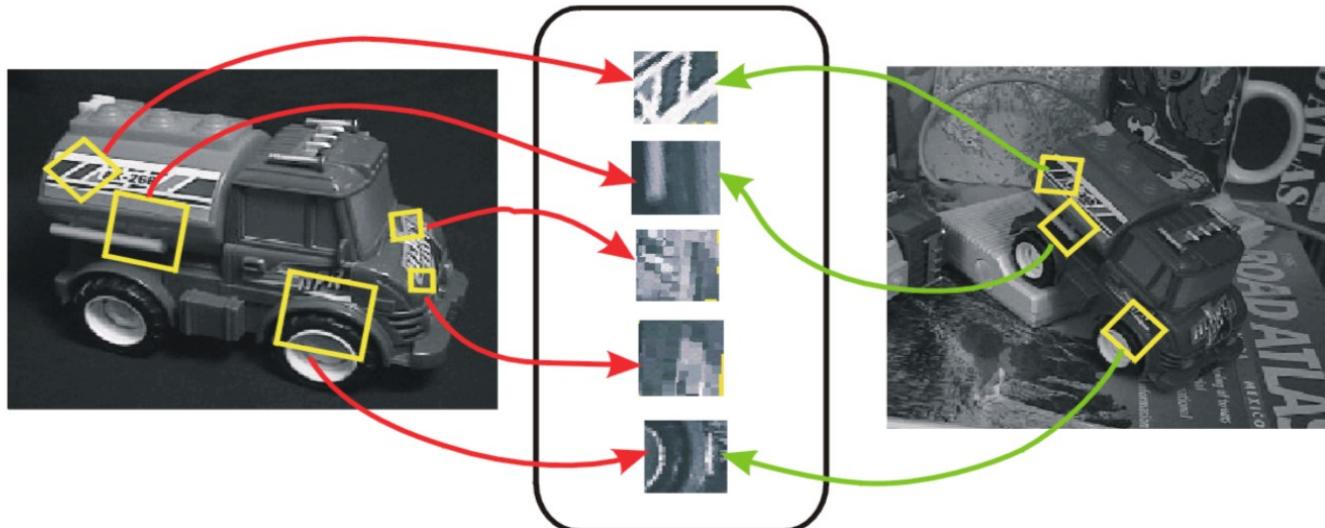




Invariance vs. Covariance



- **Invariance:**
 - $\text{features}(\text{transform}(\text{image})) = \text{features}(\text{image})$
- **Covariance:**
 - $\text{features}(\text{transform}(\text{image})) = \text{transform}(\text{features}(\text{image}))$



Covariant detection \Rightarrow invariant description



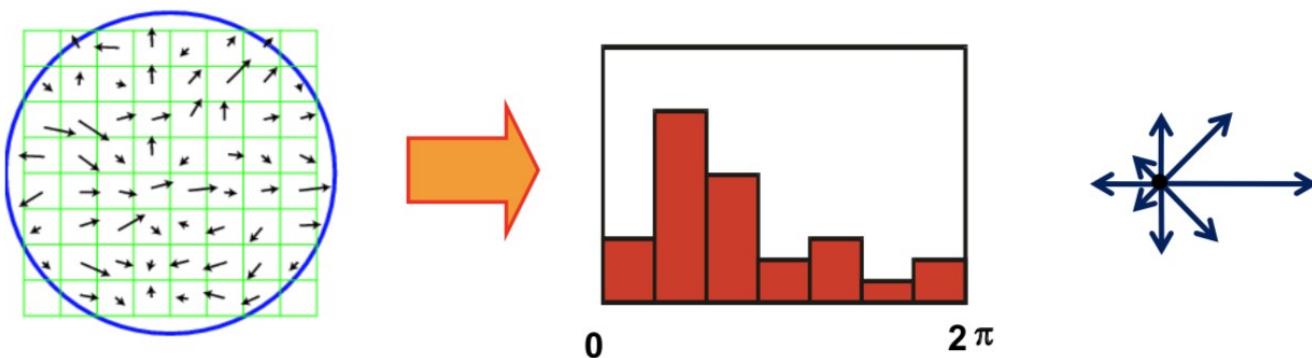
Feature Descriptor



- Disadvantage of patches as descriptors:
 - Small shifts can affect matching score a lot



- Solution: histograms

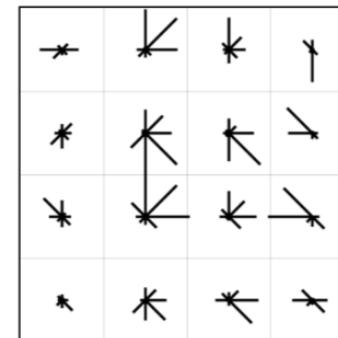
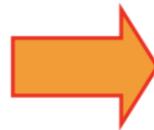
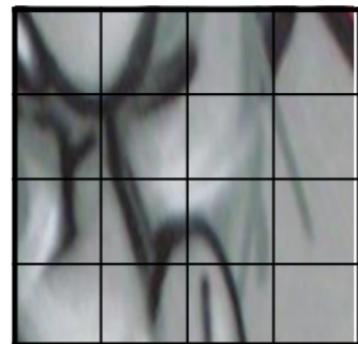




Feature Descriptor: SIFT



- **Scale Invariant Feature Transform**
- **Descriptor computation:**
 - Divide patch into 4×4 sub-patches: 16 cells
 - Compute histogram of gradient orientations (8 reference angles) for all pixels inside each sub-patch
 - Resulting descriptor: $4 \times 4 \times 8 = 128$ dimensions



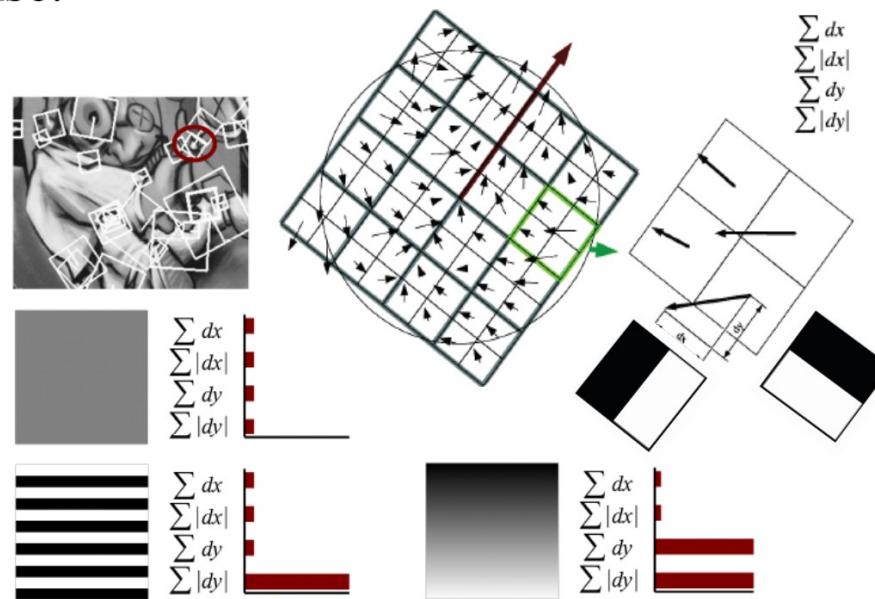
David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#)
IJCV 60 (2), pp. 91-110, 2004.



SURF: Descriptor Extraction

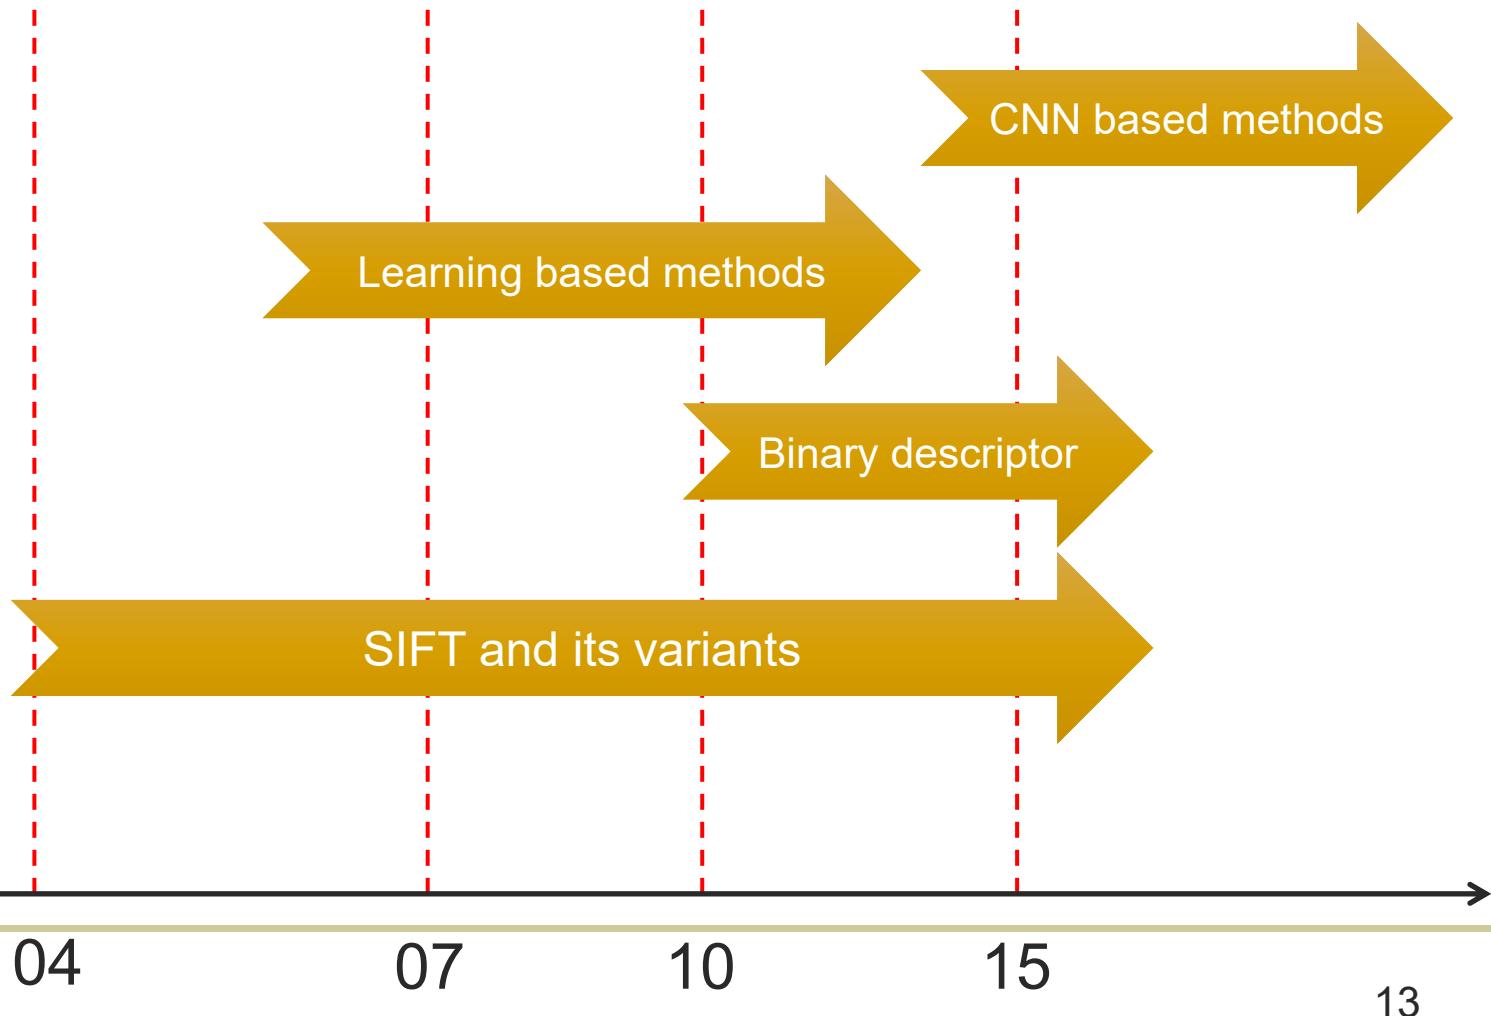


1. Split the interest region ($20s \times 20s$) into 4×4 square sub-regions.
2. Calculate Haar wavelet responses dx and dy , and weight the responses with a Gaussian kernel.
3. Sum the response over each sub-region for dx and dy , then sum the absolute value of response.
4. Concatenate summation results in all sub-regions, forming a 64D SURF descriptor.





Local Descriptors: Trend



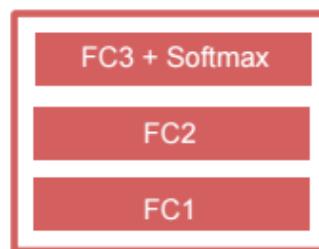


MatchNet

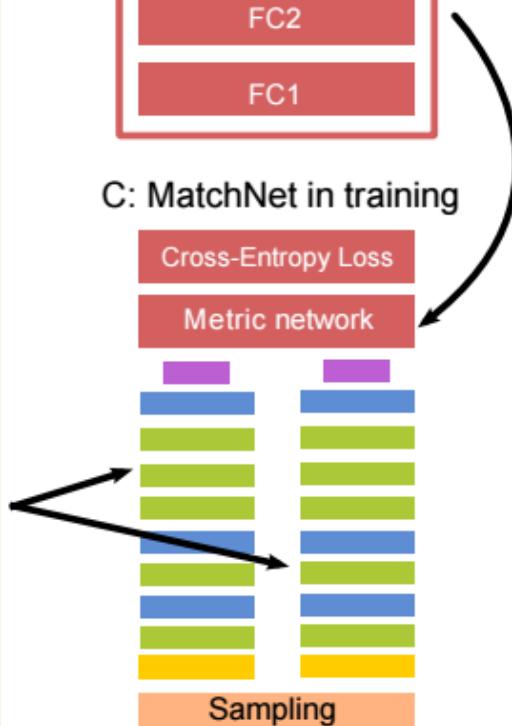
A: Feature network



B: Metric network



C: MatchNet in training



- Simultaneously learn the descriptor and the metric
- Siamese Feature descriptor network
- Metric network on top
- Cross-entropy loss, transfer matching problem to classification problem
- Train time: 1 day – 1 week



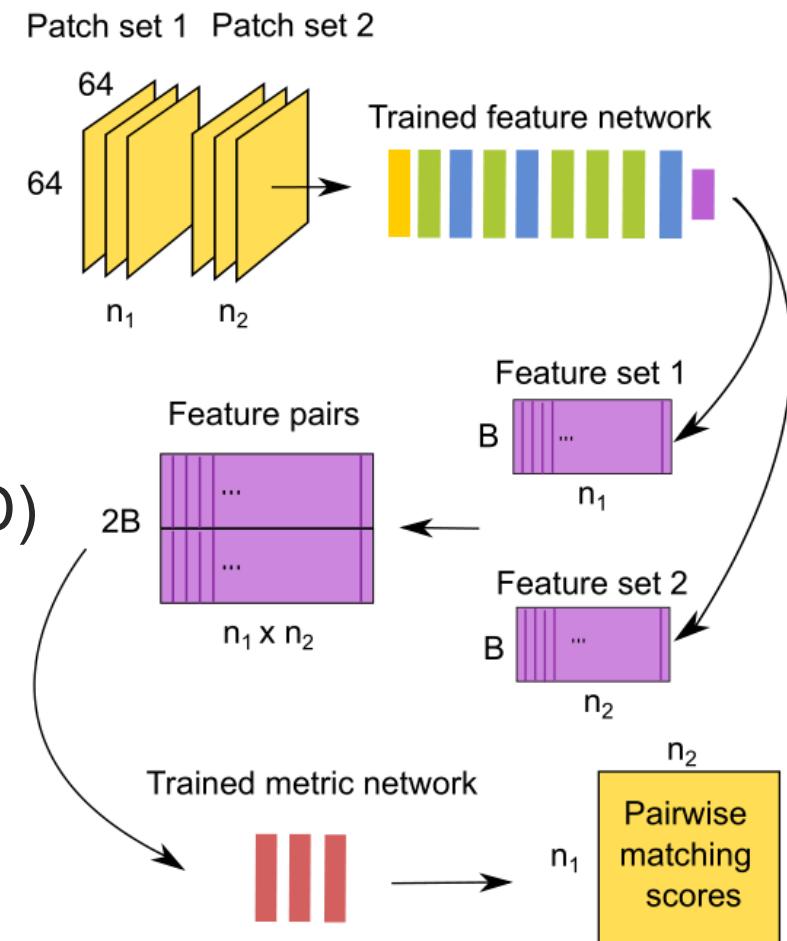
Training MatchNet

- Cross-entropy error

$$E = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

- Stochastic gradient descent (SGD)

- A special reservoir sampler for negative sampling



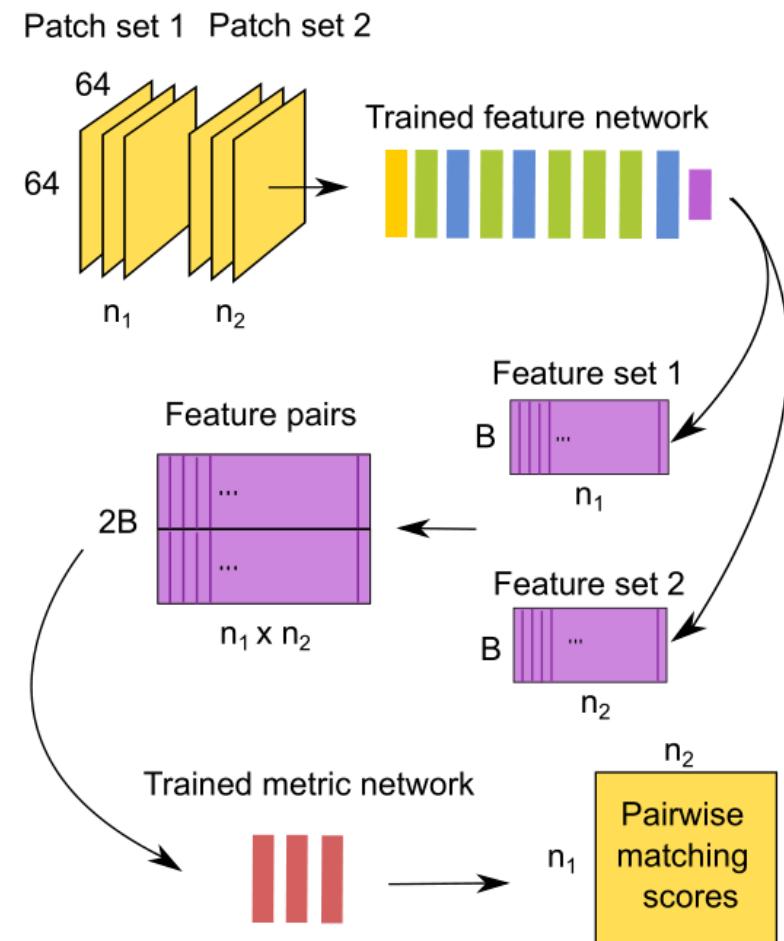


Testing MatchNet

A two-stage prediction pipeline:

1. Generate feature descriptors for all patches.

2. Pair the features and push them through the metric network to get the scores.





Today's class



- Introduction to alignment
- Alignment methods
 - Global methods
 - Hypothesize and test
- Image Transformation
 - Common image transformations
 - Examples of solving image alignment
- Homework: Mosaics



Alignment



Alignment: solving the transformation that makes two things match better

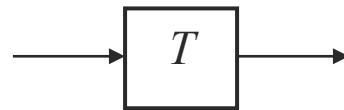




Parametric (global) warping



$$p = (x, y)$$



$$p' = (x', y')$$

Transformation T is a coordinate-changing machine:

$$p' = T(p)$$

What does it mean that T is global?

- Is the same for any point p
- can be described by just a few numbers (parameters)

For linear transformations, we can represent T as a matrix

$$p' = \mathbf{T}p$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \mathbf{T} \begin{bmatrix} x \\ y \end{bmatrix}$$



Fitting and Alignment



Fitting: find the parameters of a model that best fit the data.

Alignment: find the parameters of the transformation that best align matched points.



Fitting and Alignment



■ Design challenges

- Design a suitable **goodness of fit** measure
 - Similarity should reflect application goals
 - Encode robustness to outliers and noise
- Design an **optimization** method
 - Avoid local optima
 - Find best parameters quickly



Fitting and Alignment: Methods



- Global optimization / Search for parameters
 - Least squares fit
 - Robust least squares
 - Iterative closest point (ICP)

- Hypothesize and test
 - Generalized Hough transform
 - RANSAC



Today's class



- Introduction to alignment
- Alignment methods
 - **Global methods**
 - Hypothesize and test
- Image Transformation
 - Common image transformations
 - Examples of solving image alignment
- Homework: Mosaics

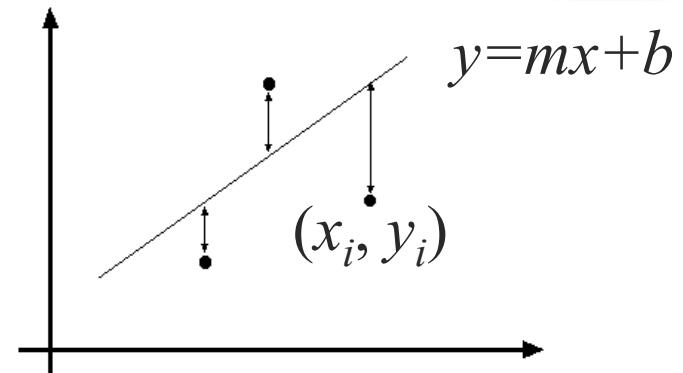


Least squares line fitting



- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = mx_i + b$
- Find (m, b) to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$



$$E = \sum_{i=1}^n \left(\begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - y_i \right)^2 = \left\| \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} - \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} \right\|^2 = \|\mathbf{Ap} - \mathbf{y}\|^2$$

$$= \mathbf{y}^T \mathbf{y} - 2(\mathbf{Ap})^T \mathbf{y} + (\mathbf{Ap})^T (\mathbf{Ap})$$

$$\frac{dE}{dp} = 2\mathbf{A}^T \mathbf{Ap} - 2\mathbf{A}^T \mathbf{y} = 0$$

$$\mathbf{A}^T \mathbf{Ap} = \mathbf{A}^T \mathbf{y} \Rightarrow \mathbf{p} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

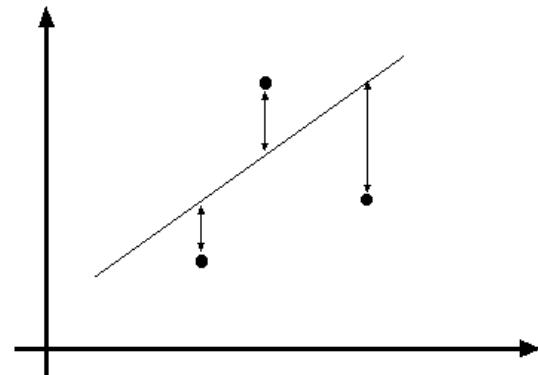
Matlab: $\mathbf{p} = \mathbf{A} \setminus \mathbf{y};$



Problem with “vertical” least squares



- Not rotation-invariant
- Fails completely for vertical lines



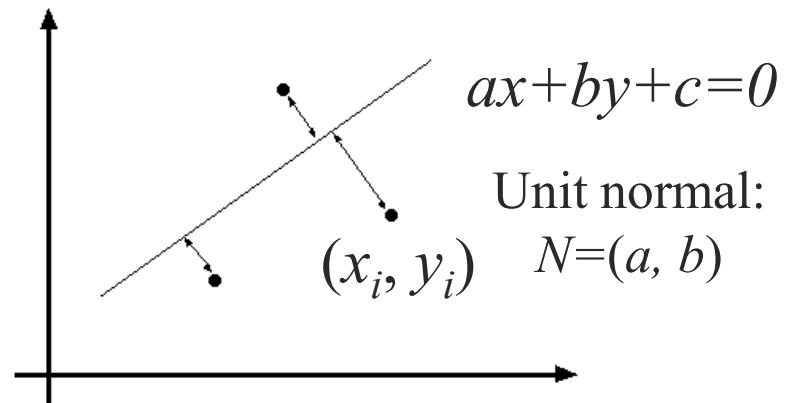


Total least squares



If $(a^2+b^2=1)$ then

Distance between point (x_i, y_i) and
line $ax+by+c=0$ is $|ax_i + by_i + c|$



proof:

<http://mathworld.wolfram.com/Point-LineDistance2-Dimensional.html>

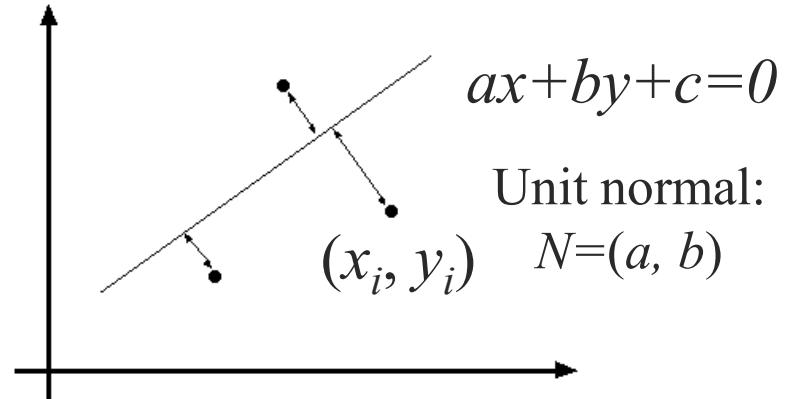


Total least squares



If $(a^2+b^2=1)$ then

Distance between point (x_i, y_i) and line $ax+by+c=0$ is $|ax_i + by_i + c|$



Find (a, b, c) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i + c)^2$$

Slide modified from S. Lazebnik



Total least squares



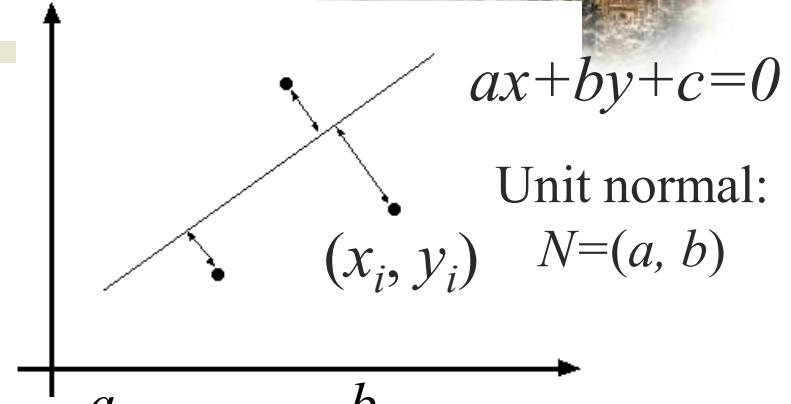
Find (a, b, c) to minimize the sum of squared perpendicular distances

$$E = \sum_{i=1}^n (ax_i + by_i + c)^2$$

$$\frac{\partial E}{\partial c} = \sum_{i=1}^n 2(ax_i + by_i + c) = 0$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = \mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p}$$

$$\text{minimize } \mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p} \quad \text{s.t. } \mathbf{p}^T \mathbf{p} = 1 \quad \Rightarrow \quad \text{minimize } \frac{\mathbf{p}^T \mathbf{A}^T \mathbf{A} \mathbf{p}}{\mathbf{p}^T \mathbf{p}}$$



$$ax + by + c = 0$$

Unit normal:
 $N = (a, b)$

$$c = -\frac{a}{n} \sum_{i=1}^n x_i - \frac{b}{n} \sum_{i=1}^n y_i = -a\bar{x} - b\bar{y}$$

Solution is eigenvector corresponding to smallest eigenvalue of $\mathbf{A}^T \mathbf{A}$

See details on Raleigh Quotient: http://en.wikipedia.org/wiki/Rayleigh_quotient



Recap: Two Common Optimization Problems



Problem statement

$$\text{minimize } \|\mathbf{Ax} - \mathbf{b}\|^2$$

least squares solution to $\mathbf{Ax} = \mathbf{b}$

Solution

$$\mathbf{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$$

$\mathbf{x} = \mathbf{A} \setminus \mathbf{b}$ (matlab)

Problem statement

$$\text{minimize } \mathbf{x}^T \mathbf{A}^T \mathbf{A} \mathbf{x} \quad \text{s.t. } \mathbf{x}^T \mathbf{x} = 1$$

$$\text{minimize } \frac{\mathbf{x}^T \mathbf{A}^T \mathbf{A} \mathbf{x}}{\mathbf{x}^T \mathbf{x}}$$

non-trivial lsq solution to $\mathbf{Ax} = 0$

Solution

$$[\mathbf{v}, \lambda] = \text{eig}(\mathbf{A}^T \mathbf{A})$$

$$\lambda_1 < \lambda_{2..n} : \mathbf{x} = \mathbf{v}_1$$



Least squares (global) optimization



Good

- Clearly specified objective
- Optimization is easy

Bad

- May not be what you want to optimize
- Sensitive to outliers
 - Bad matches, extra points
- Doesn't allow you to get multiple good fits
 - Detecting multiple objects, lines, etc.



Other ways to search for parameters



■ Line search

1. For each parameter, step through values and choose value that gives best fit
2. Repeat (1) until no parameter changes

■ Grid search

1. Propose several sets of parameters, evenly sampled in the joint set
2. Choose best (or top few) and sample joint parameters around the current best; repeat

■ Gradient descent

1. Provide initial position (e.g., random)
2. Locally search for better parameters by following gradient



Today's class



- Introduction to alignment
- Alignment methods
 - Global methods
 - **Hypothesize and test**
- Image Transformation
 - Common image transformations
 - Examples of solving image alignment
- Homework: Mosaics



Hypothesize and test



1. Propose parameters
 - Try all possible
 - Each point votes for all consistent parameters
 - Repeatedly sample enough points to solve for parameters
2. Score the given parameters
 - Number of consistent points, possibly weighted by distance
3. Choose from among the set of parameters
 - Global or local maximum of scores
4. Possibly refine parameters using inliers



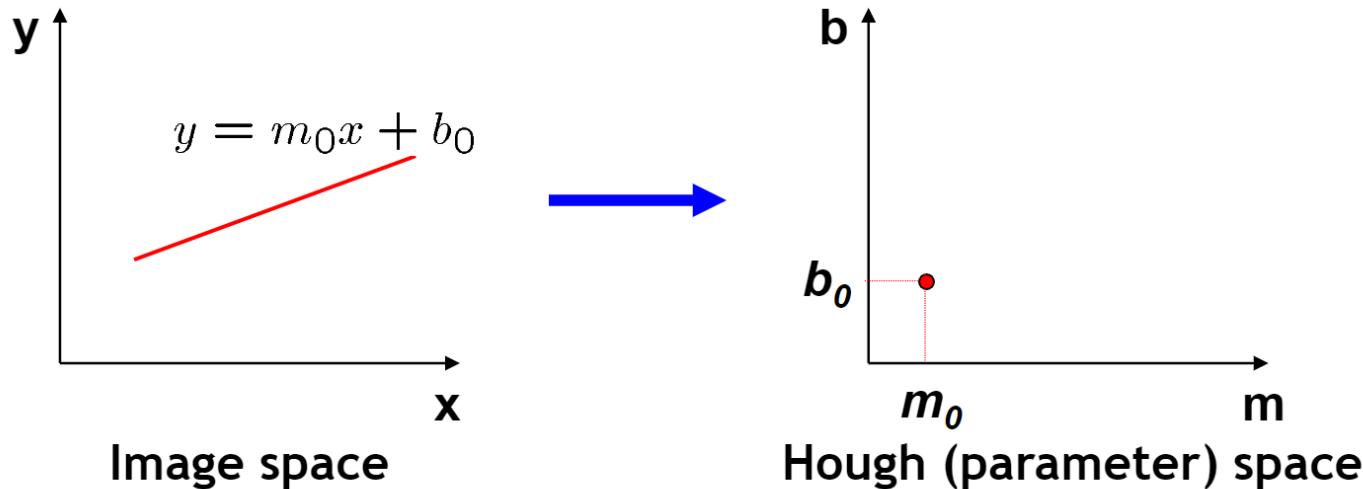
Hough Transform: Outline



1. Create a grid of parameter values
2. Each point votes for a set of parameters, incrementing those values in grid
3. Find maximum or local maxima in grid



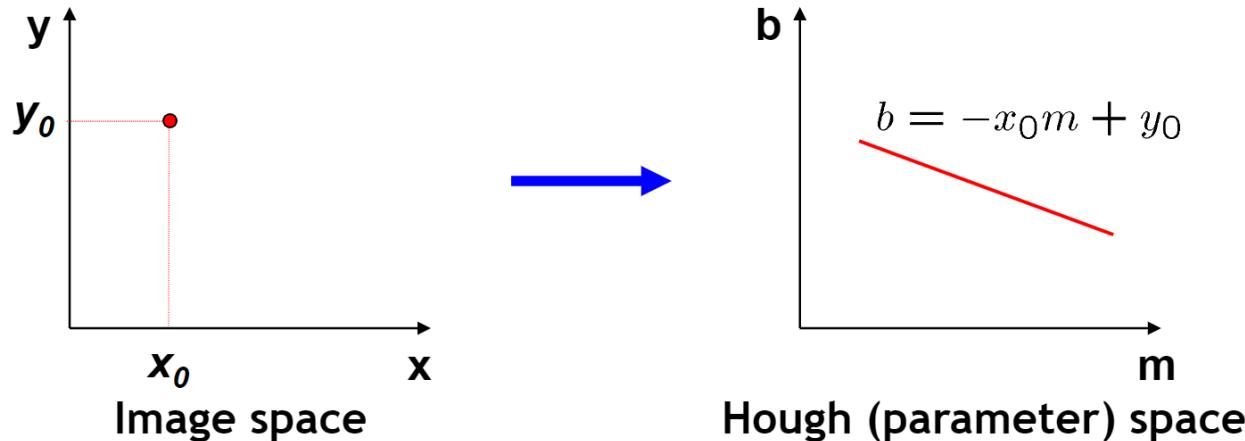
Hough Space



- **Connection between image (x,y) and Hough (m,b) spaces**
 - A line in the image corresponds to a point in Hough space.
 - To go from image space to Hough space:
 - Given a set of points (x,y) , find all (m,b) such that $y = mx + b$



Hough Space



- Connection between image (x,y) and Hough (m,b) spaces
 - A line in the image corresponds to a point in Hough space.
 - To go from image space to Hough space:
 - Given a set of points (x,y) , find all (m,b) such that $y = mx + b$
 - What does a point (x_0, y_0) in the image space map to?
 - Answer: the solutions of $b = -x_0m + y_0$
 - This is a line in Hough space.

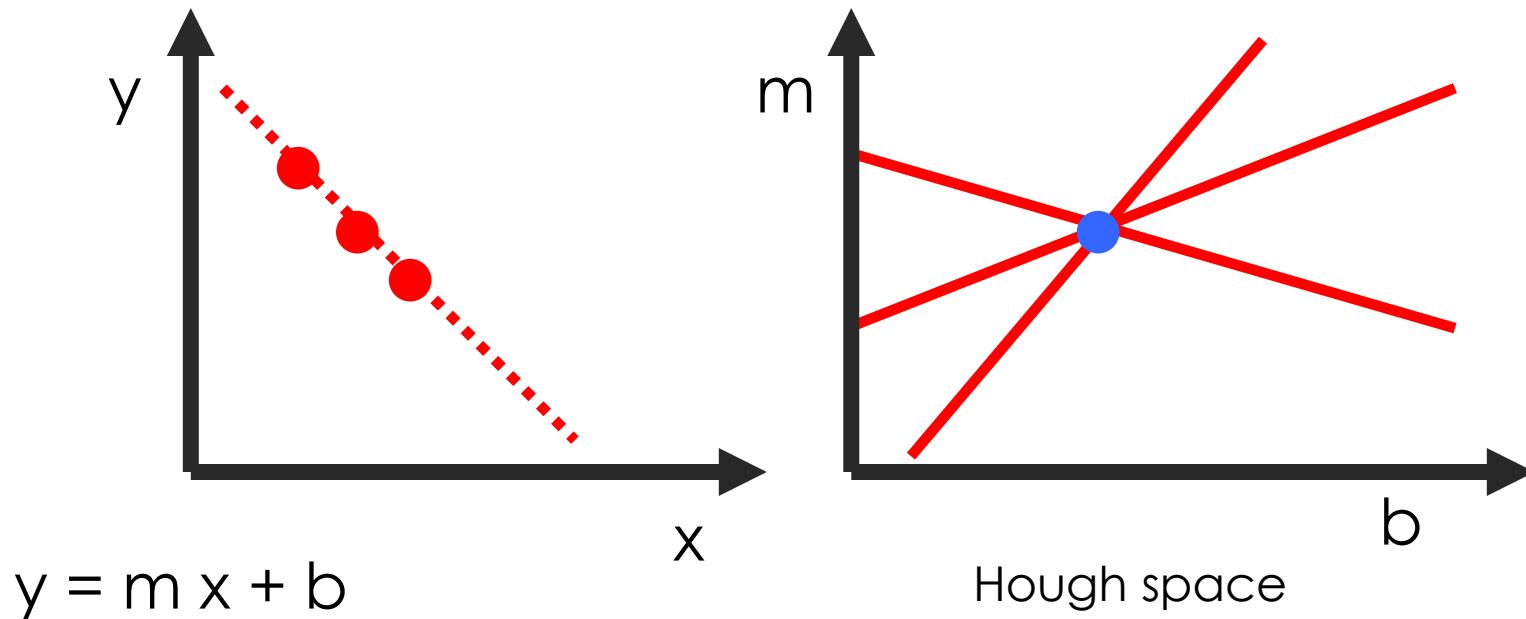


Hough transform



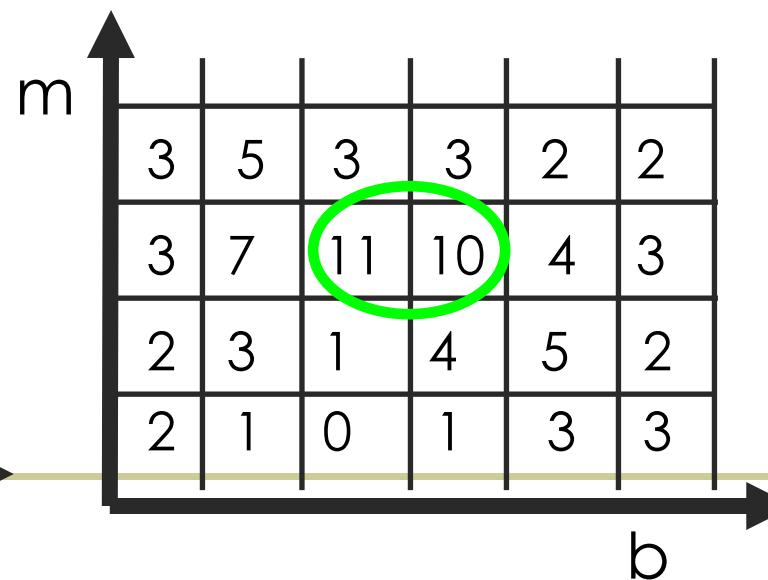
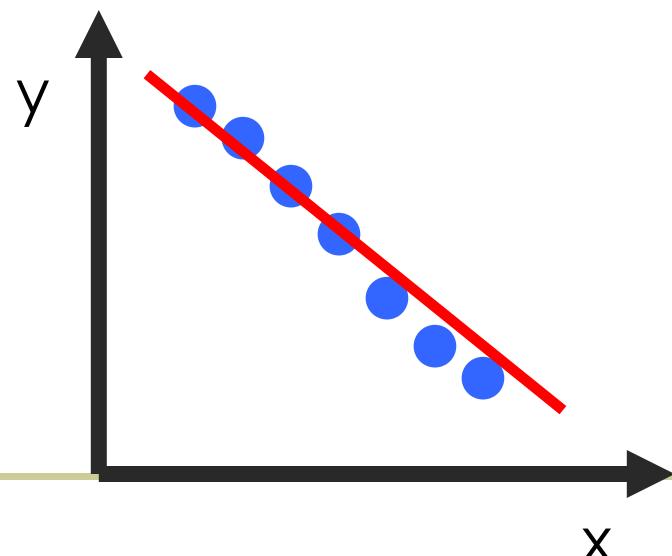
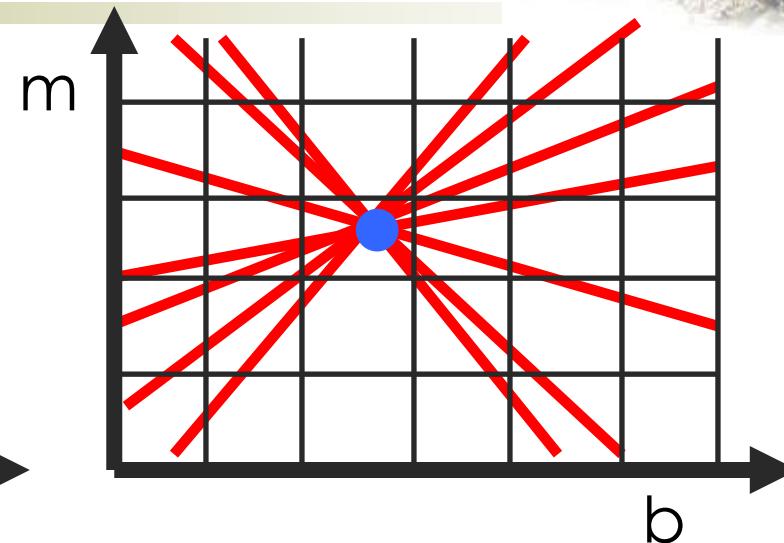
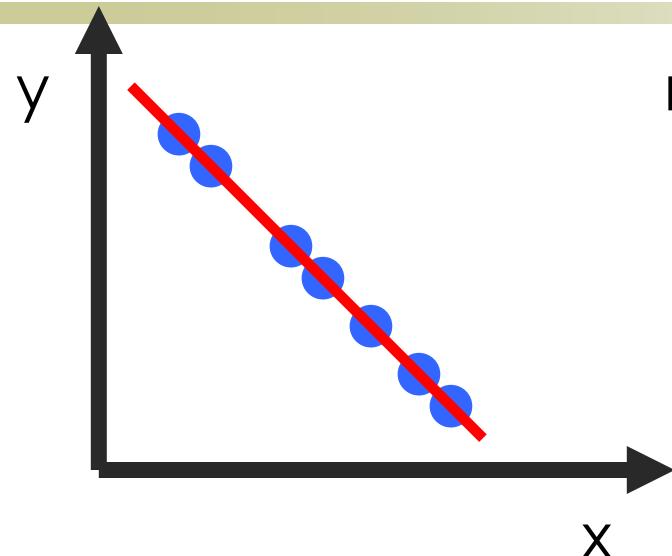
P.V.C. Hough, Machine Analysis of Bubble Chamber Pictures, Proc. Int. Conf.
High Energy Accelerators and Instrumentation, 1959

Given a set of points, find the curve or line that
explains the data points best





Hough transform

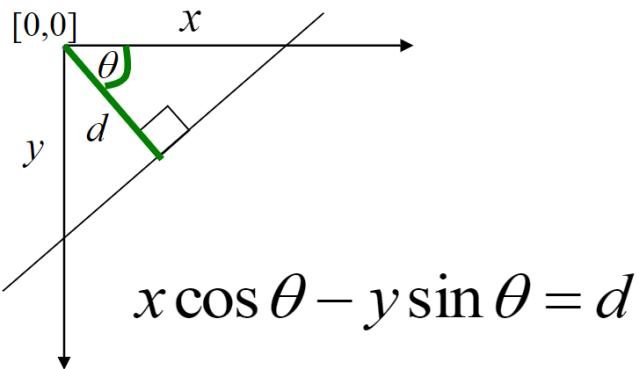




Polar Representation for Lines



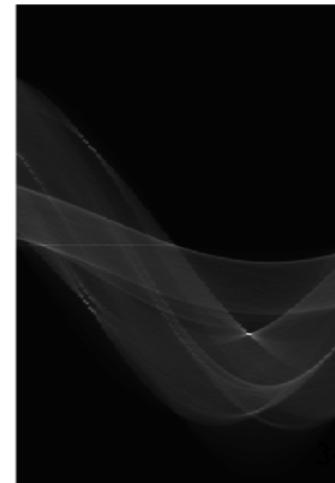
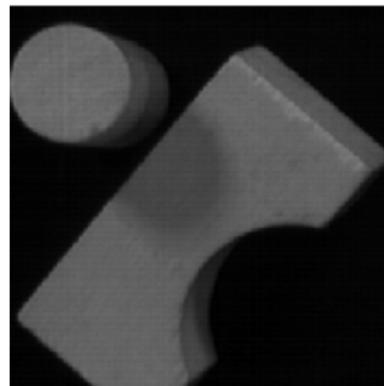
- Issues with usual (m, b) parameter space: can take on infinite values, undefined for vertical lines.



d : perpendicular distance from line to origin

θ : angle the perpendicular makes with the x-axis

- Point in image space
 \Rightarrow Sinusoid segment in Hough space



Slide adapted from Steve Seitz



Hough Transform Algorithm



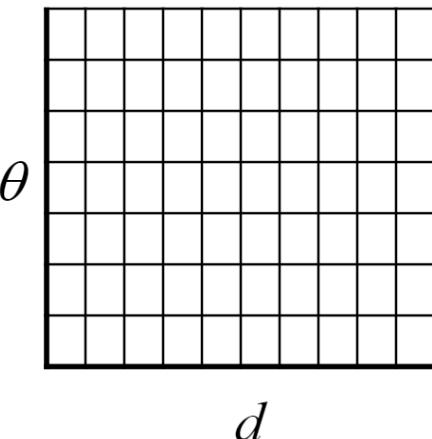
Using the polar parameterization:

$$x \cos \theta + y \sin \theta = d$$

Basic Hough transform algorithm

1. Initialize $H[d, \theta] = 0$.
2. For each edge point (x, y) in the image
 - for $\theta = 0$ to 180 // some quantization
 - $d = x \cos \theta + y \sin \theta$
 - $H[d, \theta] += 1$
3. Find the value(s) of (d, θ) where $H[d, \theta]$ is maximal.
4. The detected line in the image is given by $d = x \cos \theta + y \sin \theta$

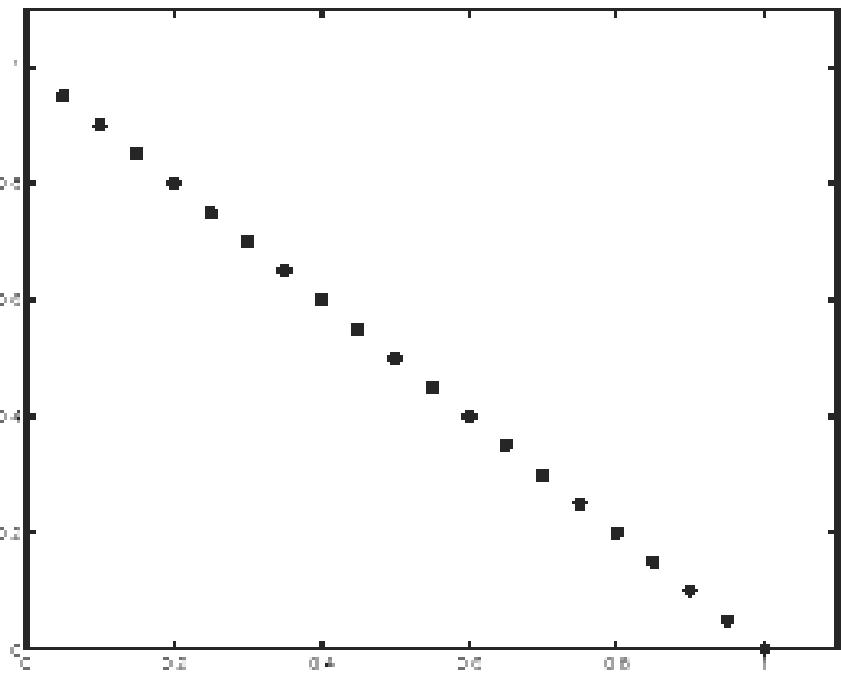
H : accumulator array (votes)



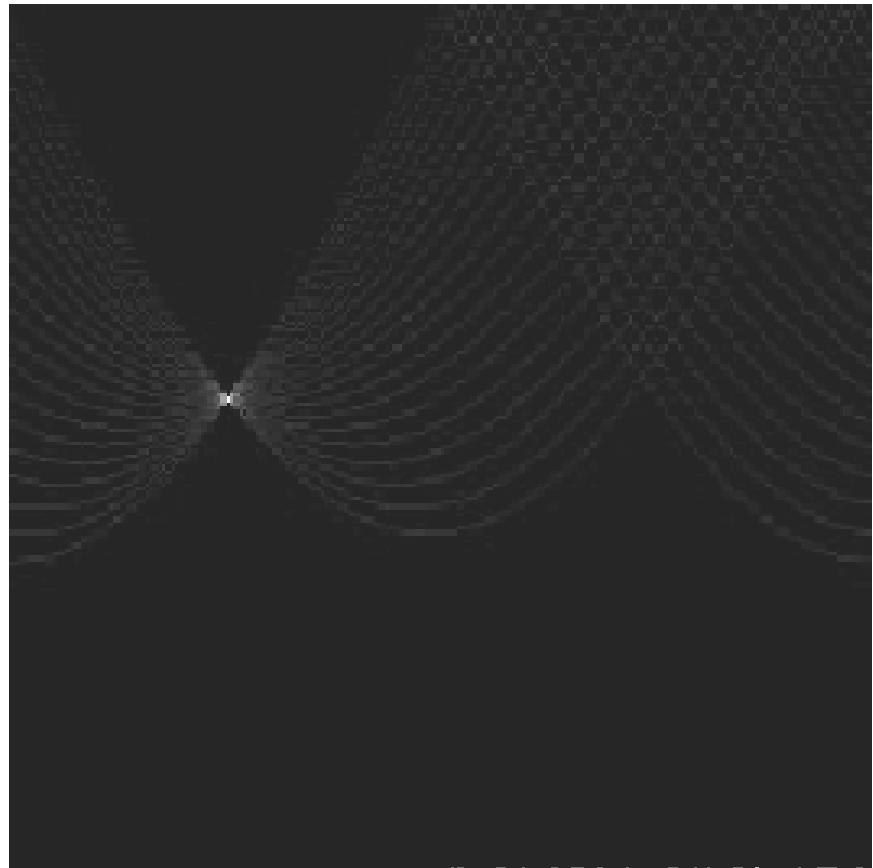
- Time complexity (in terms of number of votes)?



Hough transform - experiments



features



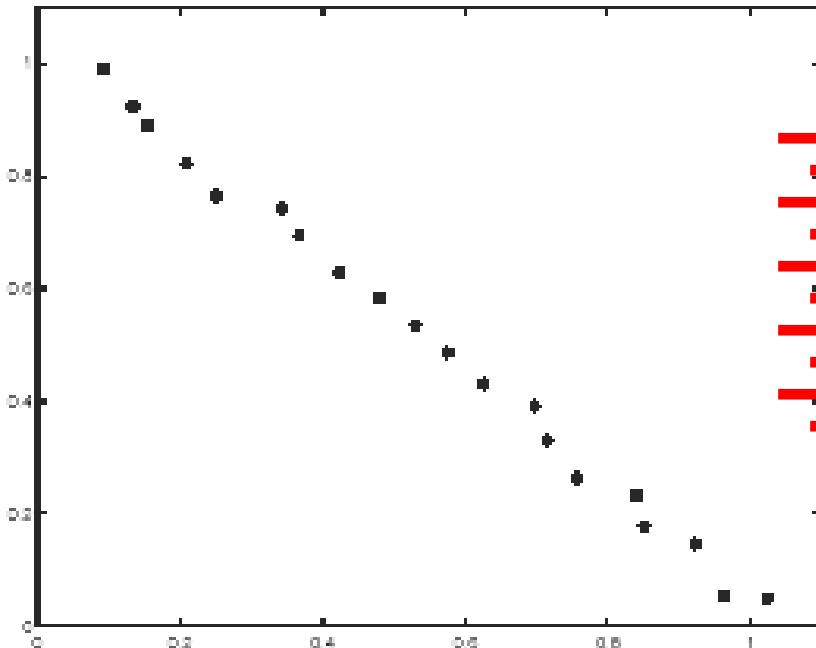
votes



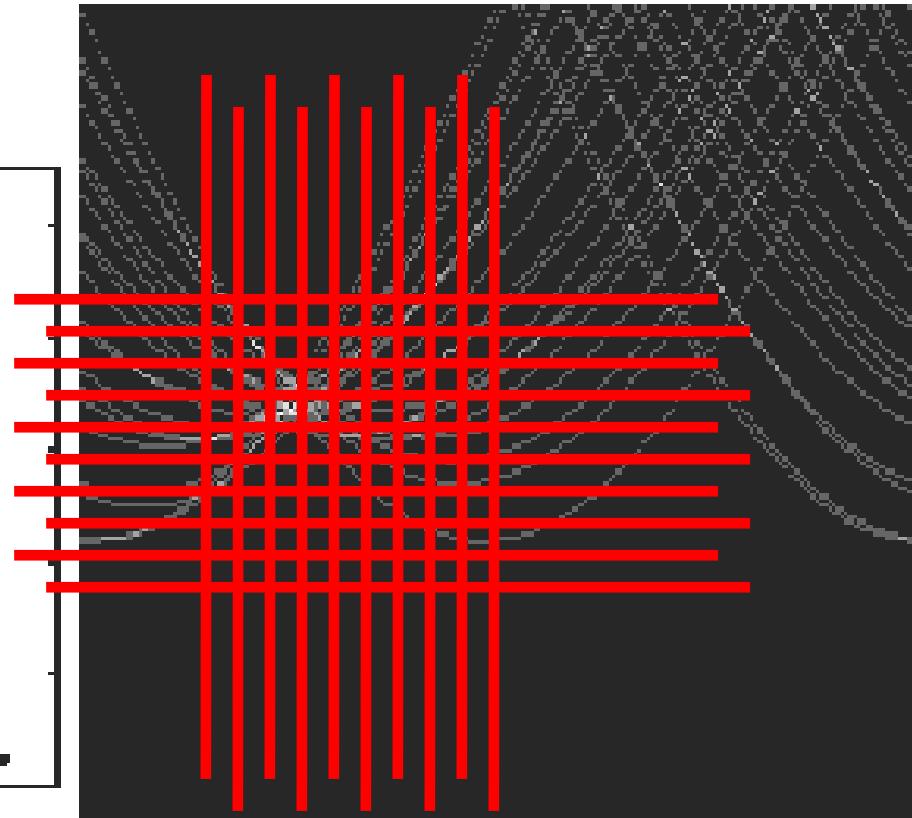
Hough transform - experiments



Noisy data



features

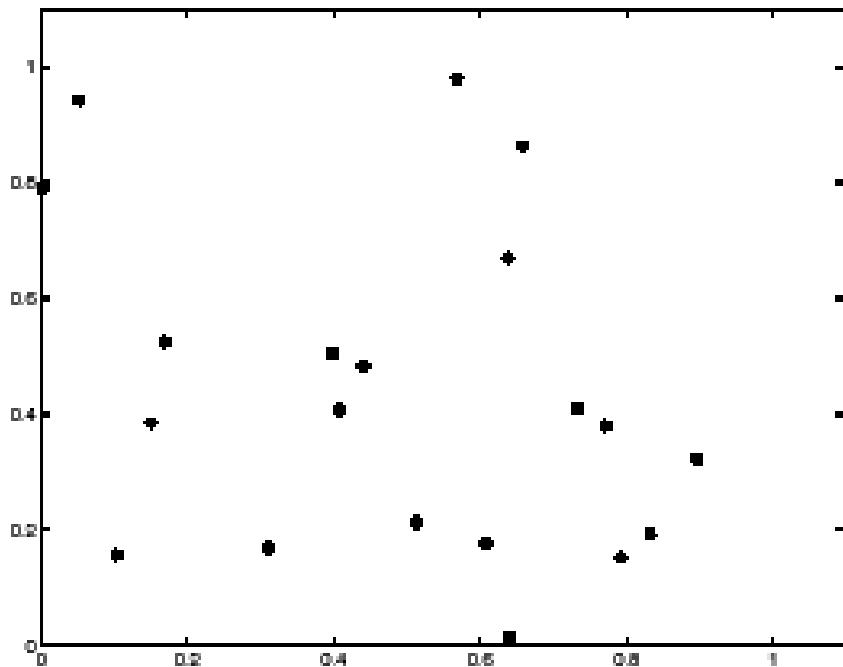


votes

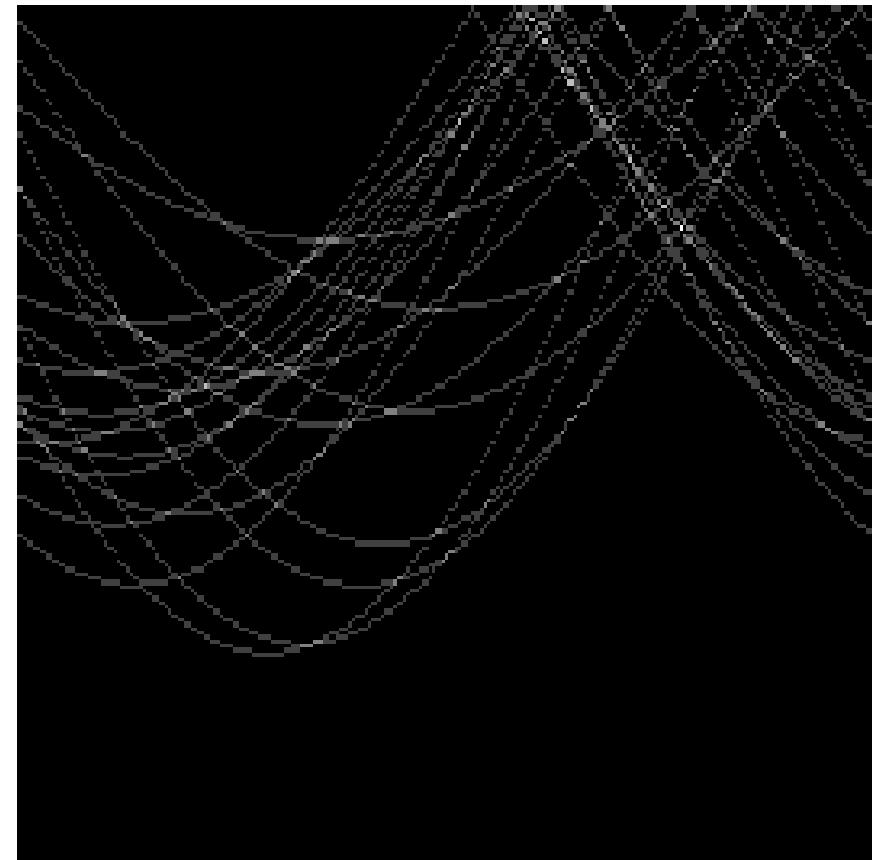
Need to adjust grid size or smooth



Hough transform - experiments



features



votes

Issue: spurious peaks due to uniform noise



1. Image → Canny





2. Canny → Hough votes

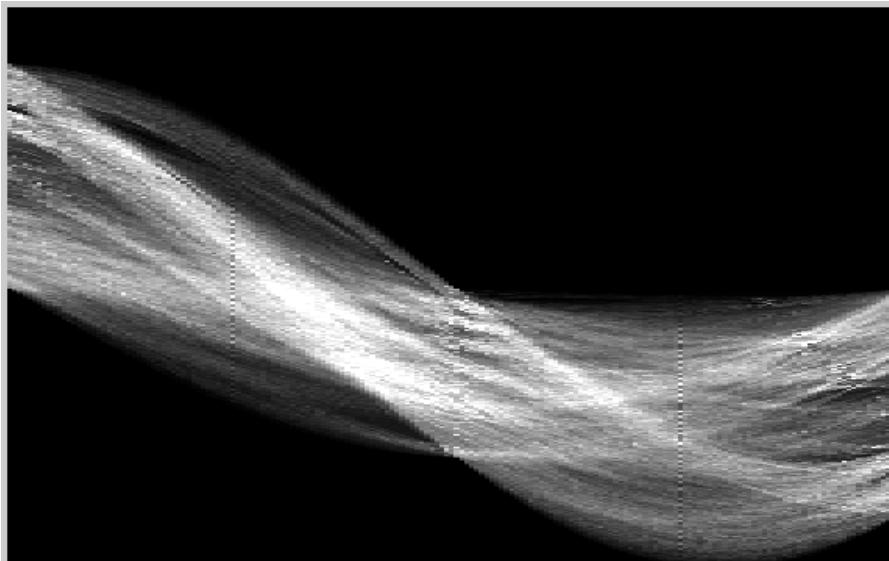




3. Hough votes → Edges

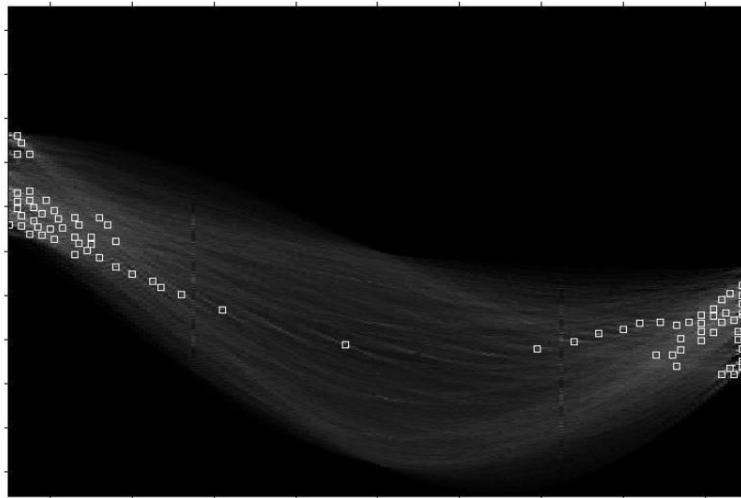
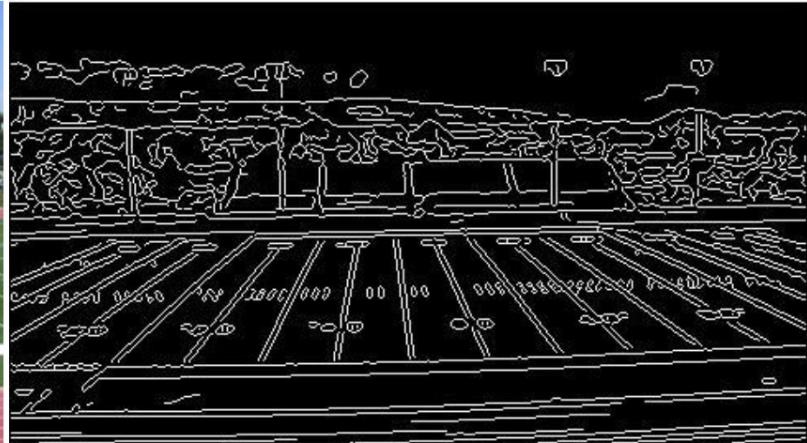


Find peaks and post-process





Real-World Example



Showing longest segments found



Hough transform: pros and cons

Pros

- All points are processed independently, so can cope with occlusion, gaps
- Some robustness to noise: noise points unlikely to contribute *consistently* to any single bin
- Can detect multiple instances of a model in a single pass

Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: can be tricky to pick a good grid size



Hough transform conclusions



Common applications

- Line fitting (also circles, ellipses, etc.)
- Object instance recognition (parameters are position/scale/orientation)
- Object category recognition (parameters are position/scale)



RANSAC



- RANdom Sample Consensus
- **Approach:** we want to avoid the impact of outliers, so let's look for “inliers”, and use those only.
- **Intuition:** if an outlier is chosen to compute the current fit, then the resulting **line** (**transformation**) won't have much support from rest of the **points** (**matches**).



RANSAC



RANSAC loop:

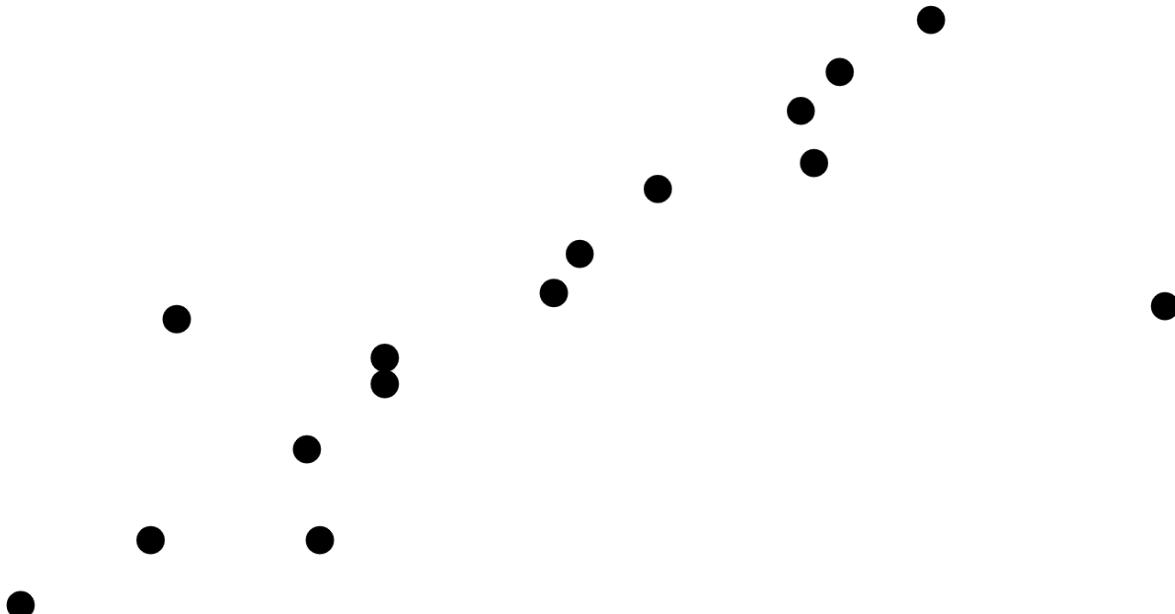
1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
 2. Compute transformation from seed group
 3. Find *inliers* to this transformation
 4. If the number of inliers is sufficiently large, recompute least-squares estimate of transformation on all of the inliers
- Keep the transformation with the largest number of inliers



RANSAC Line Fitting Example



- Task: Estimate the best line
 - *How many points do we need to estimate the line?*

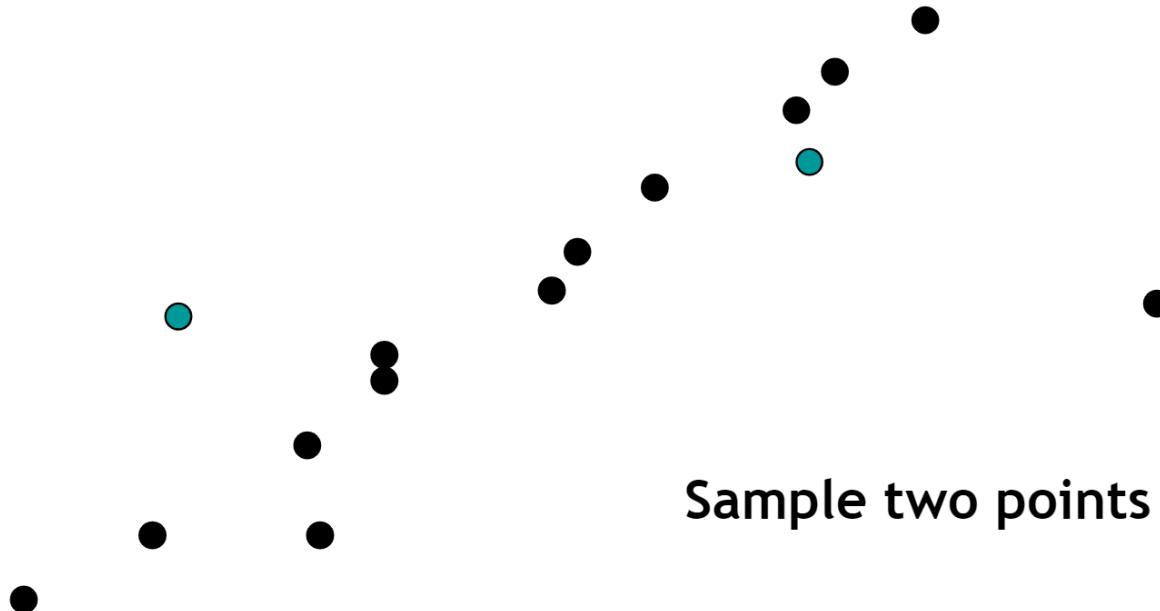




RANSAC Line Fitting Example



- Task: Estimate the best line

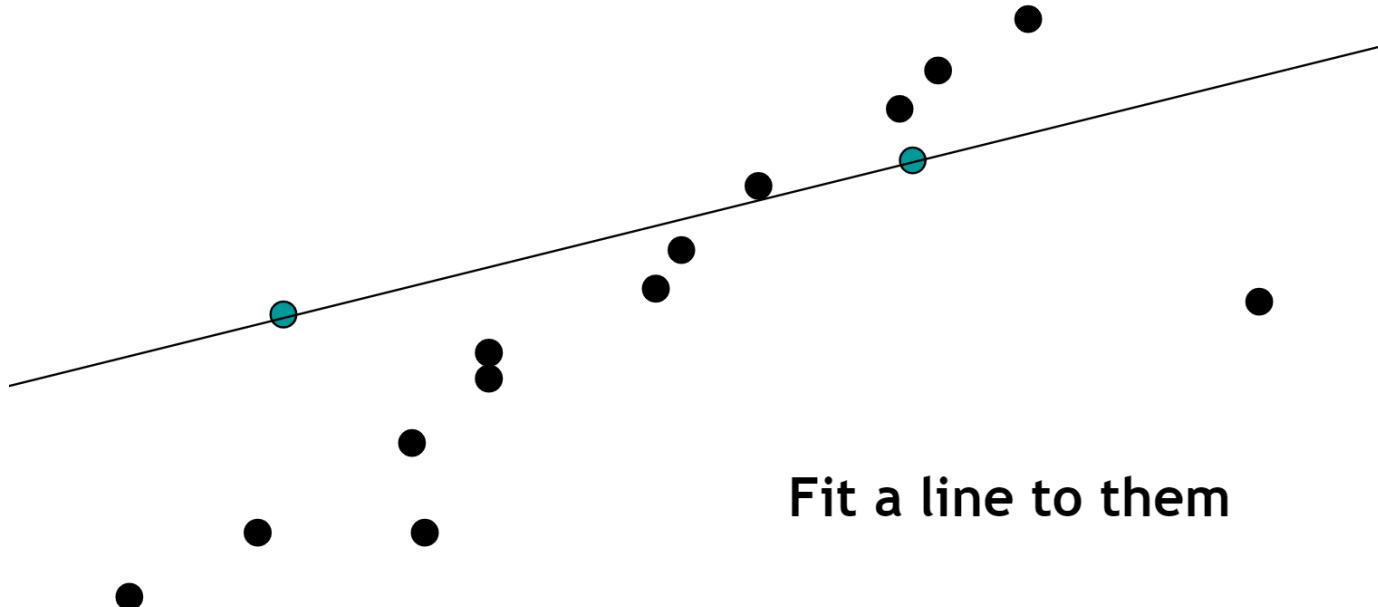




RANSAC Line Fitting Example



- Task: Estimate the best line

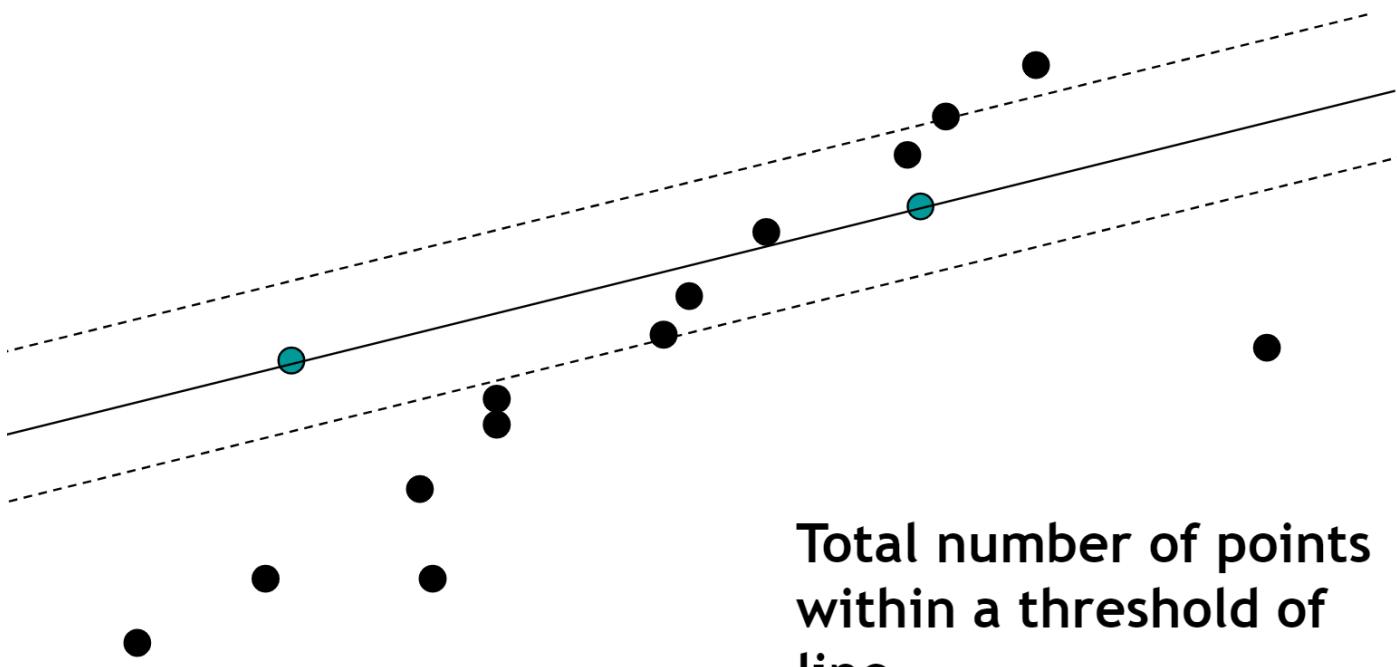




RANSAC Line Fitting Example



- Task: Estimate the best line

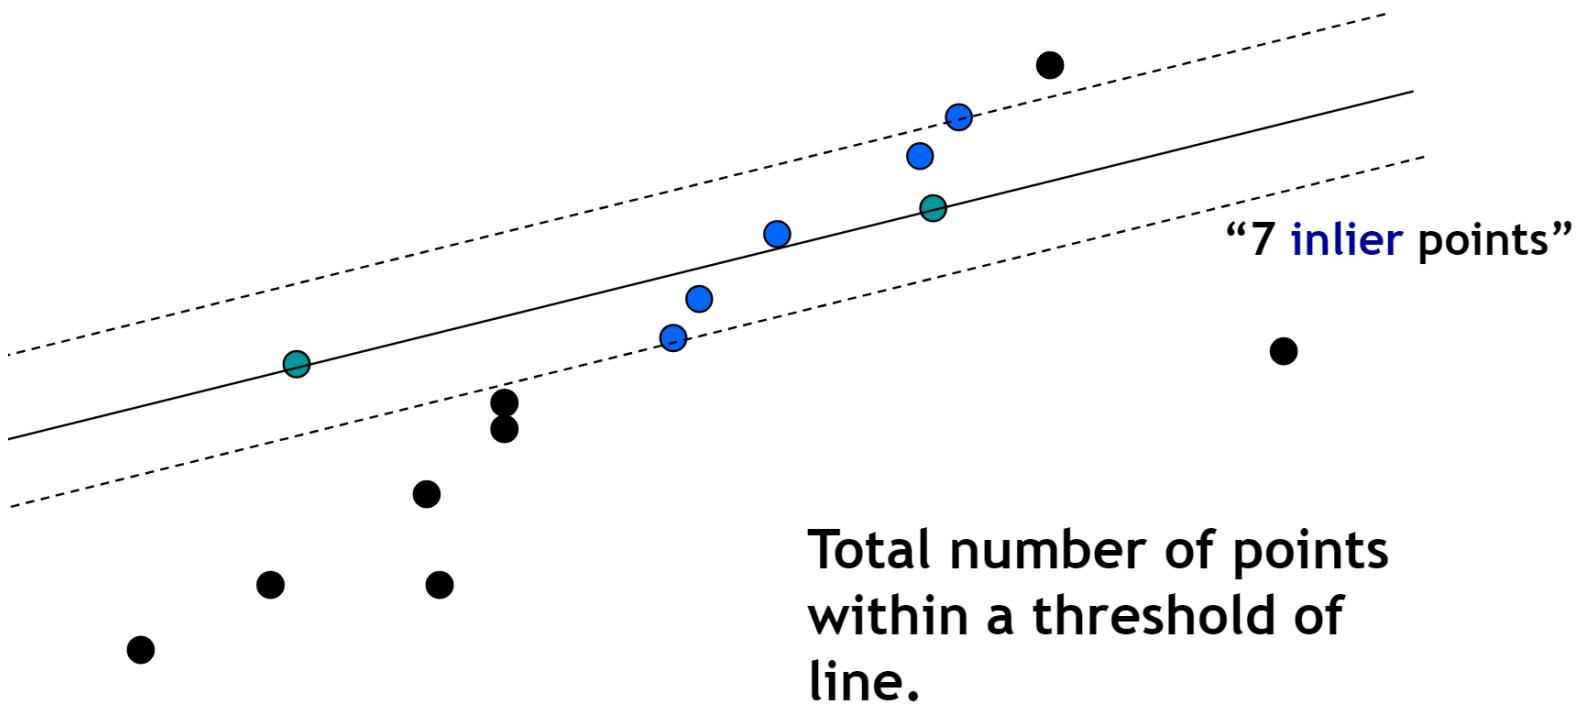




RANSAC Line Fitting Example



- Task: Estimate the best line

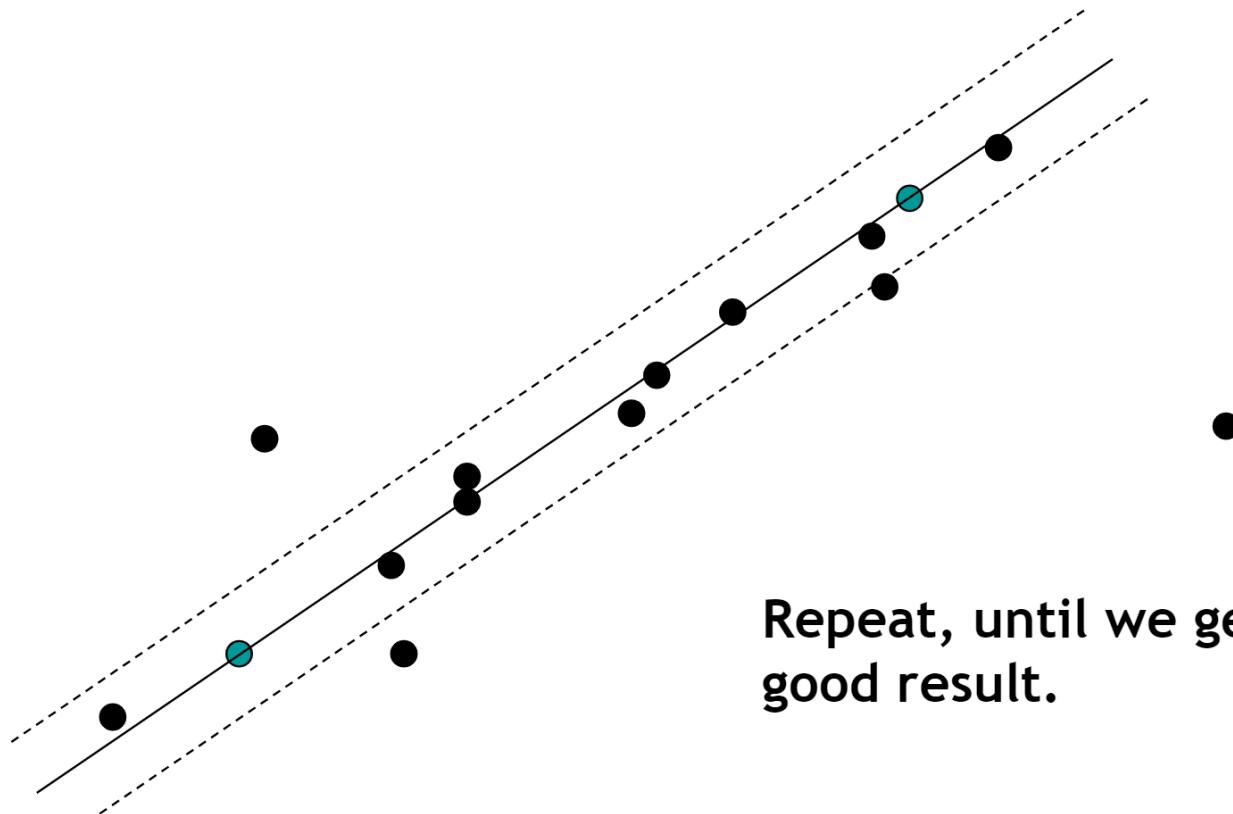




RANSAC Line Fitting Example



- Task: Estimate the best line

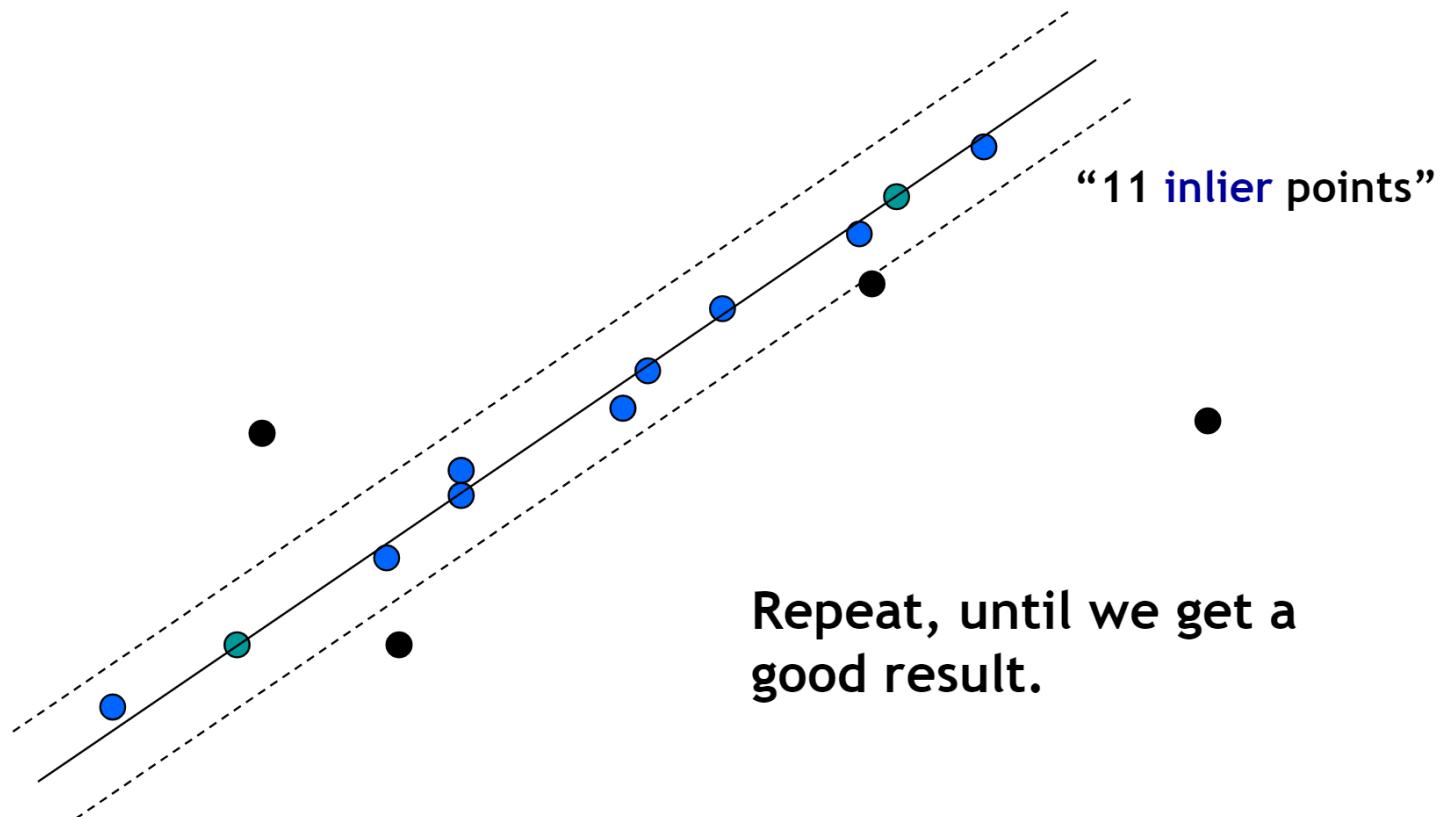




RANSAC Line Fitting Example



- Task: Estimate the best line





RANSAC: How many samples?



- **How many samples are needed?**
 - Suppose w is fraction of inliers (points from line).
 - n points needed to define hypothesis (2 for lines)
 - k samples chosen.
- Prob. that a single sample of n points is correct: w^n
- Prob. that all k samples fail is: $(1-w^n)^k$

⇒ Choose k high enough to keep this below desired failure rate.



RANSAC: Computed k (p=0.99)



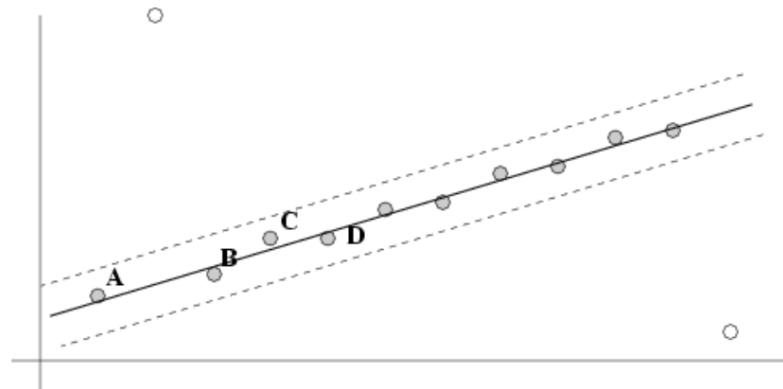
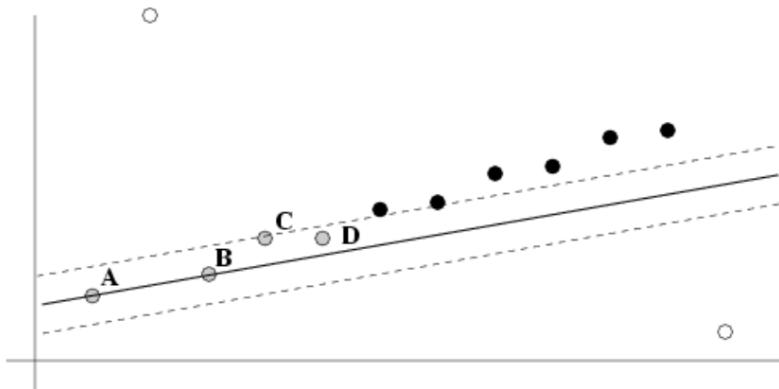
Sample size n	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177



After RANSAC



- RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.
- Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).
- But this may change inliers, so alternate fitting with reclassification as inlier/outlier.





RANSAC Example

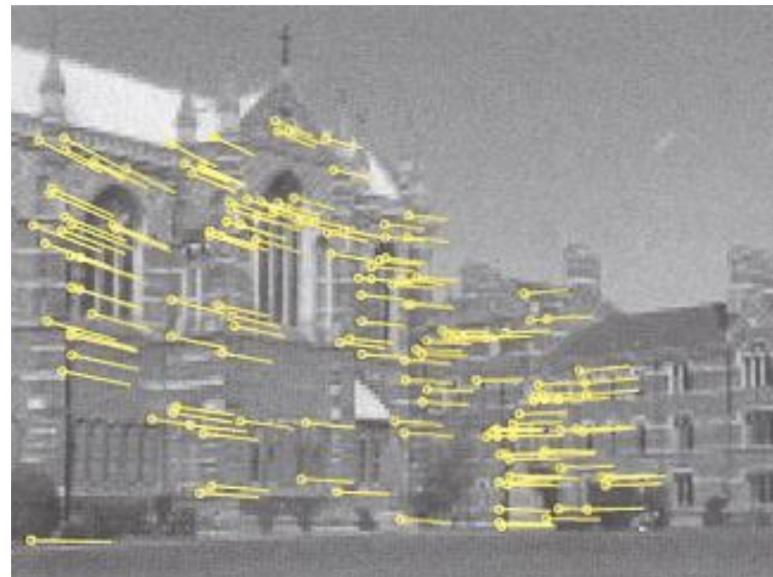


- Find best stereo match within a square search window (here 300 pixels²)
- Global transformation model: epipolar geometry

before RANSAC



after RANSAC



Images from Hartley & Zisserman



RANSAC conclusions



Good

- Robust to outliers
- Applicable for larger number of objective function parameters than Hough transform
- Optimization parameters are easier to choose than Hough transform

Bad

- Computational time grows quickly with fraction of outliers and number of parameters
- Not as good for getting multiple fits (though one solution is to remove inliers after each fit and repeat)

Common applications

- **Computing a homography (e.g., image stitching)**
- Estimating fundamental matrix (relating two views)



Today's class



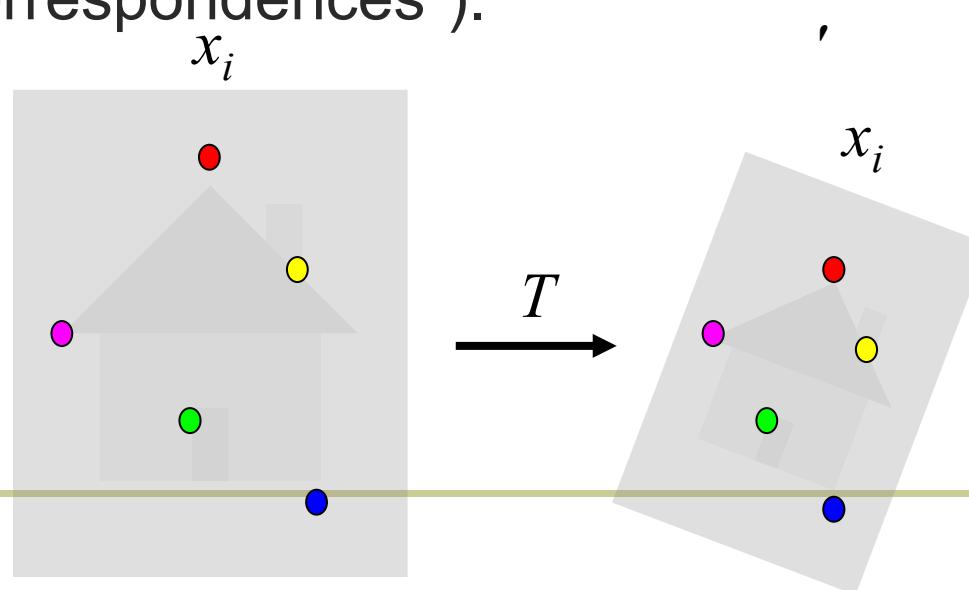
- Introduction to alignment
- Alignment methods
 - Global methods
 - Hypothesize and test
- Image Transformation
 - **Common image transformations**
 - Examples of solving image alignment
- Homework: Mosaics



Alignment problem



- We have previously considered how to **fit a model to image evidence**
 - e.g., a line to edge points
- In alignment, we will fit the parameters of some **transformation** according to a set of matching feature pairs (“correspondences”).





Common transformations



original

Transformed



translation



rotation



aspect



affine



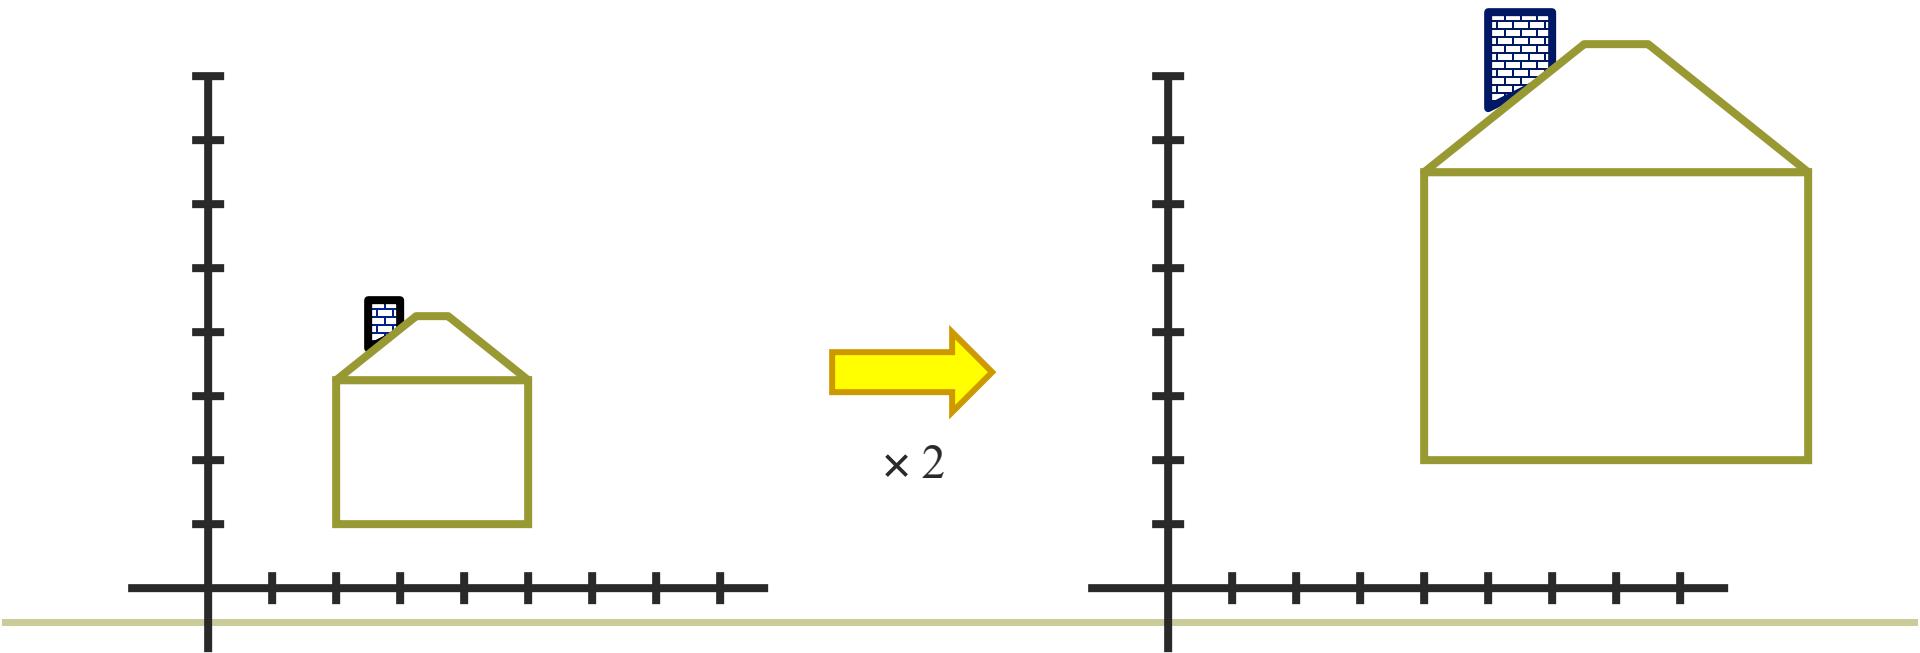
perspective



Scaling



- *Scaling* a coordinate means multiplying each of its components by a scalar
- *Uniform scaling* means this scalar is the same for all components:

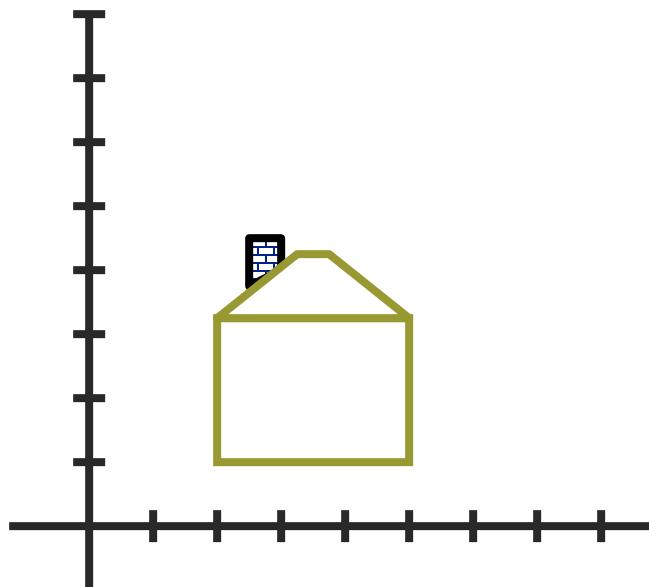




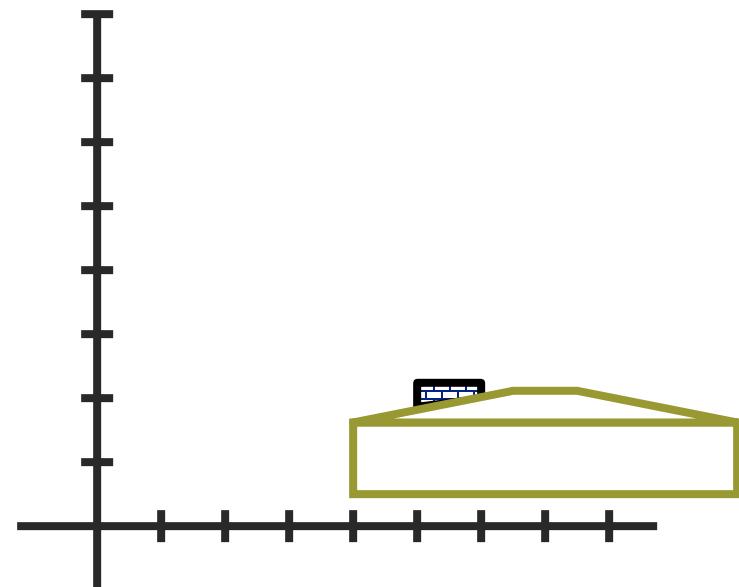
Scaling



- *Non-uniform scaling*: different scalars per component:



→
 $X \times 2$,
 $Y \times 0.5$





Scaling



- Scaling operation: $x' = ax$

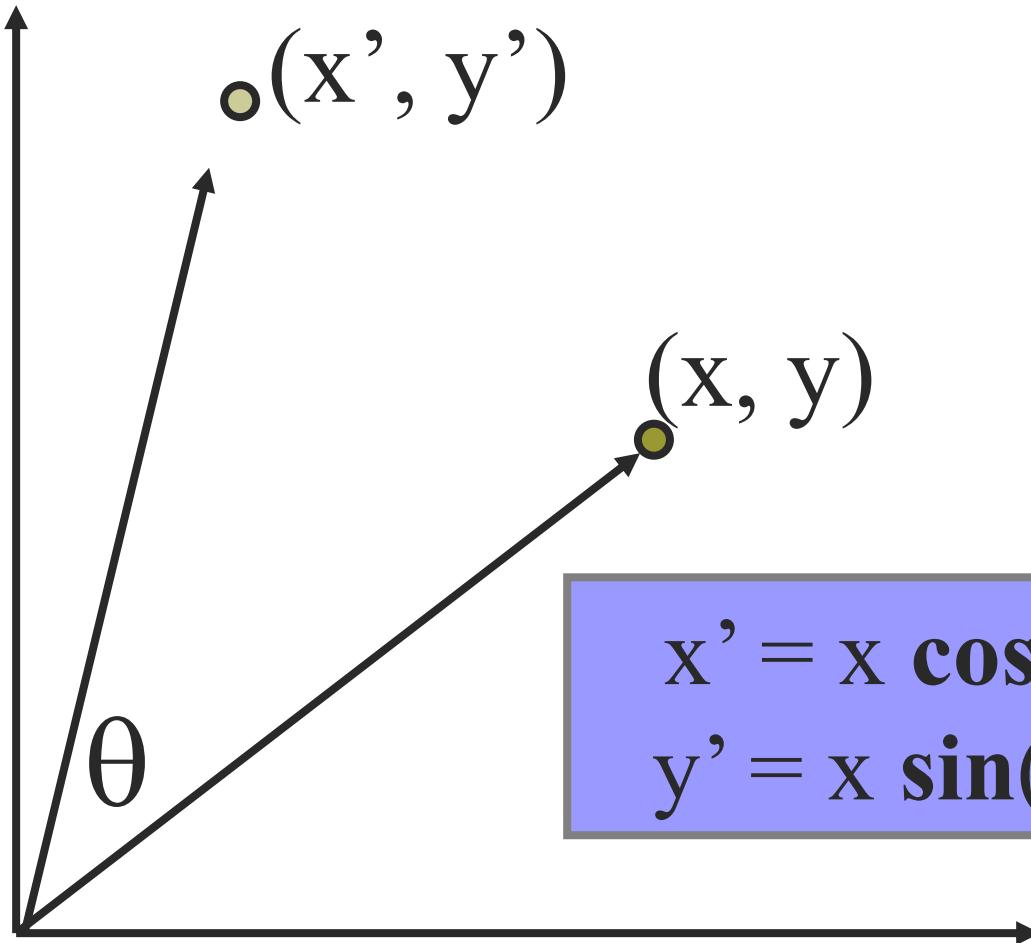
$$y' = by$$

- Or, in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}}_{\text{scaling matrix } S} \begin{bmatrix} x \\ y \end{bmatrix}$$

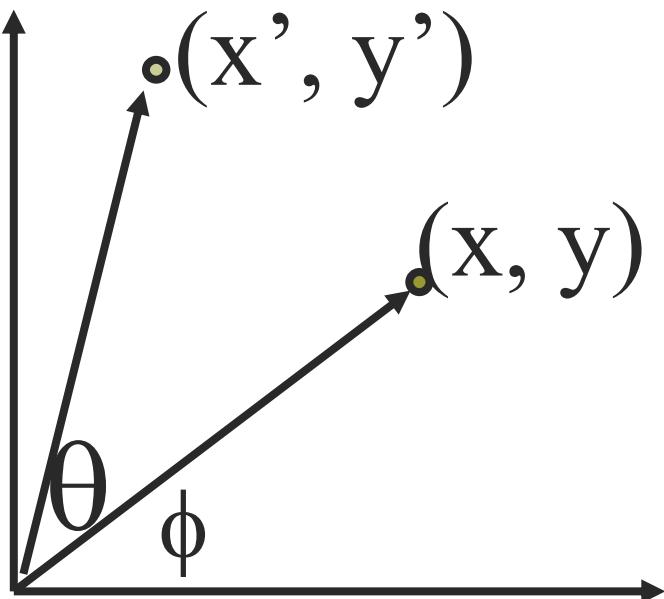


2-D Rotation





2-D Rotation



Polar coordinates...

$$x = r \cos(\phi)$$

$$y = r \sin(\phi)$$

$$x' = r \cos(\phi + \theta)$$

$$y' = r \sin(\phi + \theta)$$

Trig Identity...

$$x' = r \cos(\phi) \cos(\theta) - r \sin(\phi) \sin(\theta)$$

$$y' = r \sin(\phi) \cos(\theta) + r \cos(\phi) \sin(\theta)$$

Substitute...

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = x \sin(\theta) + y \cos(\theta)$$



2-D Rotation



This is easy to capture in matrix form:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \underbrace{\begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} x \\ y \end{bmatrix}$$

Even though $\sin(\theta)$ and $\cos(\theta)$ are nonlinear functions of θ ,

- **x' is a linear combination of x and y**
- **y' is a linear combination of x and y**

What is the inverse transformation?

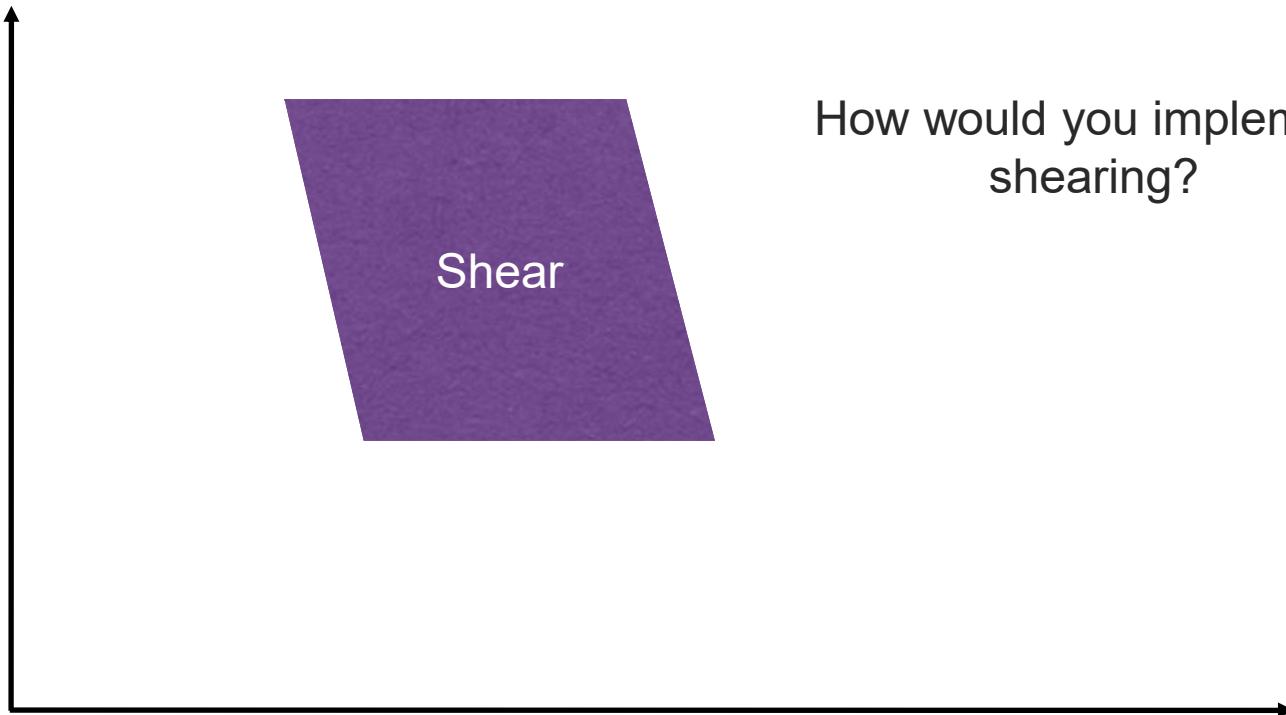
- Rotation by $-\theta$
- For rotation matrices $\mathbf{R}^{-1} = \mathbf{R}^T$



2-D Shearing



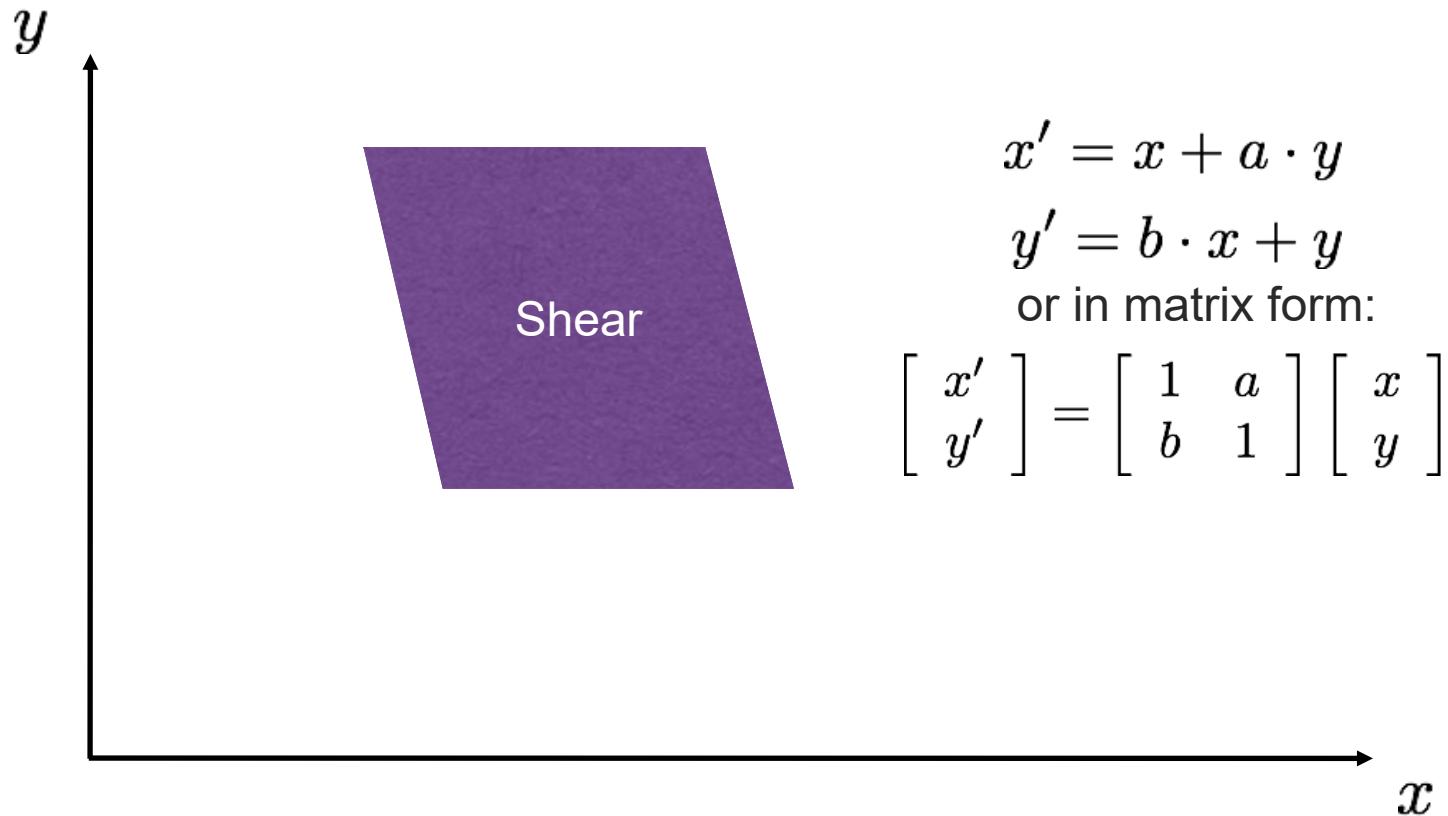
y



How would you implement
shearing?



2-D Shearing





What transformations can be represented with a 2x2 matrix?



2D Scaling?

$$x' = s_x * x$$

$$y' = s_y * y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Rotate around (0,0)?

$$x' = \cos \Theta * x - \sin \Theta * y$$

$$y' = \sin \Theta * x + \cos \Theta * y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta \\ \sin \Theta & \cos \Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Shear?

$$x' = x + sh_x * y$$

$$y' = sh_y * x + y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & sh_x \\ sh_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$



What transformations can be represented with a 2x2 matrix?



2D Mirror about Y axis?

$$x' = -x$$

$$y' = y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Mirror over (0,0)?

$$x' = -x$$

$$y' = -y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Translation?

$$x' = x + t_x$$

$$y' = y + t_y$$

NO!



2D Linear Transformations



$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

- Only linear 2D transformations can be represented with a 2×2 matrix.
- Linear transformations are combinations of ...
 - Scale,
 - Rotation,
 - Shear, and
 - Mirror



Homogeneous coordinates

To convert to homogeneous coordinates:

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

Converting *from* homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$



Homogeneous Coordinates



- Q: How can we represent 2d translation as a 3x3 matrix using homogeneous coordinates?

$$x' = x + t_x$$

$$y' = y + t_y$$

- A: Using the rightmost column:

$$\text{Translation} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

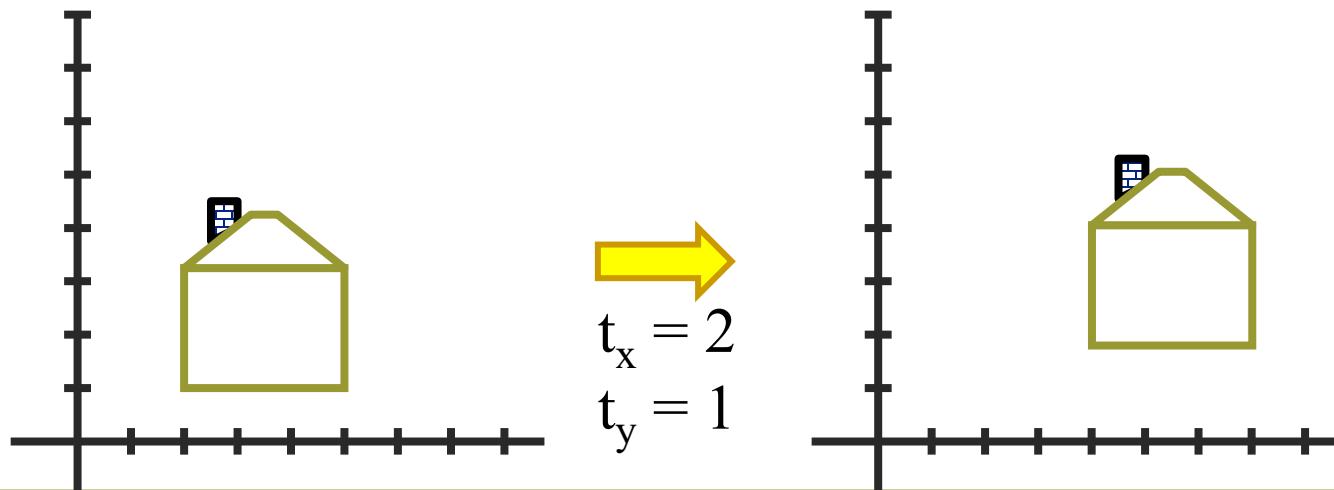


Translation



Homogeneous Coordinates

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ 1 \end{bmatrix}$$





Basic 2D Transformations



■ Basic 2D transformations as 3x3 matrices

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

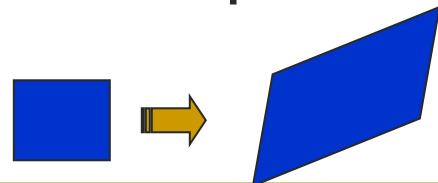


2D Affine Transformations



$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Affine transformations are combinations of ...
 - Linear transformations, and
 - Translations
- Parallel lines remain parallel





Affine Transformations



Affine transformations are combinations of

- Linear transformations, and
- Translations

Properties of affine transformations:

- Lines map to lines
- Parallel lines remain parallel
- Ratios are preserved
- Closed under composition

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

or

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$



Projective Transformations



Projective transformations are combos of

- Affine transformations, and
- Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Properties of projective transformations:

- Lines map to lines
- Parallel lines do not necessarily remain parallel
- Ratios are not preserved
- Closed under composition
- Models change of basis
- Projective matrix is defined up to a scale (8 DOF)

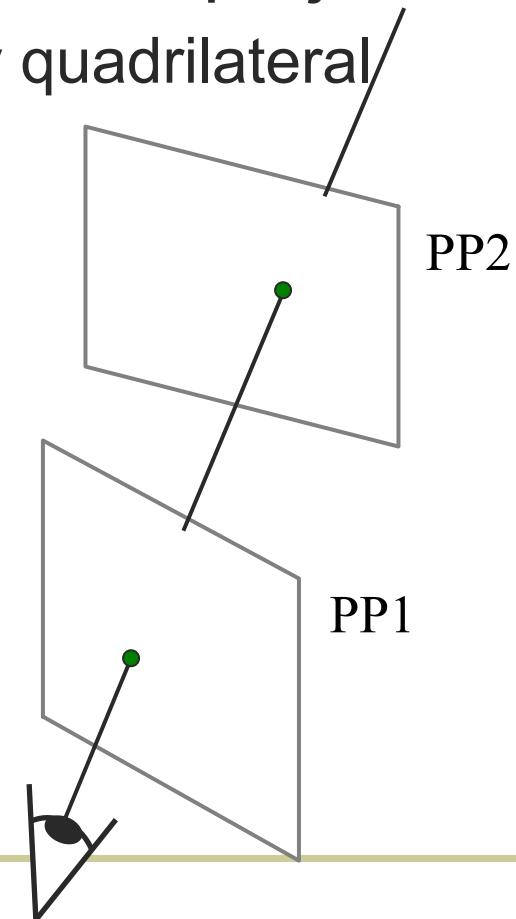


Image reprojection: Homography



- A projective transform is a mapping between any two PPs with the same center of projection
 - rectangle should map to arbitrary quadrilateral
 - parallel lines aren't
 - but must preserve straight lines
- called **Homography**

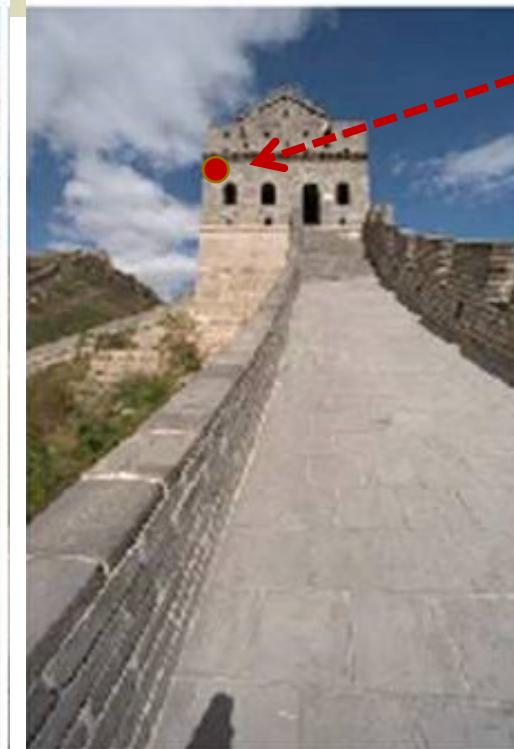
$$\begin{bmatrix} wx' \\ wy' \\ w \\ p' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ H \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \\ p \end{bmatrix}$$





Homoography

$$(x, y)$$



$$\left(\frac{wx'}{w}, \frac{wy'}{w} \right)$$

$$= (x', y')$$

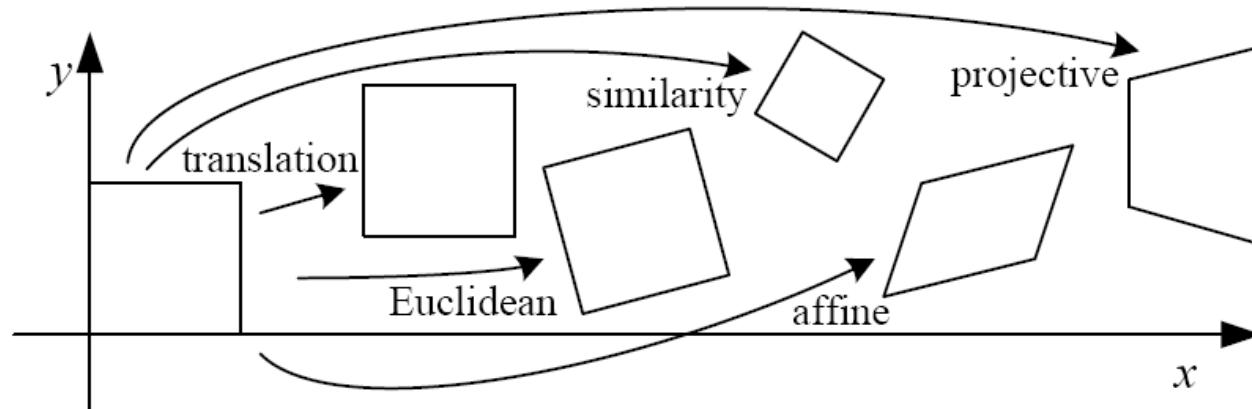
To apply a given homography \mathbf{H}

- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \mathbf{p}' \quad \mathbf{H} \quad \mathbf{p}$$



2D image transformations



Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\left[\begin{array}{c c} \mathbf{I} & \mathbf{t} \end{array} \right]_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\left[\begin{array}{c c} \mathbf{R} & \mathbf{t} \end{array} \right]_{2 \times 3}$	3	lengths + ...	
similarity	$\left[\begin{array}{c c} s\mathbf{R} & \mathbf{t} \end{array} \right]_{2 \times 3}$	4	angles + ...	
affine	$\left[\begin{array}{c} \mathbf{A} \end{array} \right]_{2 \times 3}$	6	parallelism + ...	
projective	$\left[\begin{array}{c} \tilde{\mathbf{H}} \end{array} \right]_{3 \times 3}$	8	straight lines	



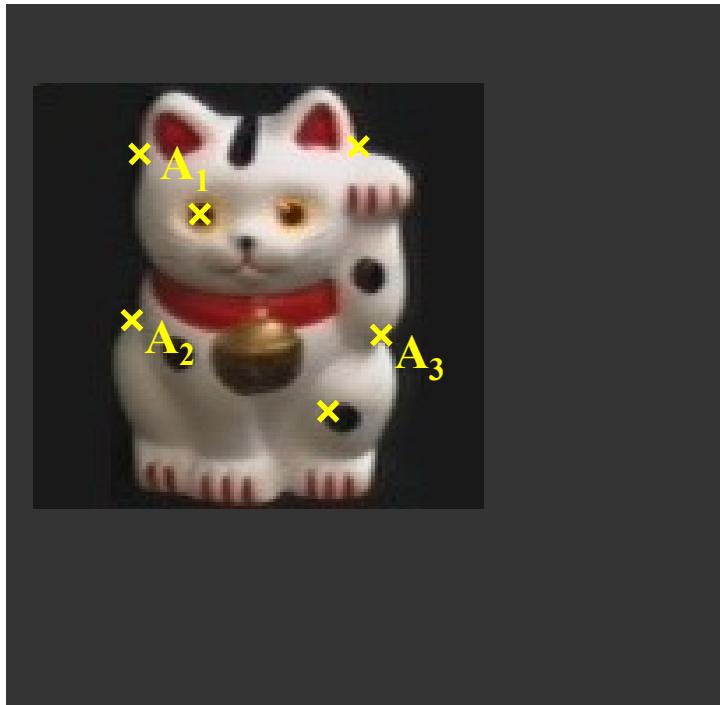
Today's class



- Introduction to alignment
- Alignment methods
 - Global methods
 - Hypothesize and test
- Image Transformation
 - Common image transformations
 - **Examples of solving image alignment**
- Homework: Mosaics



Example: solving for translation

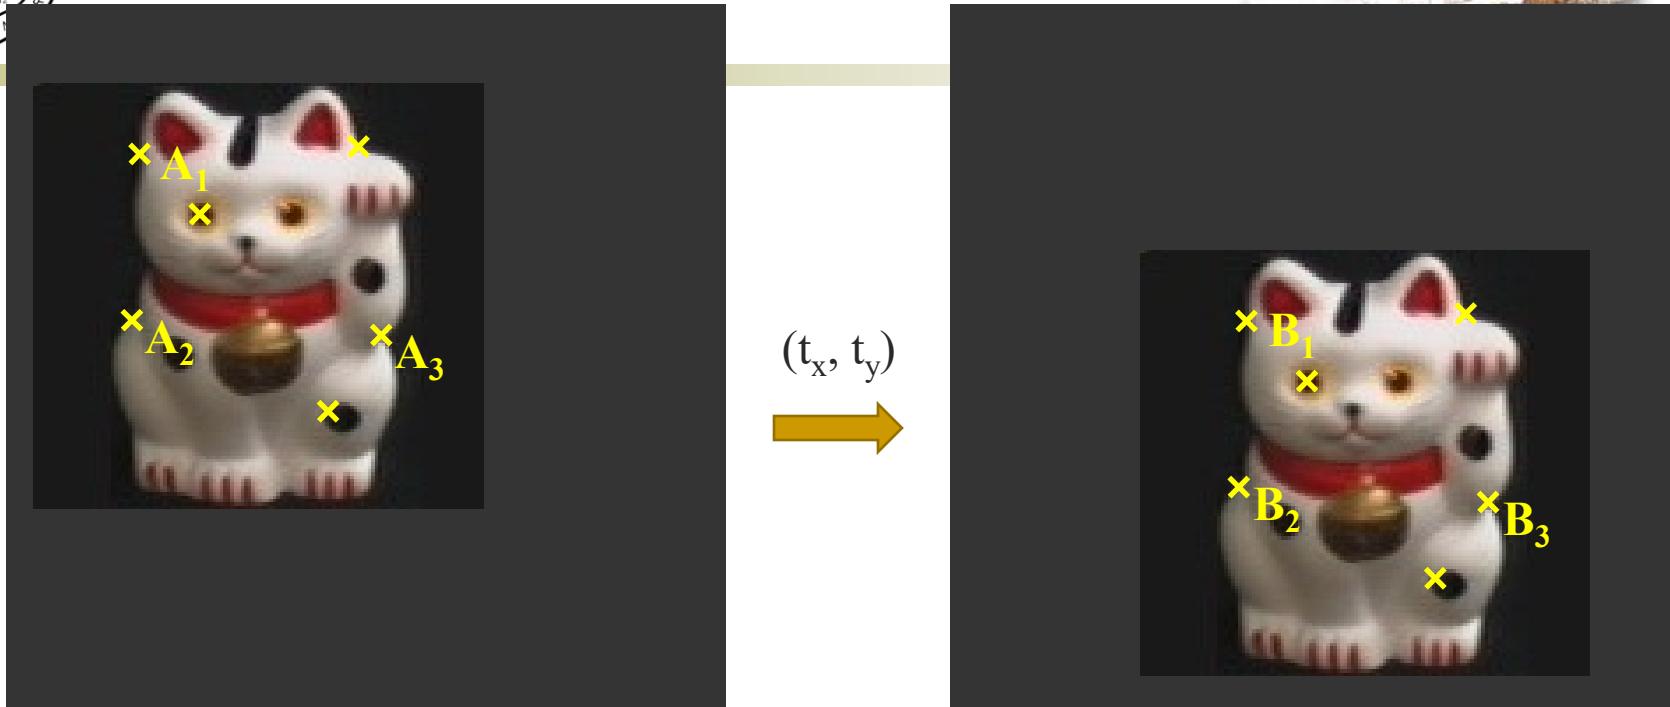


Given matched points in $\{A\}$ and $\{B\}$, estimate the translation of the object

$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$



Example: solving for translation



Least squares solution

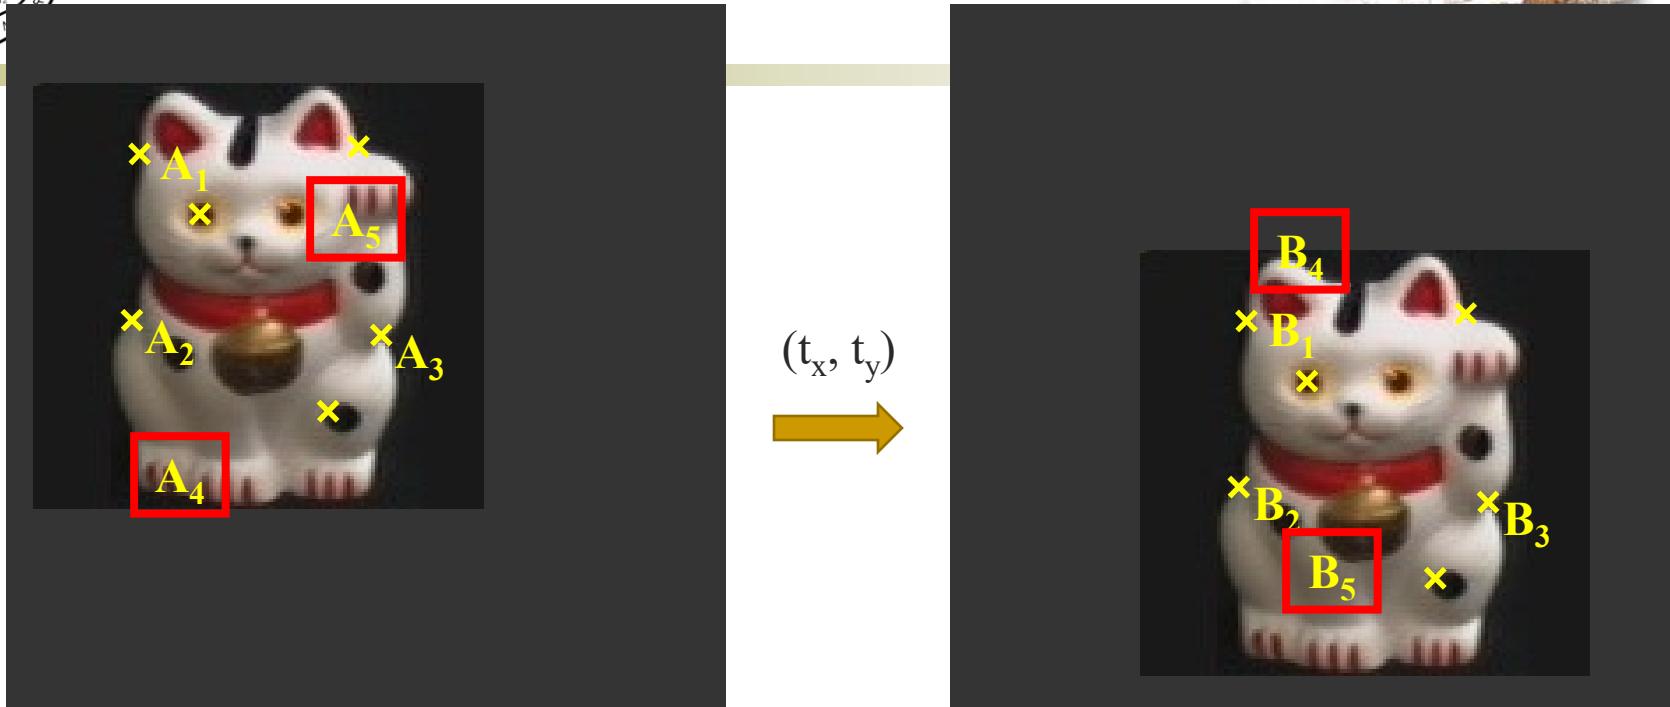
1. Write down objective function
2. Derived solution
 - a) Compute derivative
 - b) Compute solution
3. Computational solution
 - a) Write in form $Ax=b$
 - b) Solve using pseudo-inverse or eigenvalue decomposition

$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \\ \vdots & \vdots \\ 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} = \begin{bmatrix} x_1^B - x_1^A \\ y_1^B - y_1^A \\ \vdots \\ x_n^B - x_n^A \\ y_n^B - y_n^A \end{bmatrix}$$



Example: solving for translation



RANSAC solution

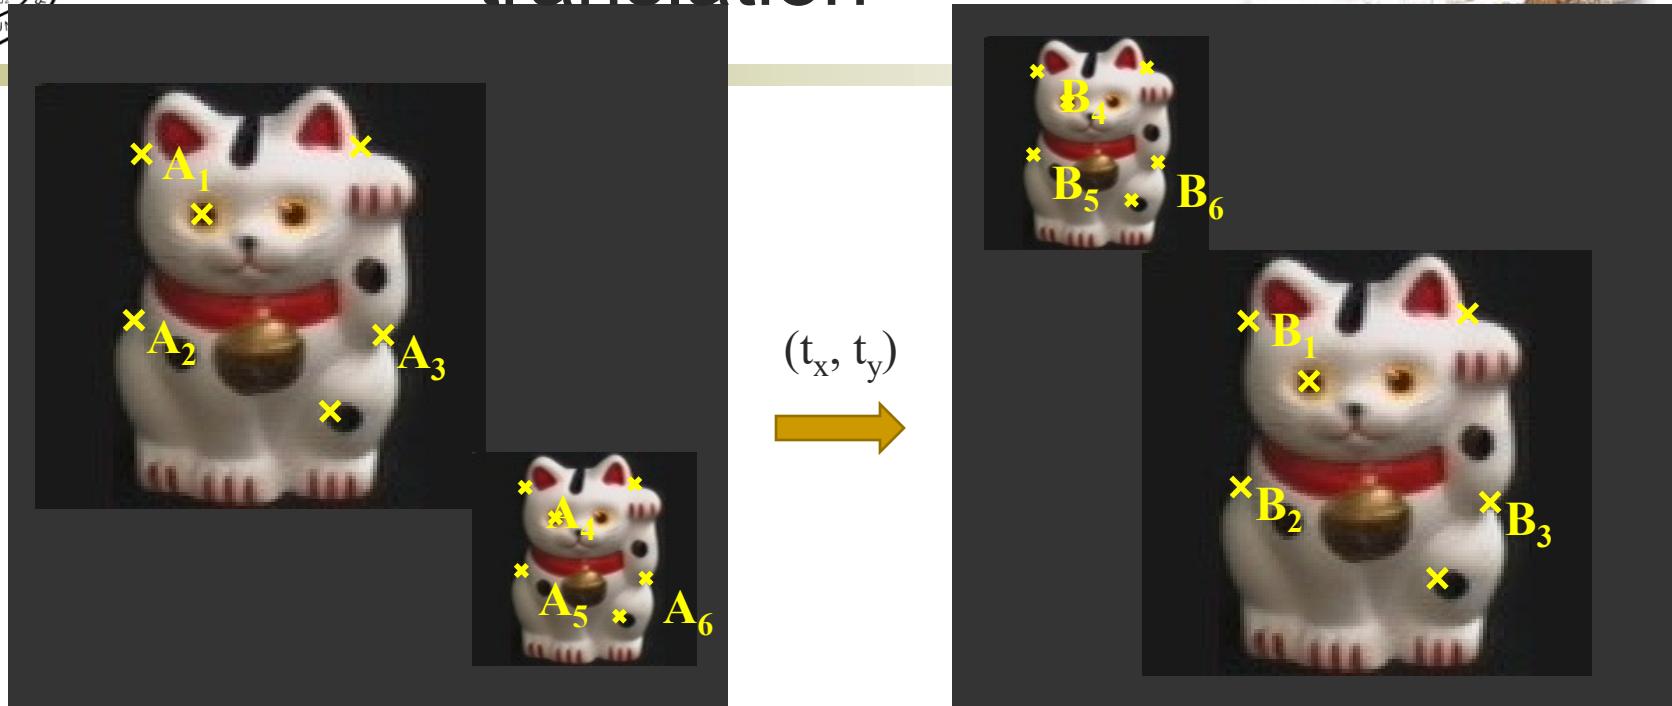
1. Sample a set of matching points (1 pair)
2. Solve for transformation parameters
3. Score parameters with number of inliers
4. Repeat steps 1-3 N times

Problem: outliers

$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$



Example: solving for translation



Problem: outliers, multiple objects, and/or many-to-one matches

Hough transform solution

1. Initialize a grid of parameter values
2. Each matched pair casts a vote for consistent values
3. Find the parameters with the most votes
4. Solve using least squares with inliers

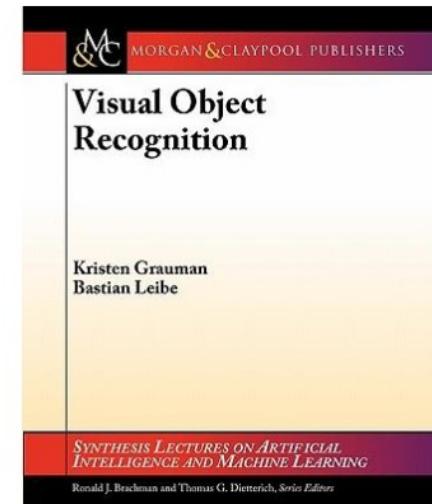
$$\begin{bmatrix} x_i^B \\ y_i^B \end{bmatrix} = \begin{bmatrix} x_i^A \\ y_i^A \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$



References

- We've created a script... for the part of the lecture on object recognition & categorization

➤ K. Grauman, B. Leibe
Visual Object Recognition
Morgan & Claypool publishers, 2011

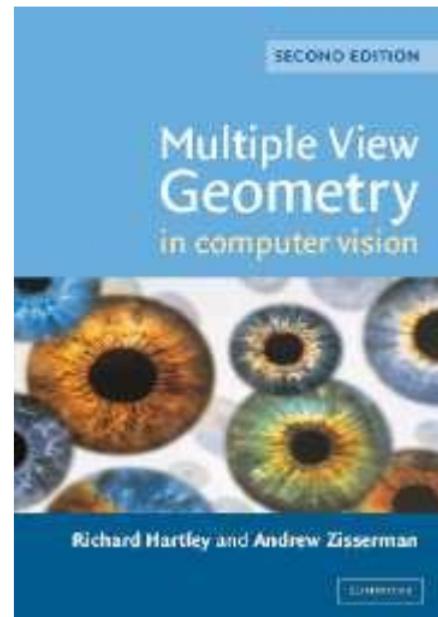


- Chapter 3: Local Feature Extraction ([Last lecture](#))
- Chapter 5: Geometric Verification ([Today](#))



References

- More details on homography estimation can be found in Chapter 4.7 of
 - R. Hartley, A. Zisserman
Multiple View Geometry in Computer Vision
2nd Ed., Cambridge Univ. Press, 2004
- Details about the DoG detector and the SIFT descriptor can be found in
 - D. Lowe, Distinctive image features from scale-invariant keypoints,
IJCV 60(2), pp. 91-110, 2004
- Try the available local feature detectors and descriptors
 - <http://www.robots.ox.ac.uk/~vgg/research/affine/detectors.html#binaries>





Today's class



- Introduction to alignment
- Alignment methods
 - Global methods
 - Hypothesize and test
- Image Transformation
 - Common image transformations
 - Examples of solving image alignment
- Homework: Mosaics



Homework: Mosaics



...

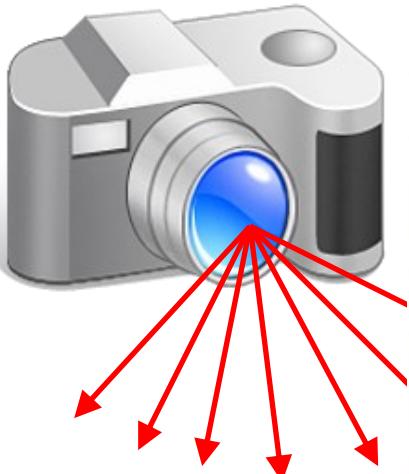
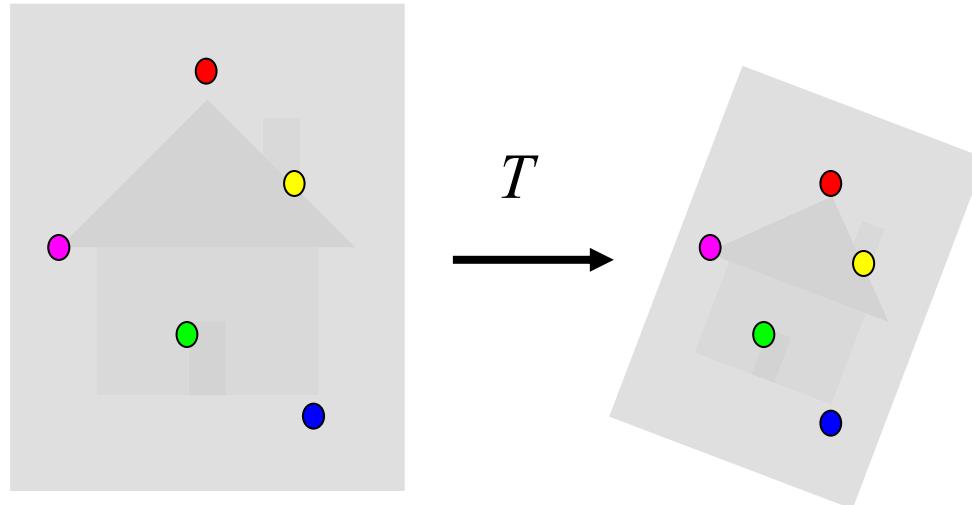


image from S. Seitz

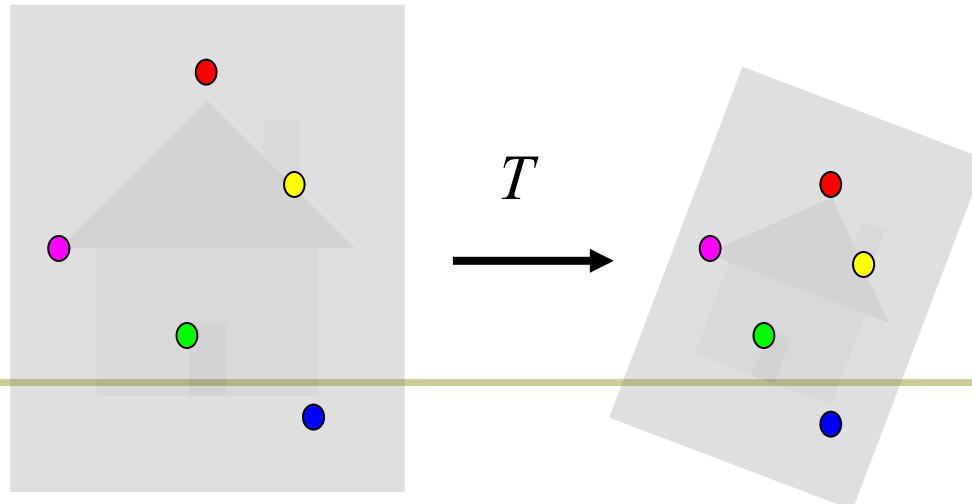
Obtain a wider angle view by combining multiple images.



Main questions



Alignment: Given two images, what is the transformation between them?



Warping: Given a source image and a transformation, what does the transformed output look like?



How to stitch together a panorama (a.k.a. mosaic)?



■ Basic Procedure

- Take a sequence of images from the same position
 - Rotate the camera about its optical center
- Compute transformation (homography) between second image and first using corresponding points.
- Transform the second image to overlap with the first.
- Blend the two together to create a mosaic.
- (If there are more images, repeat)



本科生大作业



- 综述性研究报告：针对一个专门话题，进行深入调研分析，形成一篇高质量的综述。
- 培养：查阅文献，阅读文献，撰写文献的能力，可能启发对某个方向深入研究的兴趣。
- 提交格式：CVPR论文格式，建议英文撰写。
- 调研范围：顶级会议和顶级期刊论文，以及最新的arXiv论文，
- 不要仅仅去看一些公众号、知乎、中文博客
- 提交时间：**6月20号之前（毕业班）**



话题



- Low-level vision: denoise, super resolution etc.
- Edge detection
- Grouping and segmentation
- Local descriptor and image matching
- Deep learning for vision
- Tracking, video analysis
- Vision and language
- Weakly/Self supervised learning
- Etc.